

Image Filtering



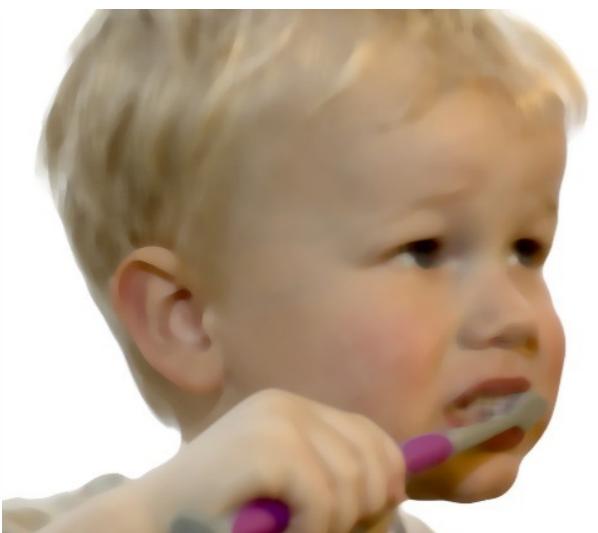
Original



Gradient Magnitude



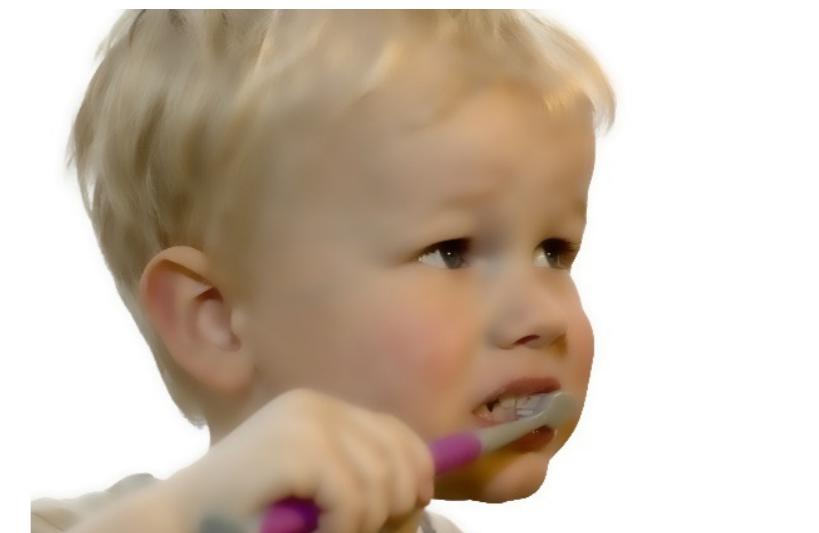
Gaussian Blur



Median



Adaptive Thresholding



Bilateral

Computer Vision
Carnegie Mellon University (Matthew O'Toole)

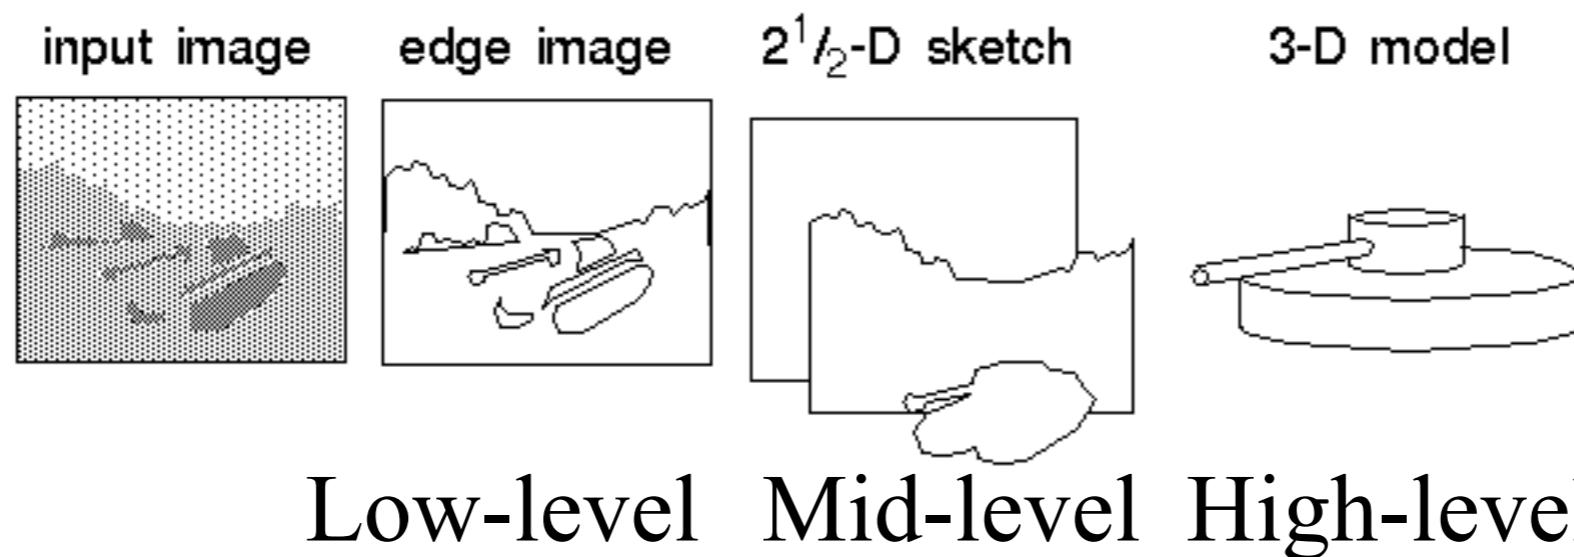
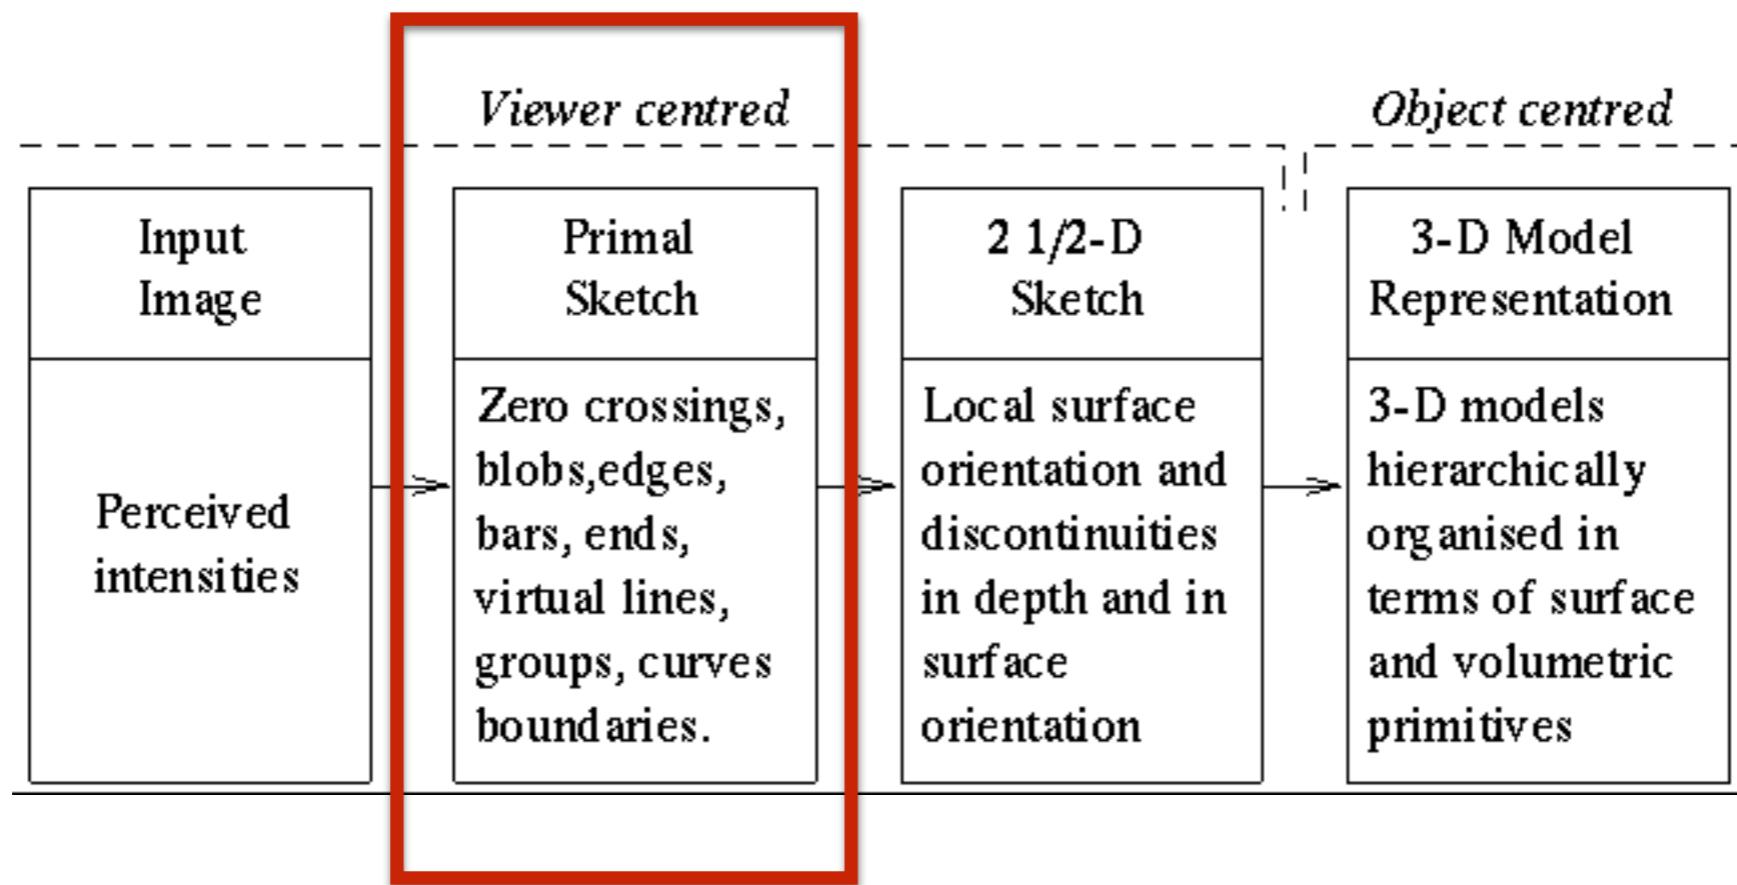
Logistics

- Class website:
 - <https://16820advancedcv.github.io/index.html>
- HW 1 will be released on Monday

Outline

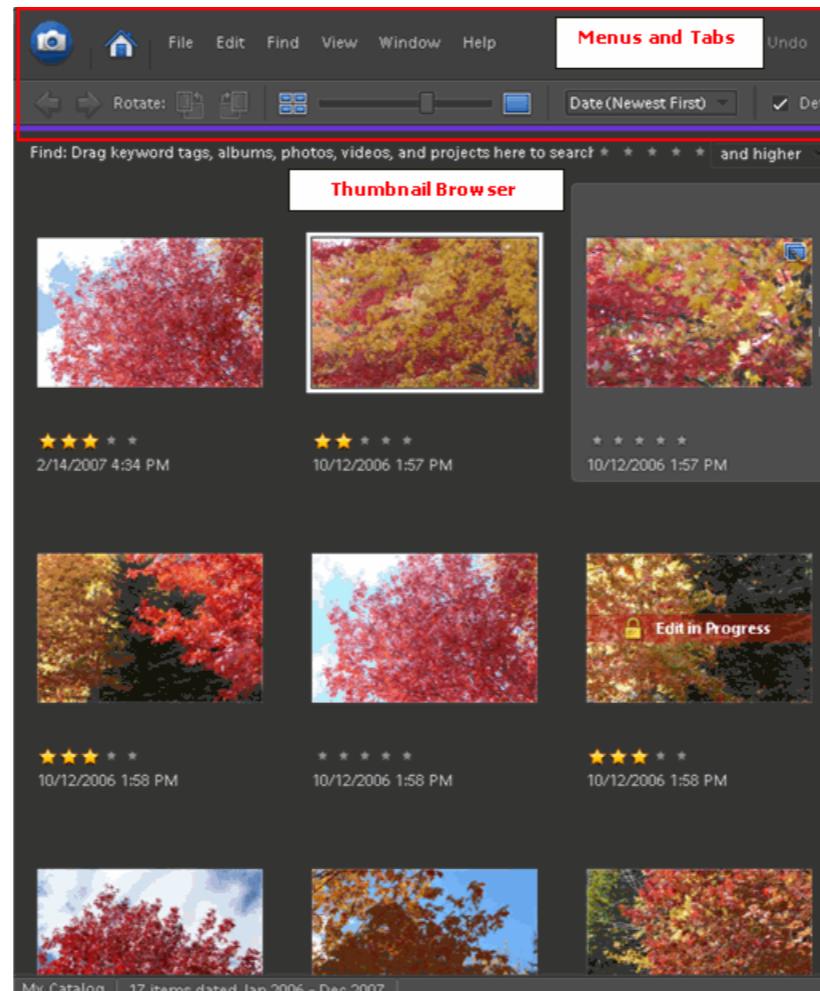
- Convolution
- What can we do with convolution?
 - Filters
 - Edge detection
 - Object detection (Template matching)
 - Convolutional Neural Networks

David Marr



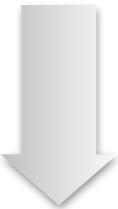
Consider family of low-level image processing operations

Photoshop / Instagram filters: blur, sharpen, colorize, etc....



What kind of image transformations can we perform?

Filtering



changes the pixel values

Warping



changes the pixel location

What kind of image transformations can we perform?

Filtering

 F 

$$G(x) = h\{F(x)\}$$

 G 

changes the **range** of image

Warping



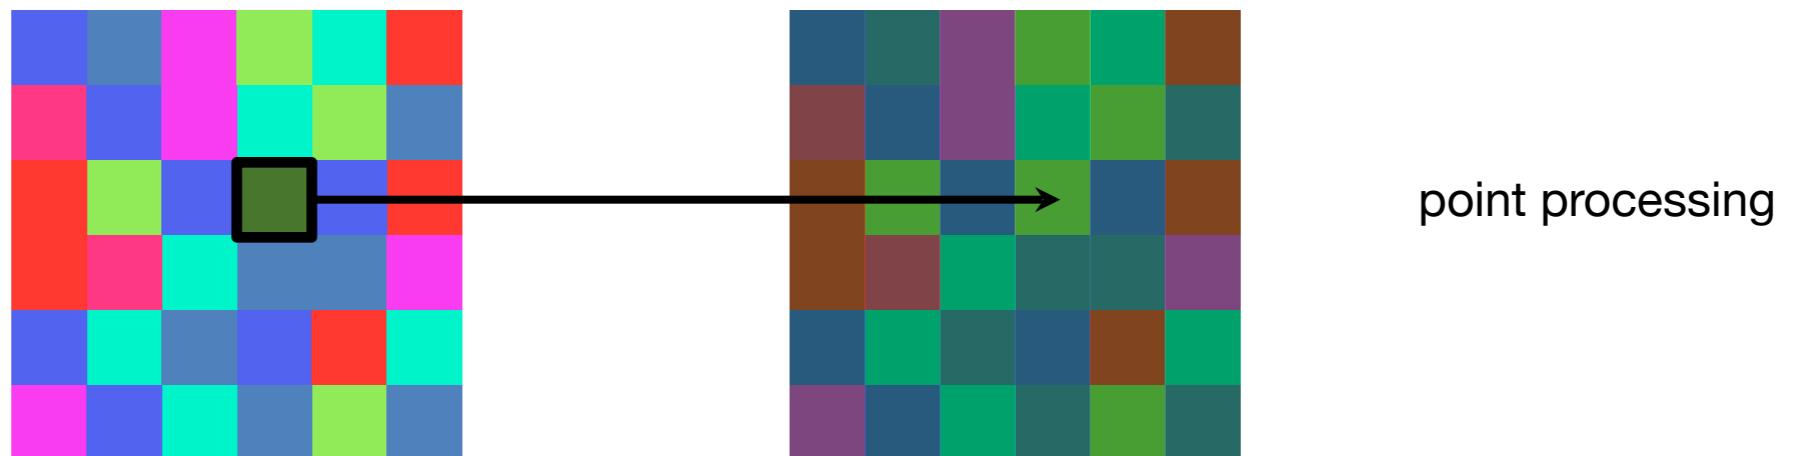
$$G(x) = F(h\{x\})$$



changes the **domain** of image

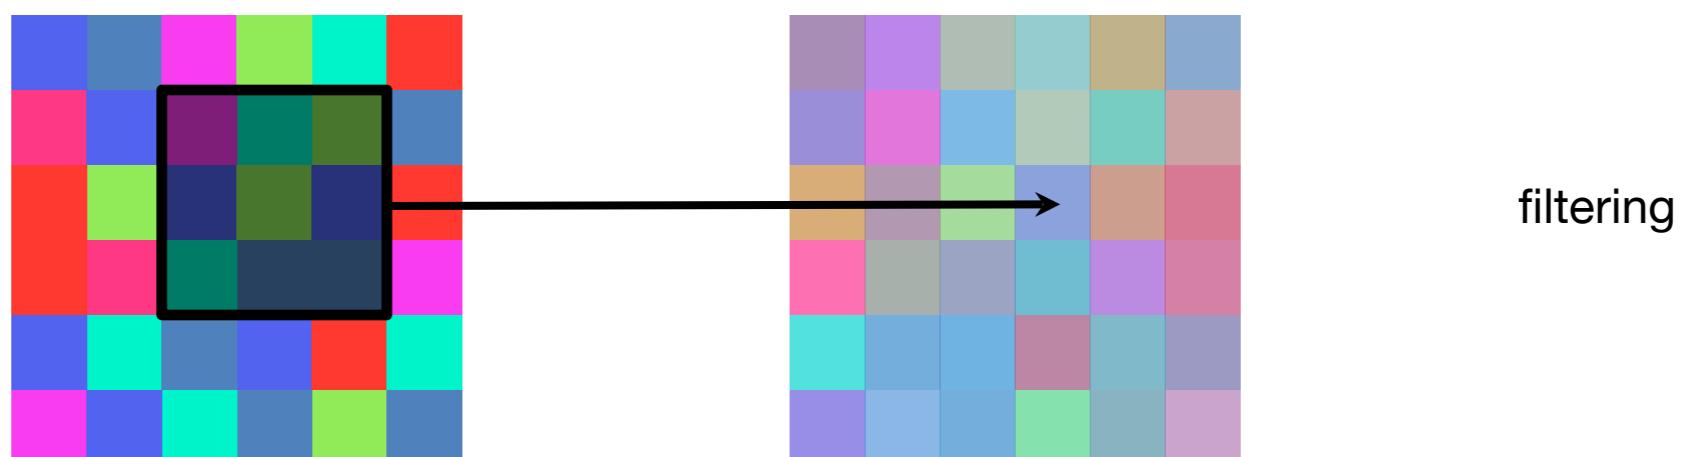
What kind of **image filtering** can we perform?

Point Operation



point processing

Neighborhood Operation



filtering

Examples of Point Processing



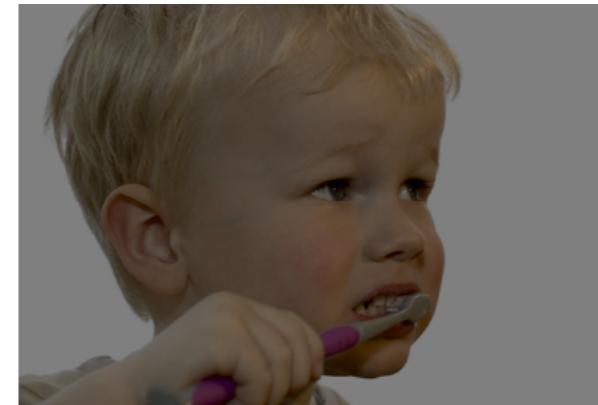
Original

$$x$$



Darken

$$x - 128$$



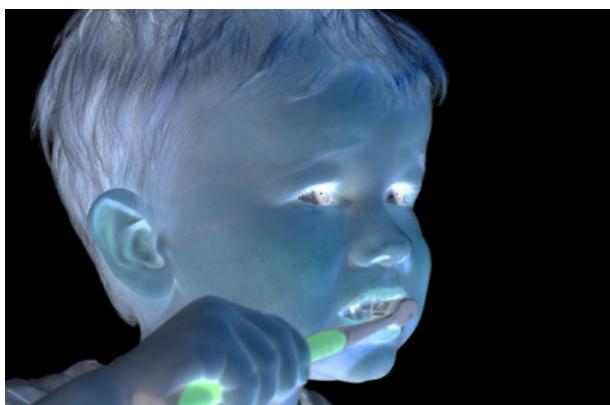
Lower Contrast

$$\frac{x}{2}$$



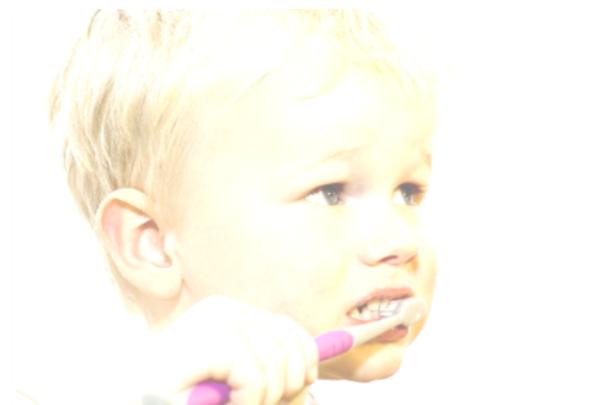
Nonlinear Lower Contrast

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$



Invert

$$255 - x$$



Lighten

$$x + 128$$



Raise Contrast

$$x \times 2$$



Nonlinear Raise Contrast

$$\left(\frac{x}{255}\right)^2 \times 255$$

Examples of filtering

Is there a mathematical way to characterize these operations?



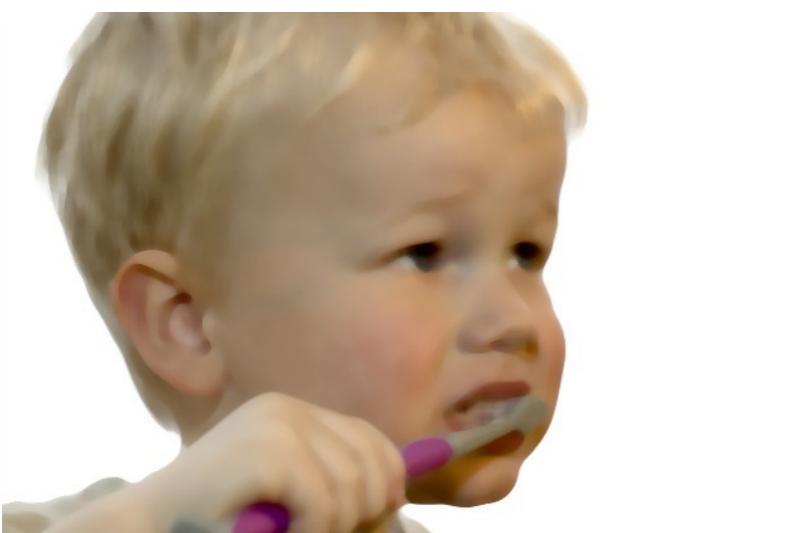
Original



Gradient Magnitude



Gaussian Blur



Median



Adaptive Thresholding



Bilateral

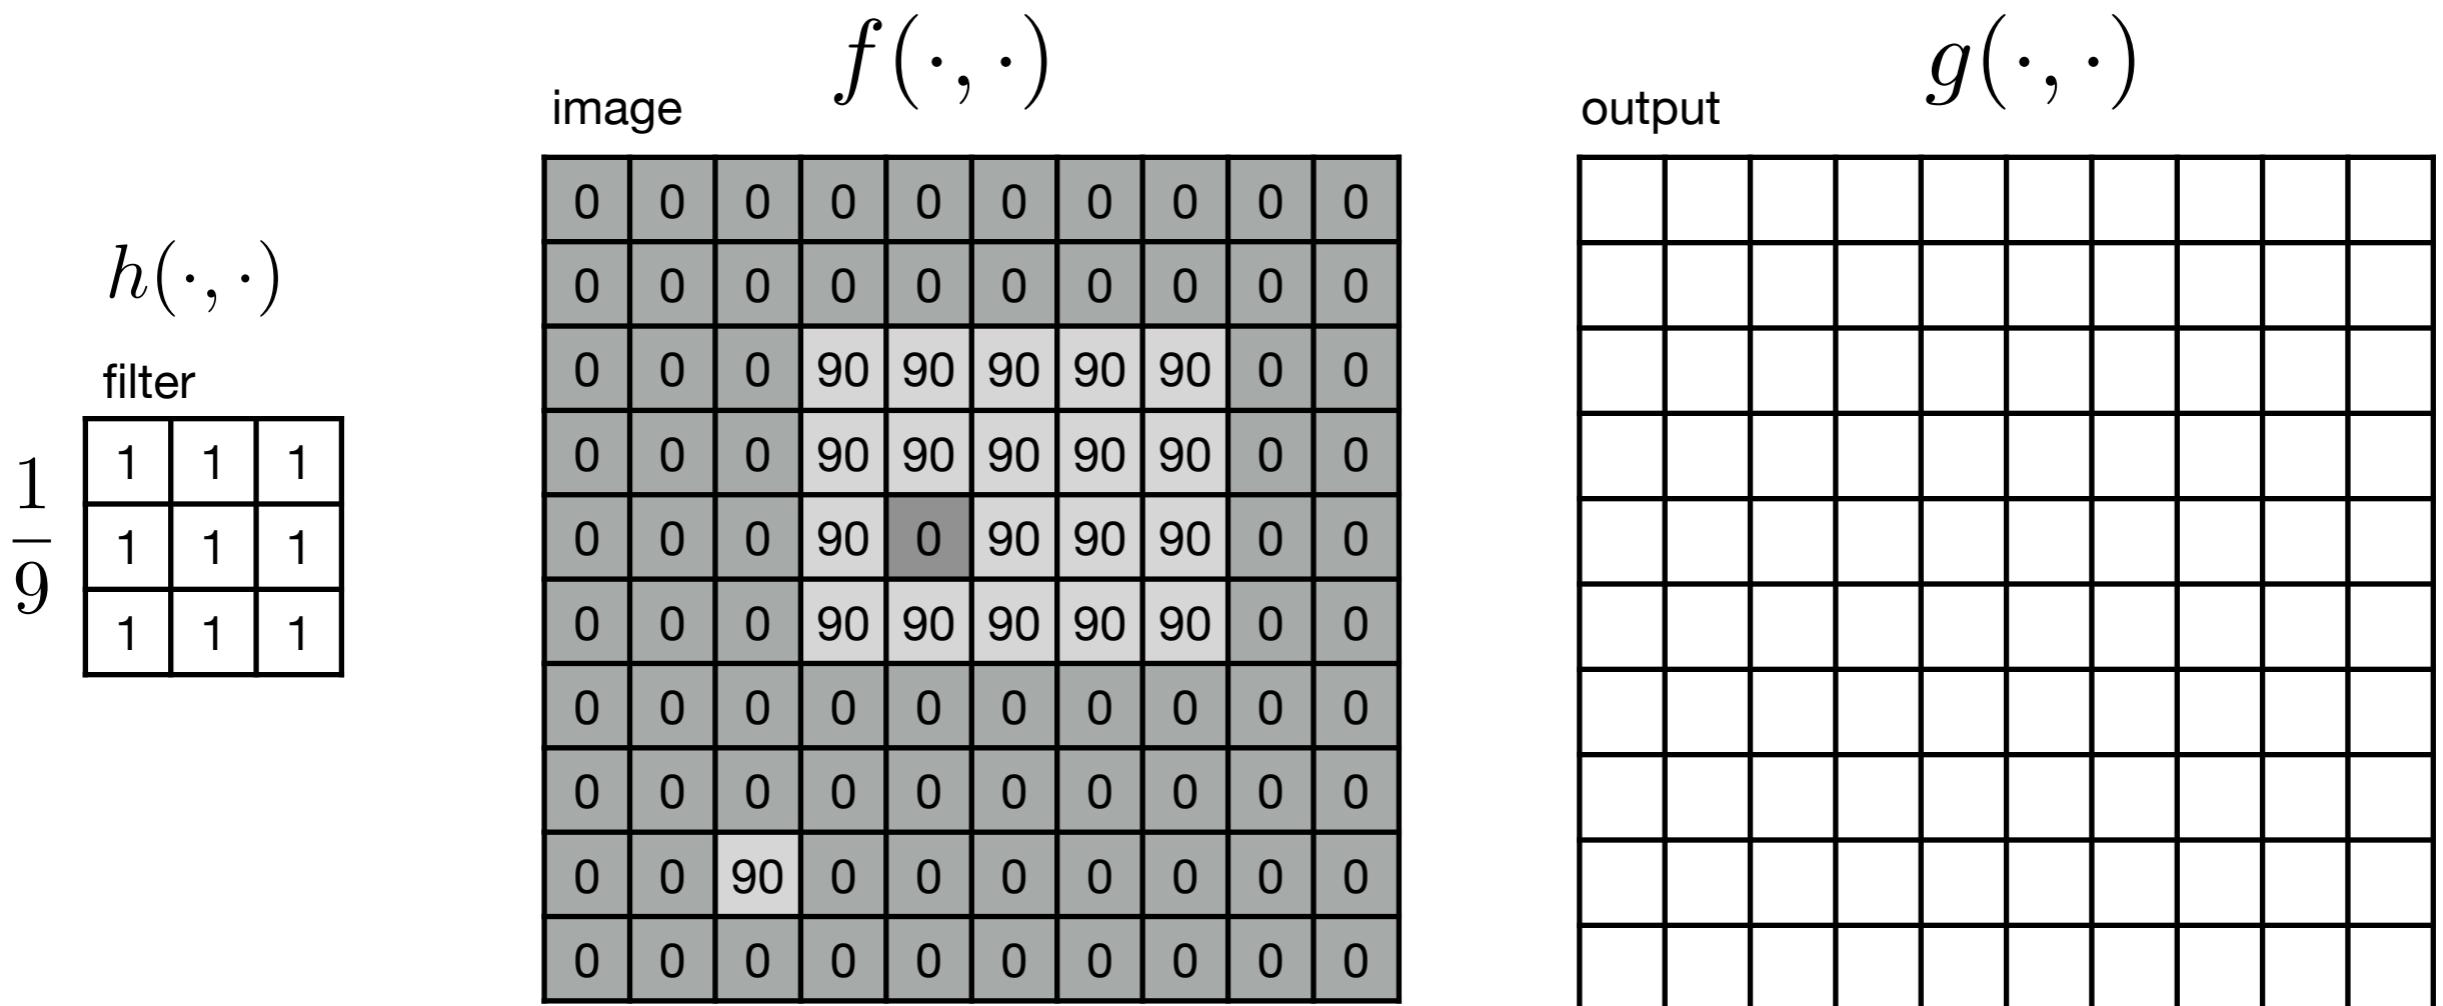
The ‘Box’ filter

$$h(\cdot, \cdot) = \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

replaces pixel with local average

has a smoothing effect

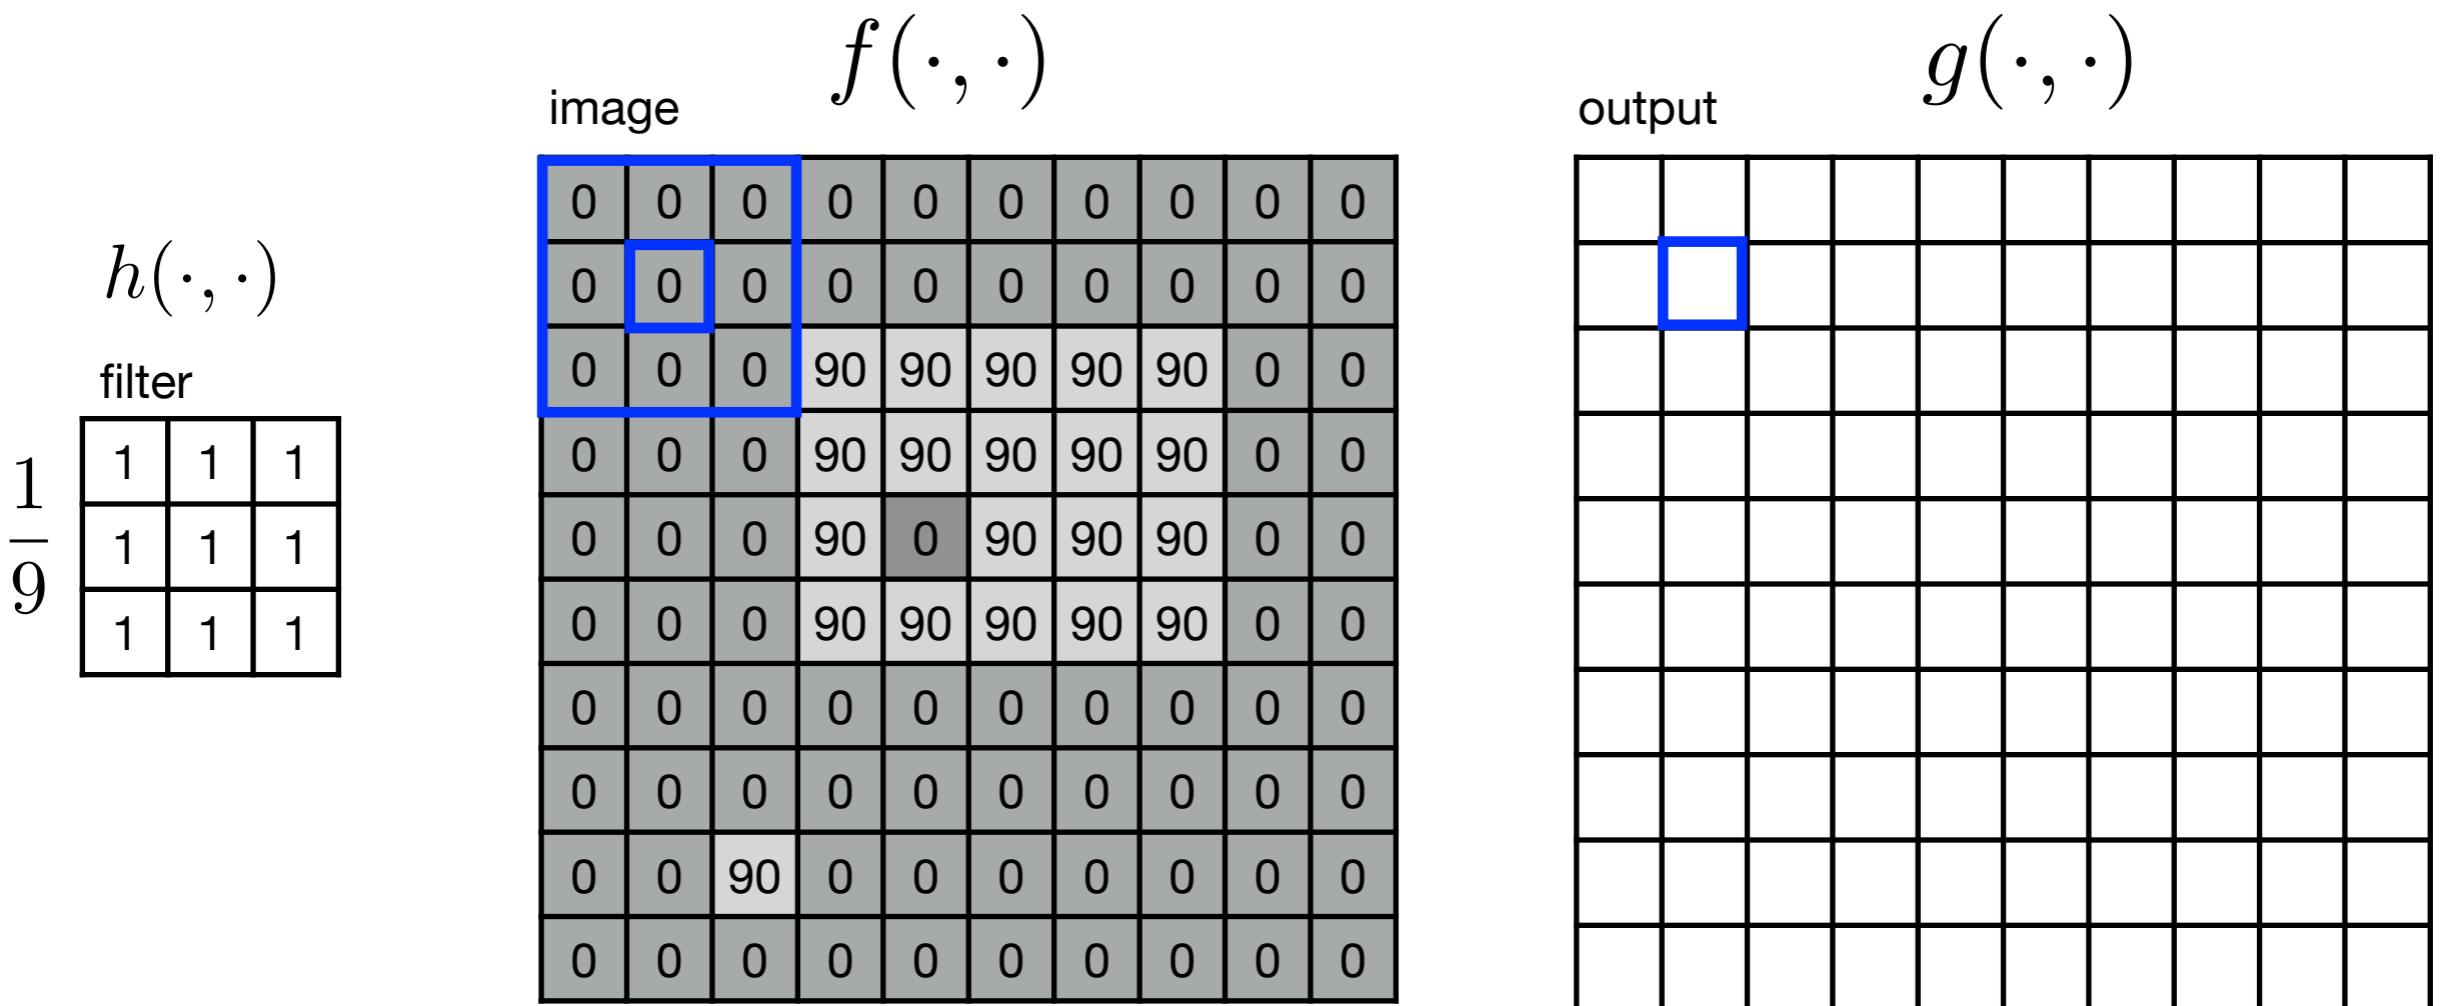
Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

image filter filter image filter image
 (flips the filter)

Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

(flips the filter)

Convolution: $F * H$
 image filter

$g(1, 1)$

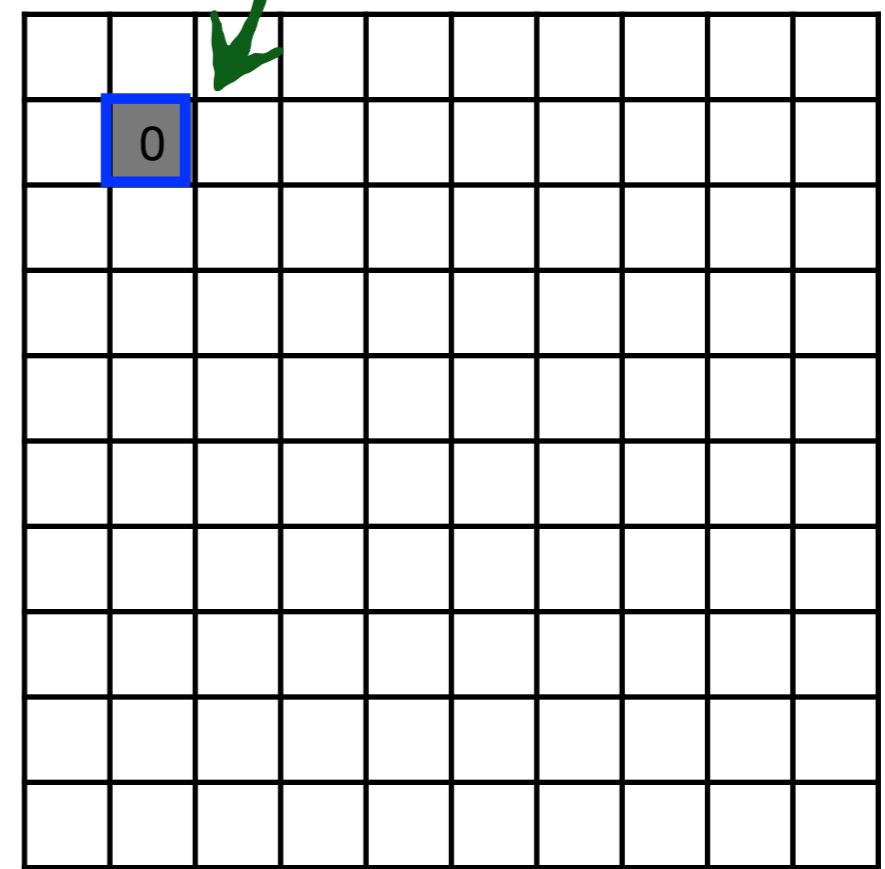
output

$g(\cdot, \cdot)$

$$h(\cdot, \cdot) \\ \text{filter} \\ \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image $f(\cdot, \cdot)$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$$G[i, j] = F * H = H * F = \sum_{u} \sum_{v} H[u, v] F[i - u, j - v]$$

output image filter filter image filter image

(flips the filter)

Convolution: $F * H$
 image filter

$$h(\cdot, \cdot)$$

filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image

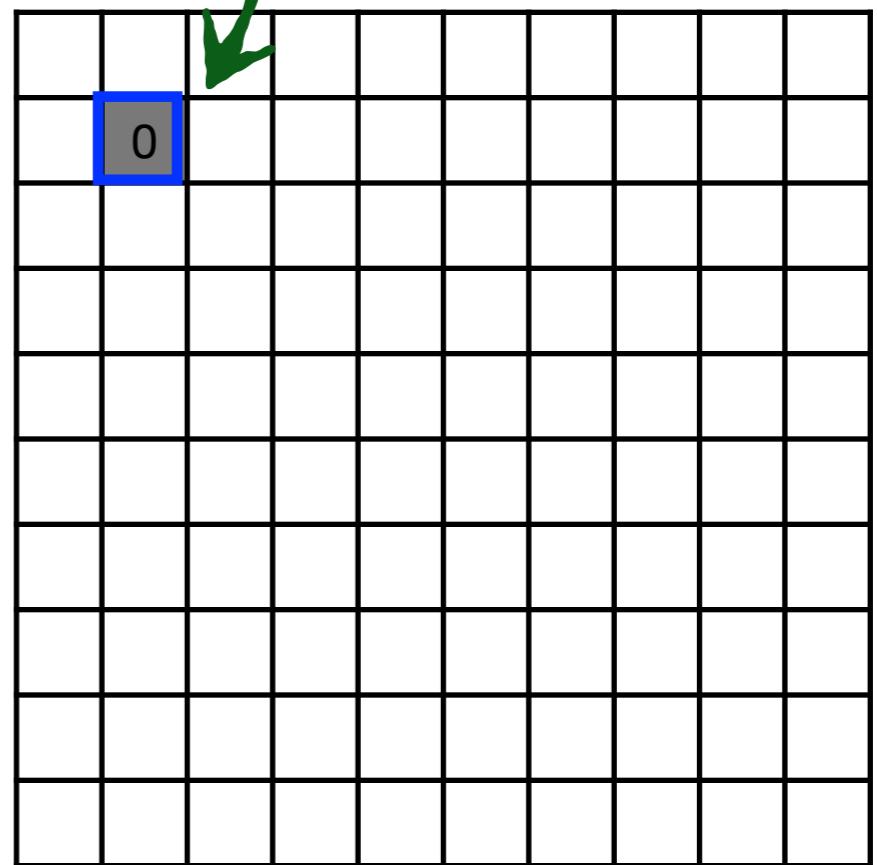
$$f(\cdot, \cdot)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g(1, 1)$

output

$g(\cdot, \cdot)$



$$G[i, j] = F * H = H * F = \sum_u \sum_v H[u, v] F[i - u, j - v]$$

output image filter filter image filter image

(flips the filter)

Convolution: $F * H$
 image filter

$$h(\cdot, \cdot)$$

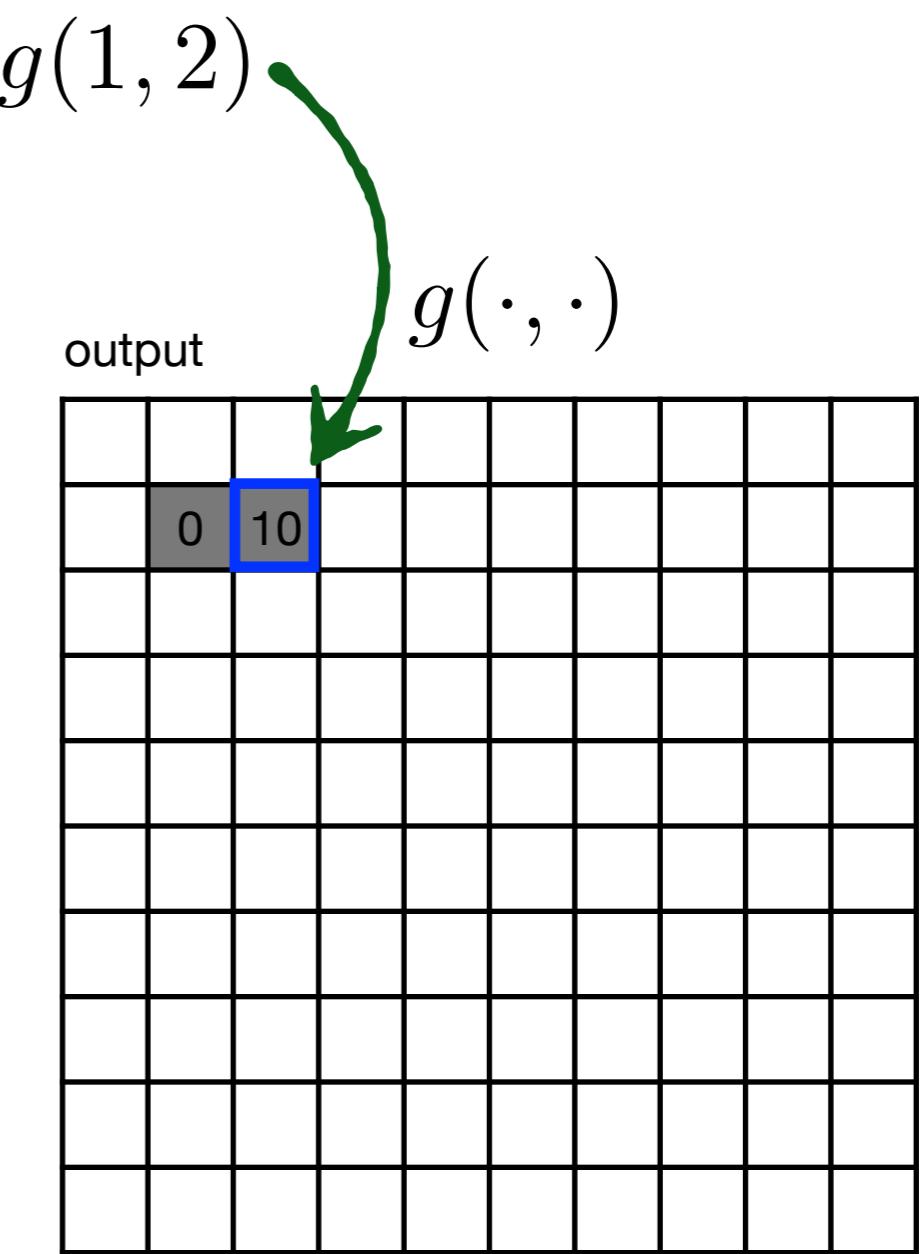
filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image

$$f(\cdot, \cdot)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

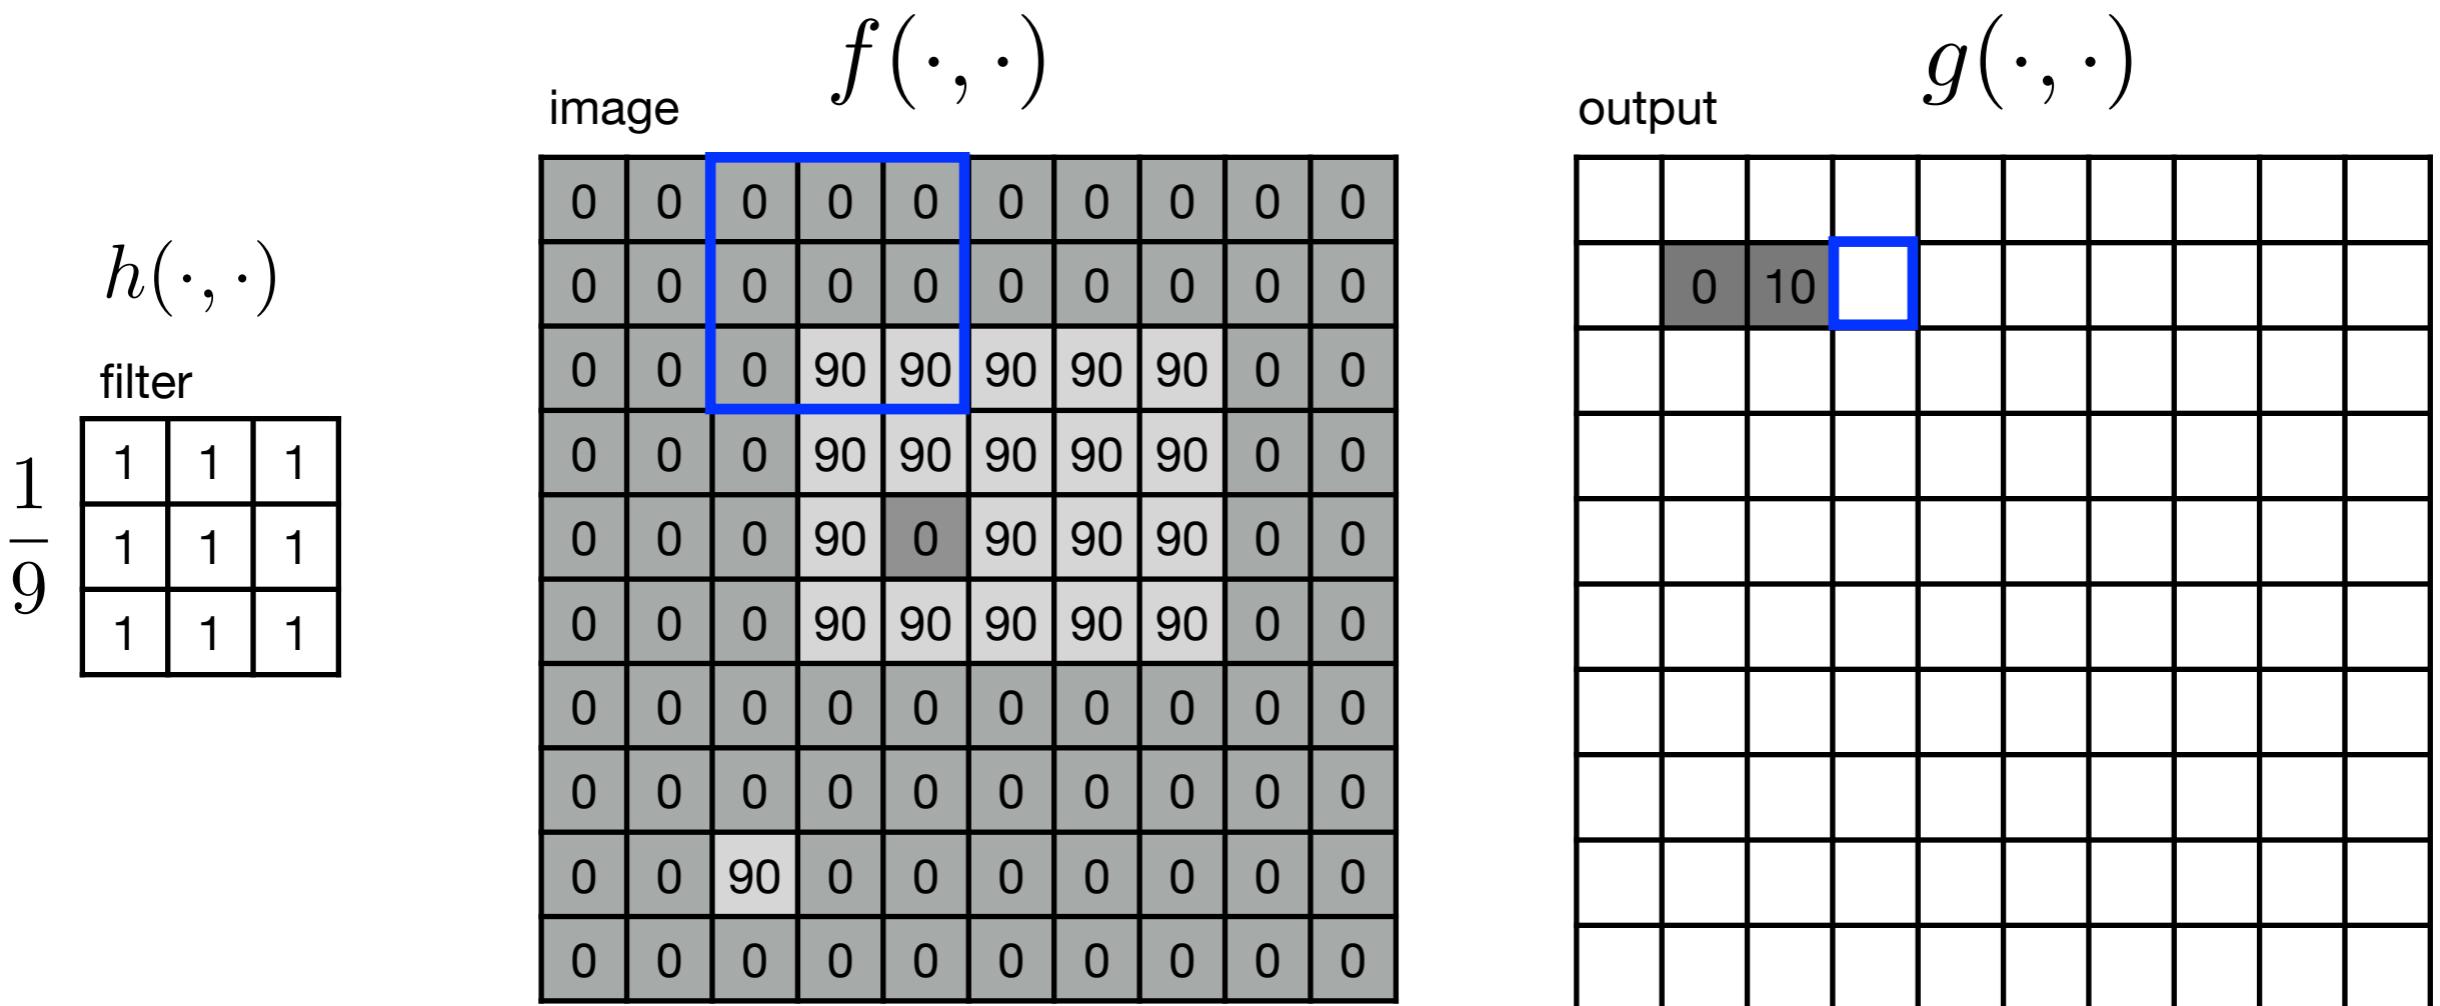


$$G[i, j] = F * H = H * F = \sum_{u} \sum_{v} H[u, v] F[i - u, j - v]$$

output image filter filter image filter image

(flips the filter)

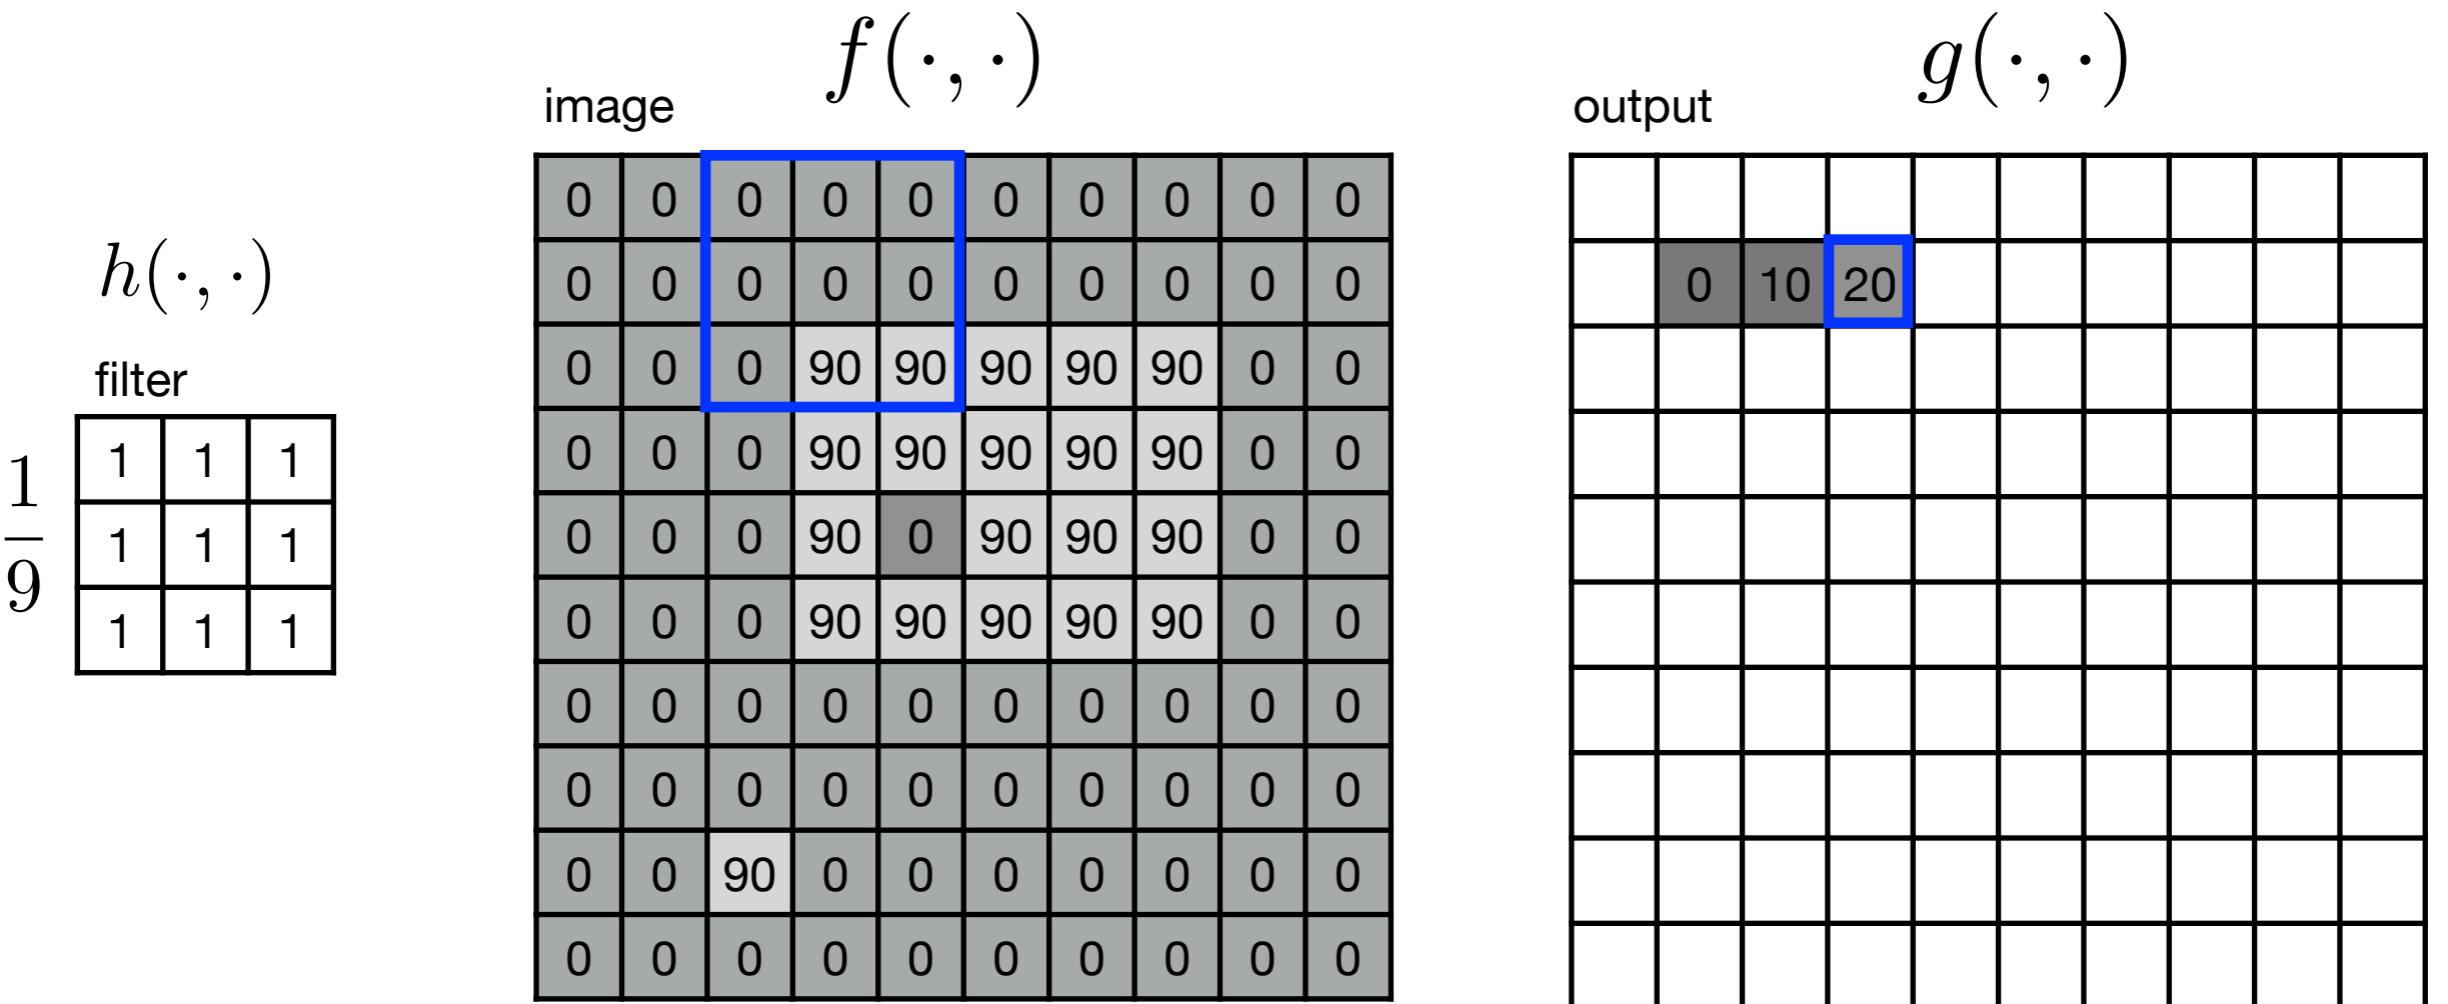
Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{u} \sum_{v} H[u, v] F[i - u, j - v]$$

output image filter filter image filter image
 (flips the filter)

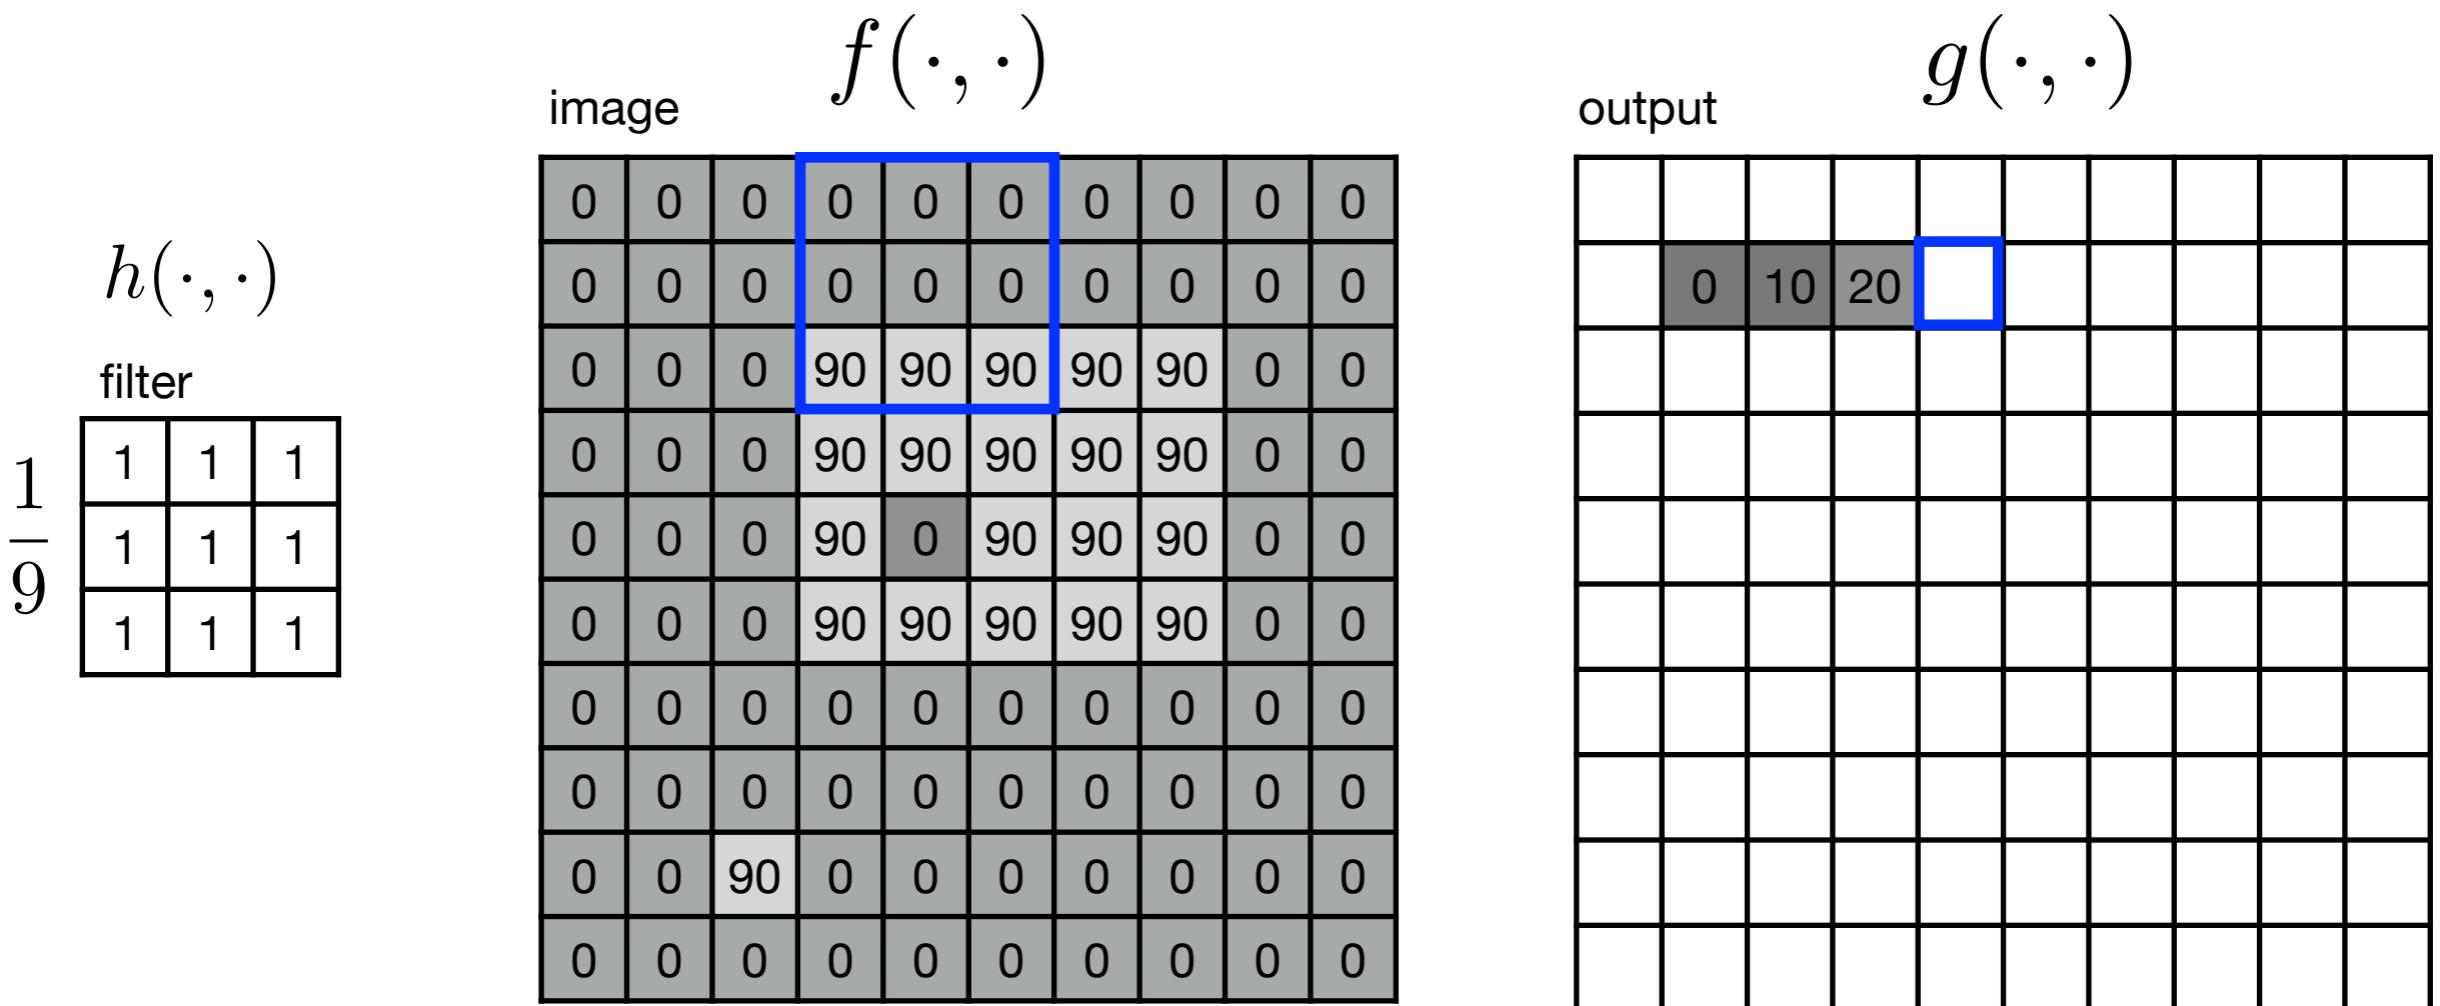
Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{u} \sum_{v} H[u, v] F[i - u, j - v]$$

output image filter filter image filter image
 (flips the filter)

Convolution: $F * H$
 image filter

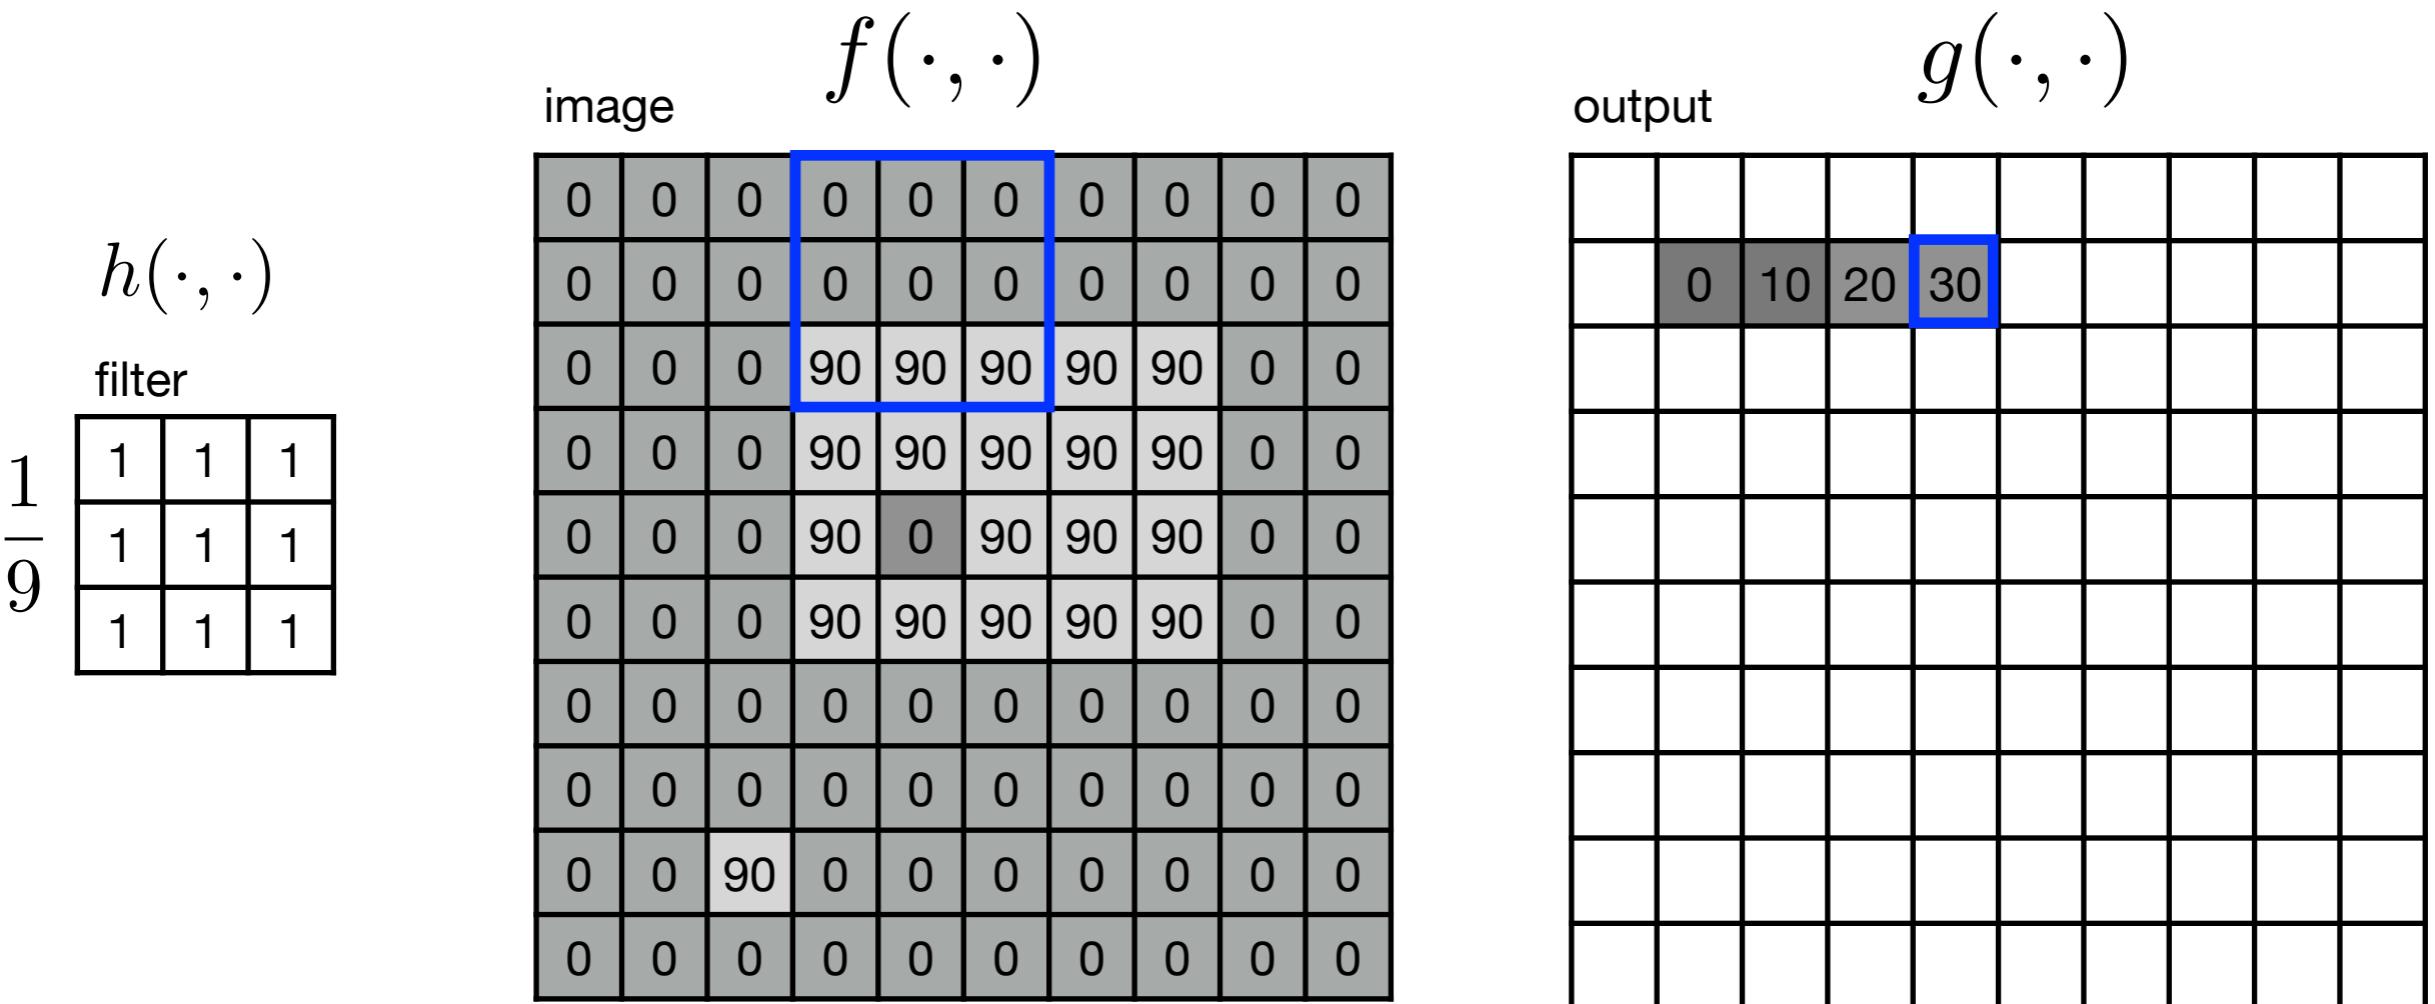


$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

image filter filter image filter image

(flips the filter)

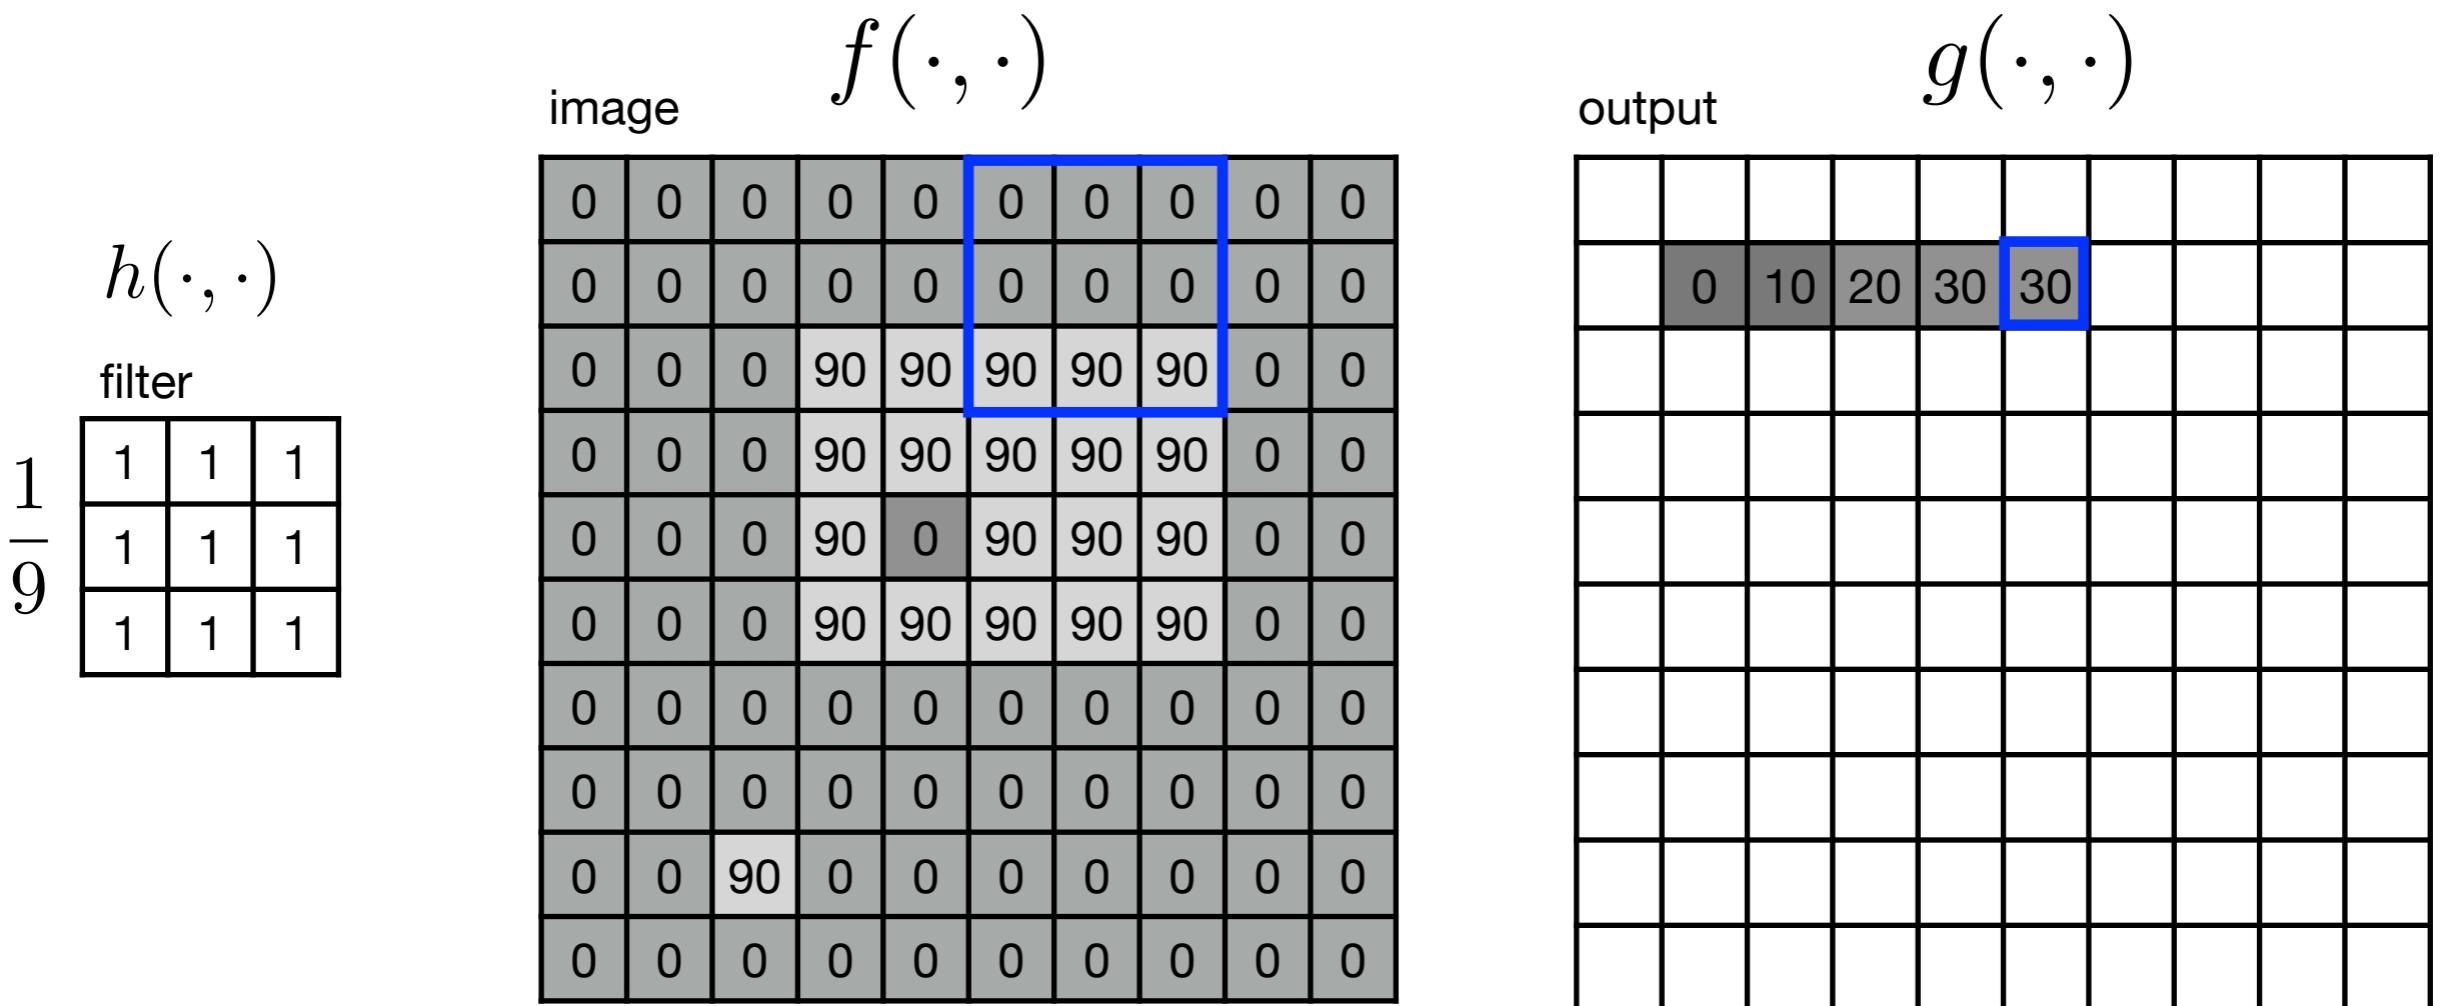
Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

image filter filter image filter image
 (flips the filter)

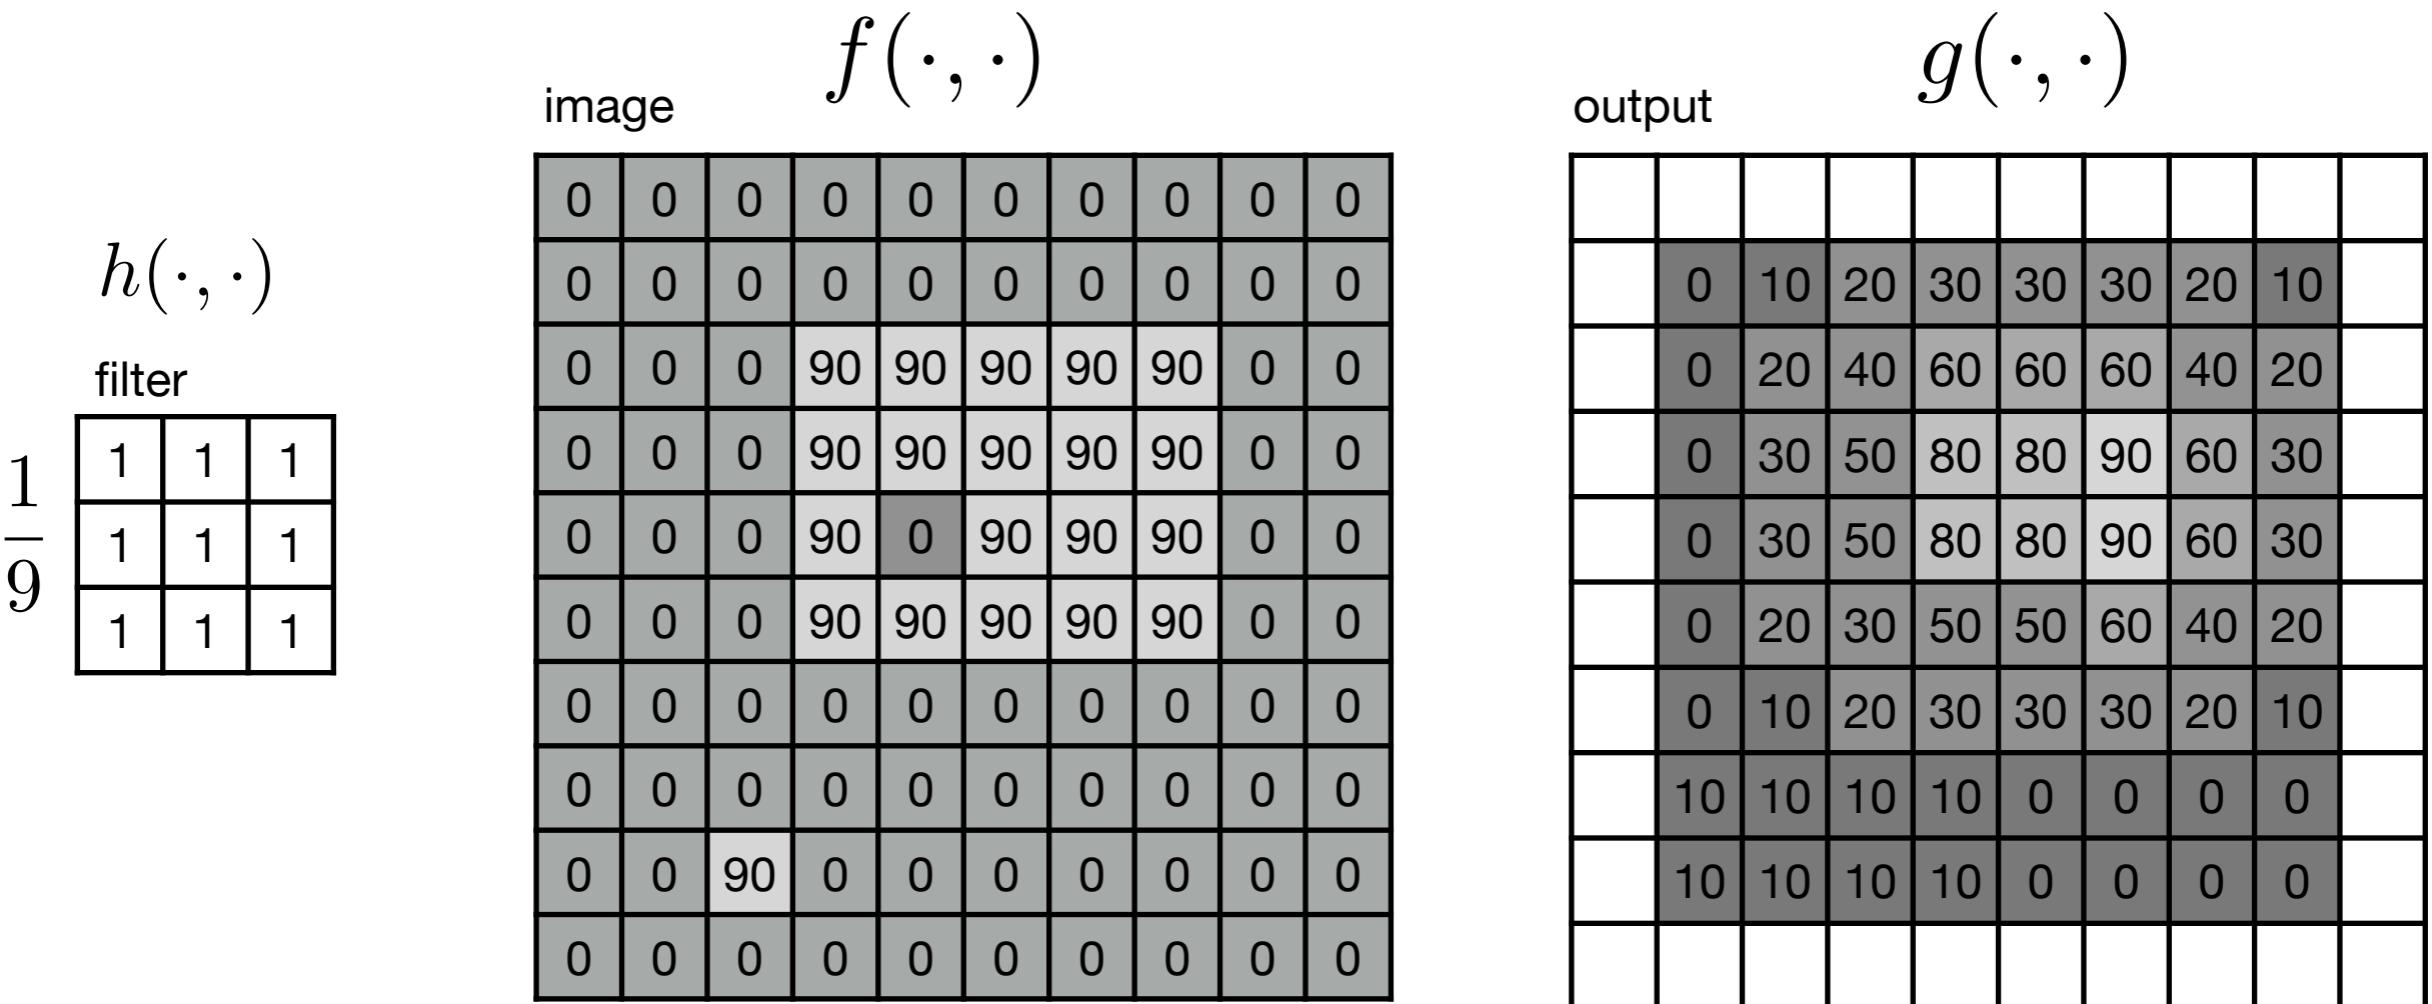
Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

filter image filter filter image filter image
 (flips the filter)

Convolution: $F * H$
 image filter

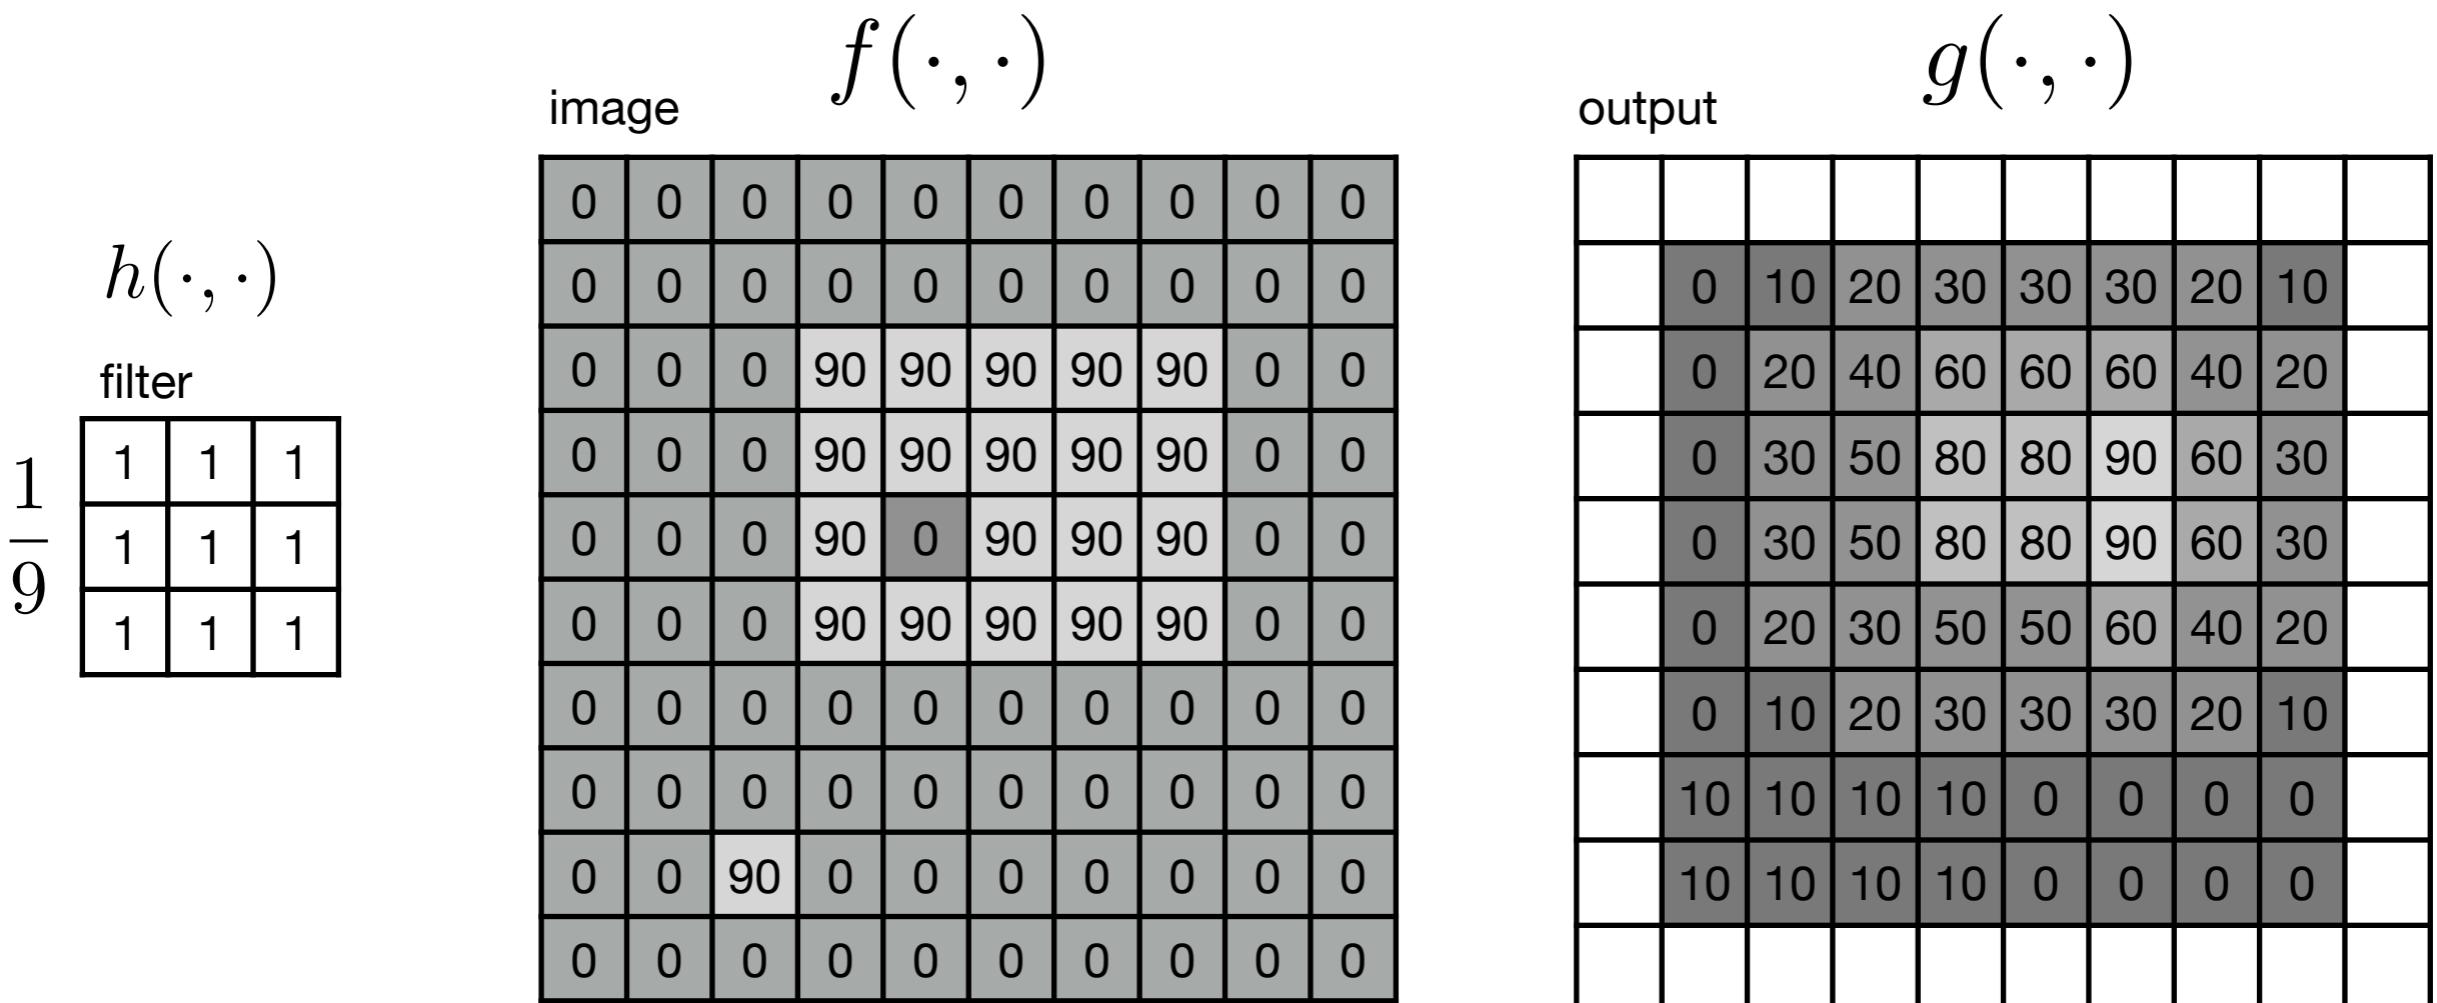


$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

image filter filter image

(flips the filter)

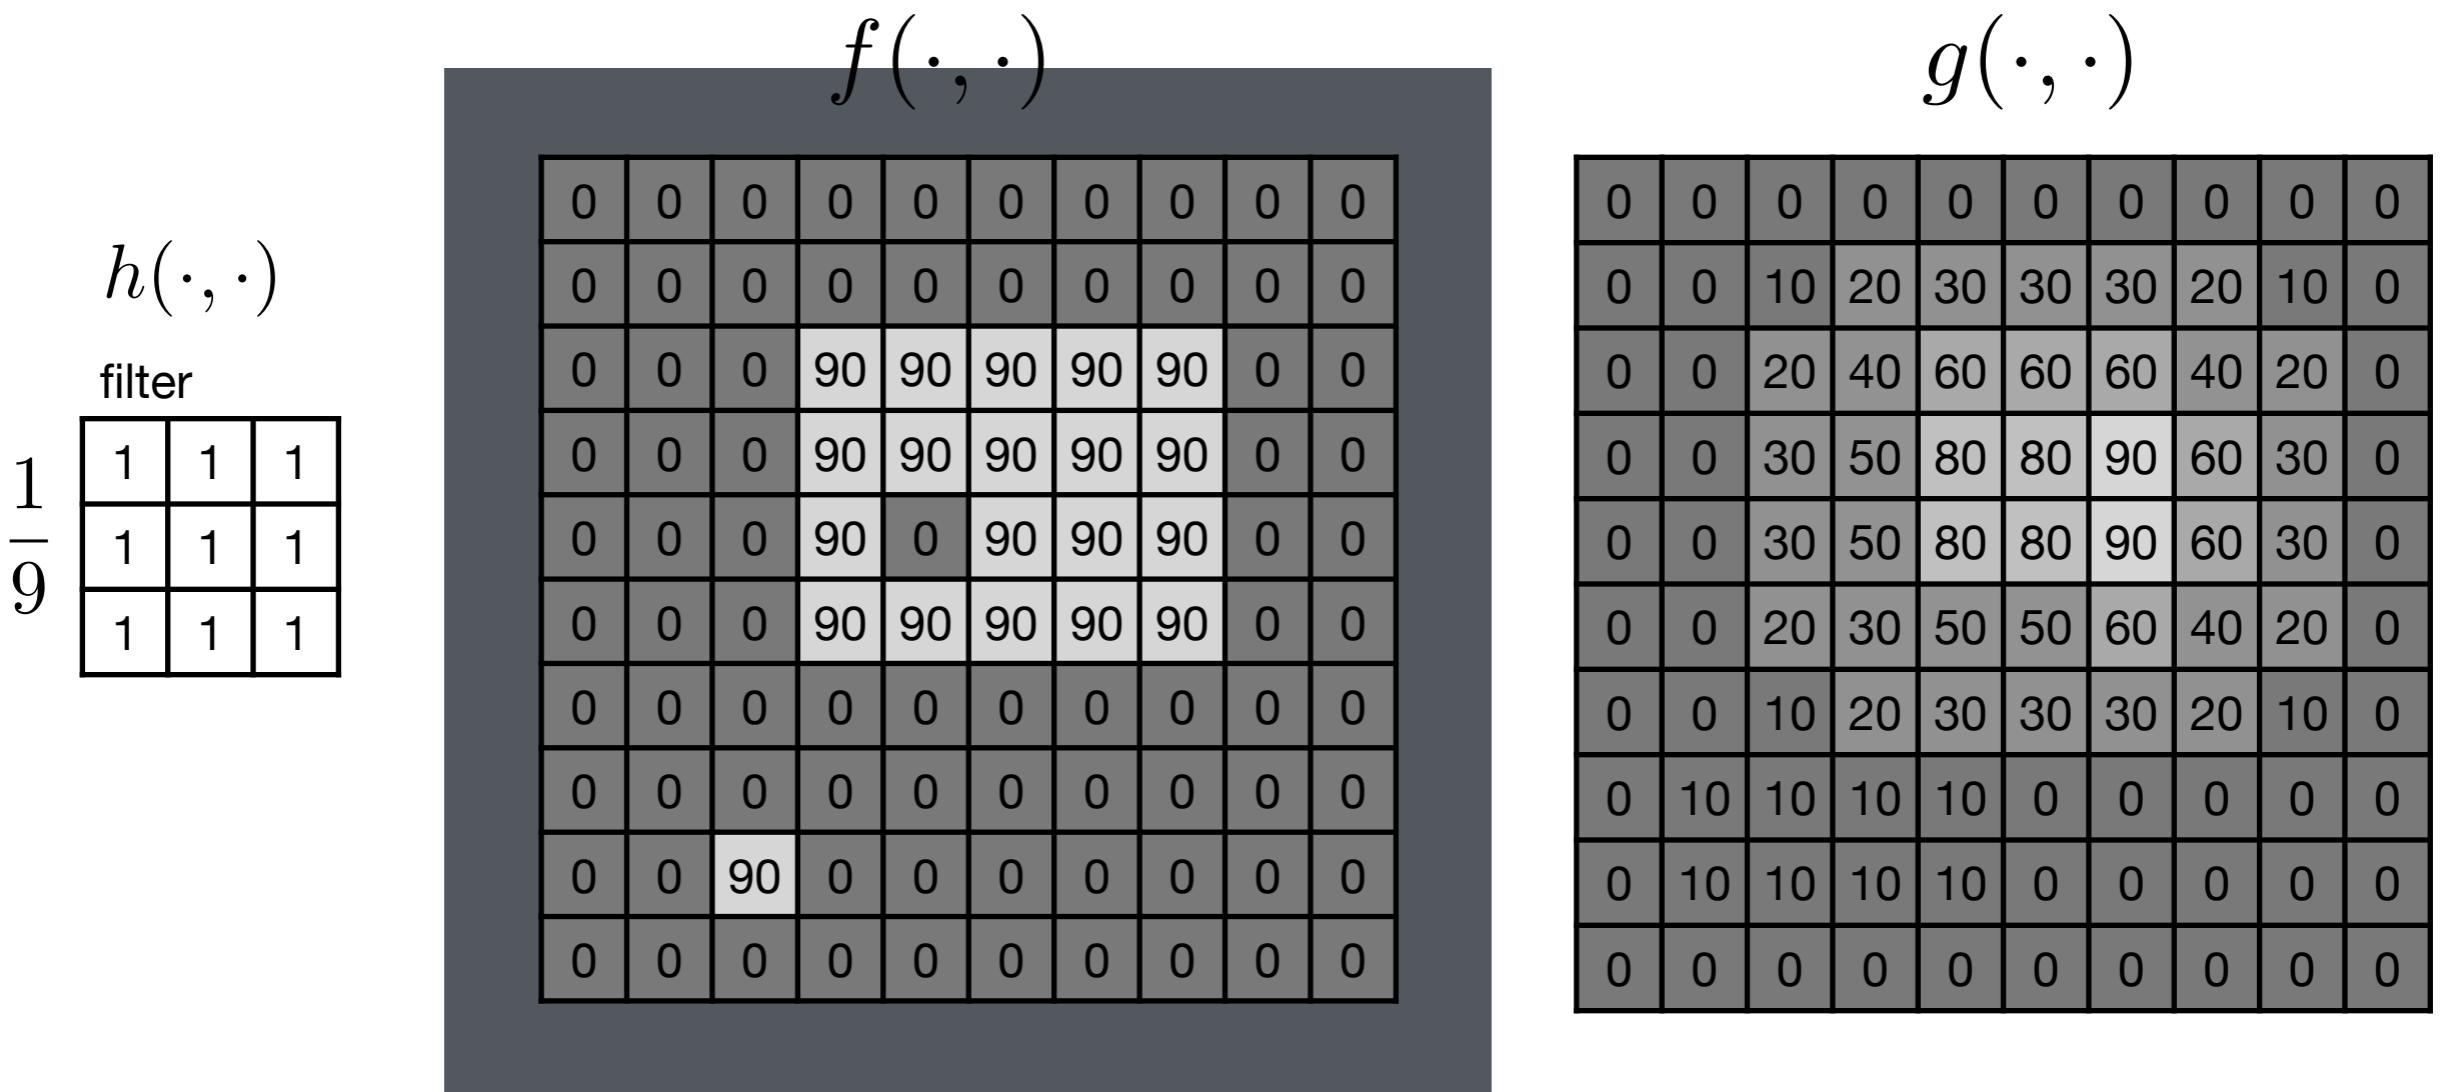
Convolution: $F * H$
 image filter



What about the borders?

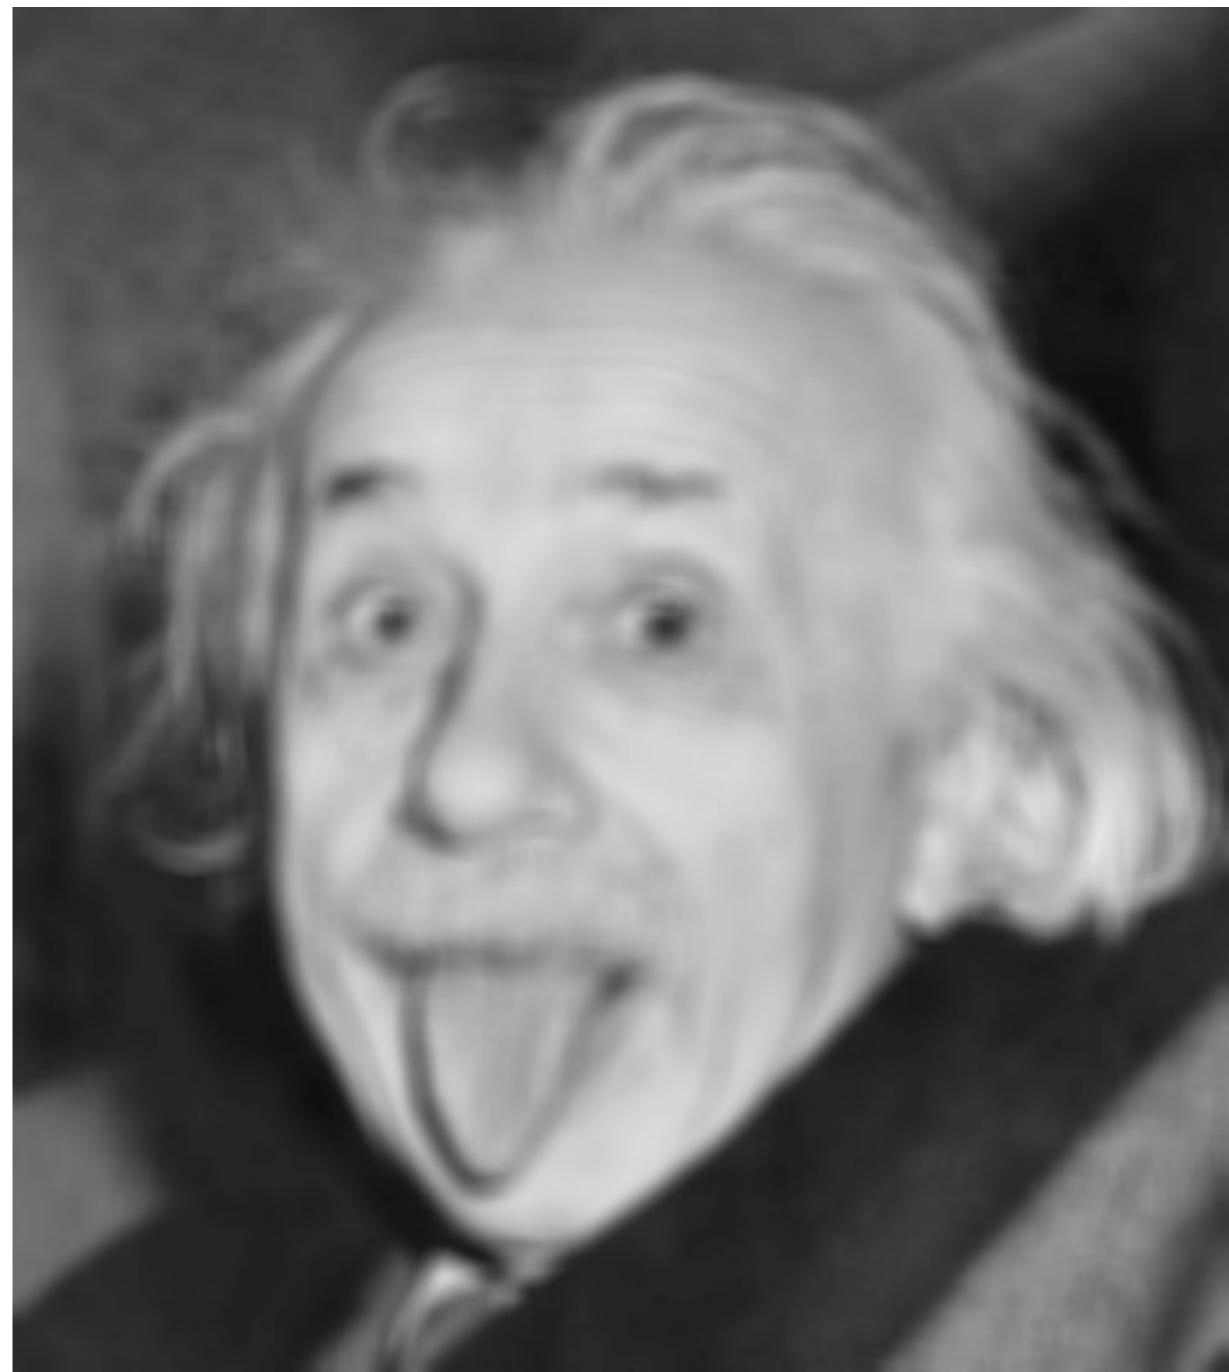
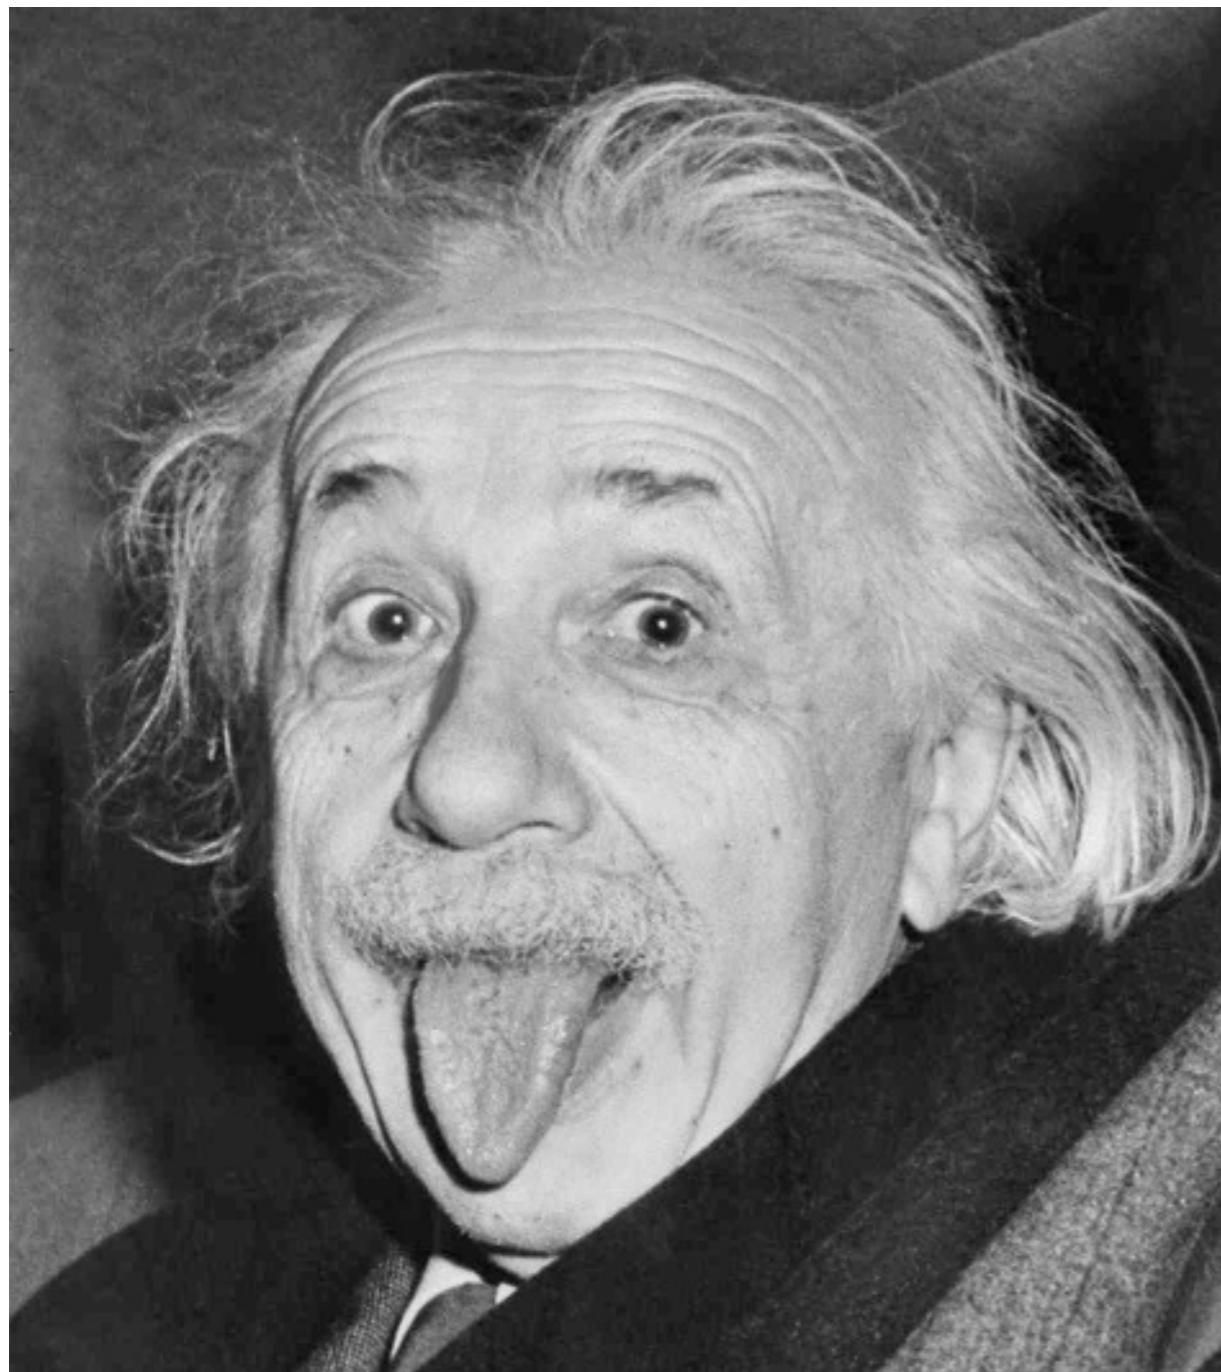
$$G[i, j] = F * H = H * F = \sum_{\text{filter}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

Convolution: $F * H$
 image filter



$$G[i, j] = F * H = H * F = \sum_{\text{output}} \sum_{\text{image}} H[u, v] F[i - u, j - v]$$

image filter filter image filter image





Properties of convolution

$$F * H = H * F$$

Commutative
(Swap the filter with the image)

$$(F * H) * G = F * (H * G)$$

First filter Second filter Combined Filter

Associative

$$(F * G) + (H * G) = (F + H) * G$$

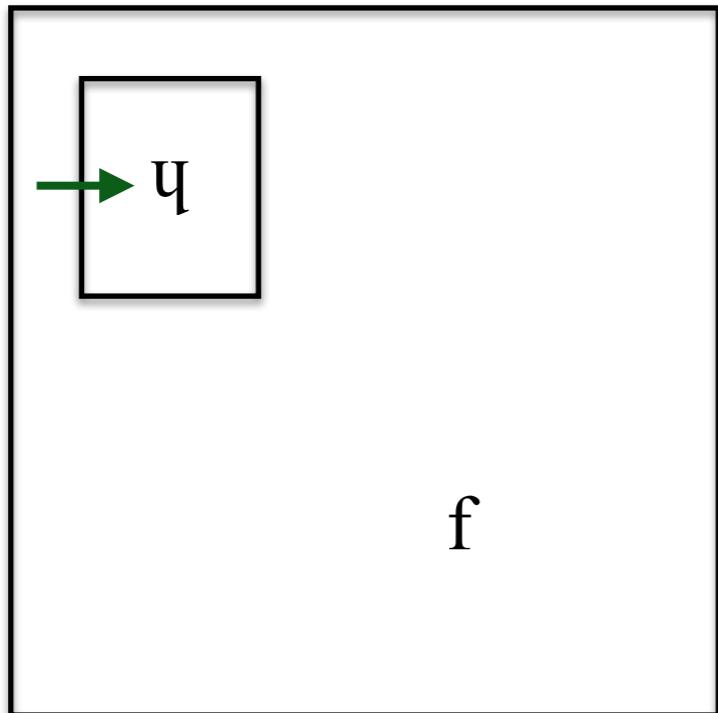
First conv Second conv Combined image One conv

Distributive

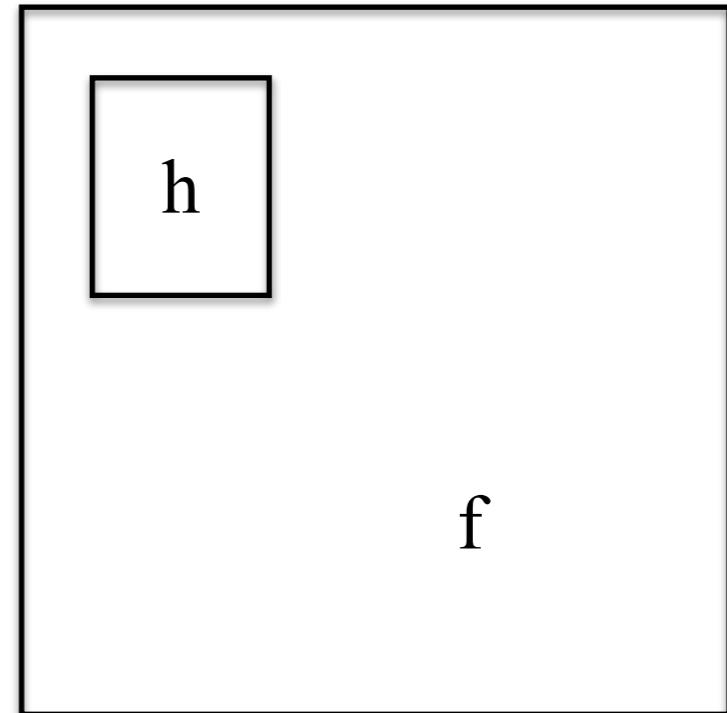
Sometimes allows us to simplify filter operations

Convolution vs Correlation

Filter / Kernel



Image



Image

Why do we have convolutional neural networks
instead of correlational neural networks?

Interesting fact: convolution is commutative and associative
Correlation is NOT!

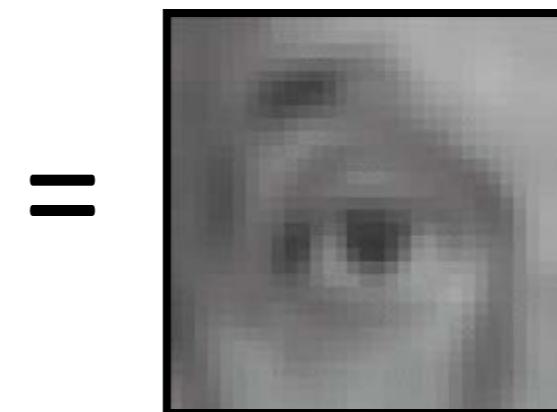
Examples of convolution filters



Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

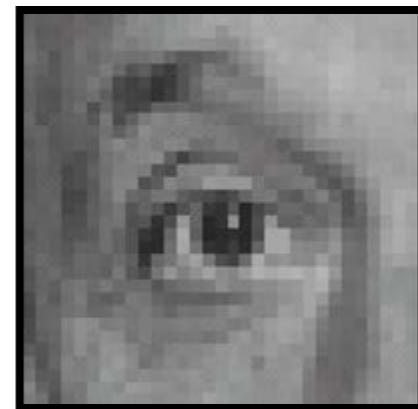
box blur /
box filter



Blur (with a mean
filter)

Why do we divide by 9?

Examples of convolution filters



Original

0	0	0
0	1	0
0	0	0

?

What do you think will happen here?

Examples of convolution filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Examples of correlation filters



0	0	0
0	0	1
0	0	0

?

Original

What do you think will happen here?

Examples of correlation filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Examples of convolution filters



1	2	1
2	4	2
1	2	1

/16



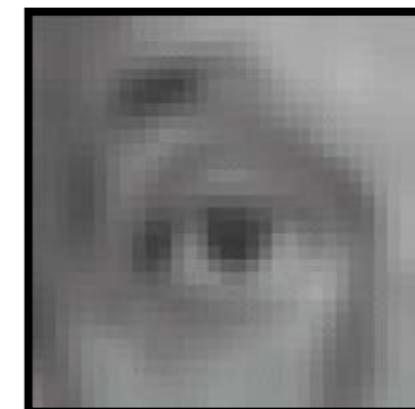
Original

Gaussian filter

How does this compare to a box filter?


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

=



Original

Blur (with a mean filter)

Examples of convolution filters



0	0	0
0	2	0
0	0	0

-

0	0	0
0	1	0
0	0	0



Original

Original

Examples of convolution filters

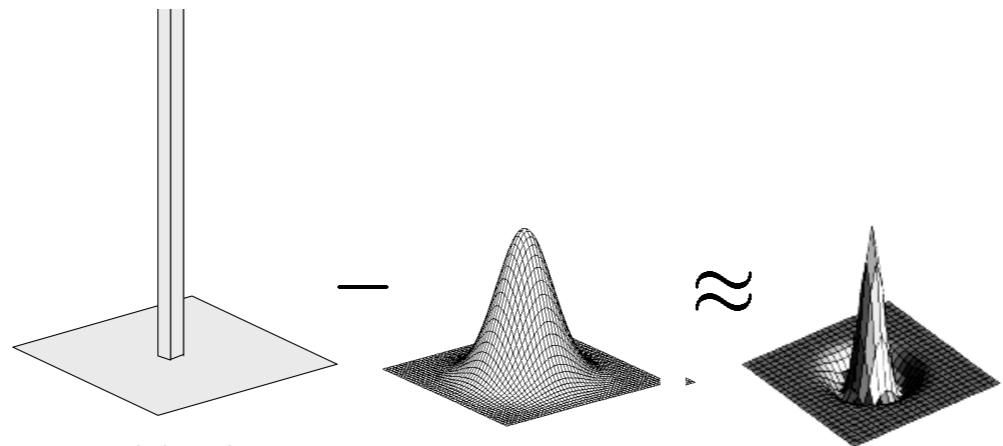


What does this do?

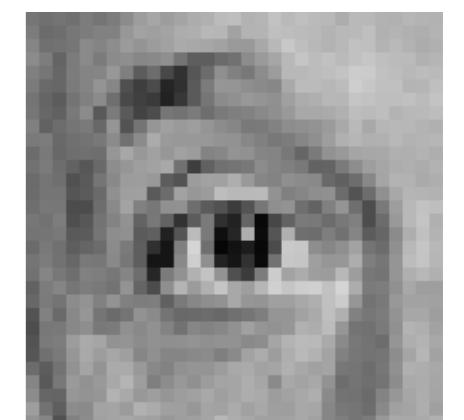
Finds high-frequency regions / details

$$(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{array}) - \begin{array}{ccc} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{array}) / 16) \alpha + \begin{array}{ccc} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{array} = \begin{array}{ccc} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{array}$$

Original



Combine into a single filter

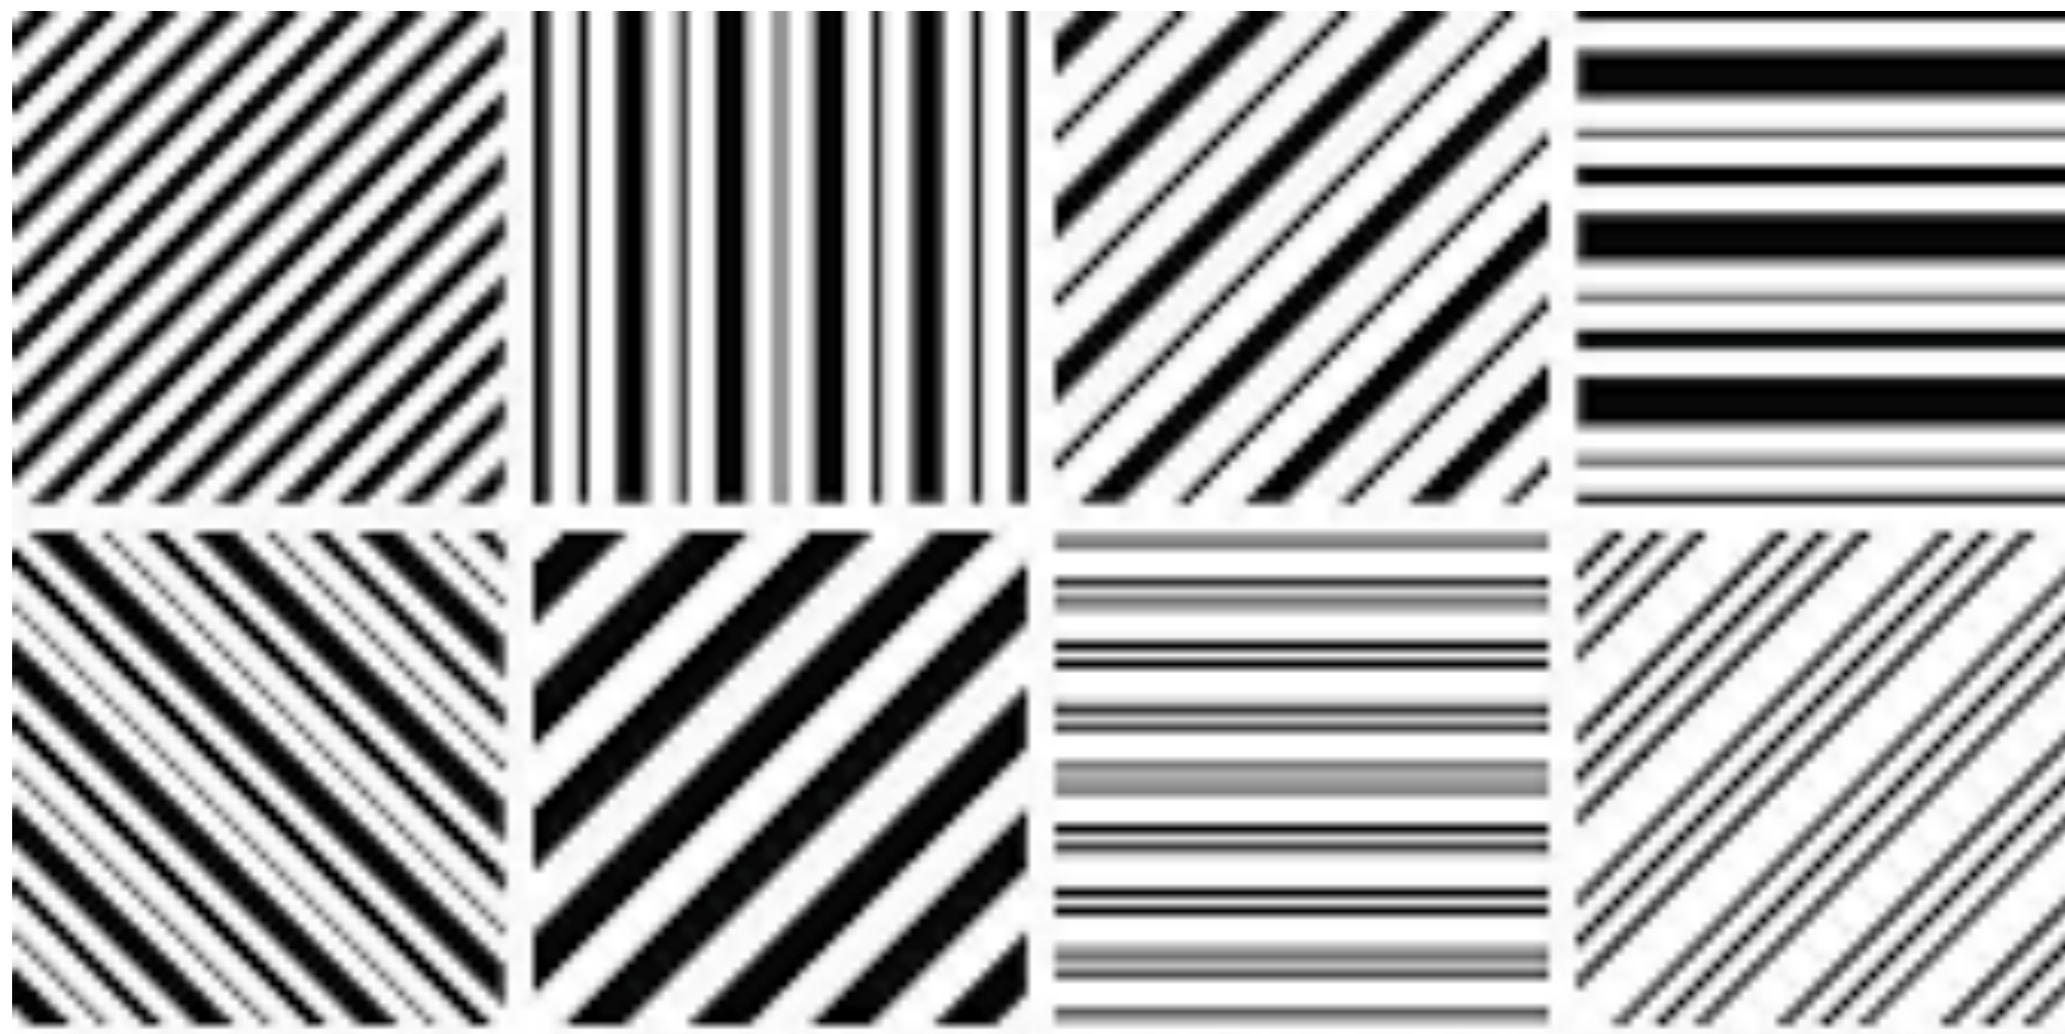


“sharper” image

Summary

- Filters
- Convolution
- Examples of Convolution

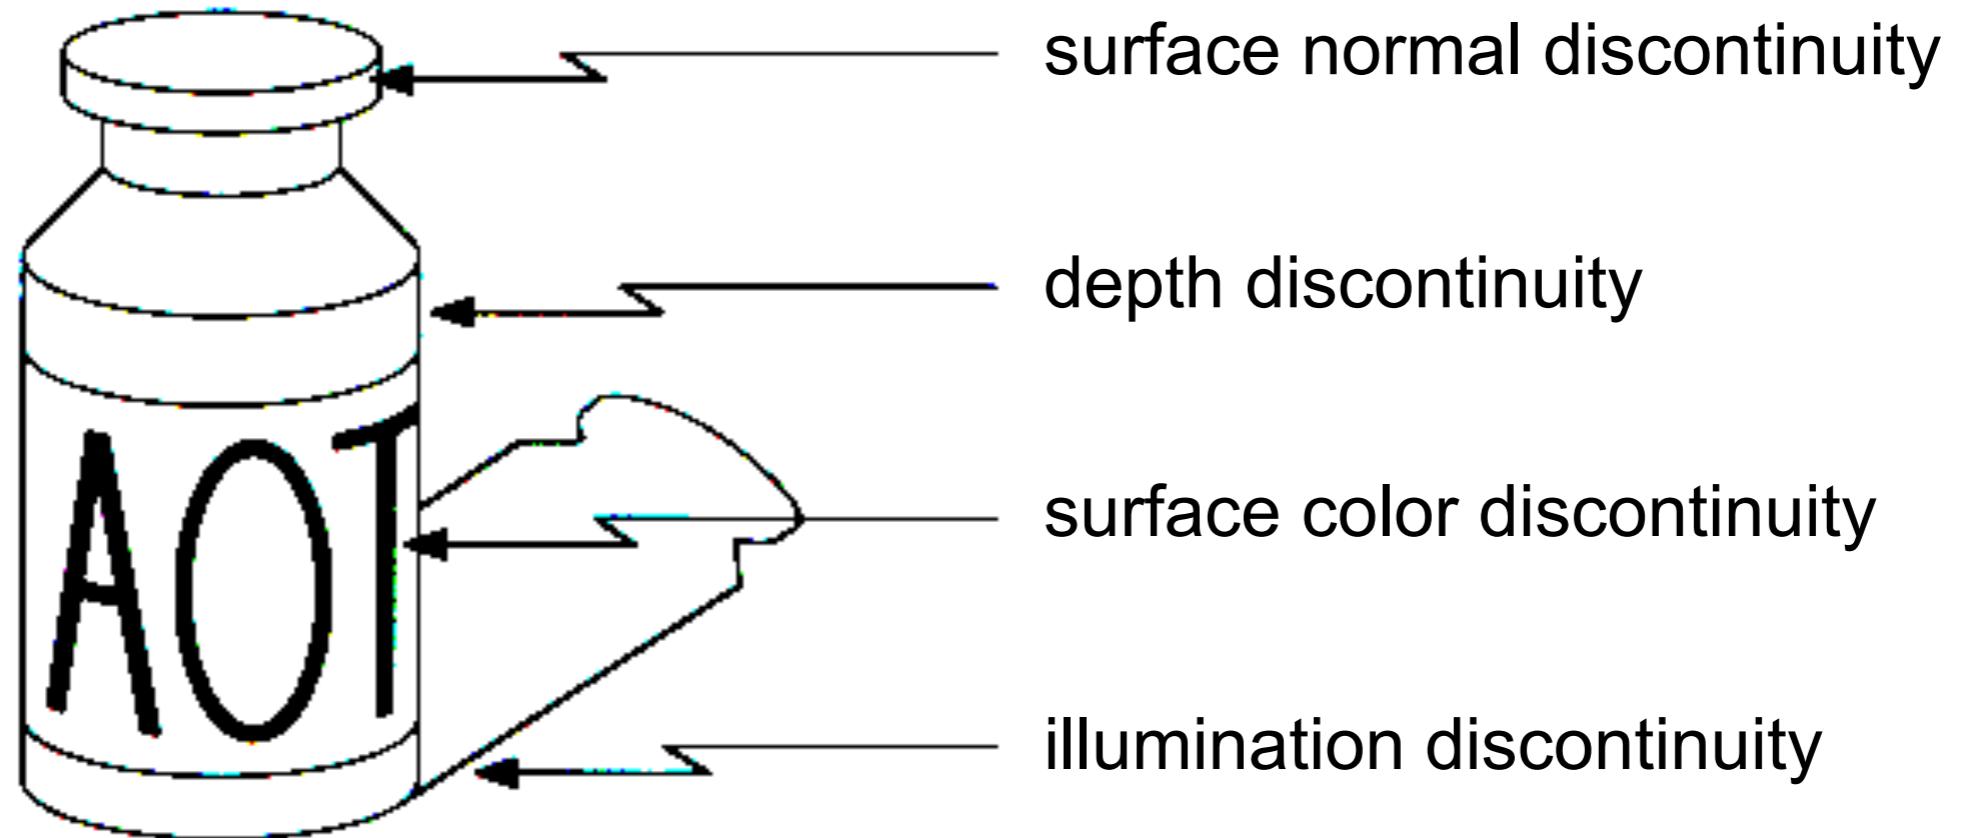
Edges



What causes edges in an image?

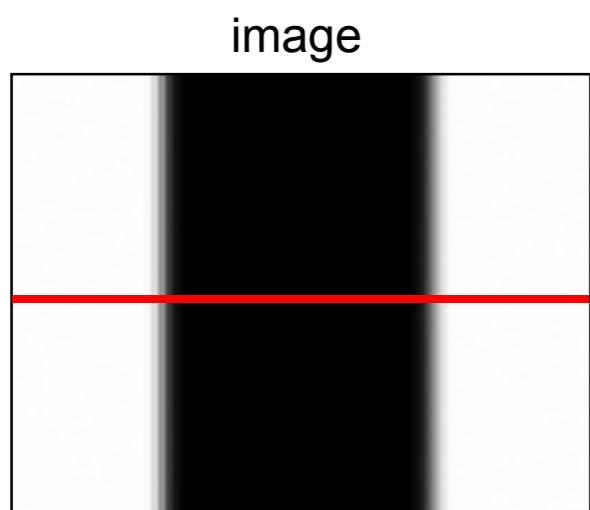


Come up with as many reasons as you can based on this image

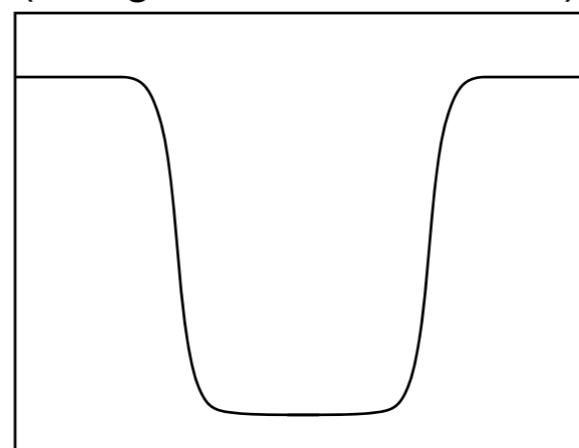


Characterizing edges

- An edge is a place of rapid change in the image intensity function

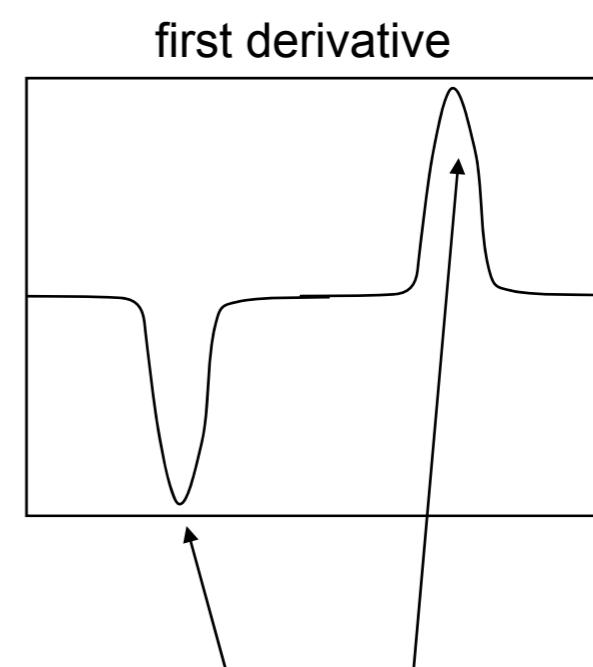


intensity function
(along horizontal scanline)



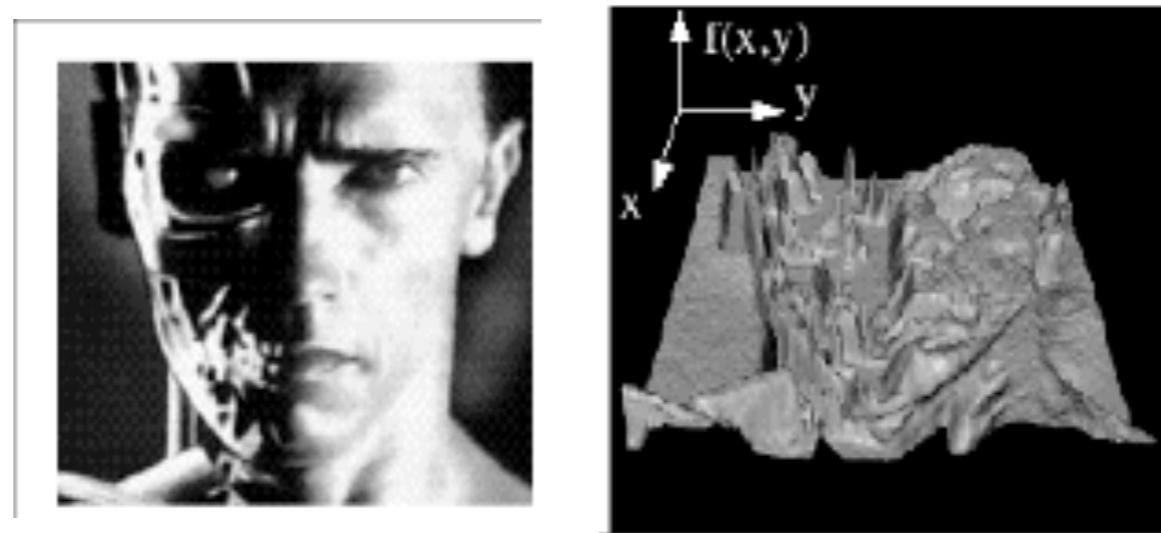
Look at pixel values along
this red line

**How do we find
image edges?**



edges correspond to
extrema of derivative

Continuous derivatives of an image



$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

Magnitude of gradient:

$$||\nabla f|| = \sqrt{f_x^2 + f_y^2}$$

Direction of gradient:

$$\alpha = \tan^{-1} \frac{f_y}{f_x}$$

Taking Derivatives by Convolution

For 2D function $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement derivatives as a convolution, what would be the associated filter?

Taking Derivatives by Convolution with an edge filter

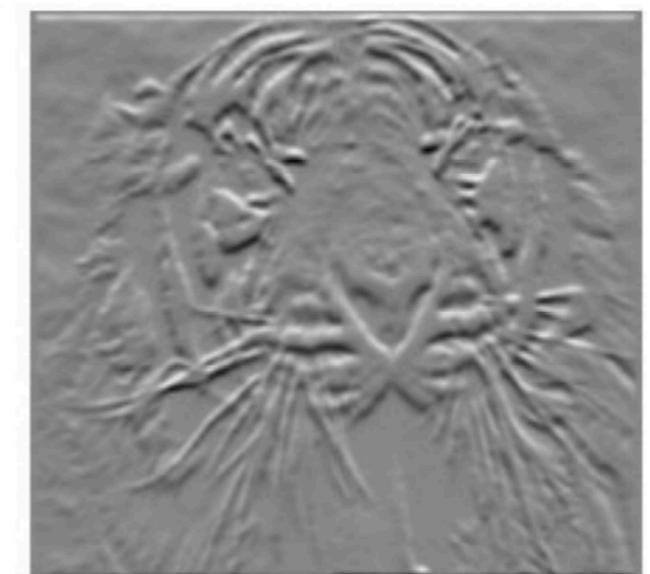
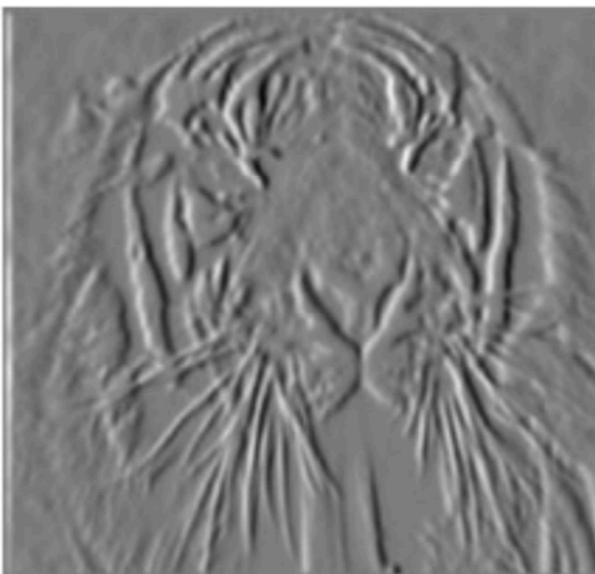
Discrete approximation of derivative
(from previous slide)

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

$\approx F \otimes [-1, 1]$ correlation

$= f * [1, -1]$ convolution

-1	1
----	---

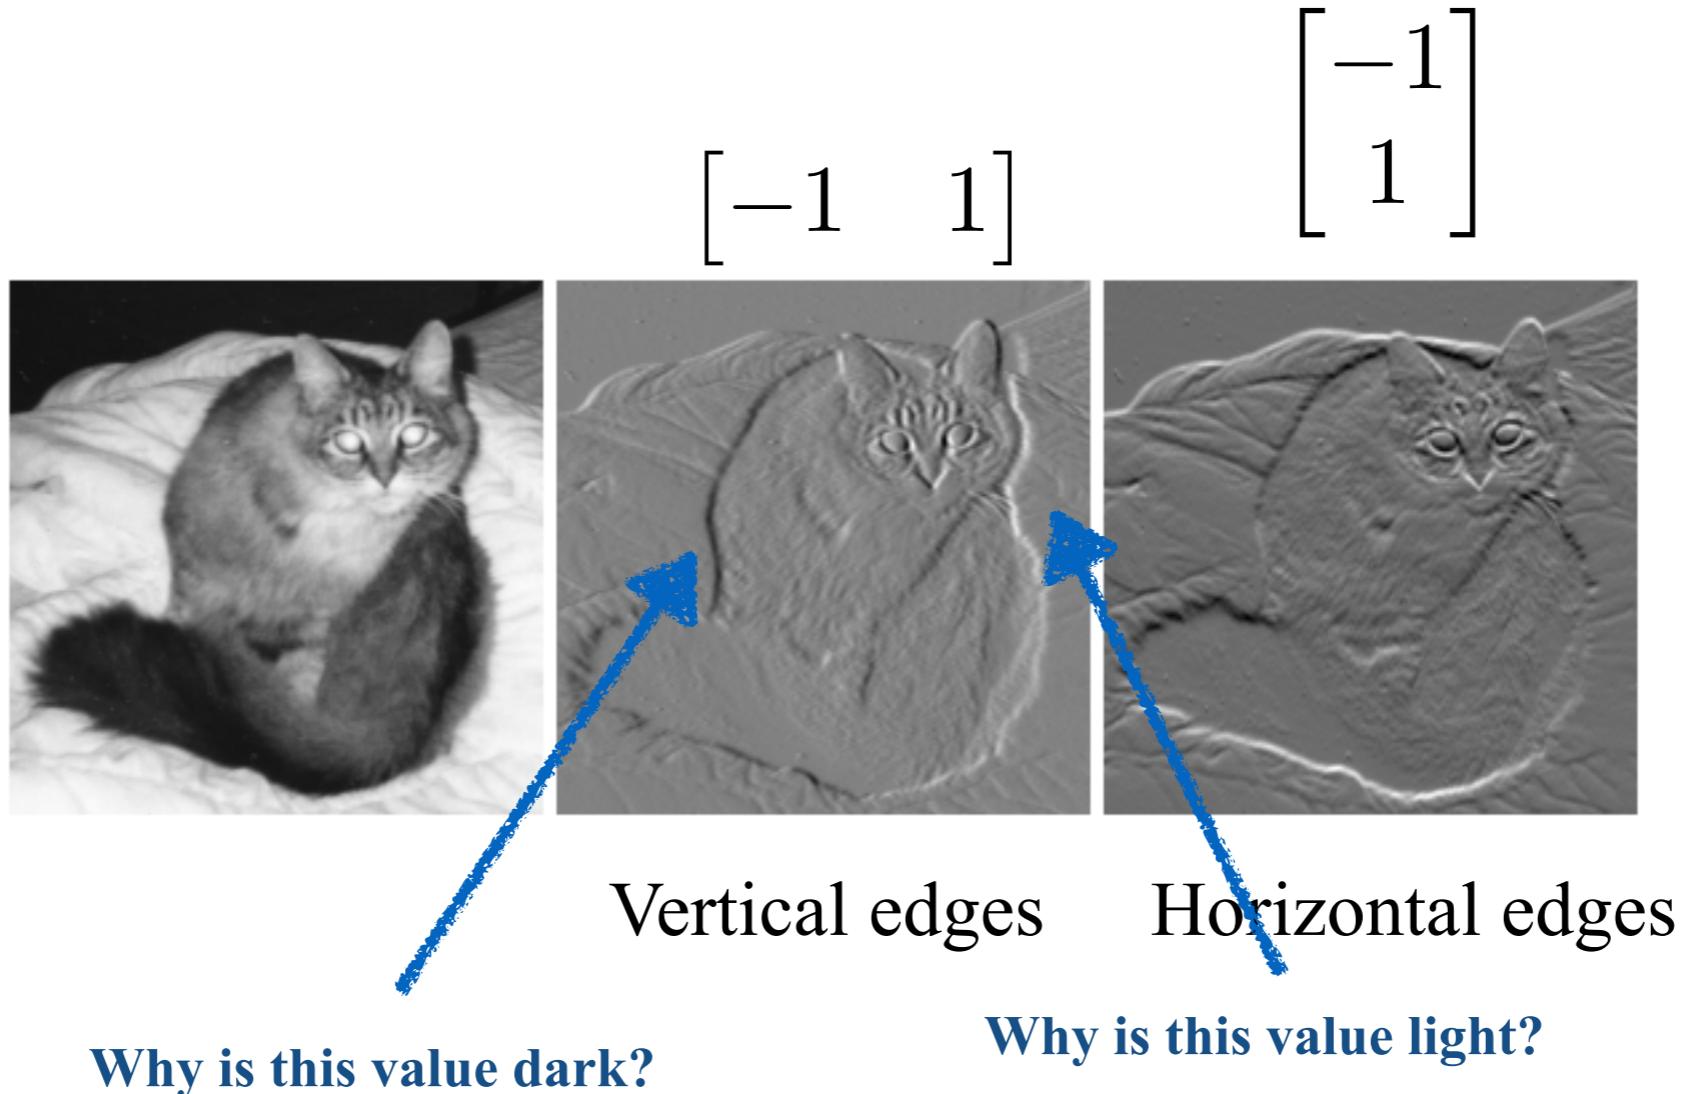


-1
1

Edge image

Derivative image

Discrete derivative filters



Other types of edge filters

Roberts: $M_x = \begin{array}{|c|c|} \hline 0 & 1 \\ \hline -1 & 0 \\ \hline \end{array}$; $M_y = \begin{array}{|c|c|} \hline 1 & 0 \\ \hline 0 & -1 \\ \hline \end{array}$

Diagonal edges

Prewitt: $M_x = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$; $M_y = \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$

Why might this work better?

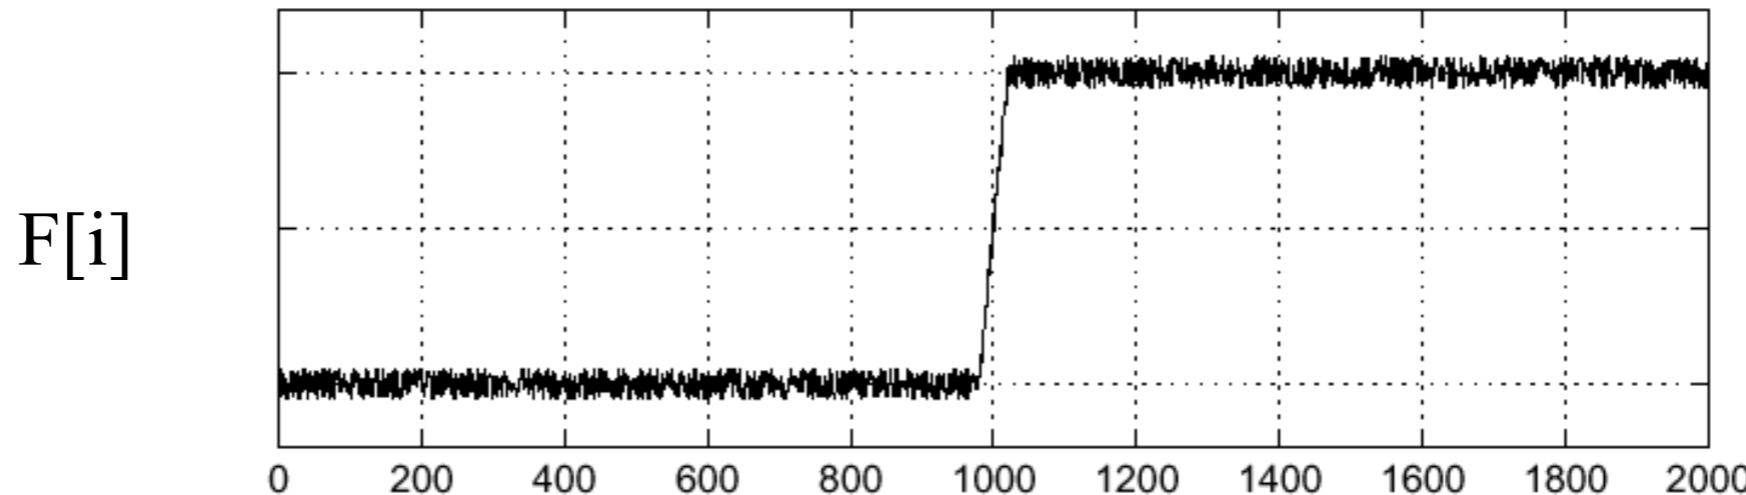
Sobel: $M_x = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$; $M_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$

Weighted average

Revisiting gradients

(let's plot a single row of image as a function)

Step edge + Gaussian noise

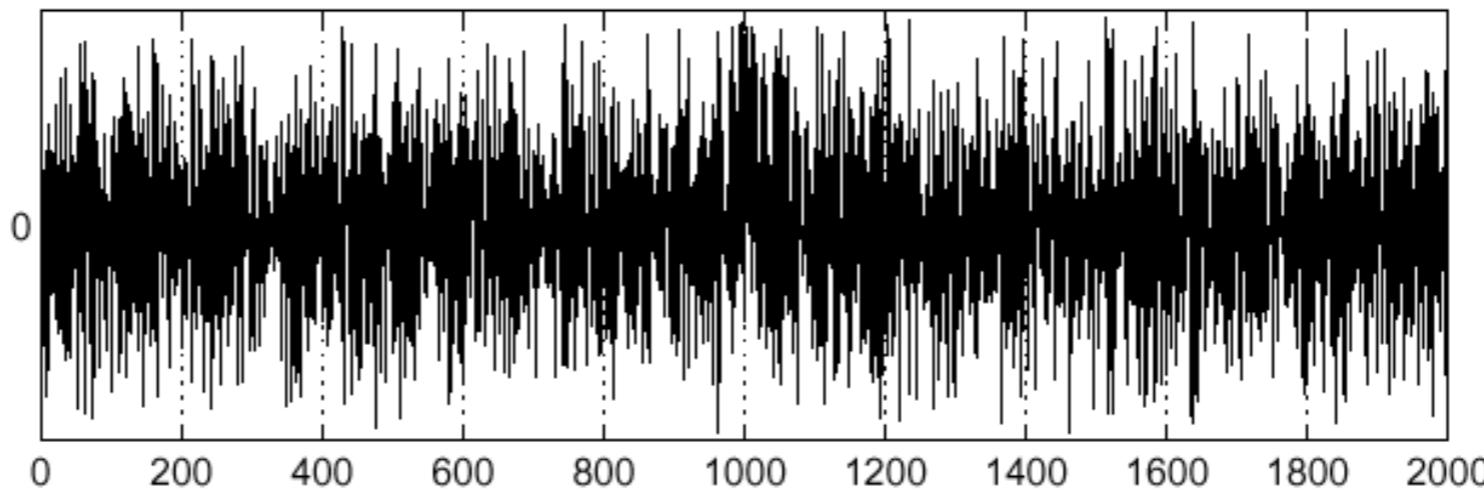


Correlate / convolve with
an edge filter

$$\begin{aligned}\frac{\partial F}{\partial i} &\approx F \otimes [-1, 1] \\ &= F * [1, -1]\end{aligned}$$

correlation

convolution



Problem: Gradient is very noisy! Where did the edge go?

How can we fix this to find the edge?

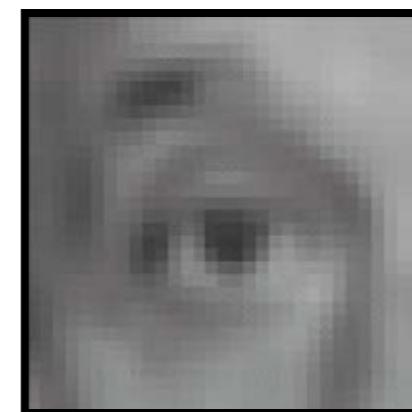
Recall the box filter

Box Filter / Mean filter



Original

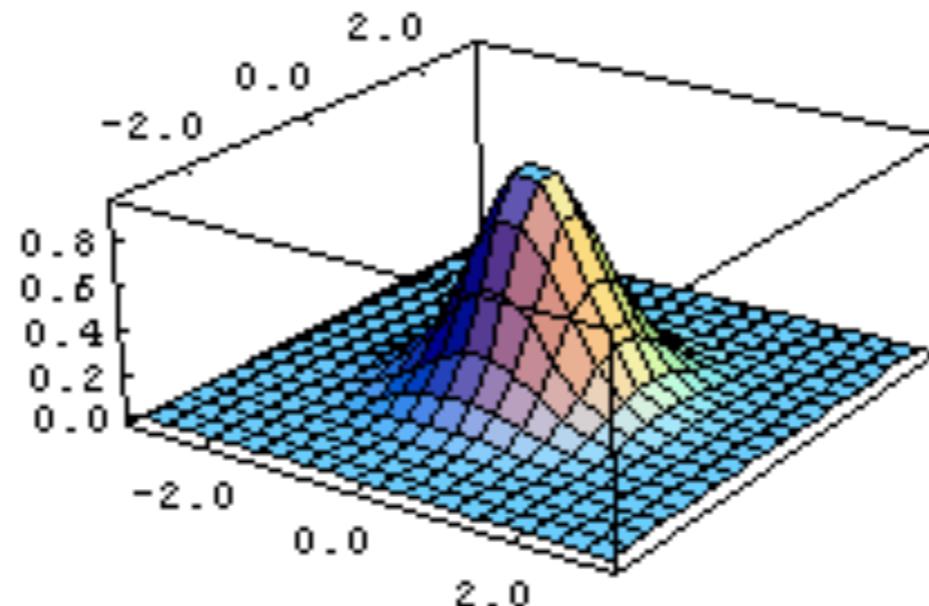
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Blur (with a mean filter)

Alternative: Gaussian Smoothing

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Discrete version
(two examples):

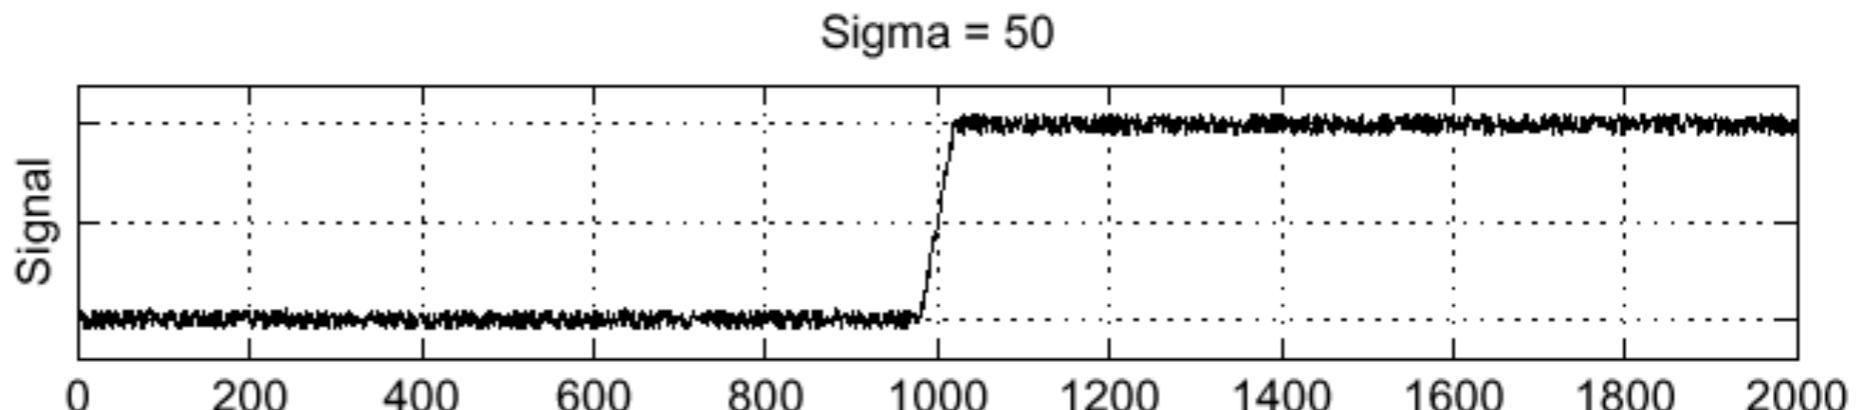
$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \quad \frac{1}{159} \begin{bmatrix} 2 & 4 & 5 & 4 & 2 \\ 4 & 9 & 12 & 9 & 4 \\ 5 & 12 & 15 & 12 & 5 \\ 4 & 9 & 12 & 9 & 4 \\ 2 & 4 & 5 & 4 & 2 \end{bmatrix}$$

Why might this be
better than a box filter?

Solution: smooth first

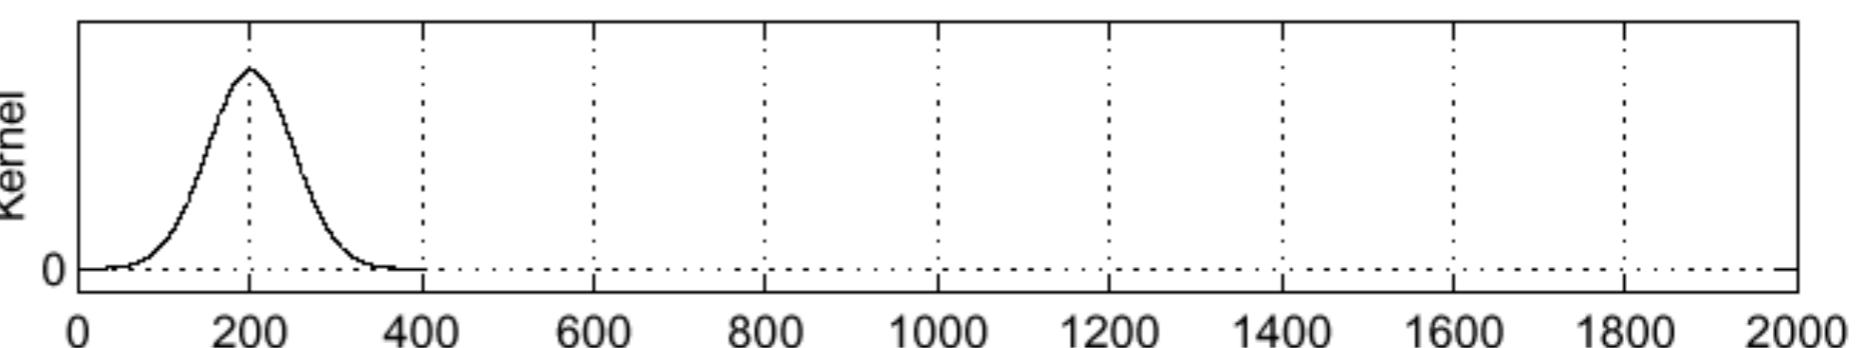
Noisy image / signal

F

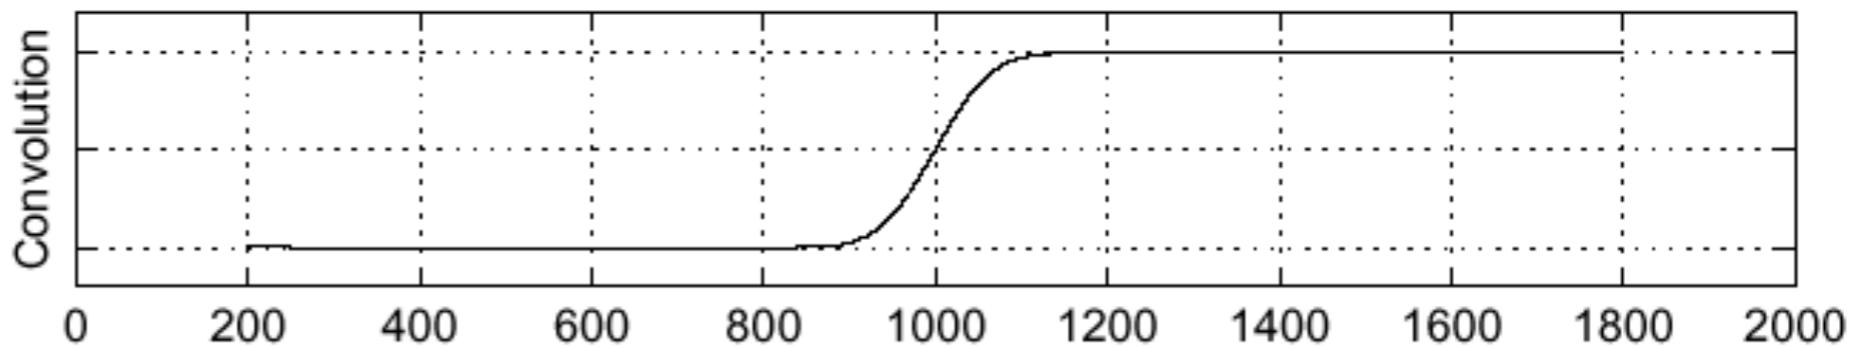


Gaussian kernel
(smooth / blur filter)

H



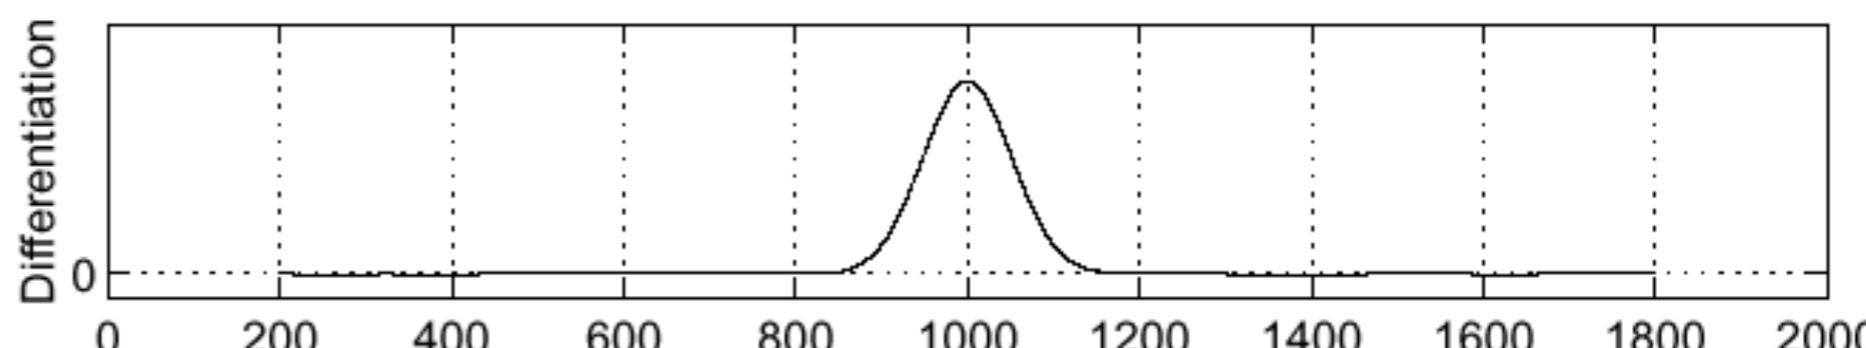
Smoothed signal $F * H$



Derivative
of smoothed
signal

$$\frac{\partial}{\partial i} (F * H)$$

$$\approx (F * H) * [1, -1]$$



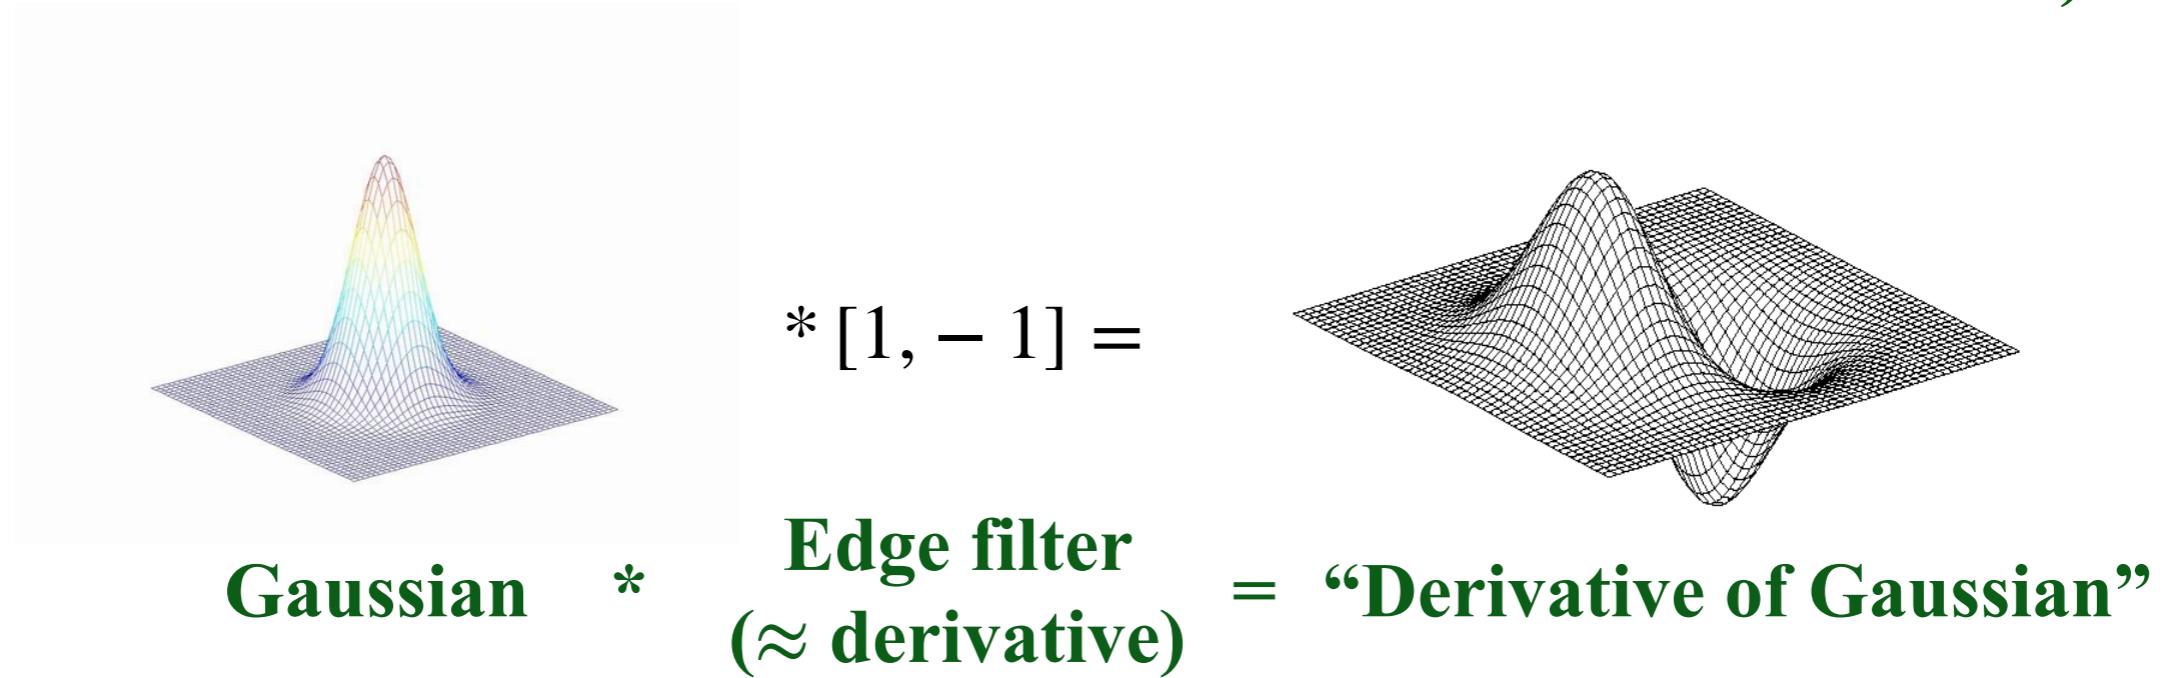
Peaks of derivative =
location of edge in original image!

Derivative of Gaussian filter

smooth derivative

Edge detection $\approx (F * \nabla) * [1, -1]$ How can we make this more efficient?

$= F * (\nabla * [1, -1])$ (Associative property of convolutions)

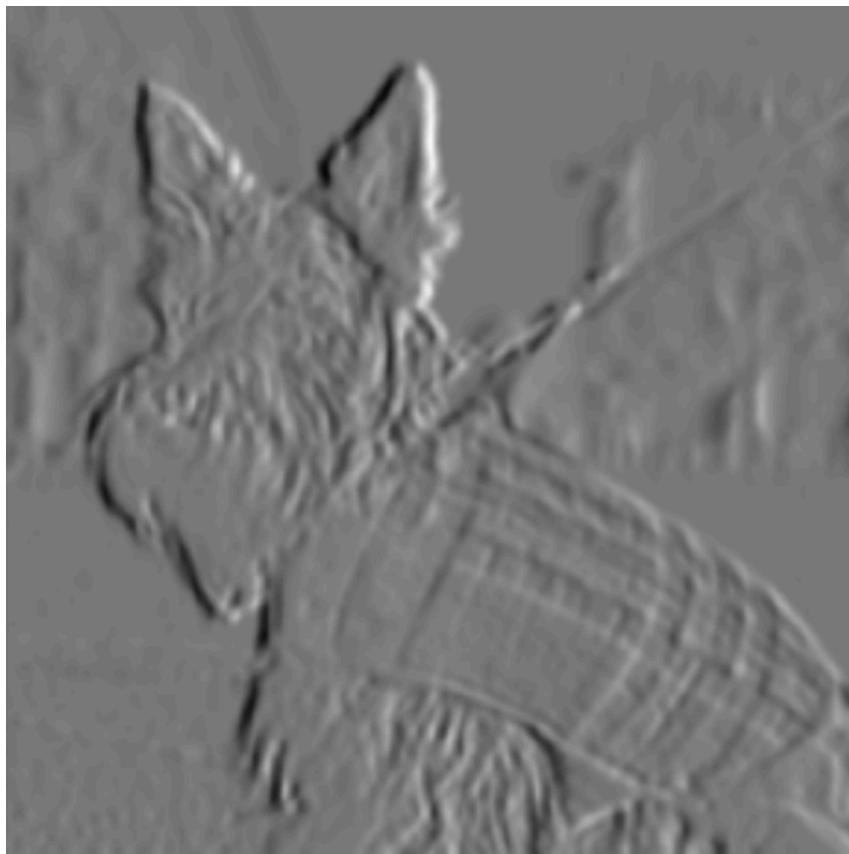


$= F * \text{“Derivative of Gaussian”}$



original image

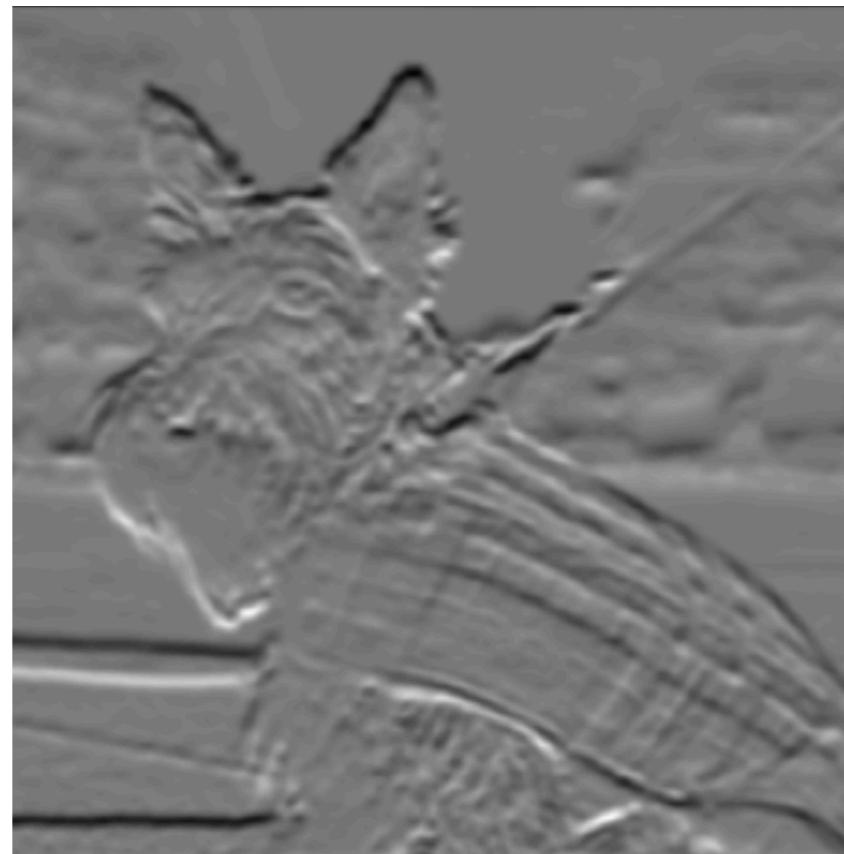
Edge Image / Gradient Image



X derivative of
Gaussian

Edges in x-direction

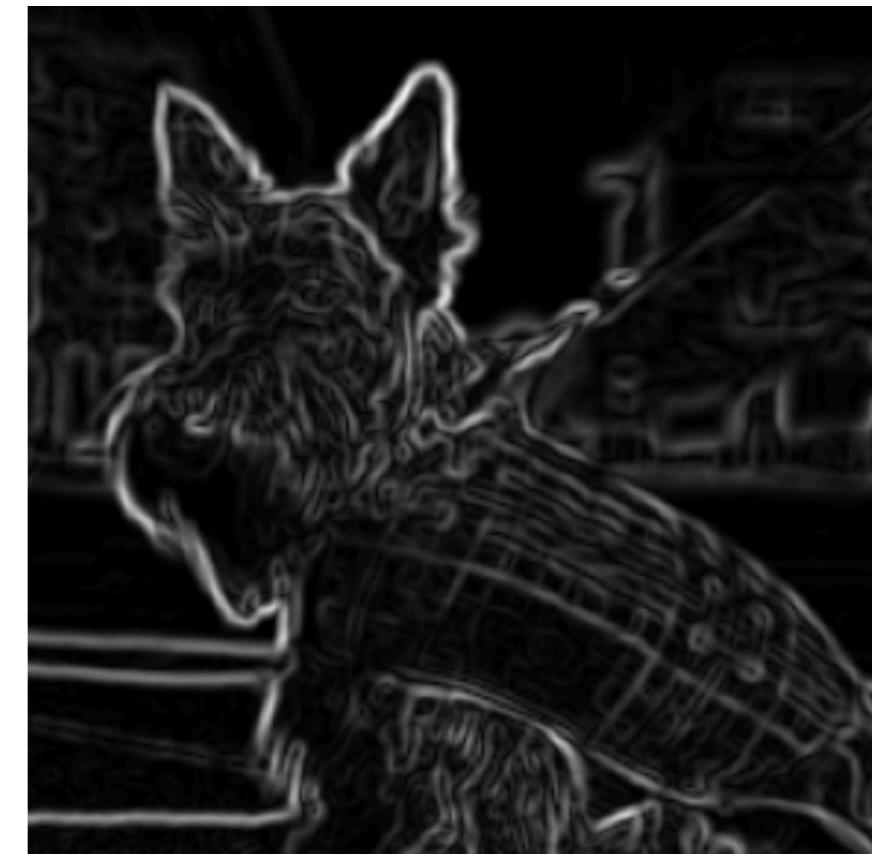
$$f_x = F * (\nabla \circledast [1, -1])$$



Y derivative of
Gaussian

Edges in y-direction

$$f_y = F * (\nabla \circledast \begin{bmatrix} 1 \\ -1 \end{bmatrix})$$



Magnitude
Gradient:

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

(For each pixel, we have 2 dimensions)

$$||\nabla f|| = \sqrt{f_x^2 + f_y^2}$$



Gradient Magnitude

Thresholding to find edge vs non-edge pixels

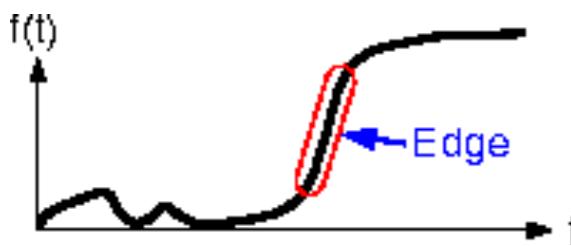


Thresholded edge image / gradient image

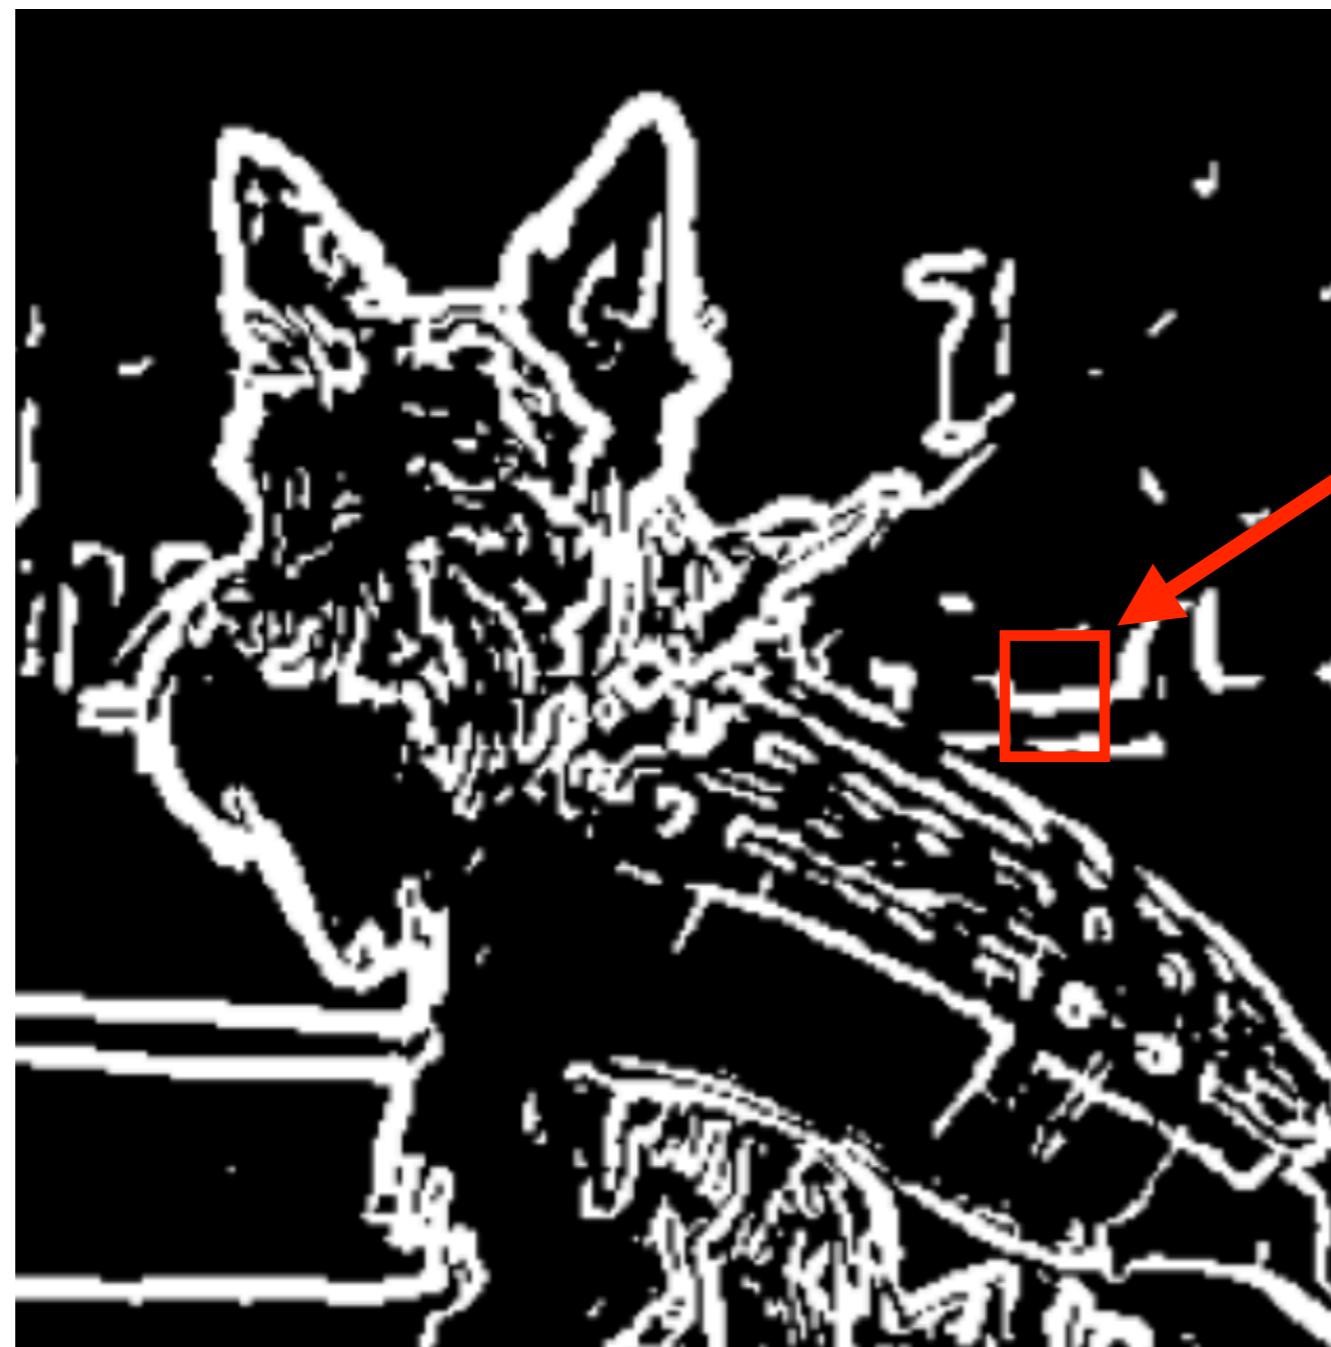
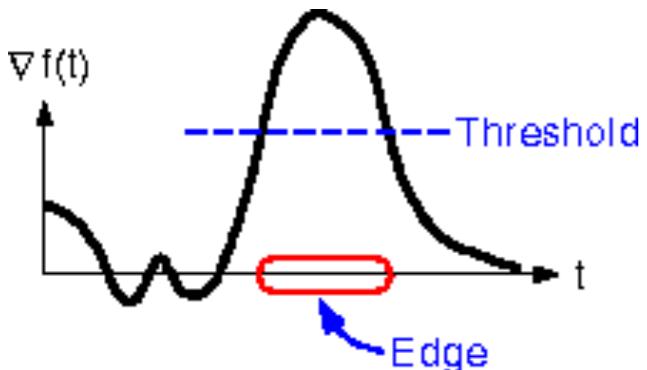
Why are these edges so thick?

Why are these edges so thick?

Original image:



Derivative



How to turn these thick regions of the thresholded gradient image into single-pixel curves?

Intuition: We need to find the center of this thick curve

Zoom-in on Thresholded Gradient Image

Intuition: we need to find the peak, not just look at all values above the threshold

Any ideas?

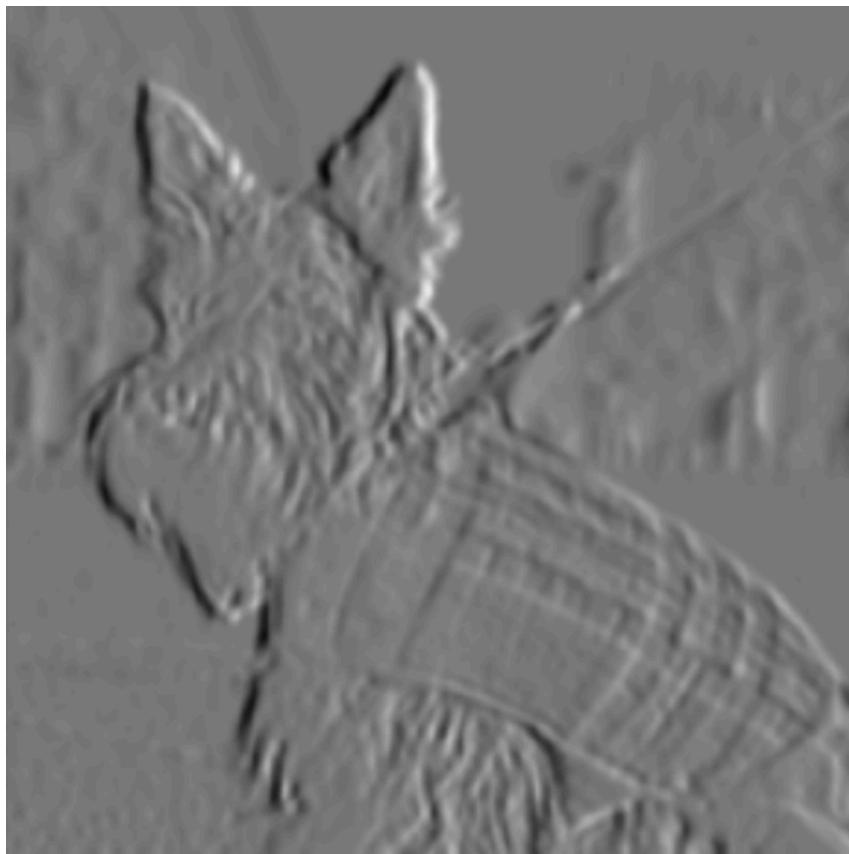
Smoothing -> edge filter -> norm of the gradient



Edge image / gradient image (before thresholding)

How can we find the edges as single-pixel edge curves ?

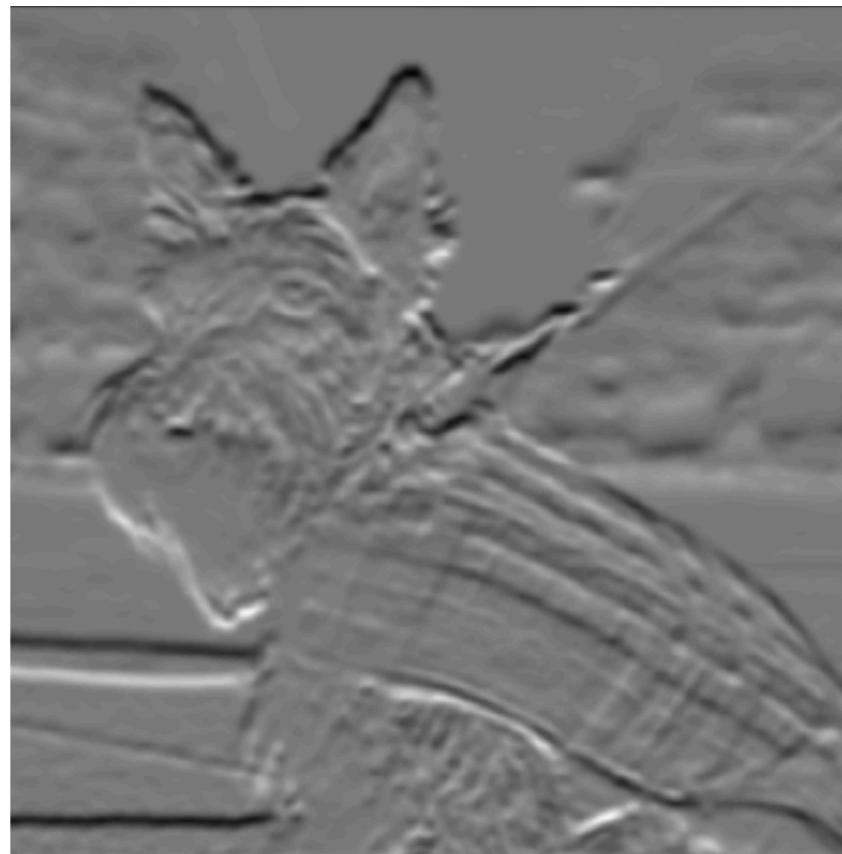
Edge Image / Gradient Image



X derivative of
Gaussian

Edges in x-direction

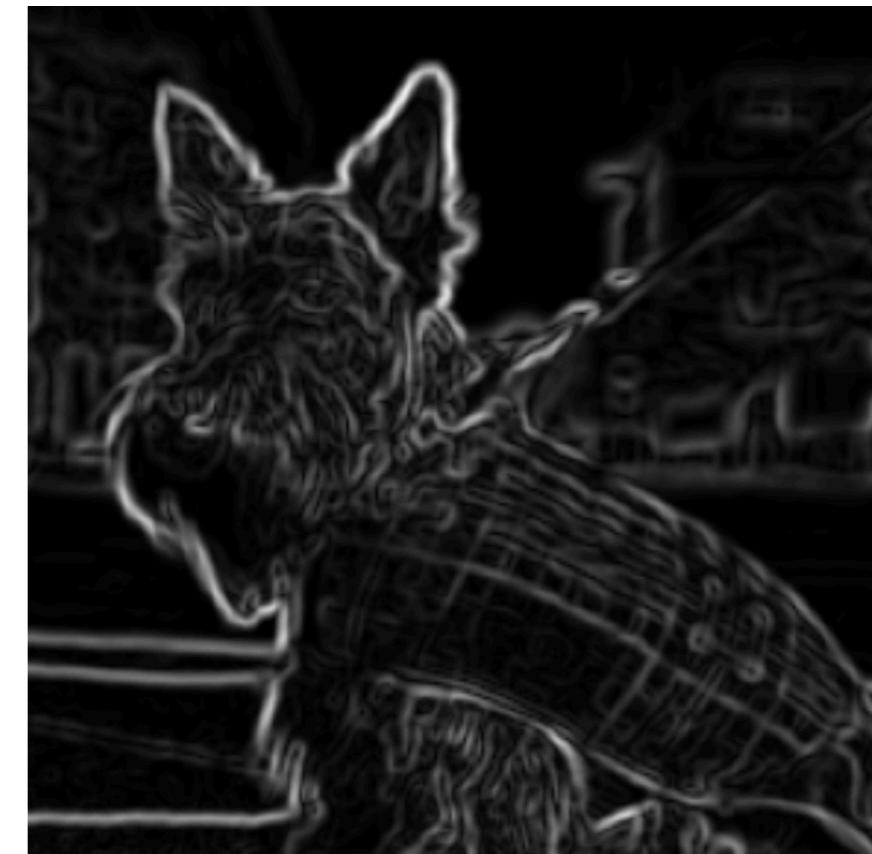
$$f_x = F * (\text{Sobel} * [1, -1])$$



Y derivative of
Gaussian

Edges in y-direction

$$f_y = F * (\text{Sobel} * \begin{bmatrix} 1 \\ -1 \end{bmatrix})$$



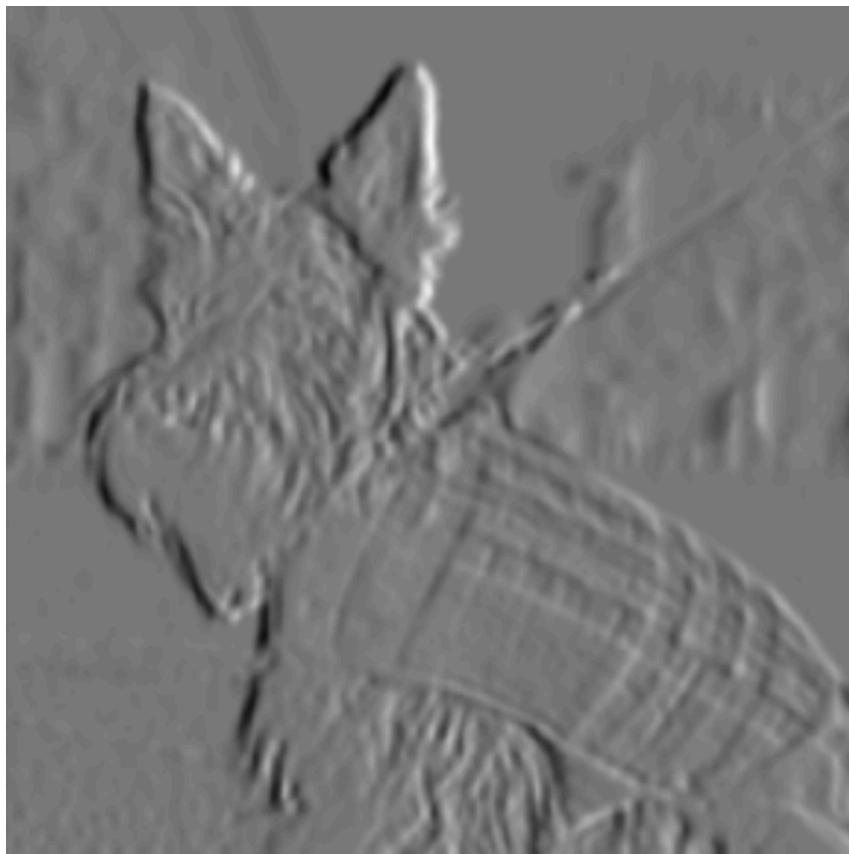
Magnitude
Gradient:

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

(For each pixel, we have 2 dimensions)

$$||\nabla f|| = \sqrt{f_x^2 + f_y^2}$$

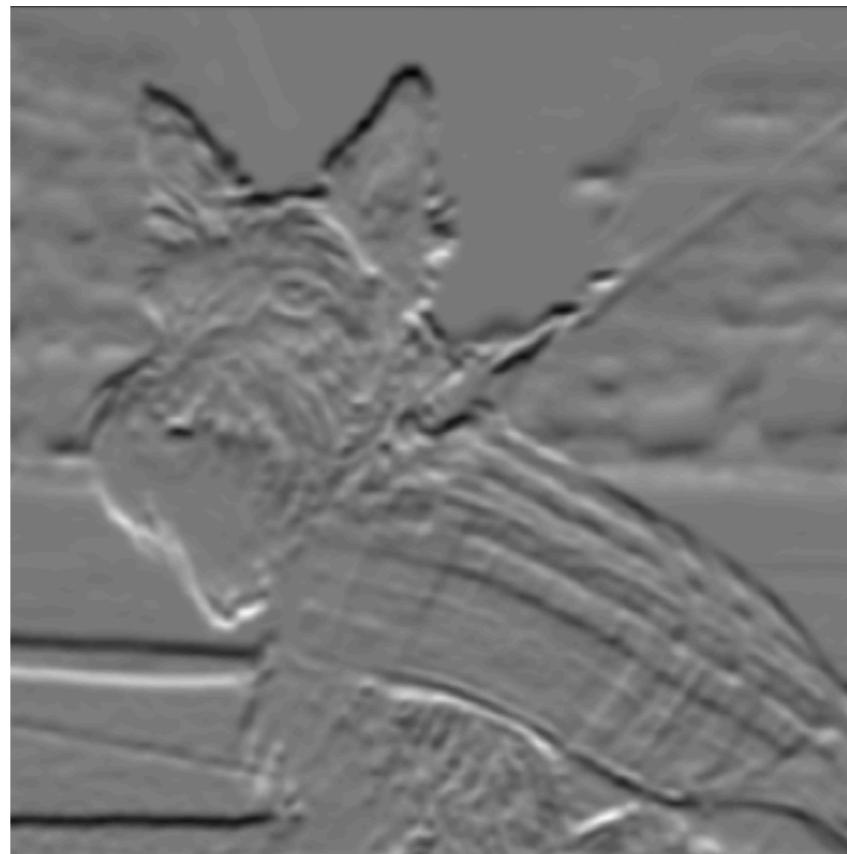
Edge Image / Gradient Image



X derivative of
Gaussian

Edges in x-direction

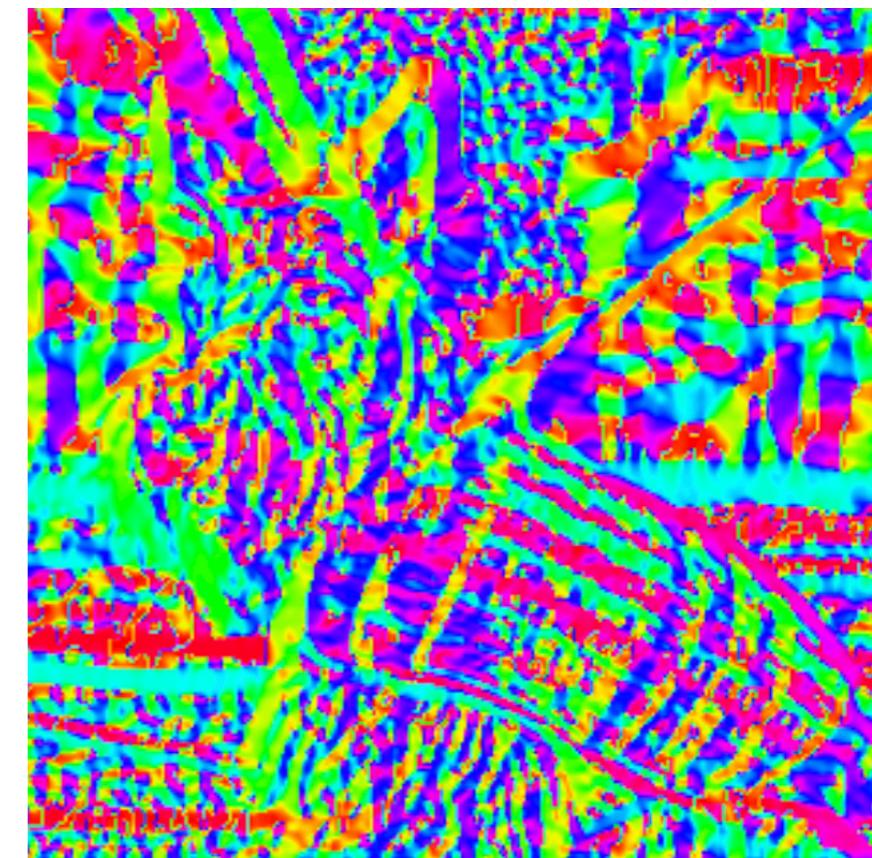
$$f_x = F * (\nabla * [1, -1])$$



Y derivative of
Gaussian

Edges in y-direction

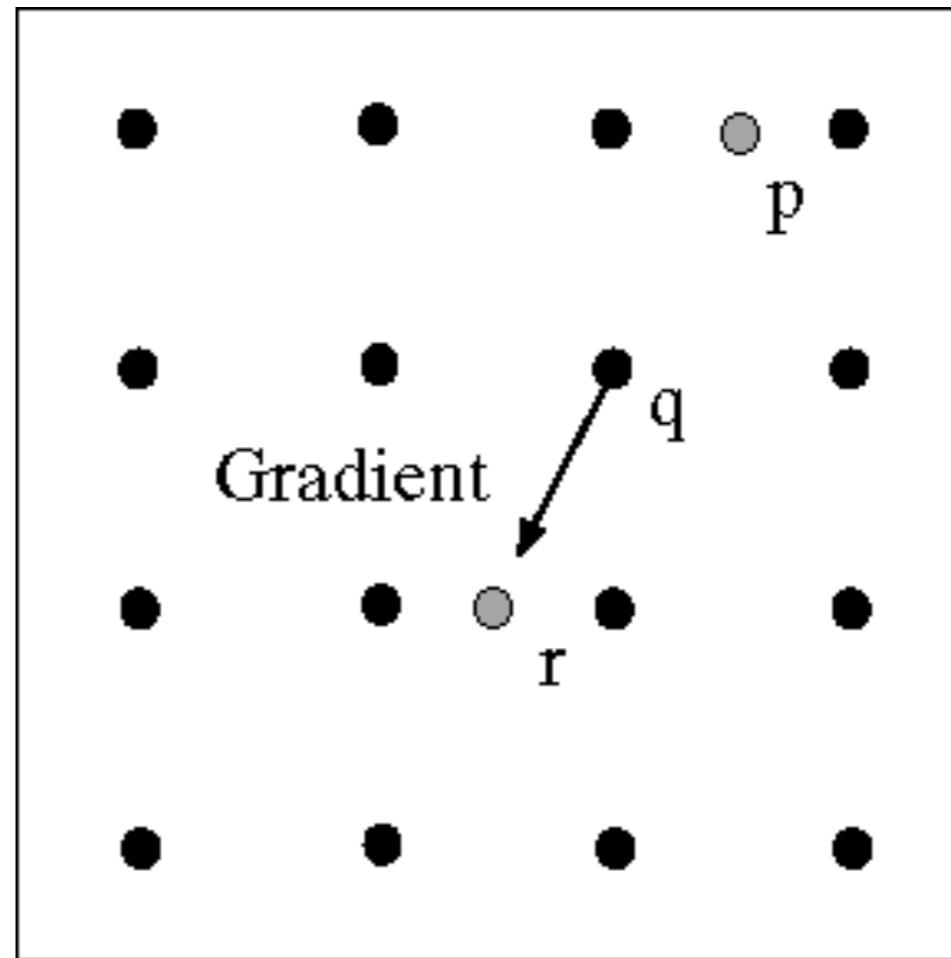
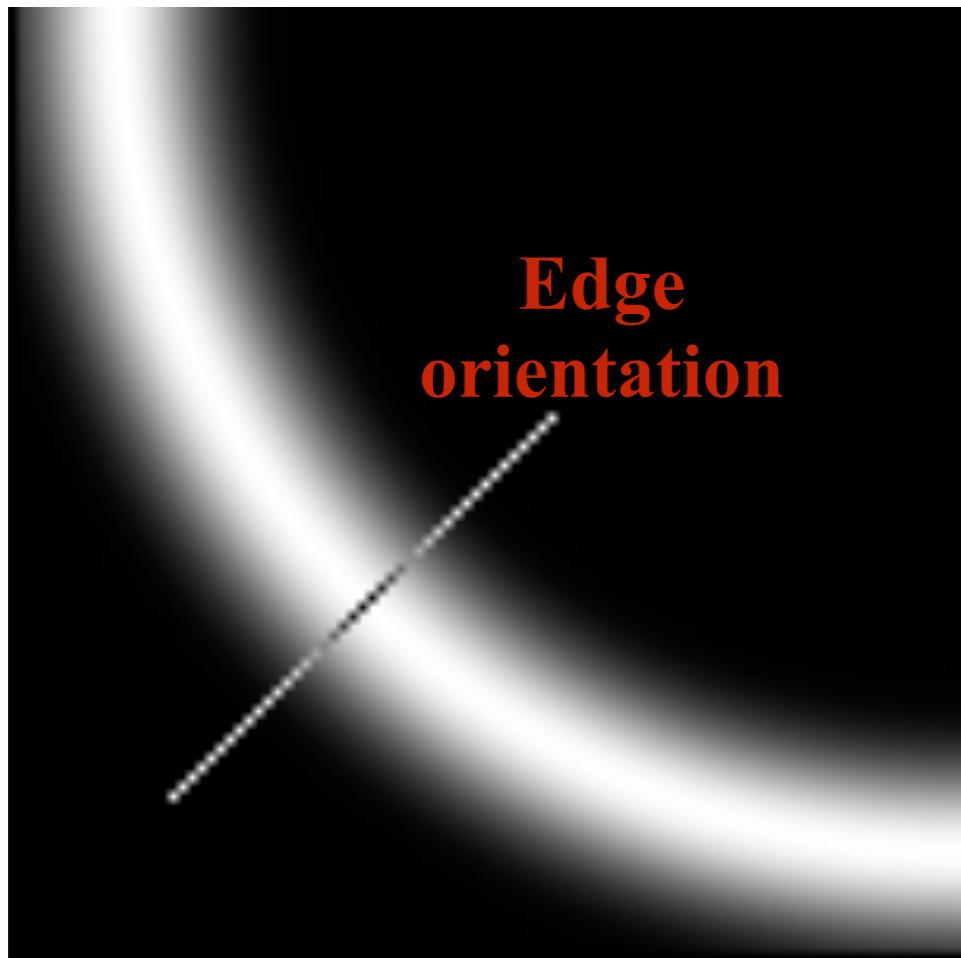
$$f_y = F * (\nabla * \begin{bmatrix} 1 \\ -1 \end{bmatrix})$$



Direction of gradient:

$$\alpha = \tan^{-1} \frac{f_y}{f_x}$$

Non-maximum suppression (NMS)



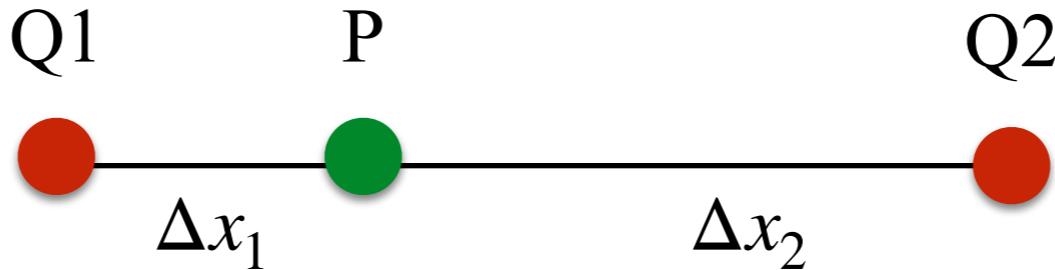
Problem: gradient direction points “in between pixels”

- Check if pixel is local maximum along gradient direction
- Select the single pixel that is max across width of the edge
- In practice: “suppress” any pixel that is not a local maximum (NMS)

How can we deal with this?

- We need to get the values “between” the pixels (p and r)

1D: Linear interpolation



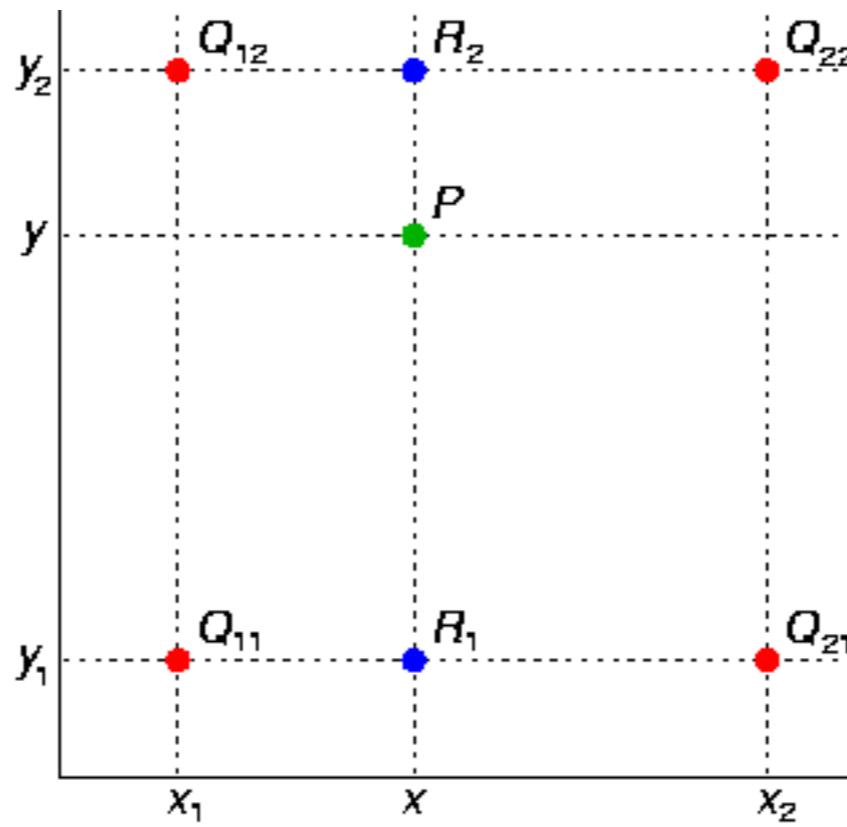
$$f(P) = \frac{\Delta x_2}{\Delta x_1 + \Delta x_2} f(Q_1) + \frac{\Delta x_1}{\Delta x_1 + \Delta x_2} f(Q_2)$$

What happens if $\Delta x_1 = 0$?

What happens if $\Delta x_2 = 0$?

What happens if $\Delta x_1 = \Delta x_2$?

2D: Bilinear interpolation

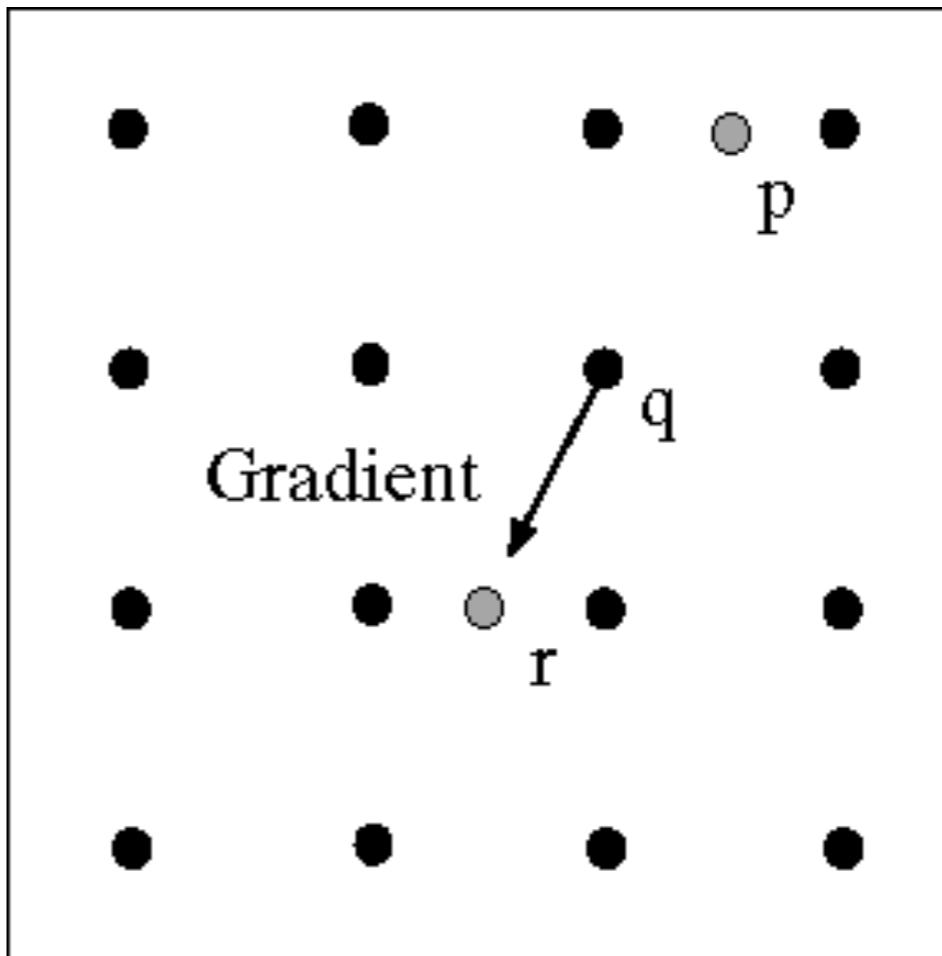
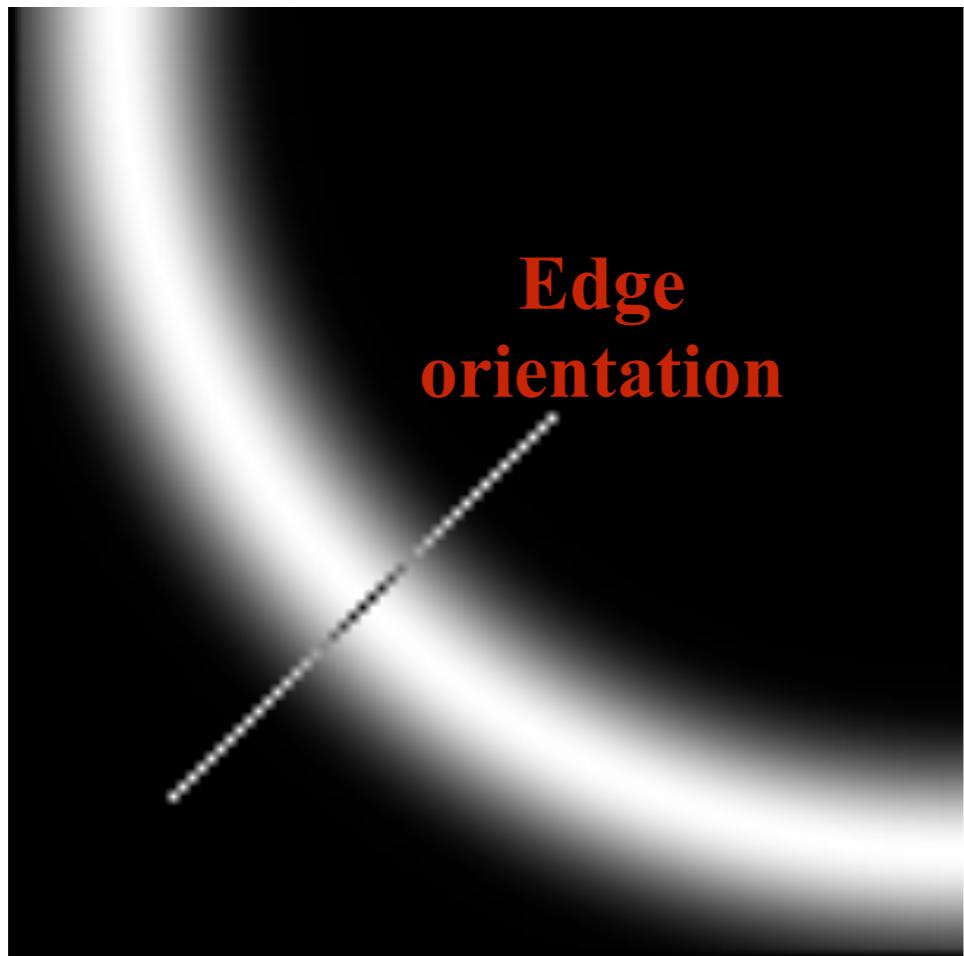


$$f(x, y_1) \approx \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21})$$

$$f(x, y_2) \approx \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22})$$

$$f(x, y) \approx \frac{y_2 - y}{y_2 - y_1} f(x, y_1) + \frac{y - y_1}{y_2 - y_1} f(x, y_2)$$

Non-maximum suppression (NMS)



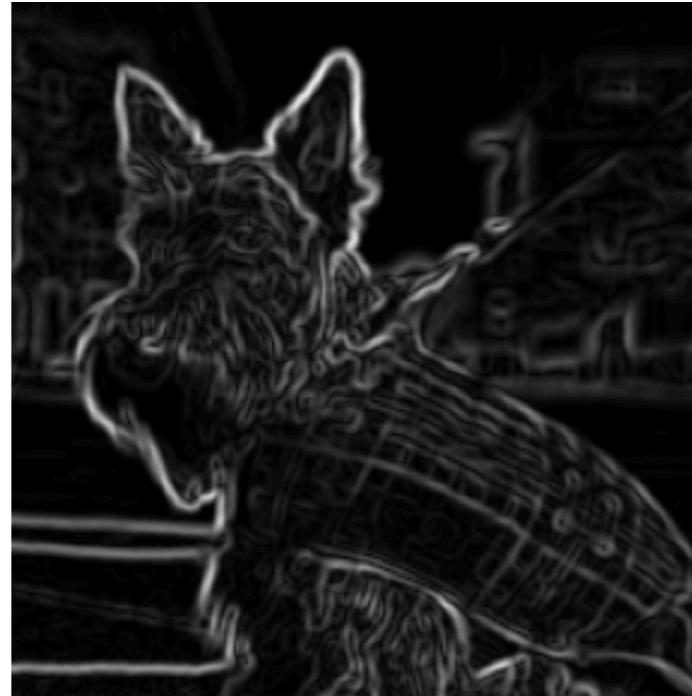
Problem: gradient direction points “in between pixels”

- Check if pixel is local maximum along gradient direction
- Select the single pixel that is max across width of the edge
- In practice: “suppress” any pixel that is not a local maximum (NMS)
 - We need to get the values “between” the pixels (p and r)
 - We can do this via **bilinear interpolation**

Edge Detector - Old version:



Original image



Edge image / gradient image



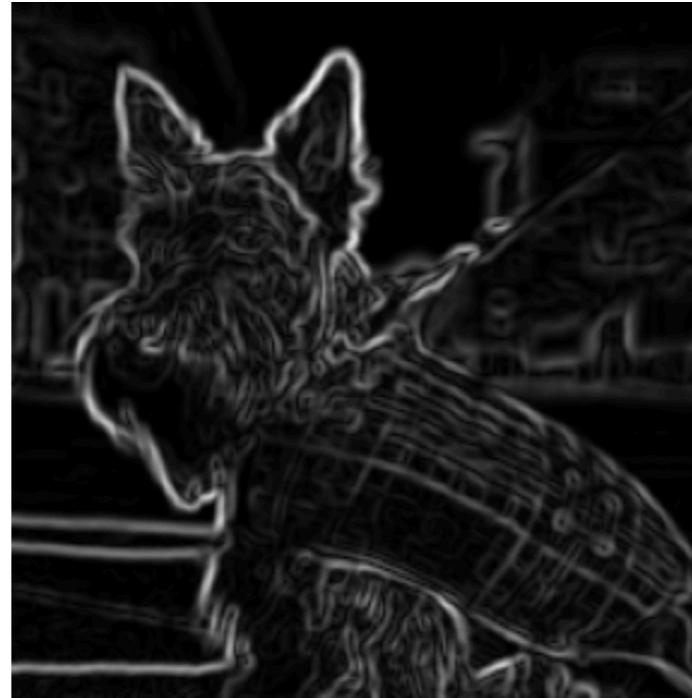
Thresholding

- Smooth the image (Gaussian filter)
- Edge / derivative filter
 - In practice combine the above 2 into a “Derivative of Gaussian” filter
- Compute norm of the gradient
- Threshold

Edge Detector - New version



Original image



Edge image / gradient image



Non-maximal suppression

- Smooth the image (Gaussian filter)
- Edge / derivative filter
 - In practice combine the above 2 into a “Derivative of Gaussian” filter
- Compute norm of the gradient
- Threshold
- **Non-maximum suppression to find the single-pixel edges**

The Canny edge detector



After non-maximum suppression

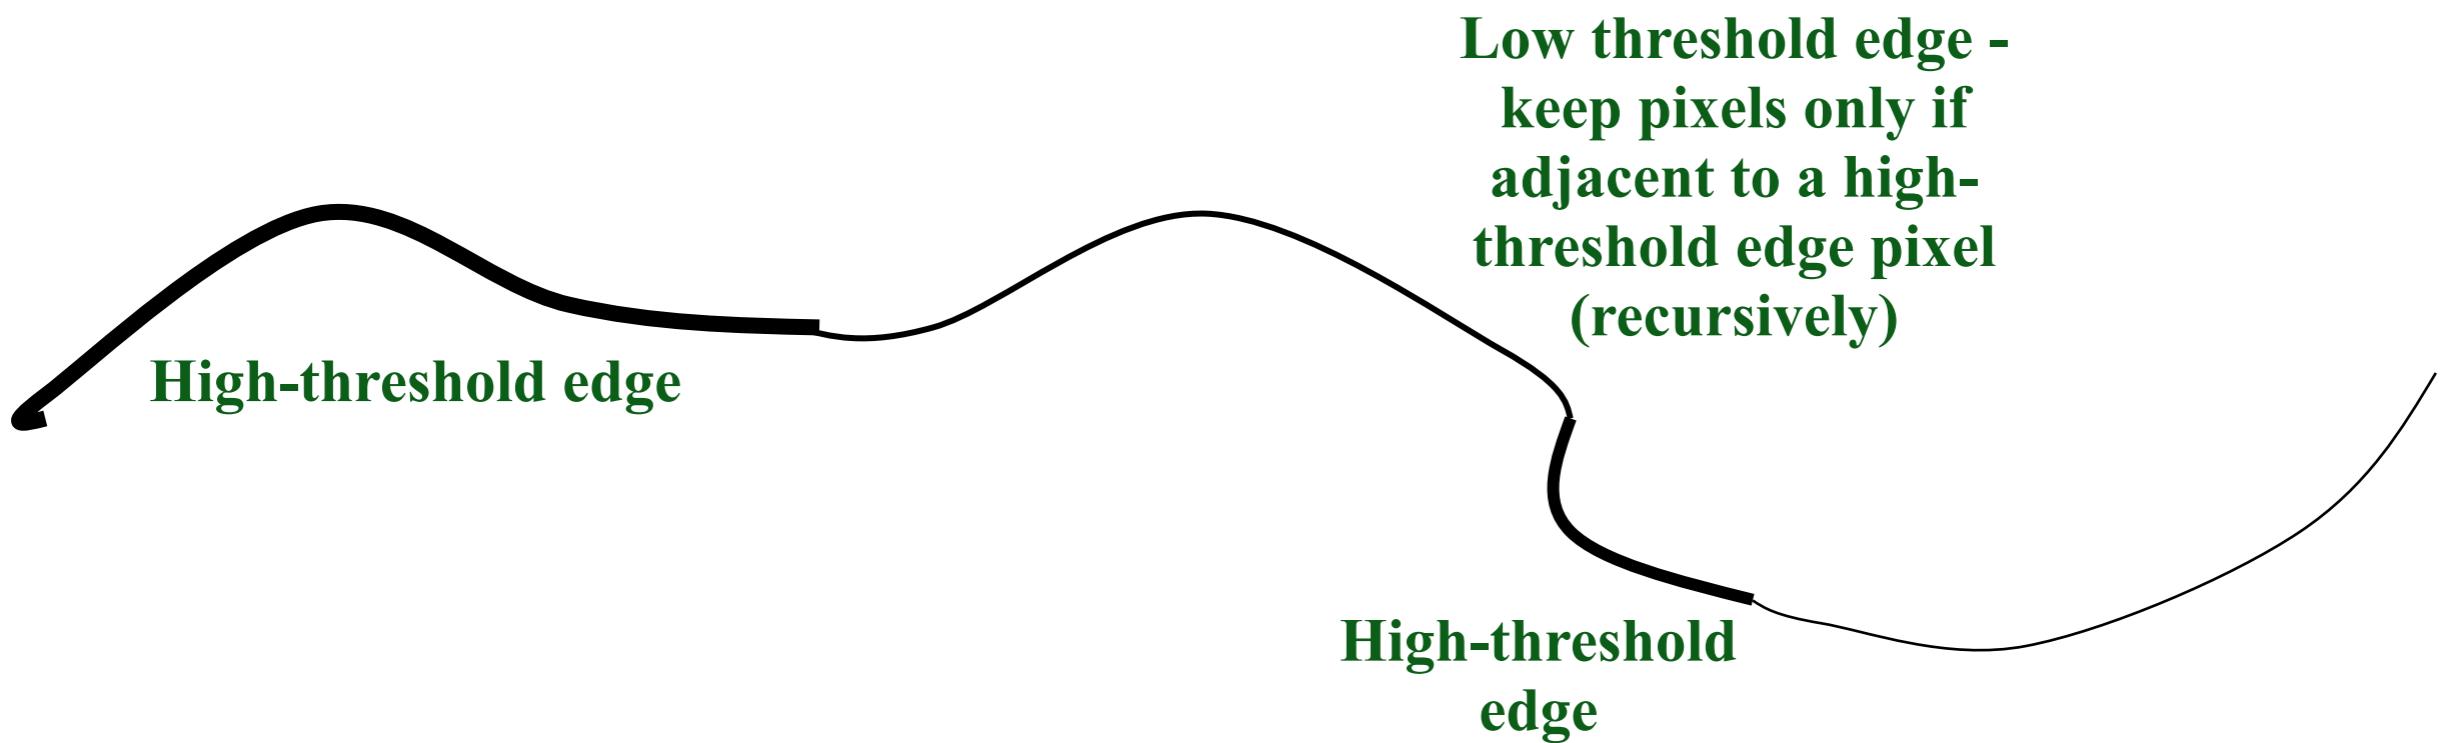
Problem: some pixels along this edge didn't survive the thresholding

(gaps in the edge)

How can we fix this?

Hysteresis thresholding

- How to fill in the gaps in an edge? Use **hysteresis**
 - Use a **high** threshold to **start** edge curves
 - Use **low** threshold to **continue** them

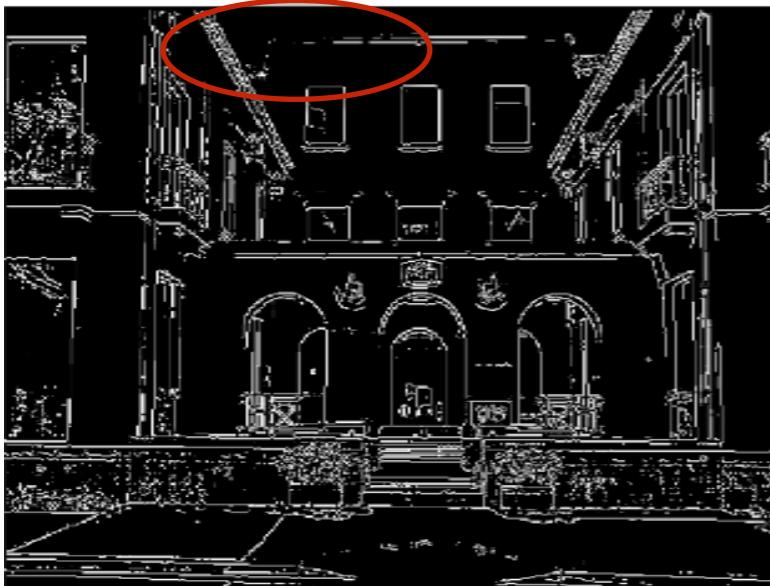


- This is a generally useful trick anytime we want to find curves in images / videos (e.g., tracking as spacetime curve fitting!)

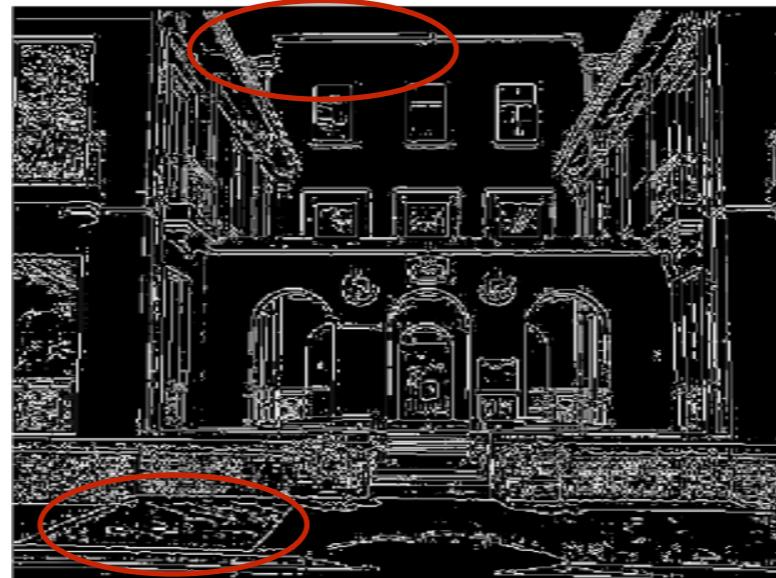
Hysteresis thresholding



original image



high threshold
(strong edges)



low threshold
(weak edges)



hysteresis threshold
68

Edge Detector



Original image



Edge image / gradient image



Non-maximum suppression

- Smooth the image (Gaussian filter)
- Edge / derivative filter
 - In practice combine the above with “Derivative of Gaussian” filter
- Compute norm of the gradient
- Threshold
- Non-maximum suppression to find the single-pixel edges
- **Hysteresis thresholding**

This is the “Canny Edge Detector”

IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. PAMI-8, NO. 6, NOVEMBER 1986

A Computational Approach to Edge Detection

JOHN CANNY, MEMBER, IEEE

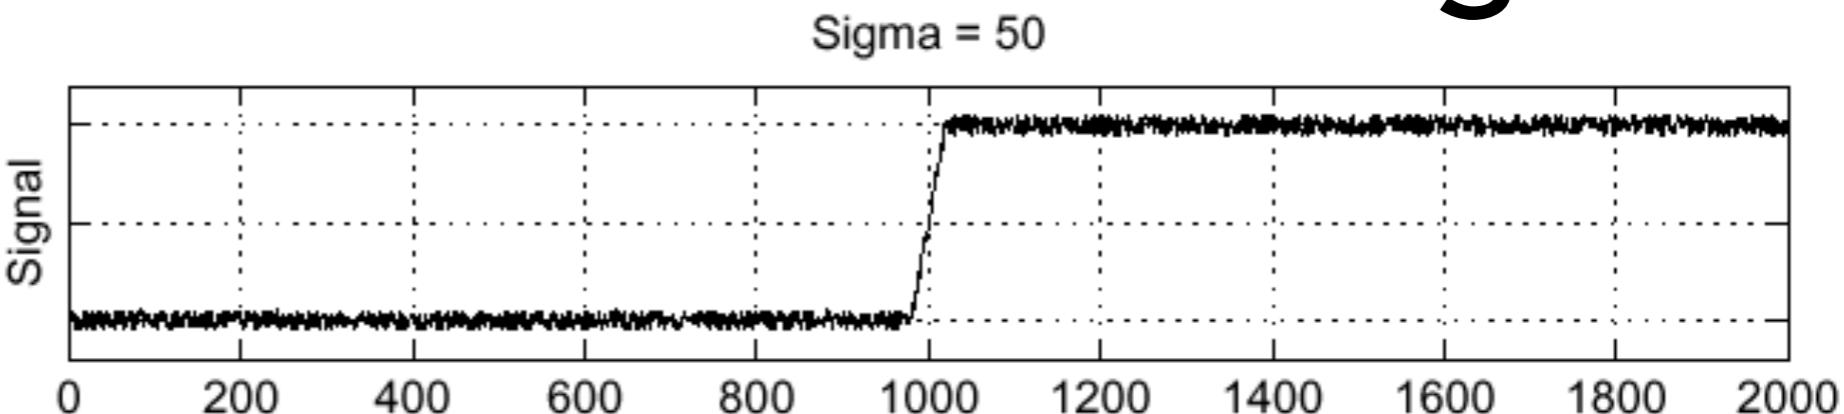
Simple and very robust way to find edges in images

Intuition can help with many other problems:
NMS, hysteresis thresholding

Recall how we found the edges:

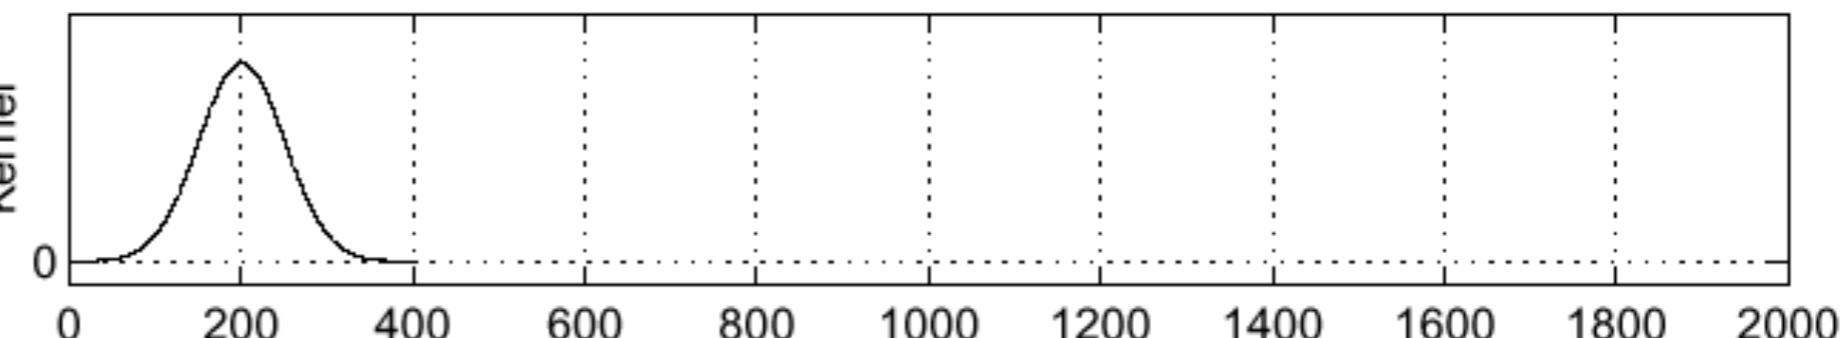
Noisy image / signal

F

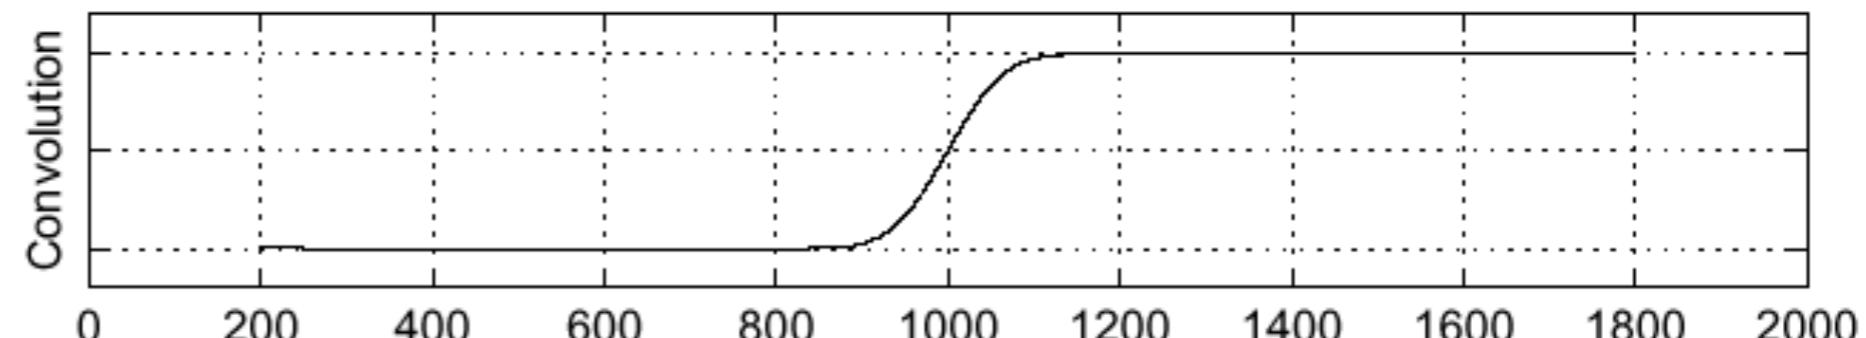


Gaussian kernel
(smooth / blur filter)

H



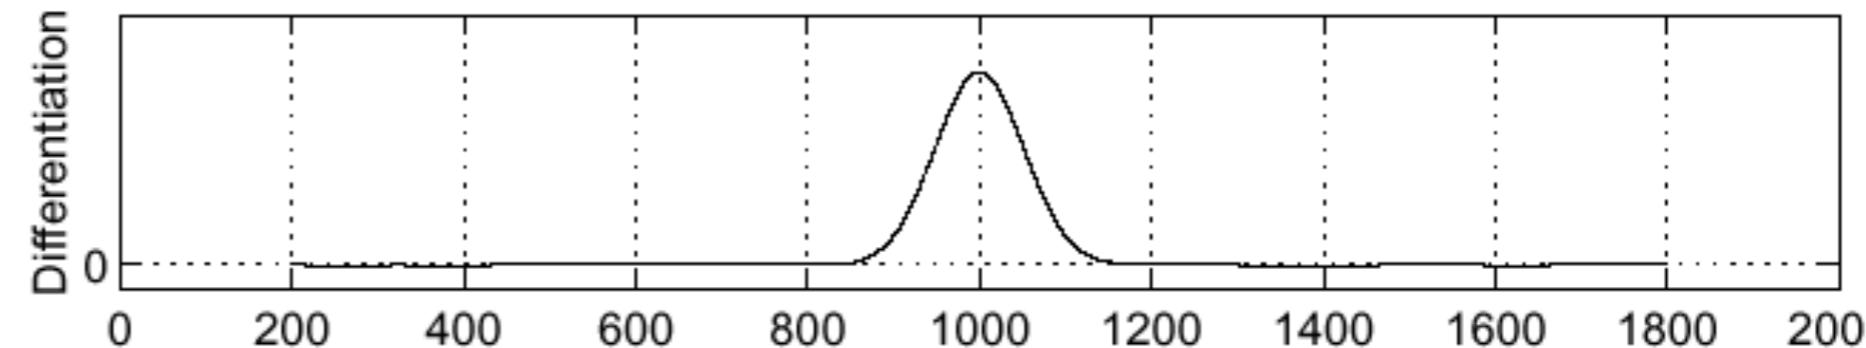
Smoothed signal $F * H$



Derivative
of smoothed
signal

$$\frac{\partial}{\partial i} (F * H)$$

$$\approx (F * H) * [1, -1]$$



We want to find peaks in this output
Canny: threshold + non-max suppression
What else can we do to find peaks?

Alternative: take the **second** derivative and set equal to 0

Computing a Second-Derivative Filter

$$F * H \quad (\text{first derivative filter}) \quad H = [1 \ -1]$$

How do we compute the 2nd derivative?

$$F * (H * H) \quad (\text{second derivative filter})$$

Computing the second derivative filter $G = H * H$:

$$G = [1 \ -2 \ 1]$$

Discrete approximation of a second derivative

Problem: Derivatives are noisy; second derivatives are even more noisy!

This is even more sensitive to noise than before!

As before, we will need to do smoothing (but this is still a problematic issue)

Recall: Derivative of Gaussian

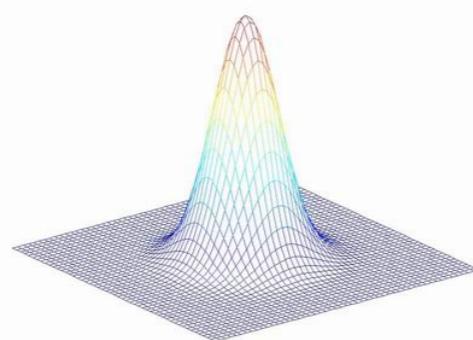
smooth derivative

$$\text{Edge detection} \approx (F * \nabla) * [1, -1]$$

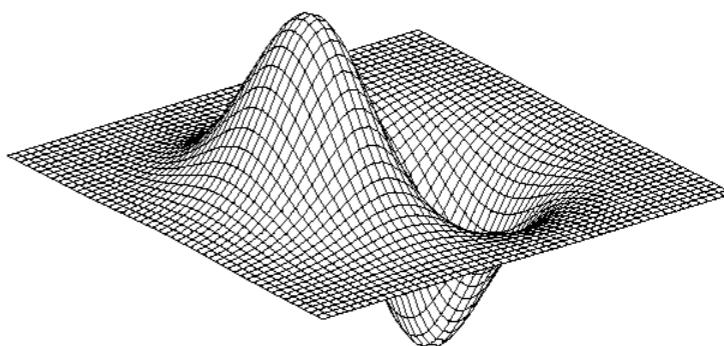
How can we make this
more efficient?

$$= F * (\nabla * [1, -1])$$

(Associative property of
convolutions)



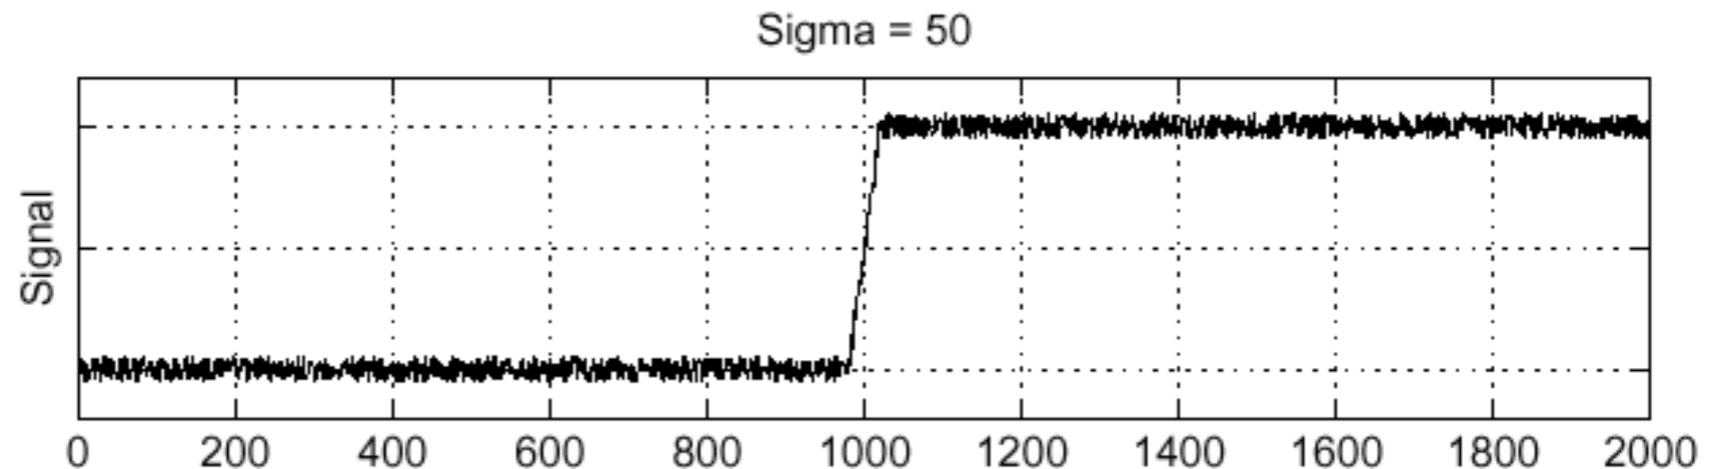
$$* [1, -1] =$$



Gaussian * **Edge filter**
 \approx derivative) = “Derivative of Gaussian”

$$= F * \text{“Derivative of Gaussian”}$$

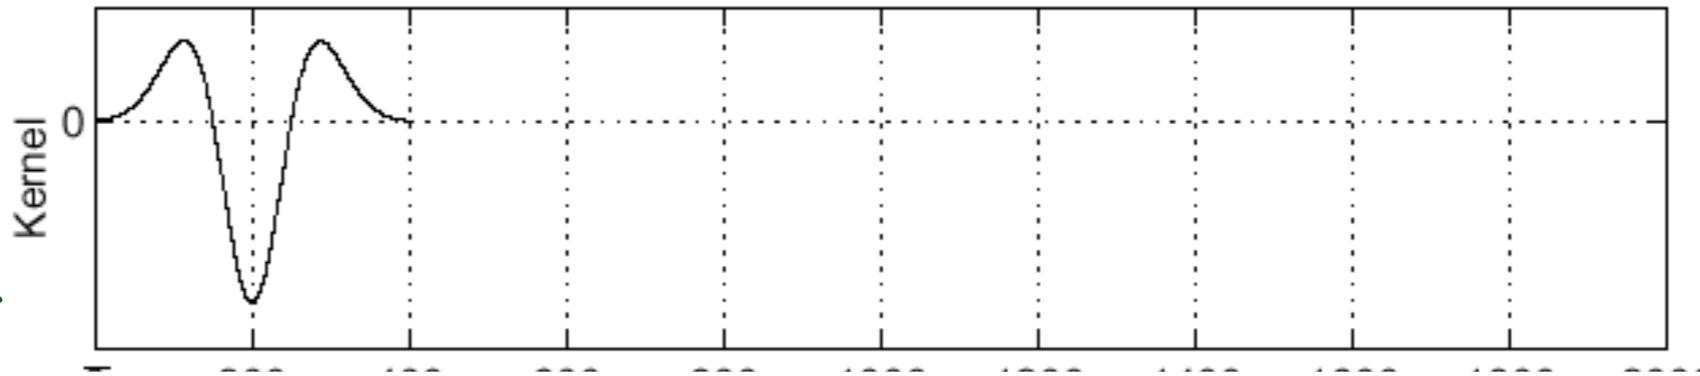
Laplacian of Gaussian (LoG)



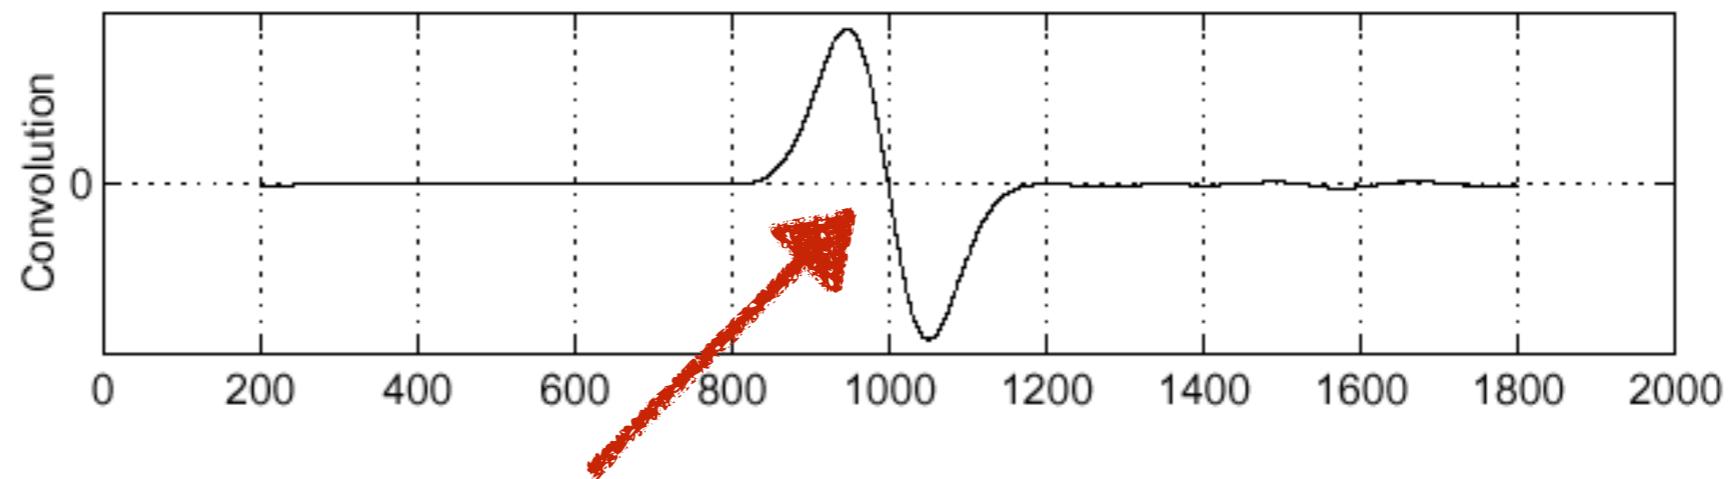
Gaussian smoothing *
second-derivative filter

$$[1 \ -2 \ 1]$$

This is called “Laplacian of
Gaussian” (LoG)

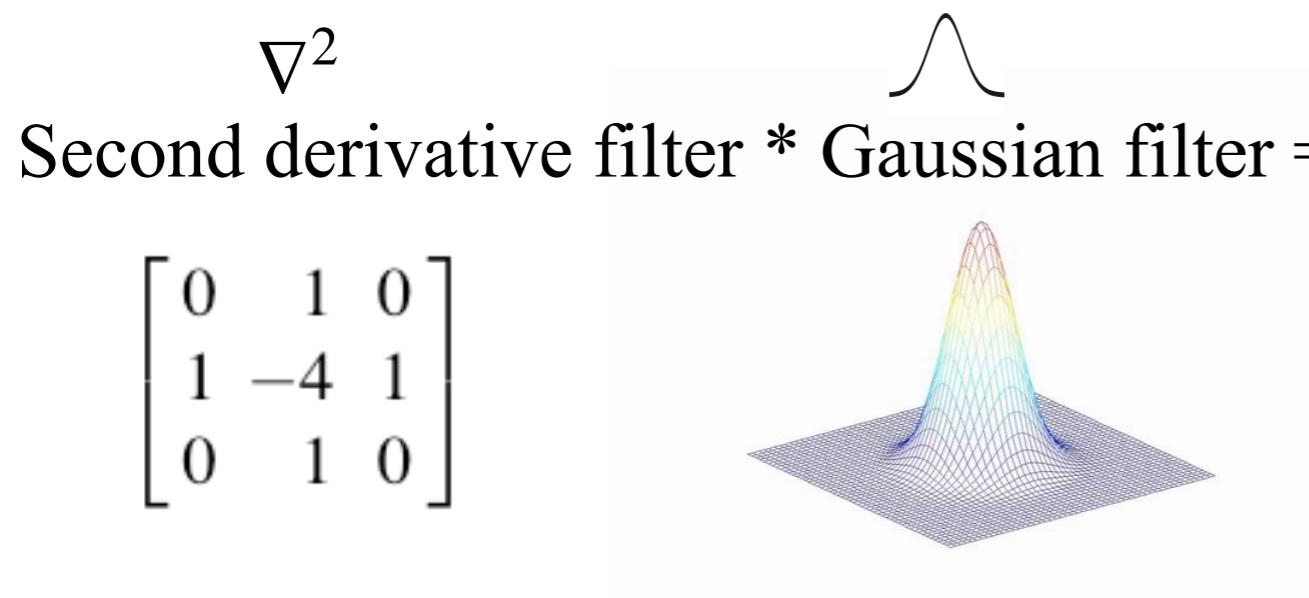


Look for zero-crossings
of second derivative:



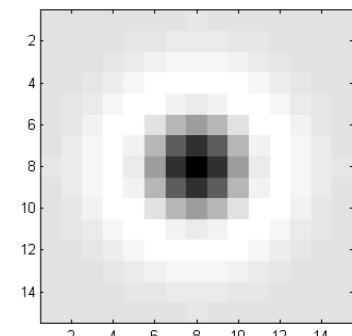
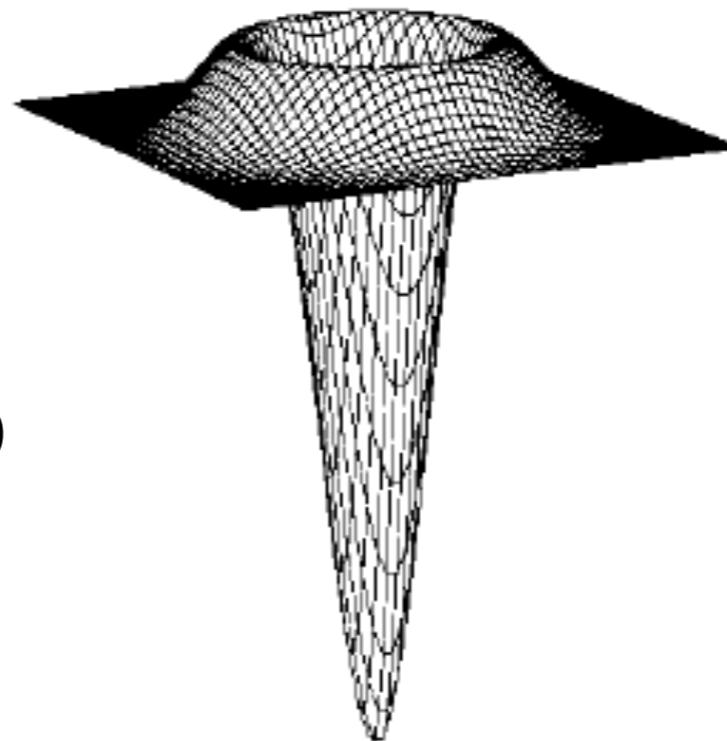
Laplacian of Gaussian (LoG)

∇^2 is the **Laplacian** operator: $\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$



$$\nabla^2 \curvearrowleft = \text{LoG}(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Laplacian of Gaussian



Convolve image with an LoG

Then find where the output is equal to 0

Pro: Fast

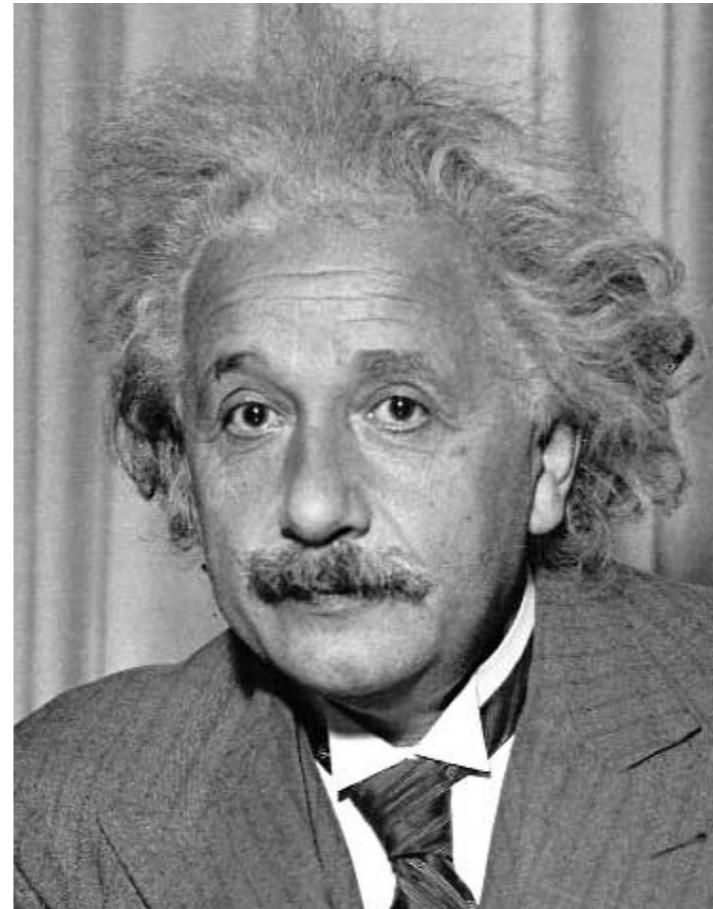
Con: Noisy (1st derivative is noisy; 2nd derivative is more noisy)

Review

- Edges
 - Derivative-of-gaussians
 - Canny edge detection (hysterisis)
 - Laplacian-of-Gaussians (LoG)

Can we use filtering to build detectors?

$F[i,j]$



$H[i,j]$



We want to find this template...

... in this image

Let's try to do template matching with correlation...

Attempt 1: correlate with eye patch

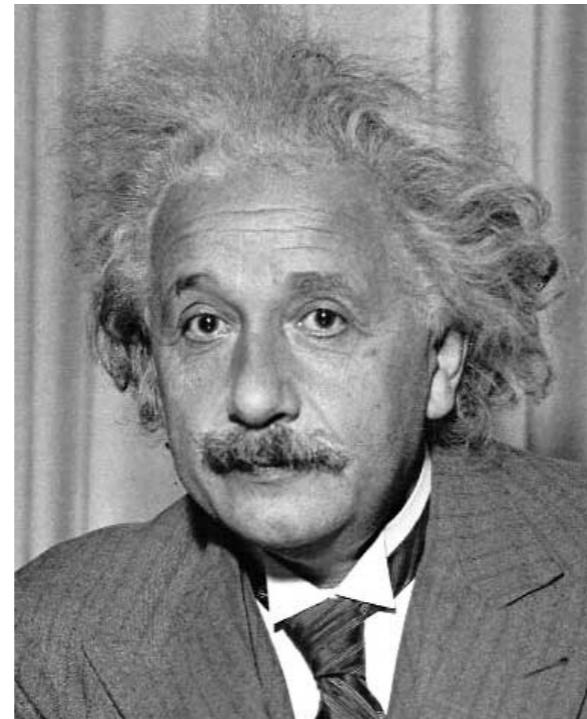
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

Why are the highest values
not all near the eye?

How can we fix this to have the
highest values only near the eyes?



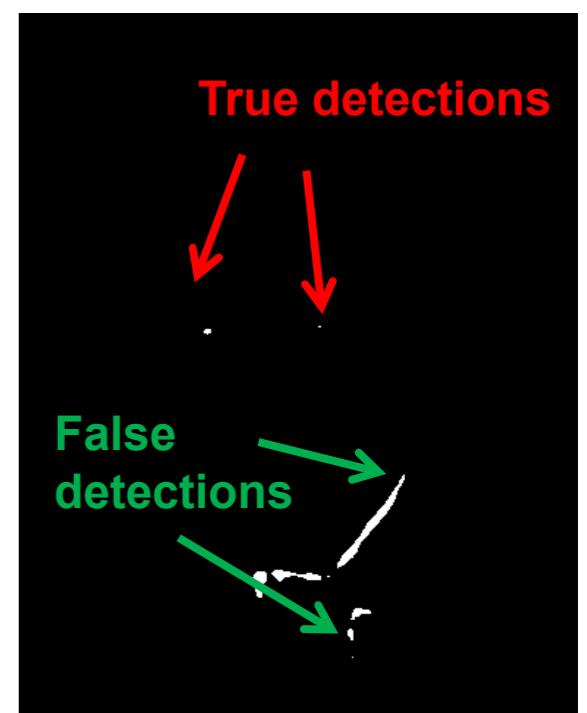
H



Input



Filtered Image (scaled)



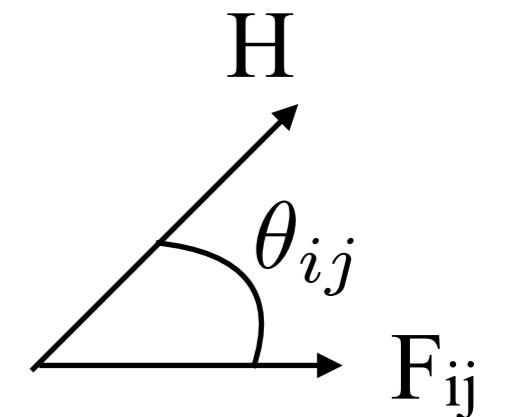
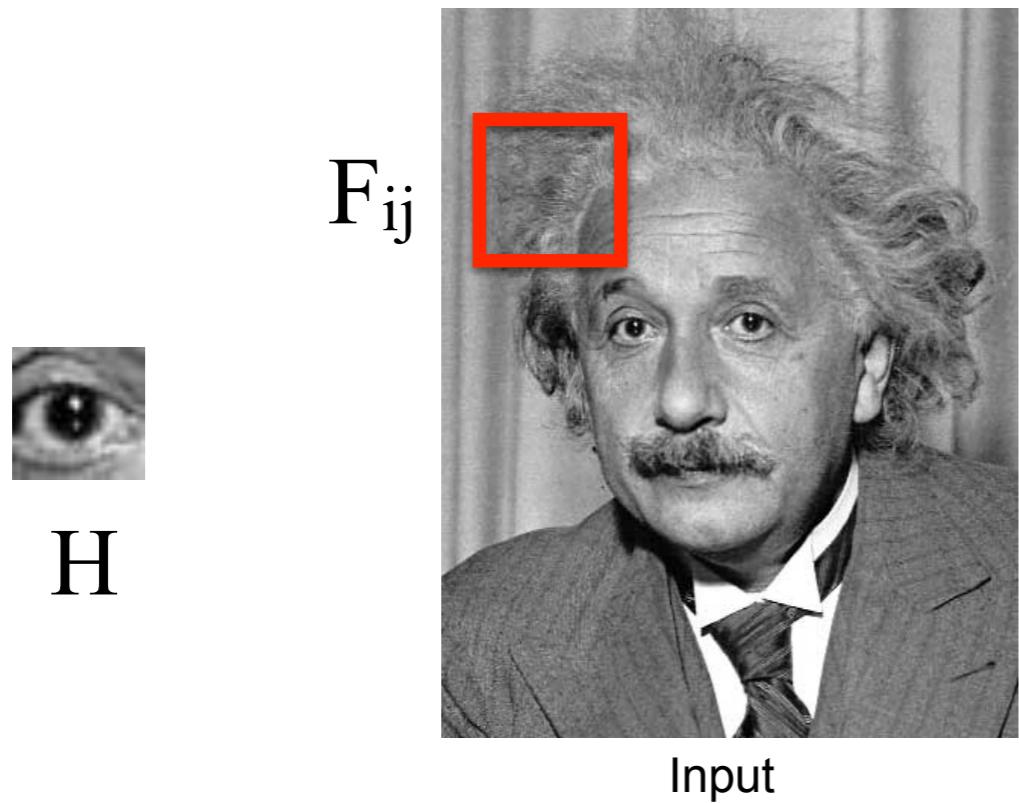
Thresholded Image

Attempt 1: correlate with eye patch

$$\begin{aligned} G[i, j] &= \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v] \\ &= H^T F_{ij} = \|H\| \|F_{ij}\| \cos \theta \end{aligned}$$

When will the inner product be high?

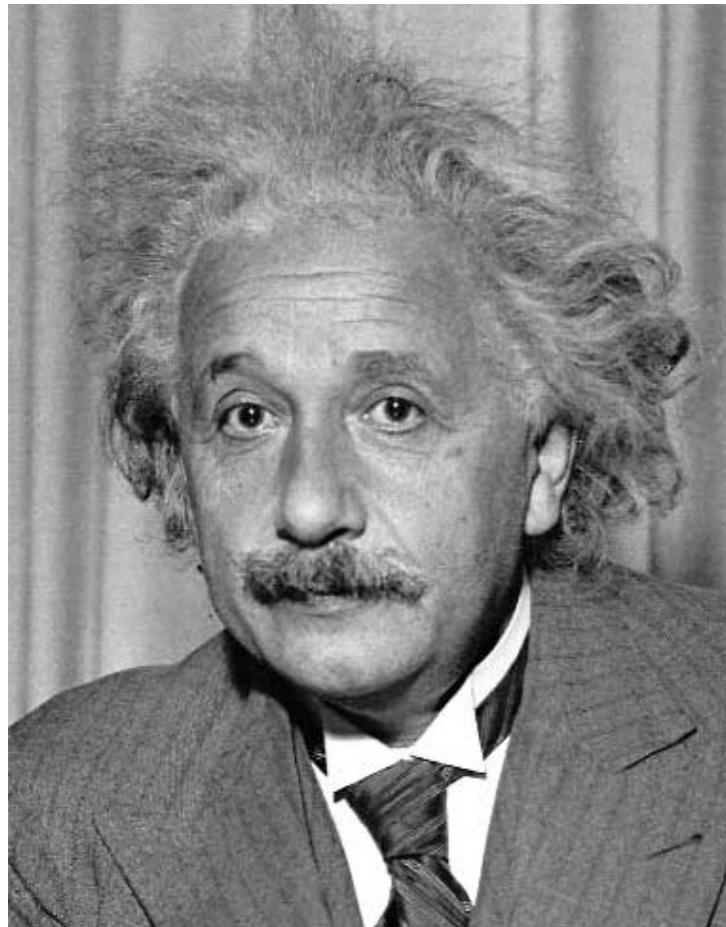
Useful to think about correlation as **inner product of vectors $H[:]$ and $F_{ij}[:]$**



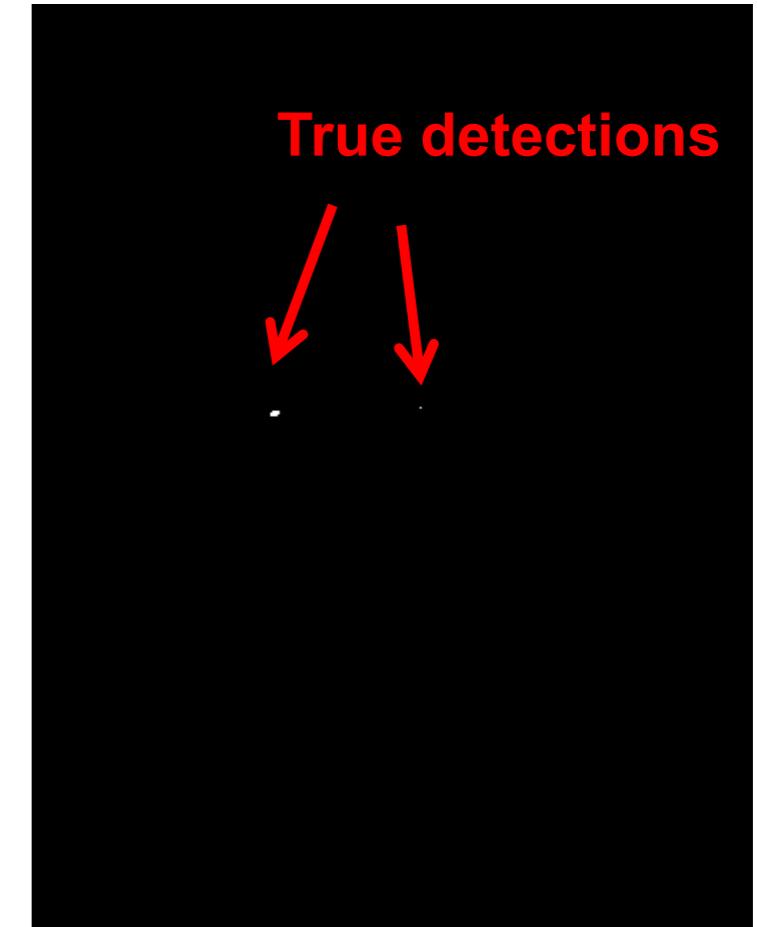
Attempt 2: sum-of-squared-difference (SSD)

$$\begin{aligned} SSD[i, j] &= \|H - F_{ij}\|^2 \quad (\text{flatten the matrices into vectors}) \\ &= (H - F_{ij})^T (H - F_{ij}) \end{aligned}$$

Can this be implemented with filtering? (correlation or convolution)



-SSD(patch, image)



Thresholded image

Attempt 2: sum-of-squared-difference (SSD)

$$\begin{aligned} SSD[i, j] &= \|H - F_{ij}\|^2 \quad (\text{flatten the matrices into vectors}) \\ &= (H - F_{ij})^T (H - F_{ij}) \end{aligned}$$

Can this be implemented with filtering? (correlation or convolution)

$$= H^T H - 2H^T F_{ij} + F_{ij}^T F_{ij}$$

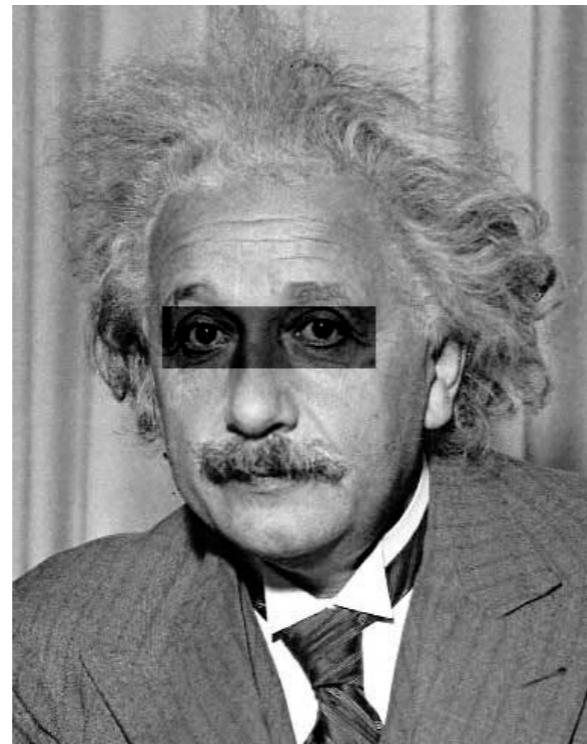
Sum of squared elements in filter (filter norm)

Cross-correlation of filter H with image F

Cross-correlation of each image patch with itself

Trick: Compute a “squared image” (element-wise square), then convolve with a “ones” filter to add up elements within each region

What will SSD find here?



Eyes are in shadow
(eyes have been
darkened by .5 scale
factor)



-SSD(patch, image)

SSD has strongest
response on the shirt
(not the eyes)

SSD is not robust to changes in lighting!

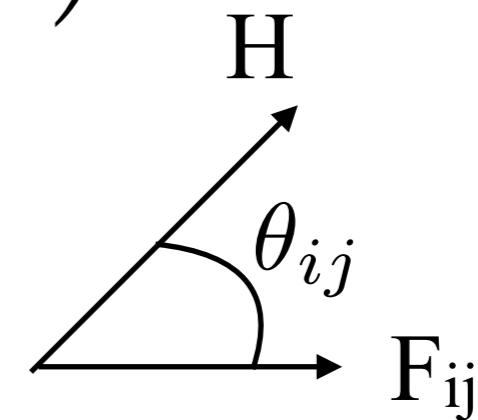
SSD is not “lighting invariant”

Attempt 3: Normalized cross correlation (NCC)

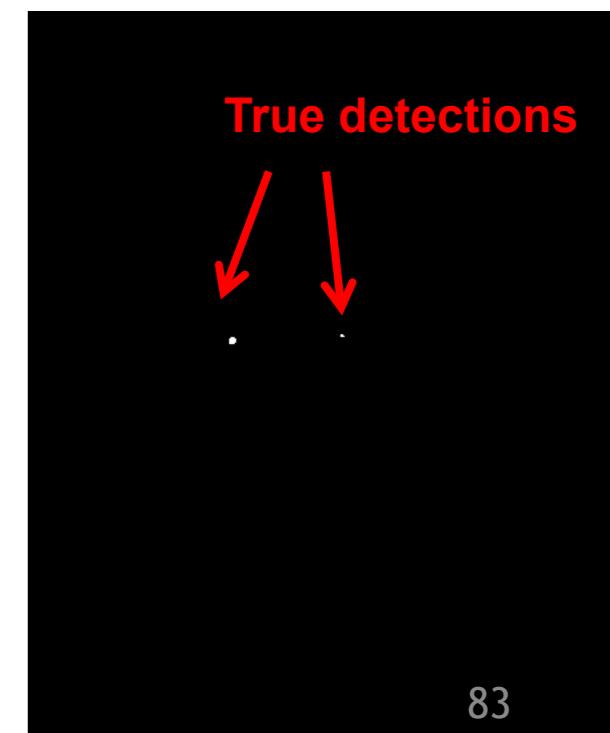
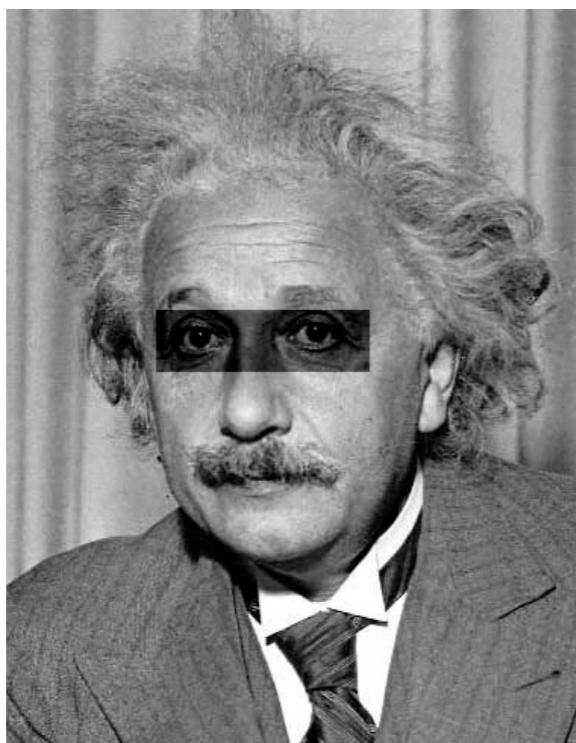
NCC is lighting / scale invariant!

$$\begin{aligned} NCC[i, j] &= \frac{H^T F_{ij}}{\|H\| \|F_{ij}\|} = \left(\frac{H}{\|H\|} \right)^T \left(\frac{F_{ij}}{\|F_{ij}\|} \right) & \|X\| = \sqrt{X^T X} \\ &= \frac{(\|H\| \|F_{ij}\| \cos \theta)}{\|H\| \|F_{ij}\|} \\ &= \cos \theta \end{aligned}$$

NCC



NCC thresholded



Summary: CC vs SSD vs NCC

$$CC_{ij} = H^T F_{ij}$$

**Simple cross-correlation;
Sensitive to “bright” patches**

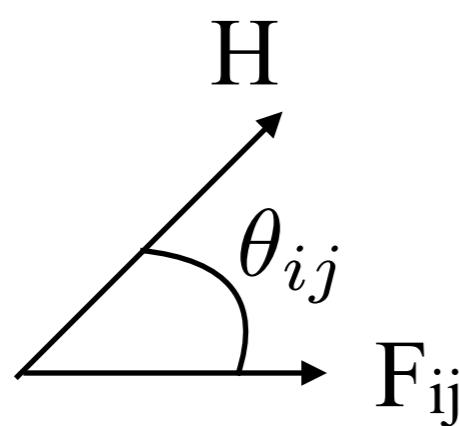
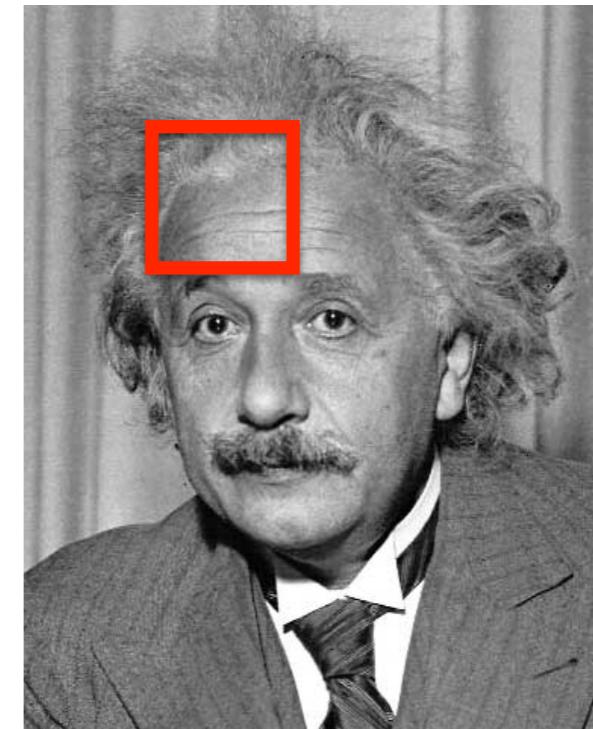
$$SSD_{ij} = \|H - F_{ij}\|^2$$

**Looks for closest direct match;
sensitive to lighting changes**

$$NCC_{ij} = \cos(\theta_{ij})$$

**Invariant to lighting changes;
very common and powerful approach!**

$$F_{ij}$$



$$H$$