

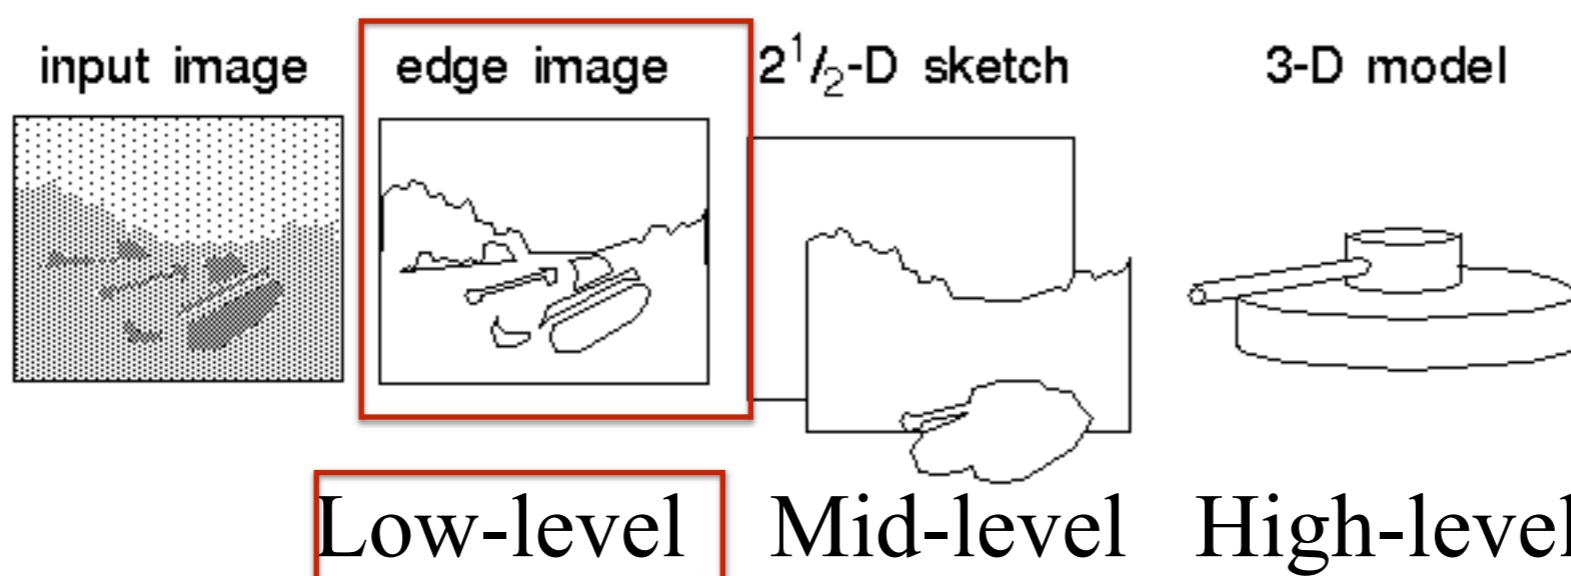
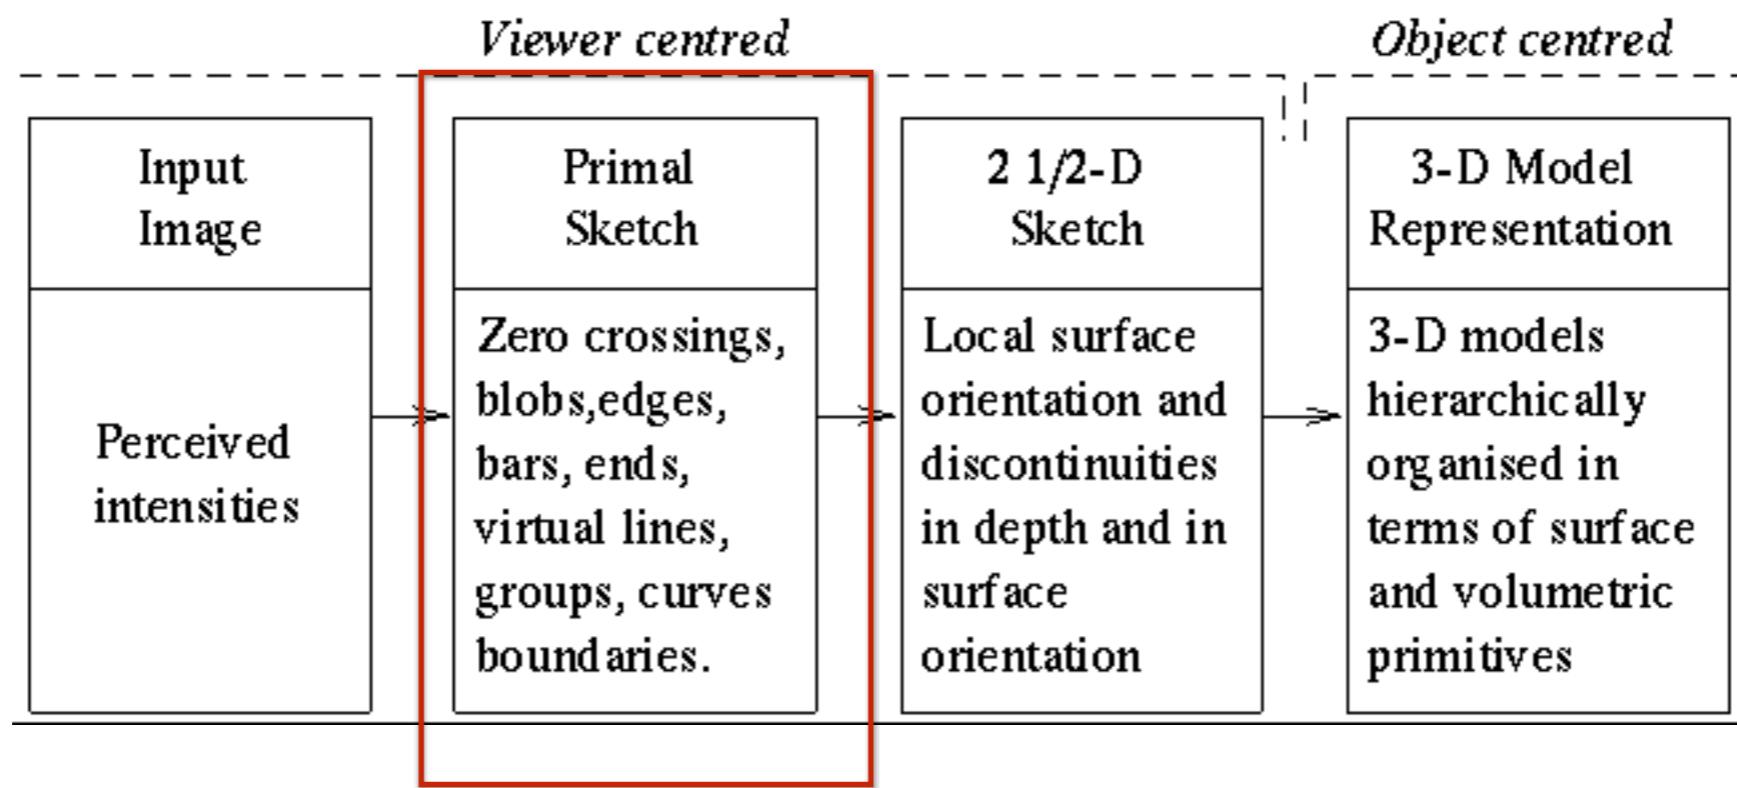
Motion and flow



Logistics

- HW 2 has been released
- It is due on Oct 2 (in 2 weeks).

David Marr's Taxonomy of Vision



Outline

- Egomotion
 - Time-to-Contact, Parallax, Focus-of-Expansion
- Optical Flow

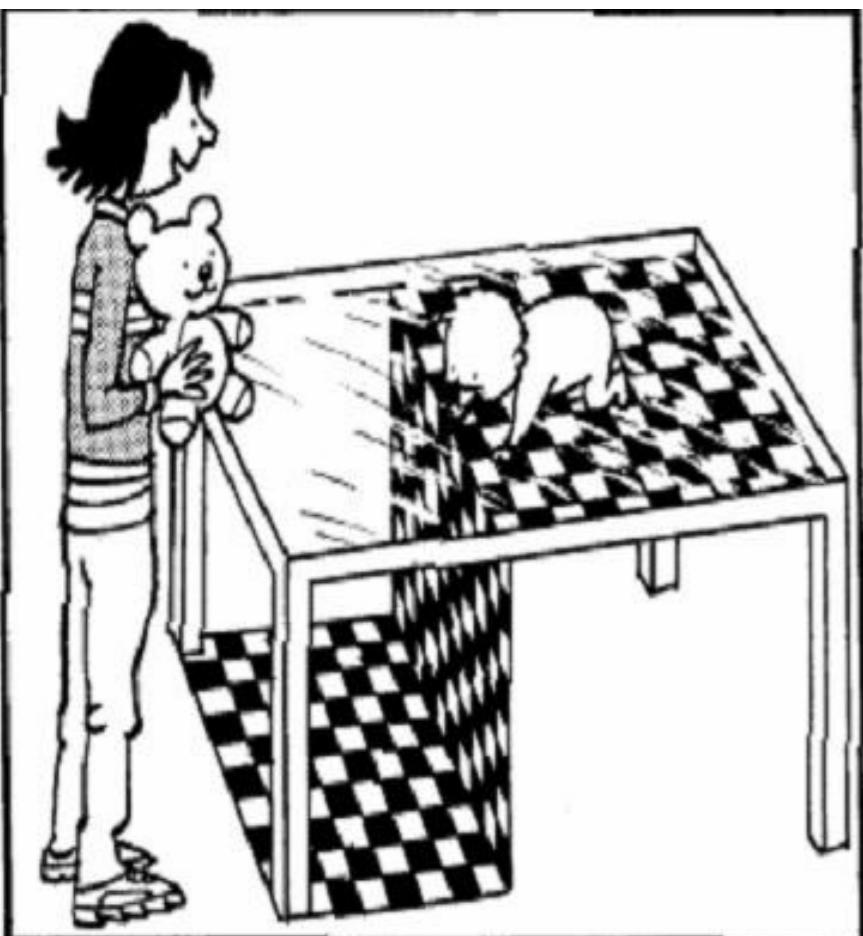
Moving is a part of life!

Sea squirt:

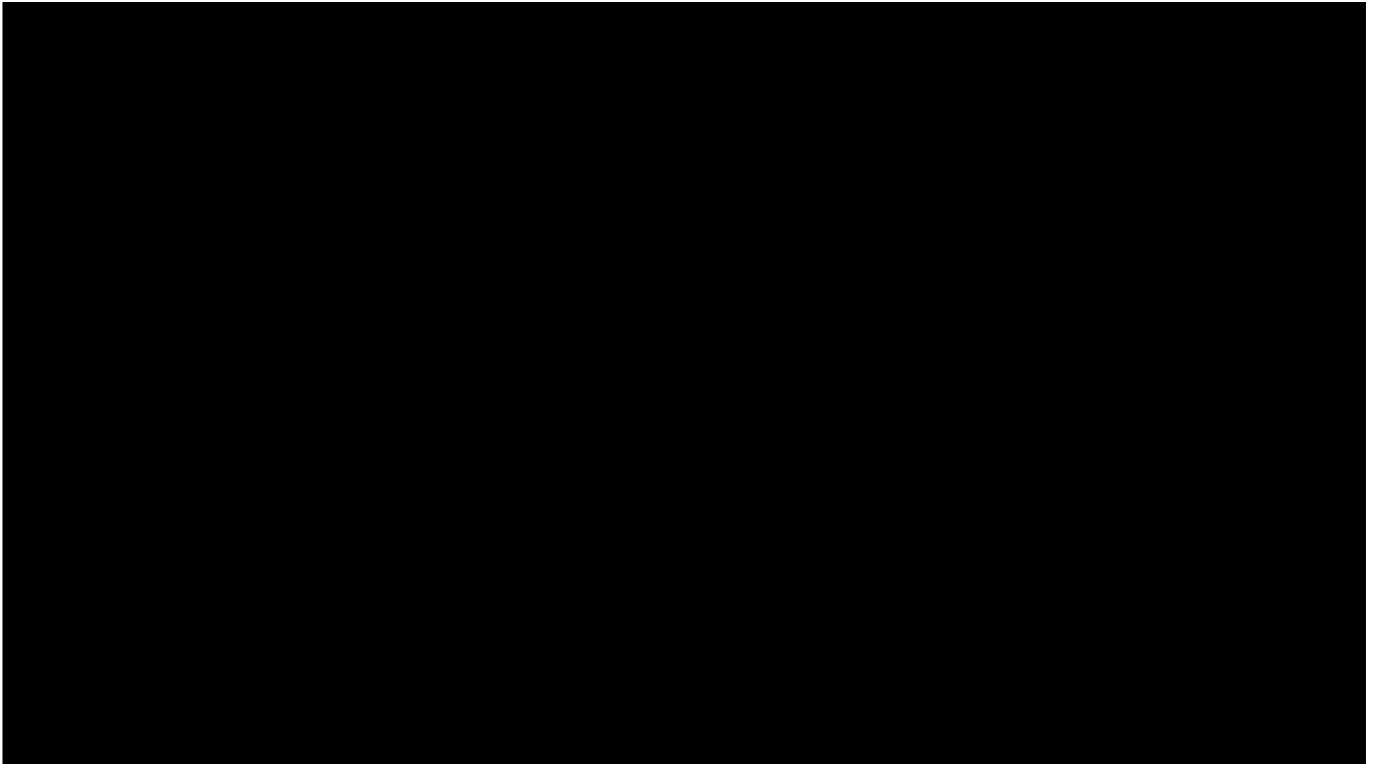


- The sea squirt moves through water and attaches itself to a rock or other stable structure.
- Once it finds a suitable place to attach itself, it “eats” its own brain (its brain is absorbed by its body).
- Being permanently attached to a home makes the sea squirt’s brain that control its locomotion unnecessary.
- This is an example of the connection between movement and intelligence in animals!

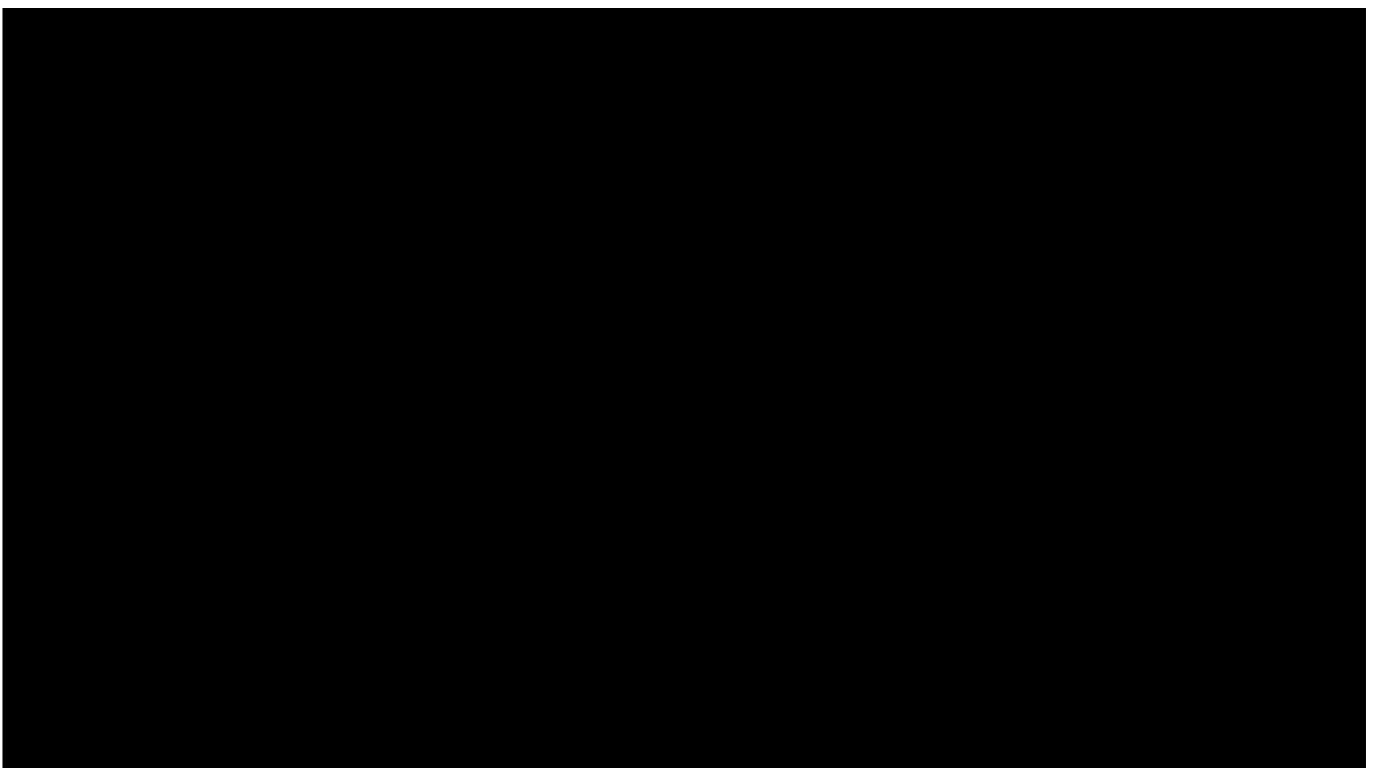
Human development: Crawling teaches babies depth perception



After 1
week of
crawling:



After
several
weeks of
crawling:



Today's goal

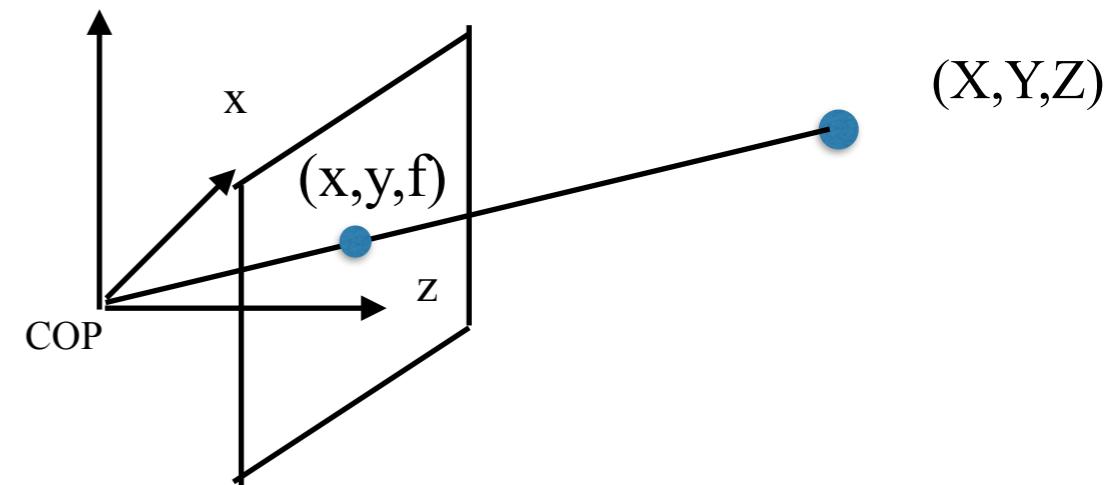


Understand the space of image transformations that we see when we move
egomotion - the task of estimating our own motion (e.g. from a video)

Optical Flow



Motion



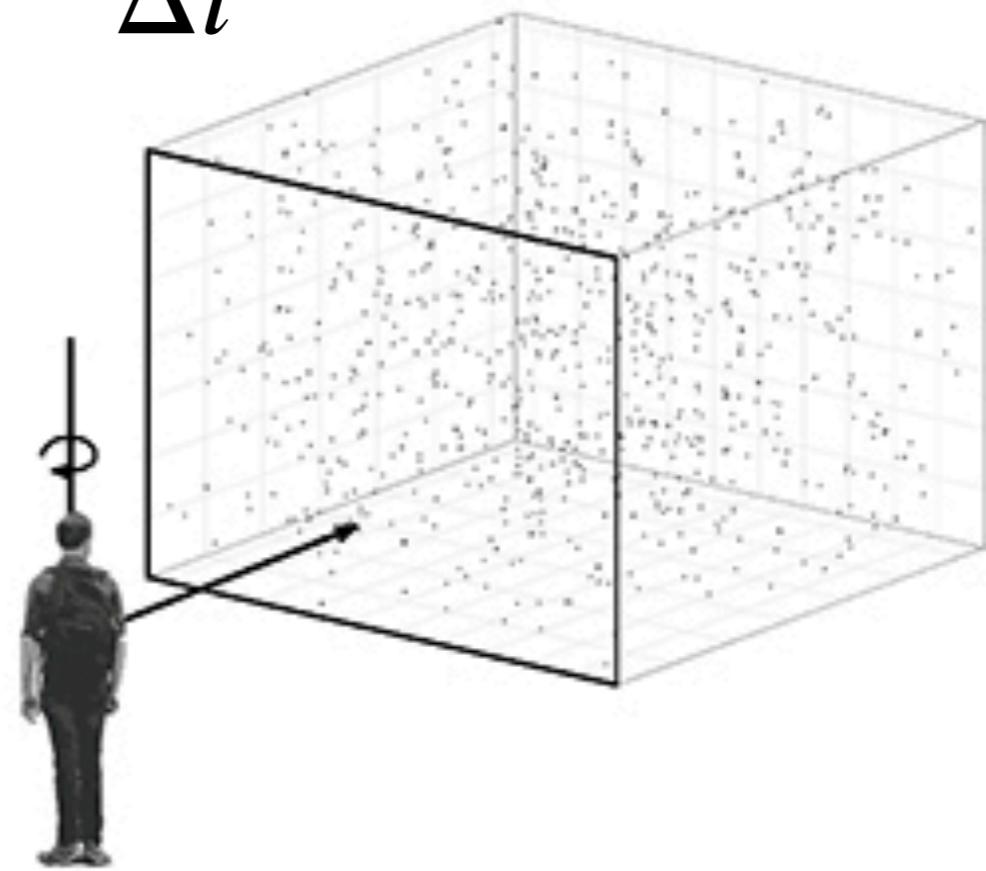
When a point (X, Y, Z) moves relative to the camera, its projection on the image (x, y) moves as well

This movement in the image plane is called “optical flow”

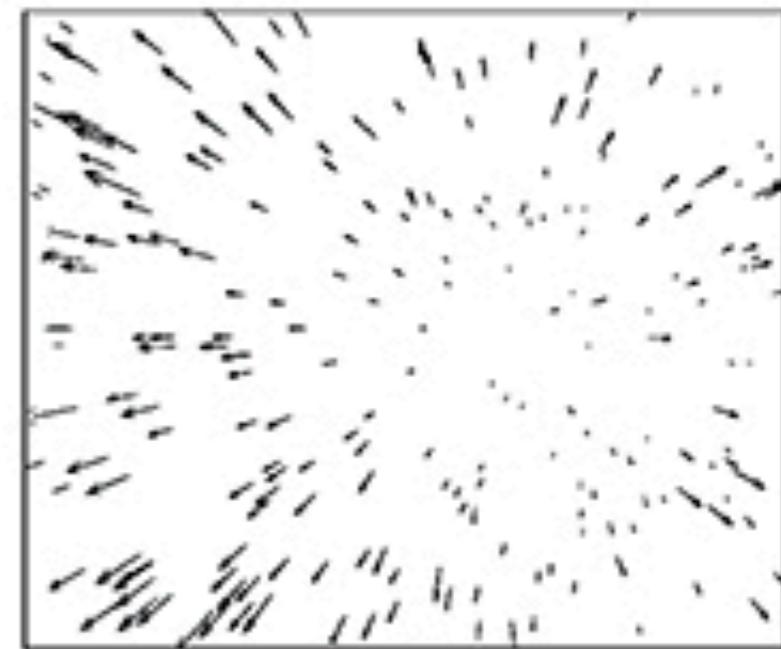
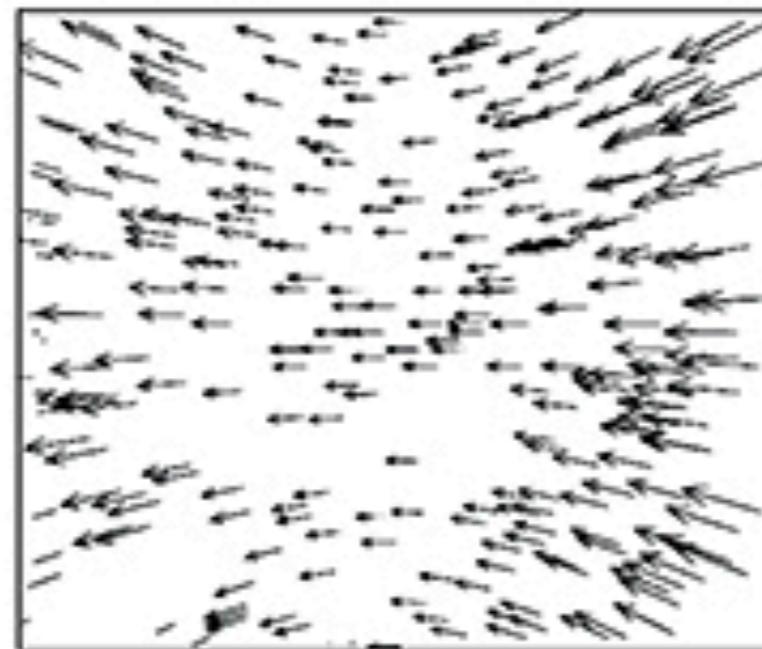
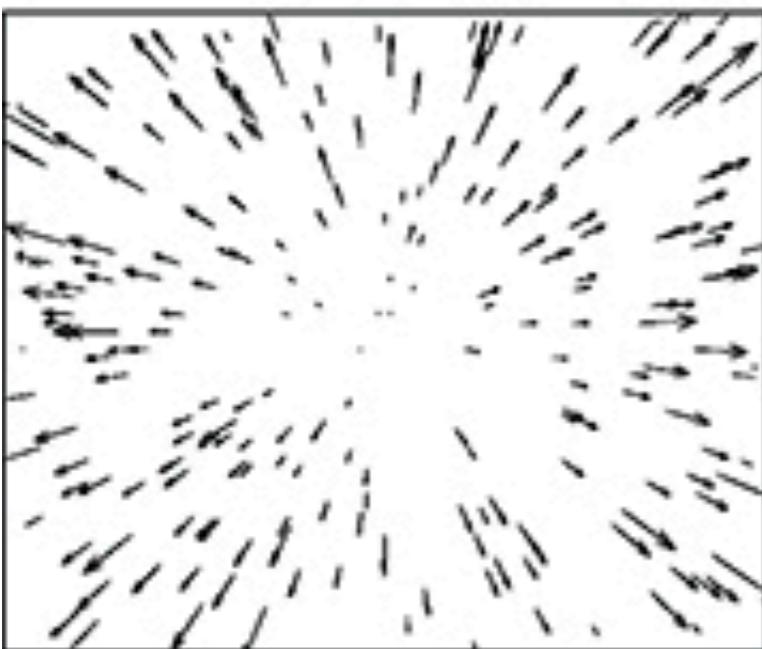
Suppose the point at pixel (x, y) moves to the pixel $(x + \Delta x, y + \Delta y)$ in time Δt , then the two components of the optical flow at (x, y) are:

$$u = \frac{\Delta x}{\Delta t}, \quad v = \frac{\Delta y}{\Delta t}$$

We can visualize $u = \frac{\Delta x}{\Delta t}$, $v = \frac{\Delta y}{\Delta t}$ as a “flow field”

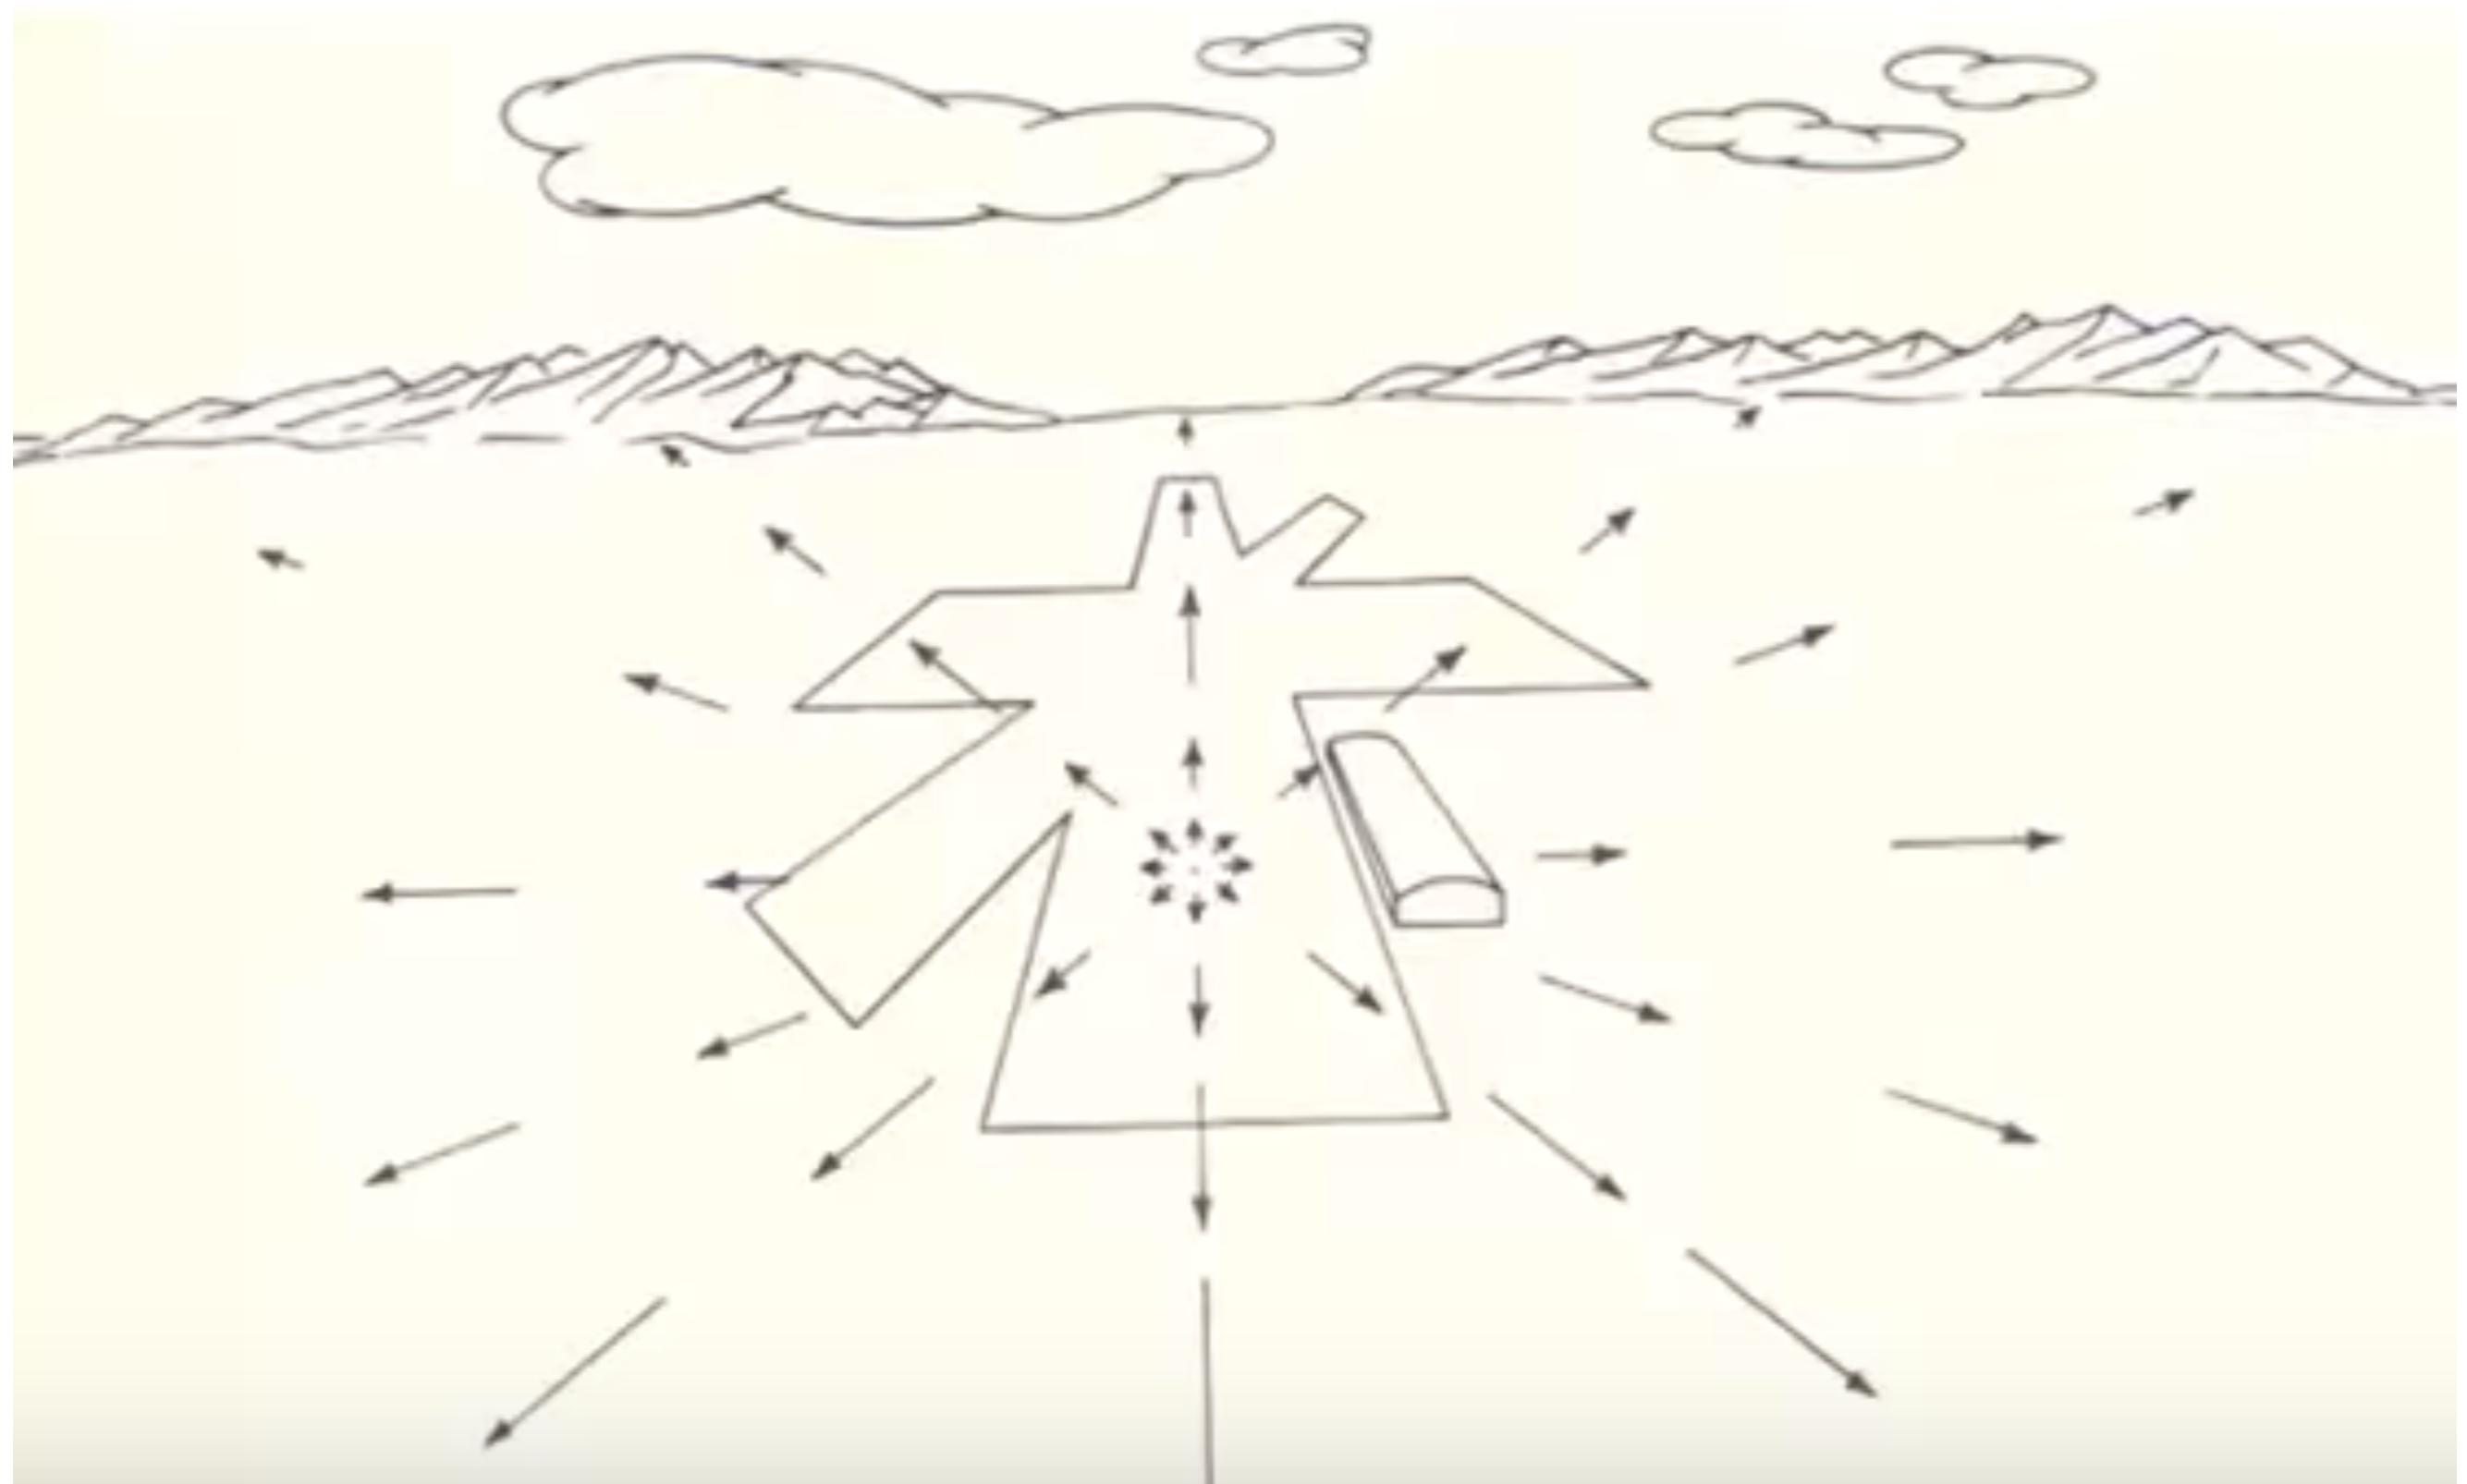


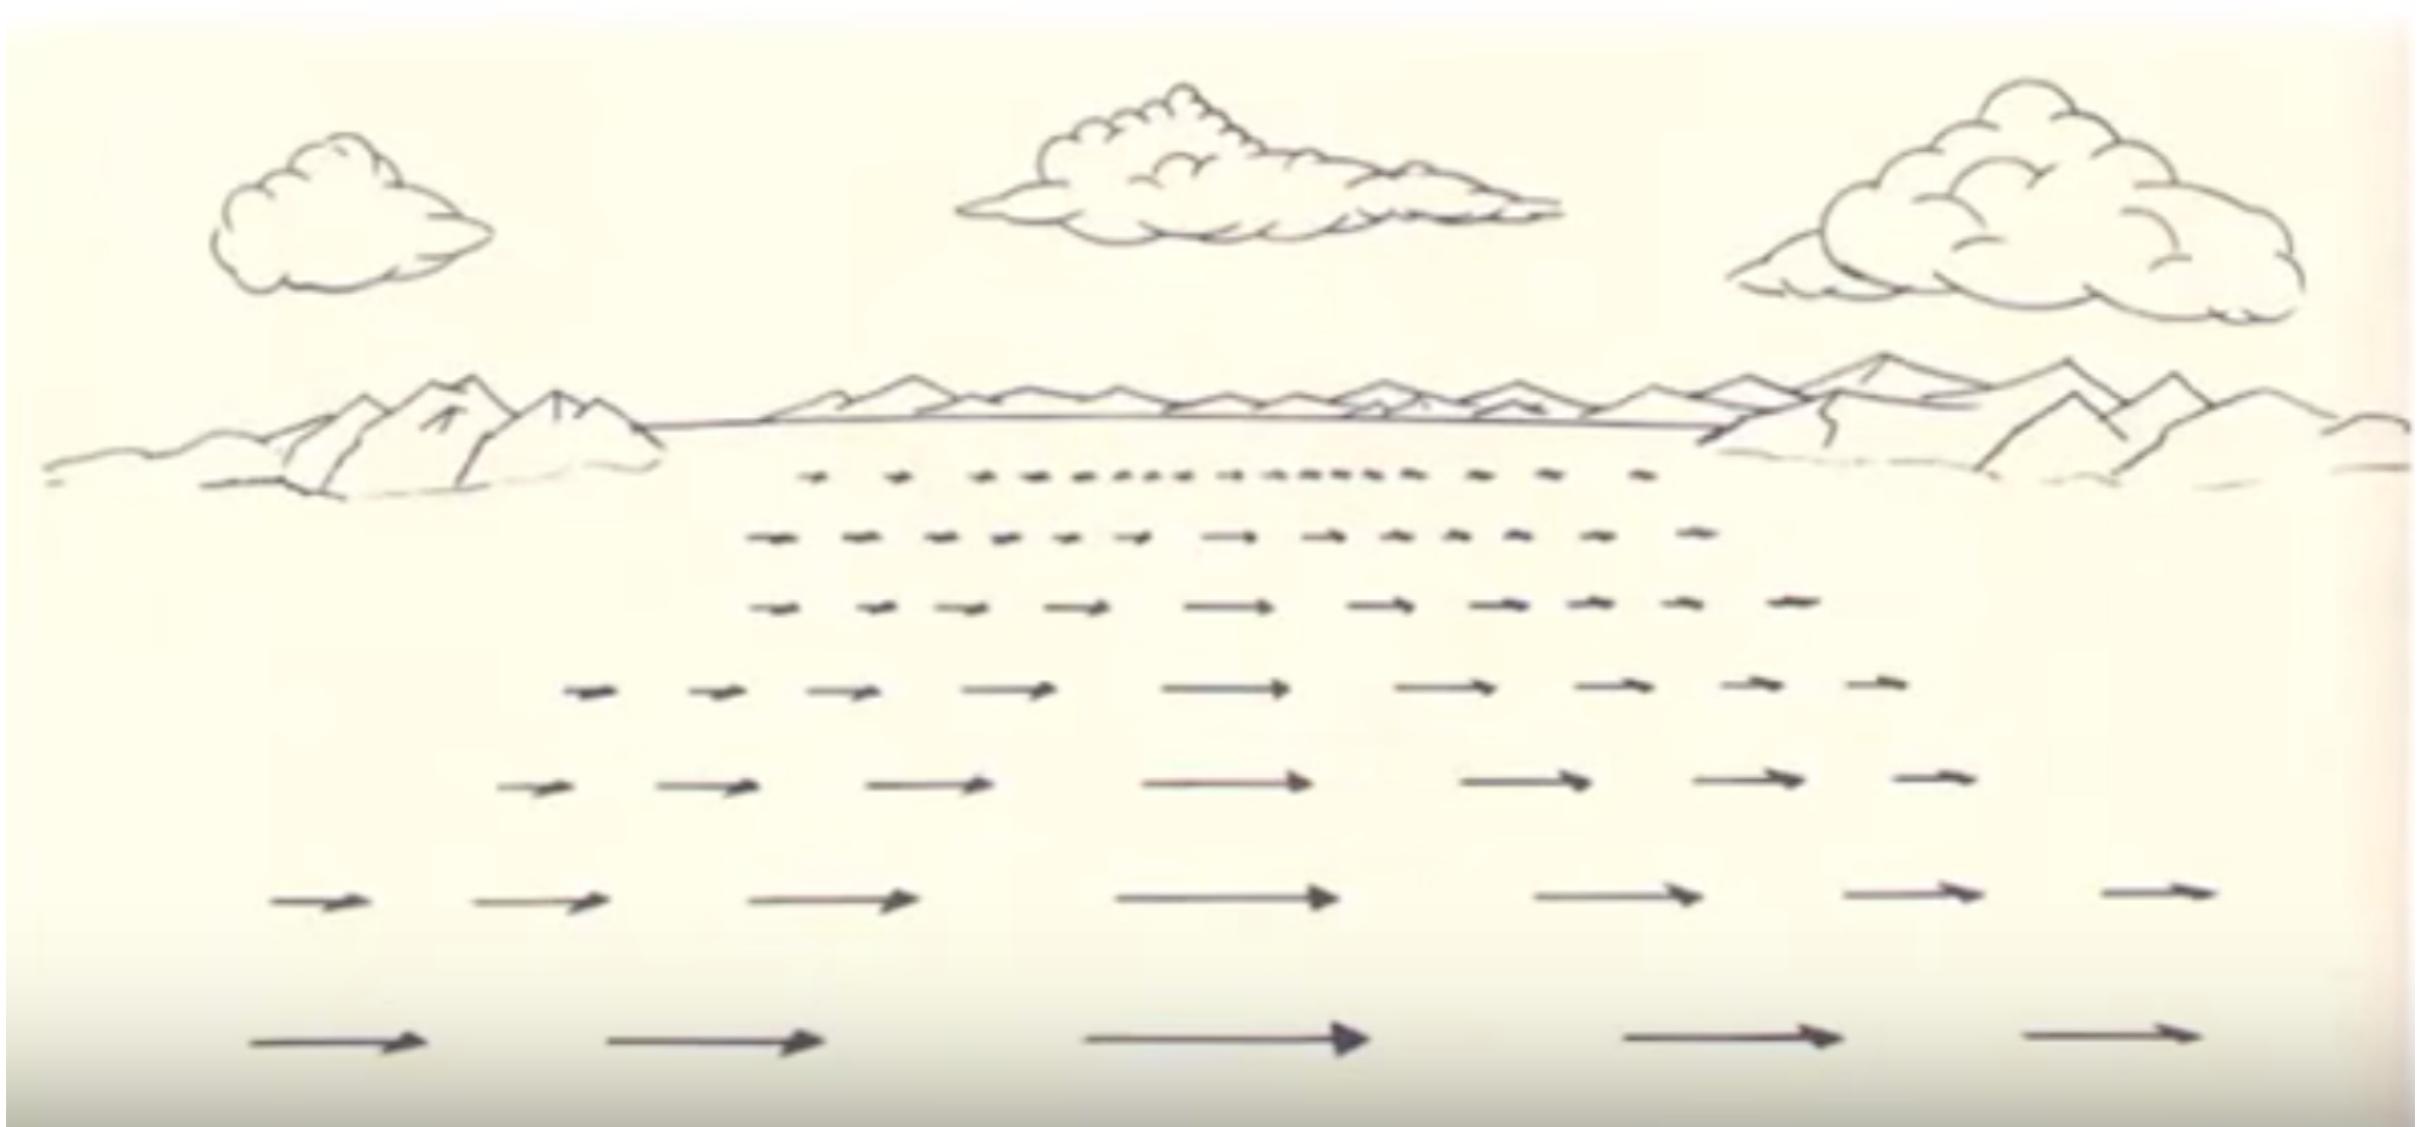
What direction do you think each of these flow fields correspond to?



Flow field: Moving closer towards a point

We can determine the direction of motion by finding the “center of flow”





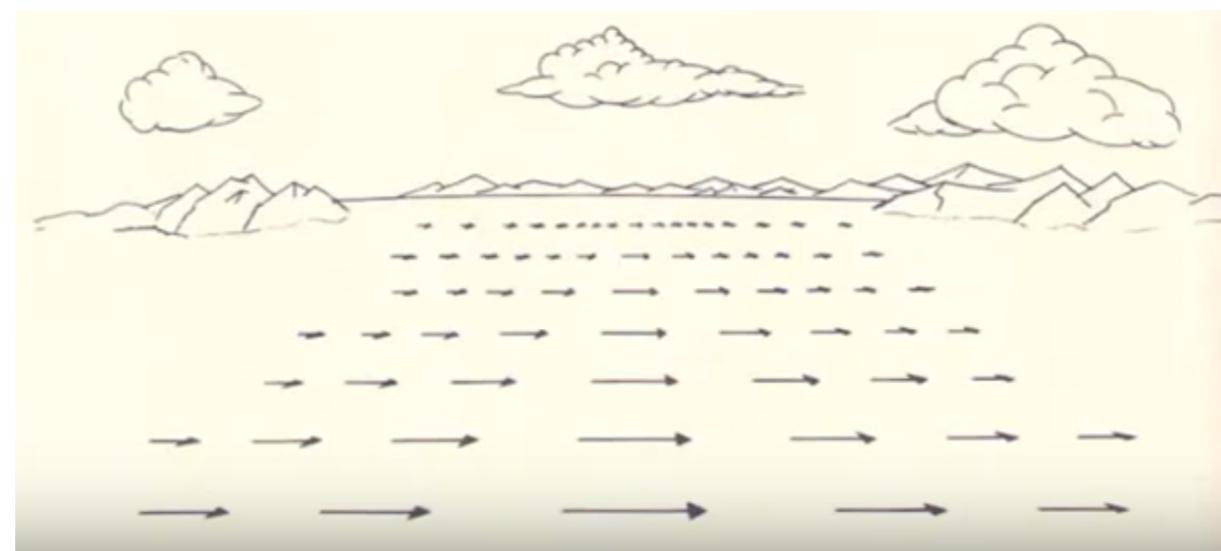
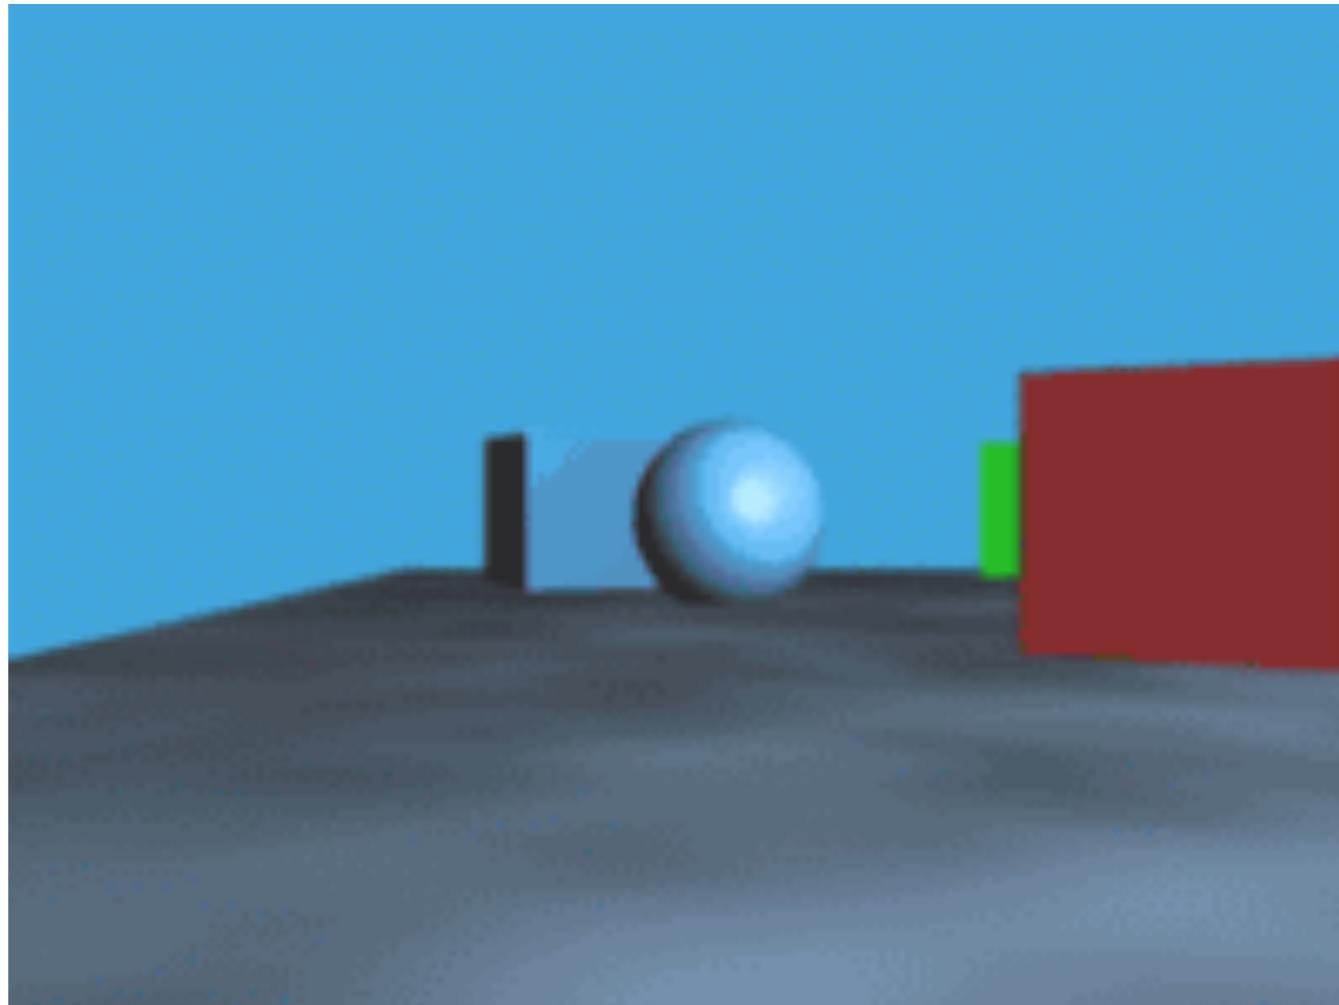
Why is the flow smaller for more distant objects?

Apparent motion on the image is smaller due to perspective projection
(similar to how more distant objects appear smaller in the image)

Thus, we can use the size of the flow field of a static world / moving camera to estimate depth!

Flow due to translating cameras reveals depth about scene

This pixel motion is known as “parallax”



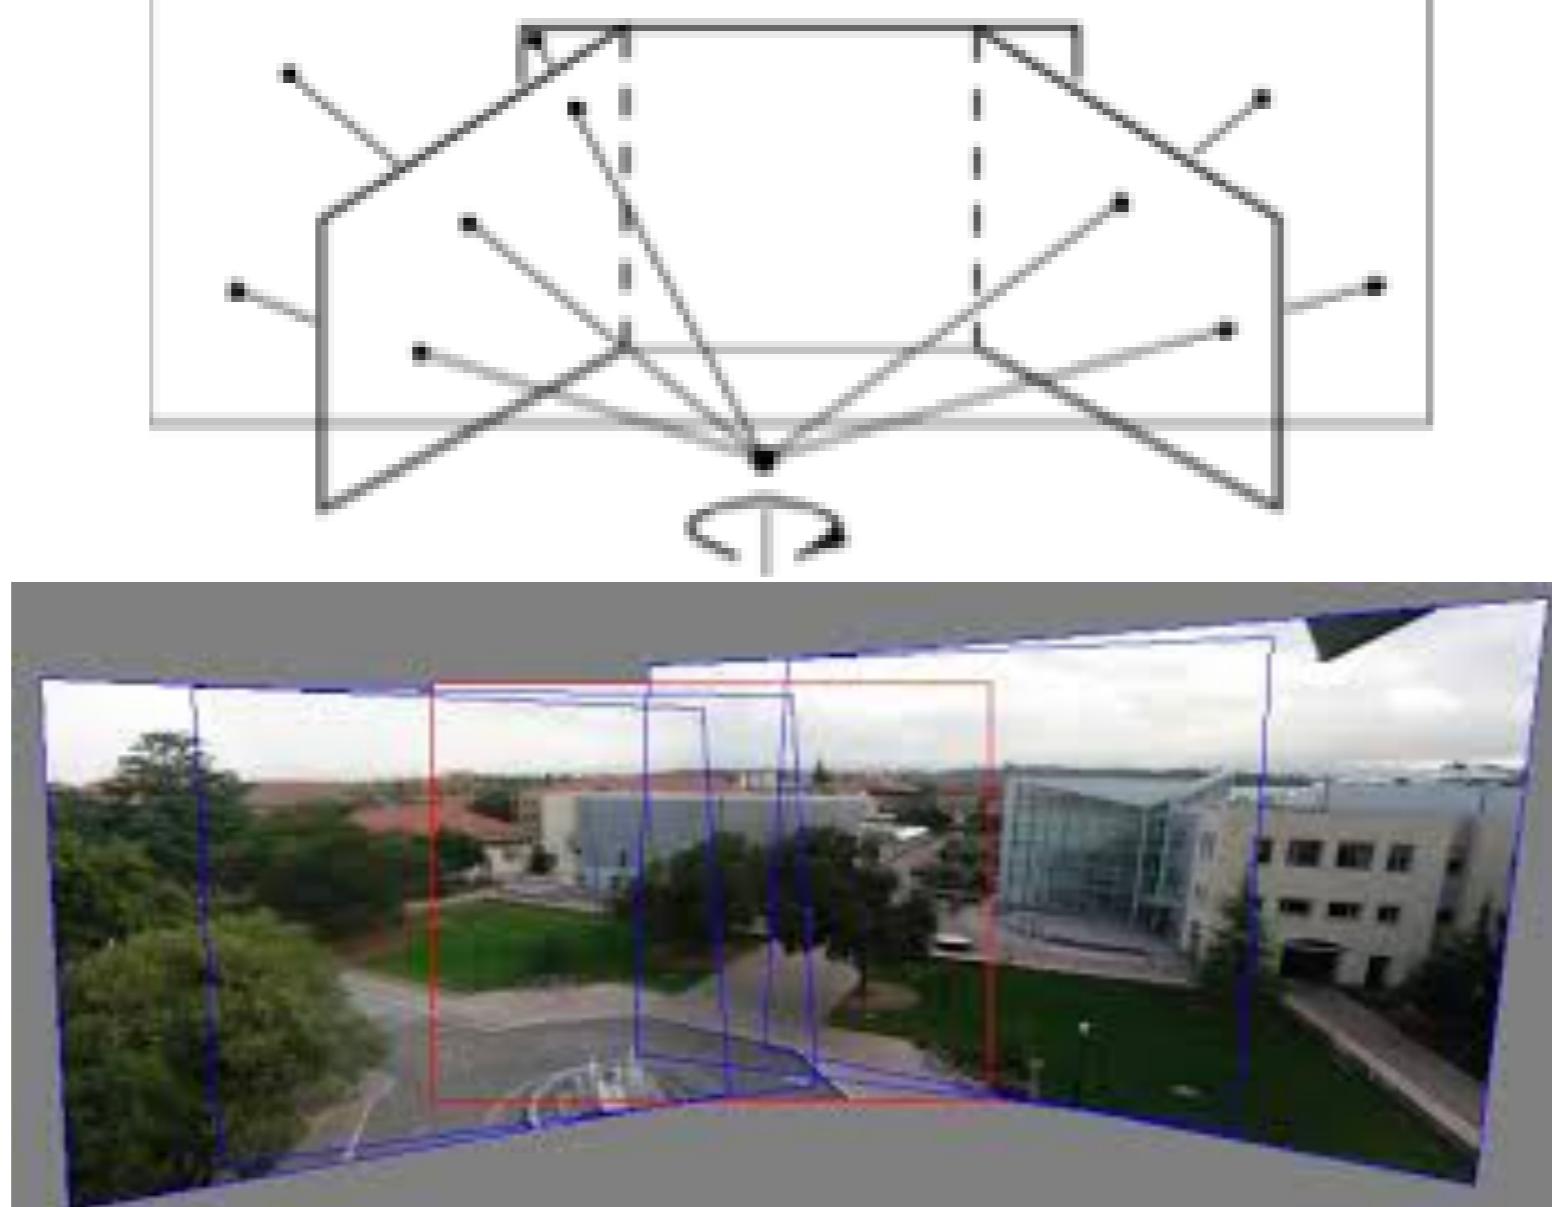
Rotating cameras

What does a rotating camera tell us about scene depth?

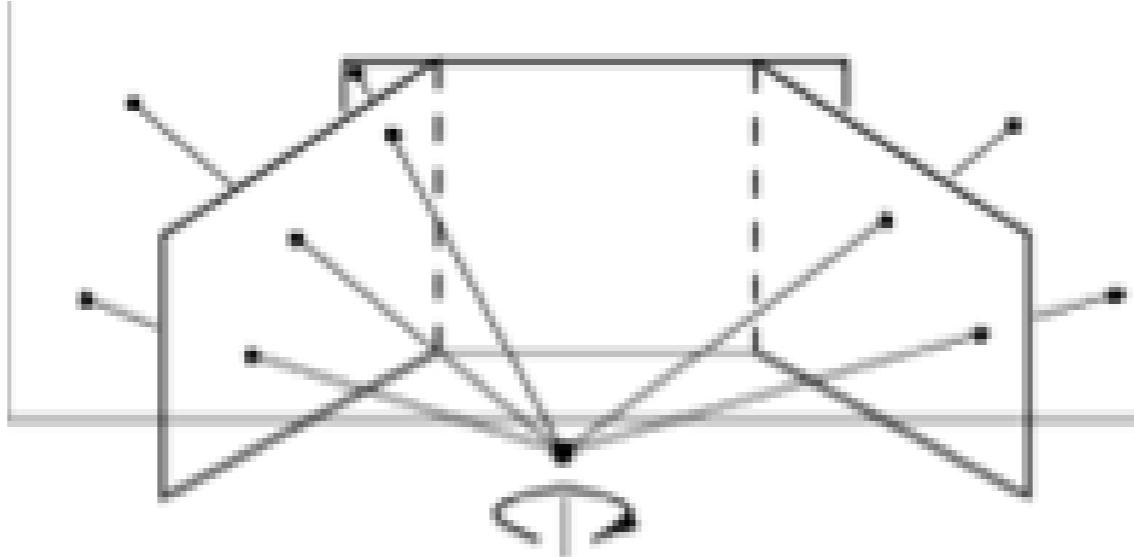
How do pixels change if we just rotate (not translate) the camera?

Since pixels are related by a homography transformation,
this reveals **nothing about the scene depth!**

The motion is determined only by the camera rotation, not the scene depth



Rotations about camera center



Relation between 3D camera coordinates:

$$\begin{bmatrix} X_2 \\ Y_2 \\ Z_2 \end{bmatrix} = R \begin{bmatrix} X_1 \\ Y_1 \\ Z_1 \end{bmatrix}$$

**Rotation
matrix**

3D->2D projection:

$$\lambda_2 \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} = \begin{bmatrix} f_2 & 0 & 0 \\ 0 & f_2 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_2 \\ Y_2 \\ Z_2 \end{bmatrix}$$

**Intrinsic camera
matrix K**

Combining both:

$$\lambda \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} = K R K^{-1} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

**This is a 3x3
homography!**

**Scene depth is not
part of the
equation**

**This is why 360
degree cameras work**

Summary

- Camera rotation causes apparent object movement, **independent** of the depth of the scene
- Camera translation causes apparent object movement, **dependent** of the depth of the scene

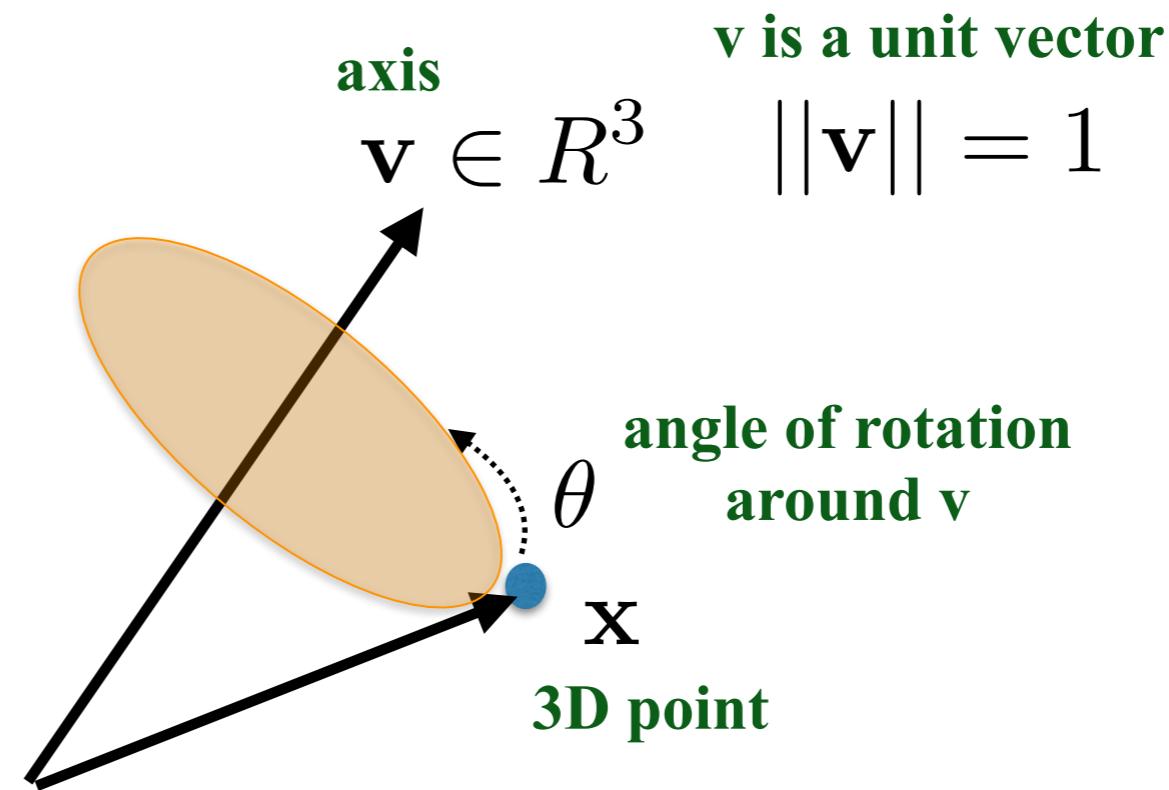
Roadmap

We want to describe the relationship between camera movement, depth, and flow.

Here's how we will do this:

- Let's parameterize rigid-body camera movement by translation t and a rotation
- We will look at flow-fields produced by a small camera movement:
 - Flow fields: $u(x,y)$ and $v(x,y)$ for every pixel (x,y)
 - u is the horizontal component of flow; v is the vertical component of flow
- We'll explicitly derive an equation that shows how flow depends on camera movement and the depth of each point

Axis-angle representation of rotations

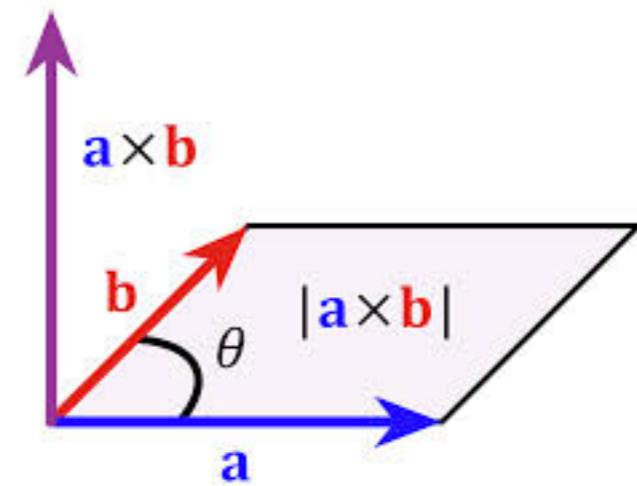


Can represent with 3 numbers as: θv

Review: cross products

Cross product: $\mathbf{a} \times \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \sin \theta \mathbf{n}$

$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} a_2 b_3 - a_3 b_2 \\ b_1 a_3 - a_1 b_3 \\ a_1 b_2 - a_2 b_1 \end{bmatrix}$$



(Cross product
as a linear
operation)

$$= \hat{\mathbf{a}}\mathbf{b} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix} \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

Cross product
matrix **Cross product**
matrix

A green arrow points from the label "Cross product matrix" to the first column of the matrix.

Angular velocity

https://en.wikipedia.org/wiki/Angular_velocity

Notation switch: let's call the vector of the rotation axis as “angular velocity” ω

We can approximate the change in position due to a small rotation as:

$$\dot{\mathbf{X}} = \omega \times \mathbf{X} = \hat{\omega} \mathbf{X}$$

**linear motion from
angular velocity**

Recall: exponentials are solutions of linear differential equations

$$\dot{x}(t) = ax(t)$$

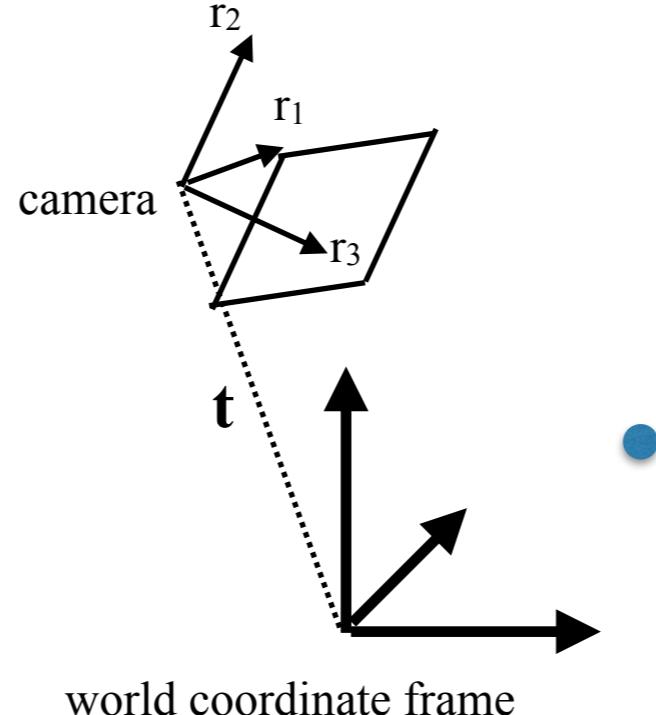
$$x(t) = e^{at}x(0)$$

Matrix exponentials are solutions of matrix linear differential equations

$$\dot{\mathbf{X}}(t) = \hat{\omega} \mathbf{X}(t)$$

$$\mathbf{X}(t) = e^{\hat{\omega}t} \mathbf{X}(0)$$

How does a fixed scene point X move with respect to the camera?



Suppose a camera is moving at translational velocity of \mathbf{t} and rotating with angular velocity $\boldsymbol{\omega}$ in the world frame

Then with respect to the camera frame, the point is moving with a linear velocity of:

$$\dot{\mathbf{X}} = -\mathbf{t} - \boldsymbol{\omega} \times \mathbf{X}$$

Point's velocity in camera frame	linear camera velocity	linear motion from camera's angular velocity	negative sign since the point appears to move in the opposite direction from the camera motion
--	------------------------------	--	--

$$\begin{bmatrix} \dot{X} \\ \dot{Y} \\ \dot{Z} \end{bmatrix} = - \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} - \begin{bmatrix} \omega_y Z - \omega_z Y \\ \omega_z X - \omega_x Z \\ \omega_x Y - \omega_y X \end{bmatrix}$$

expand out cross-product

$$\begin{bmatrix} \dot{X} \\ \dot{Y} \\ \dot{Z} \end{bmatrix} = - \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} - \begin{bmatrix} \omega_y Z - \omega_z Y \\ \omega_z X - \omega_x Z \\ \omega_x Y - \omega_y X \end{bmatrix}$$

If we assume $f=1$, $x(t) = \frac{X(t)}{Z(t)}$ and $y(t) = \frac{Y(t)}{Z(t)}$ (Projection equations)

**x image coordinate
over time** **y image coordinate
over time**

Then what is $\frac{dx}{dt}, \frac{dy}{dt}$?

change in image
coordinates over time

$$\dot{x} = \frac{\dot{X}Z - \dot{Z}X}{Z^2}, \quad \dot{y} = \frac{\dot{Y}Z - \dot{Z}Y}{Z^2}$$

**image velocity in x-
direction, "u"** **image velocity in y-
direction, "u"** **quotient rule**

Plug in equations for $\dot{X}, \dot{Y}, \dot{Z}, \frac{X}{Z} = x, \frac{Y}{Z} = y$

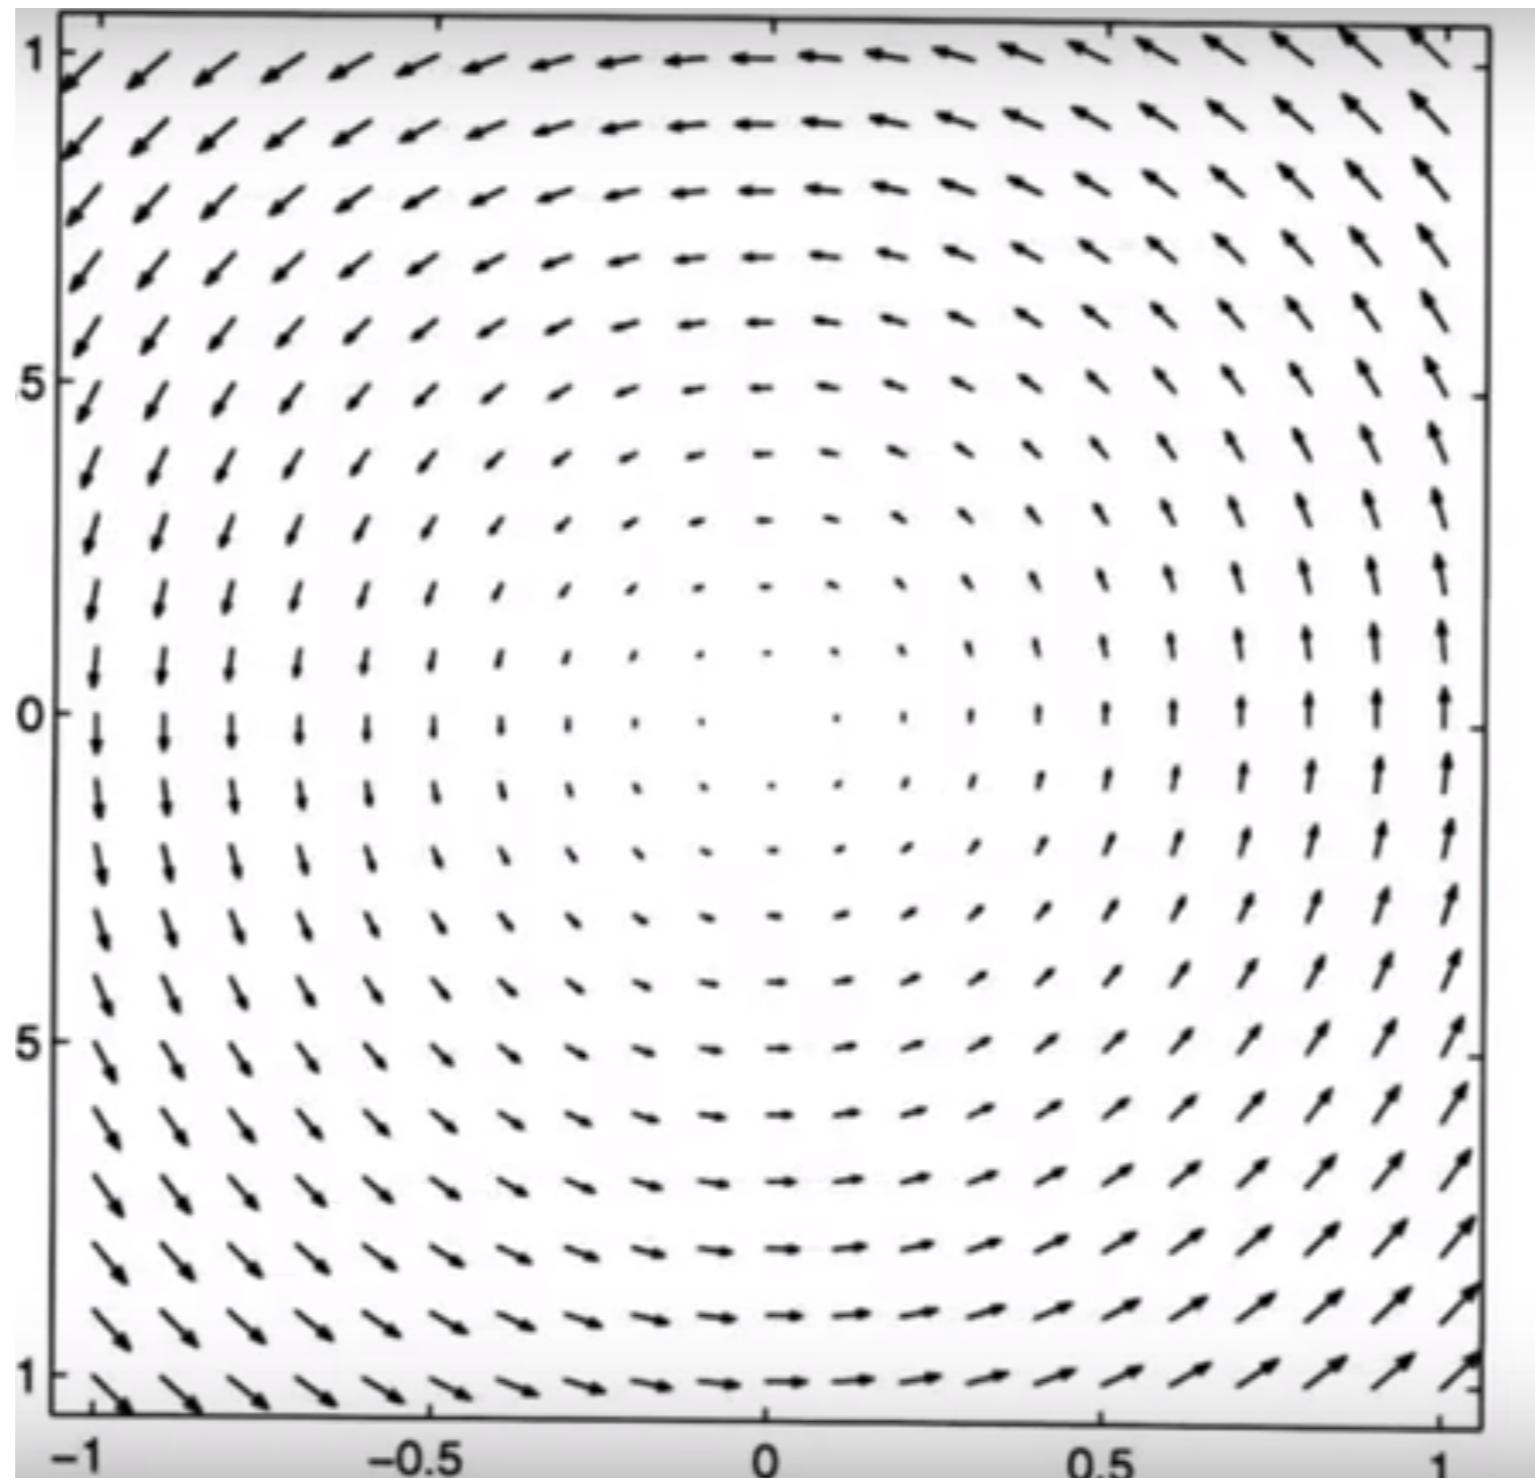
Final result:

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ 1+y^2 & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

**velocity of image
coordinates** **depth** **camera
translation** **(no depth)** **camera
rotation**

This is the important part!

Example: What does the flow look like for rotating about the z-axis?



Flow does not depend on object depth

Egomotion optical flow

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ 1+y^2 & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

velocity of image coordinates depth camera translation (no depth) camera rotation

velocity of image coordinates	depth	translation part	rotation part
	$\frac{1}{Z(x,y)}$		
$u(x,y) =$		$(-t_x + xt_z) + xy\omega_x - (1+x^2)\omega_y + y\omega_z$	
$v(x,y) =$		$(-t_y + yt_z) + (1+y^2)\omega_x - xy\omega_y - x\omega_z$	

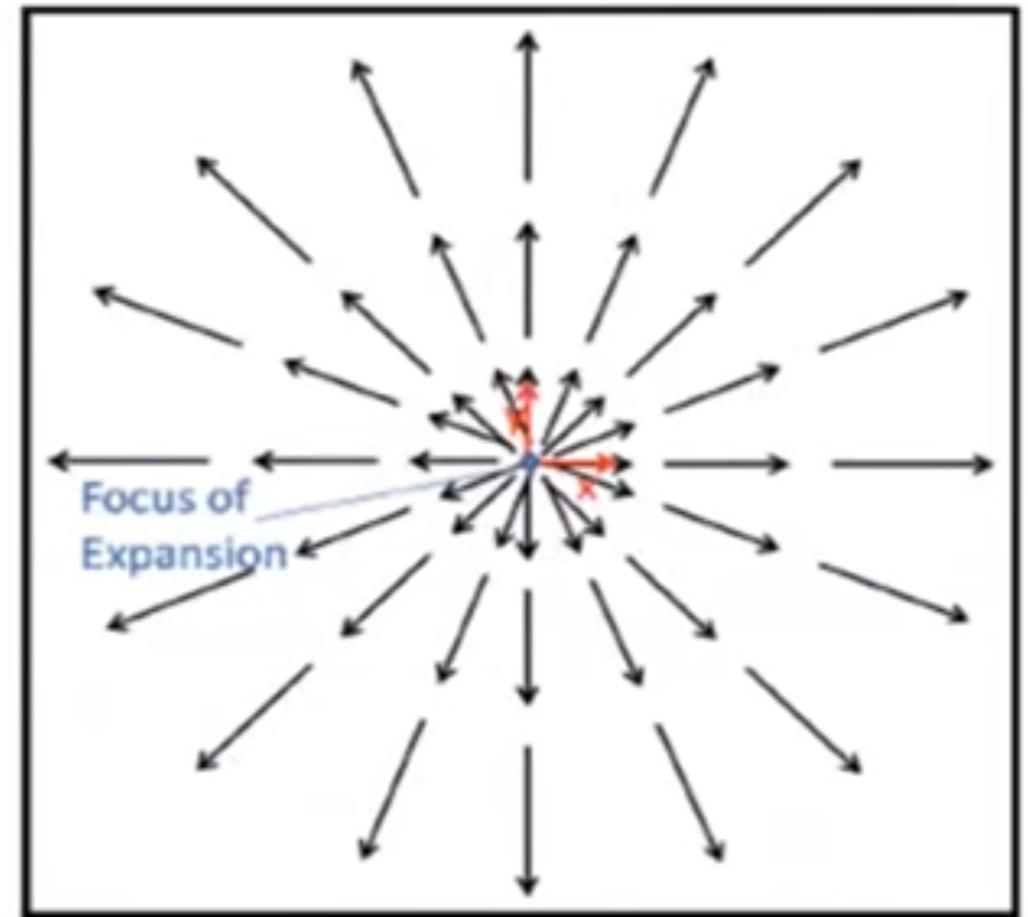
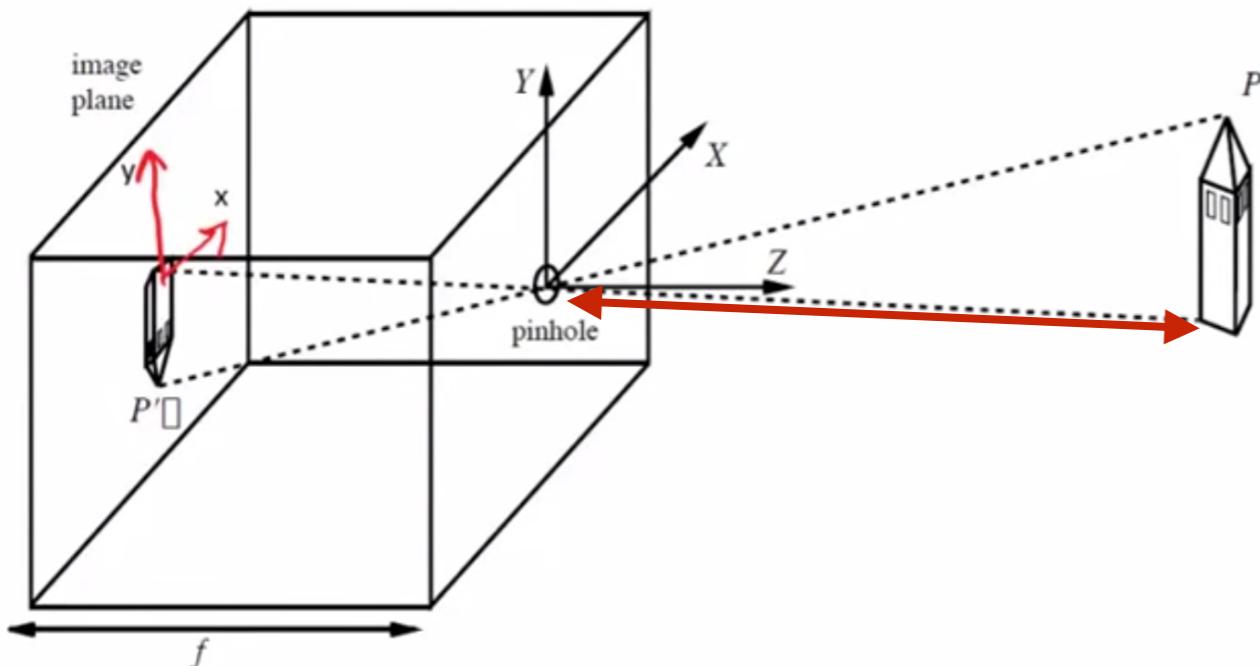
What would happen if there is no camera rotation (pure translation)?

What would happen if we translate along the optical axis ($t_x = 0, t_y = 0$)? i.e. move camera forward

$$u(x,y) = \frac{xt_z}{Z(x,y)} \quad v(x,y) = \frac{yt_z}{Z(x,y)} \quad \begin{bmatrix} u(x,y) \\ v(x,y) \end{bmatrix} = \frac{t_z}{Z(x,y)} \begin{bmatrix} x \\ y \end{bmatrix}$$

What does this flow field look like?

Optical flow for translation along camera's z-axis



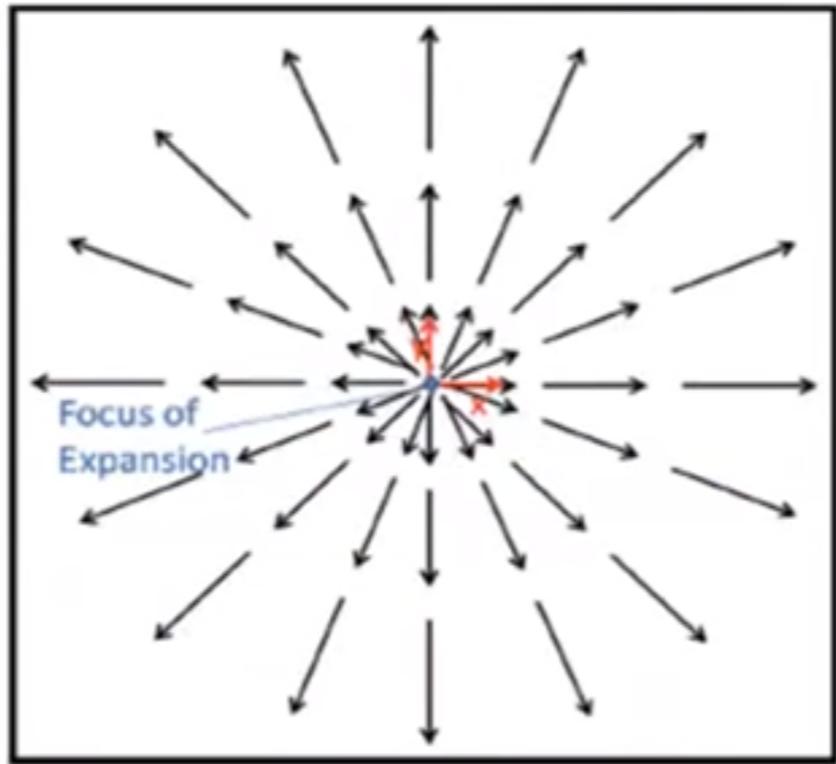
$$\begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix} = \frac{t_z}{Z(x, y)} \begin{bmatrix} x \\ y \end{bmatrix}$$

Optical flow vector is a scalar multiple of the position vector
(points in the direction of x,y) - expands out!

(the scalar depends upon the speed of camera translation *and* the depth of the scene point)

What is the flow in the center of the image?

Moving toward a wall (constant distance to all points Z)



$$\begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix} = \frac{t_z}{Z} \begin{bmatrix} x \\ y \end{bmatrix}$$

If speed (t_z) and distance (Z) to wall doubles, what happens to flow?

Nothing => fundamental scale ambiguity

If I am walking to a wall that is close by, or I am running towards a wall that is far away, I will get the same flow field!

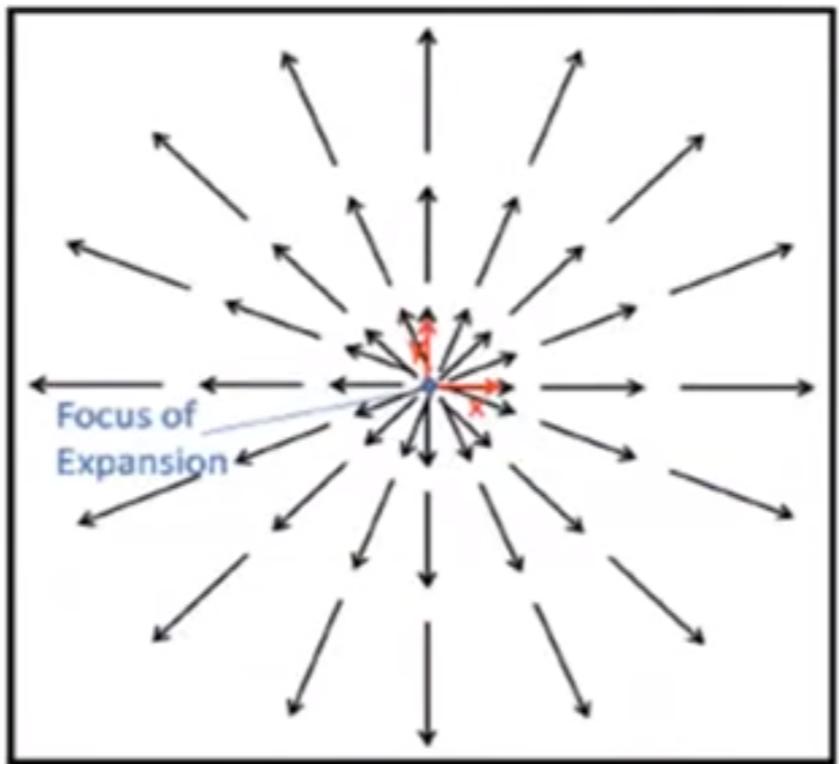
From a video, we cannot determine if the person holding the camera is walking or running if we do not know the depth of the scene!

We cannot determine the depth of the scene unless we know the speed of the camera!

Of course, with priors on depth and speed, we could guess both, but we won't know for sure

If we have a robot with a sensor to measure its own speed, we can estimate depth

Moving toward a wall (constant distance to all points Z)



$$\begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix} = \frac{t_z}{Z} \begin{bmatrix} x \\ y \end{bmatrix}$$

If we are traveling at 2 ft/sec, and the wall is 4 ft away, what's the time-to-contact?

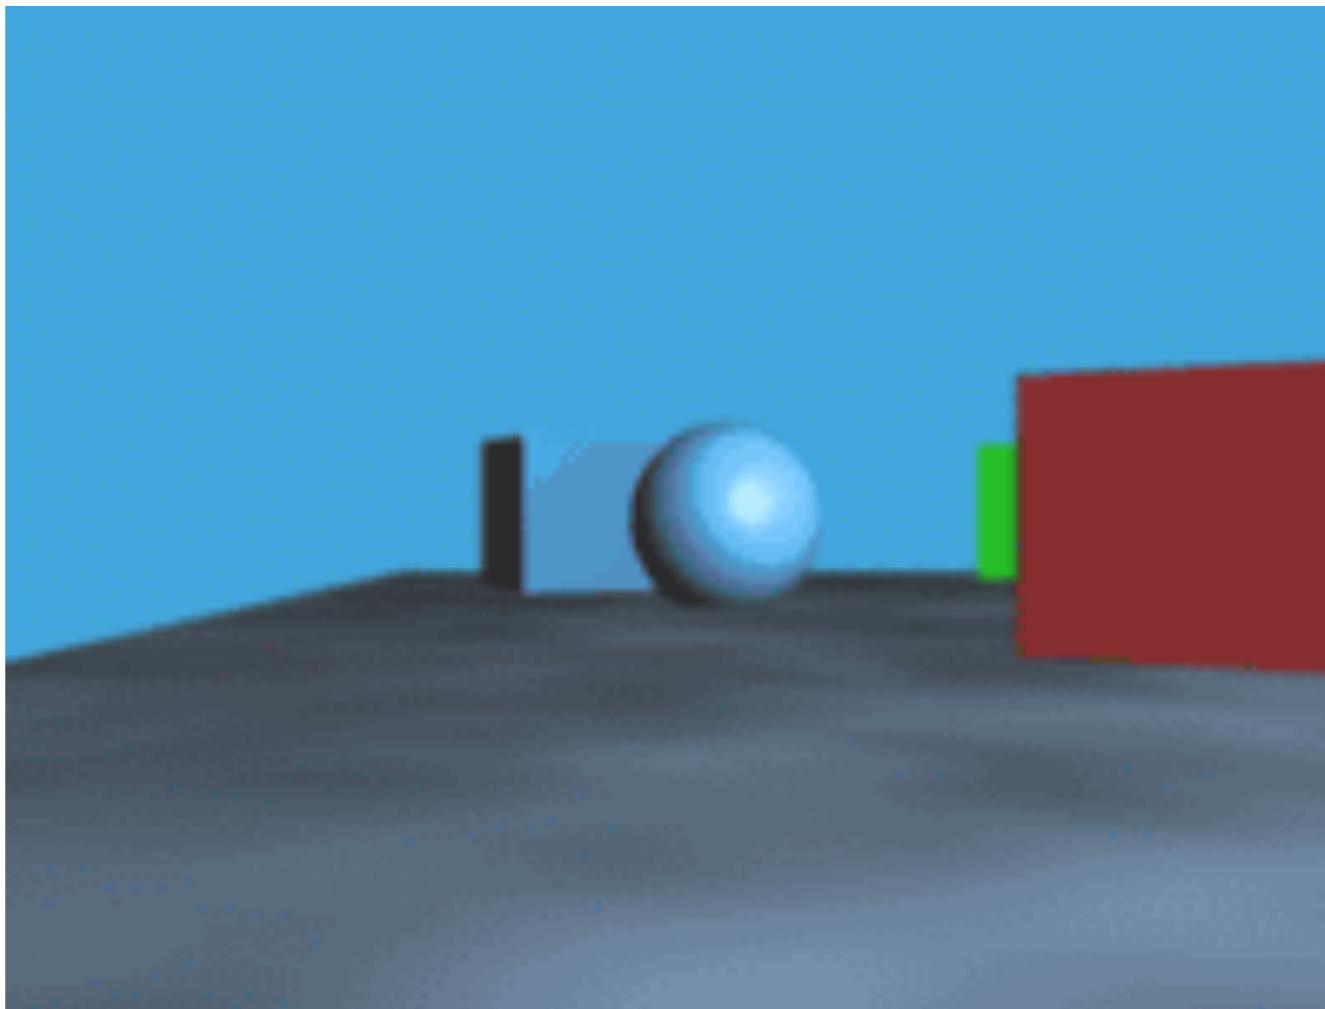
$$\text{Time to contact} = Z / t_z$$

Implies that time-to-contact (assuming constant depth) can be computed from 2D flow!
(Even though we don't know the depth or speed)

Example: walking to a close wall or running to a far wall

Thus we can avoid hitting things without estimating the depth of the scene,
just by estimating the flow!
(Primitive animals can do this)

Optical flow for translation along camera's x-axis ($t_y, t_z = 0$)



$$u(x, y) = \frac{-t_x + xt_z}{Z(x, y)}, v(x, y) = \frac{-t_y + yt_z}{Z(x, y)}.$$

$$v(x, y) = 0$$

$$u(x, y) = -t_x / Z(x, y)$$

Flow depends on depth and translational velocity!

Optical flow for general translation

$$u(x, y) = \frac{-t_x + xt_z}{Z(x, y)}, v(x, y) = \frac{-t_y + yt_z}{Z(x, y)}.$$

When is the flow $(u, v) = (0, 0)$?

$$x = \frac{t_x}{t_z} \qquad y = \frac{t_y}{t_z}$$

If the camera is translating at (t_x, t_y, t_z) , then there will be some point where the flow is 0:

$$\left(\frac{t_x}{t_z}, \frac{t_y}{t_z} \right)$$

If I find the point (x, y) where the flow is 0, then I know the ratio of the flows!

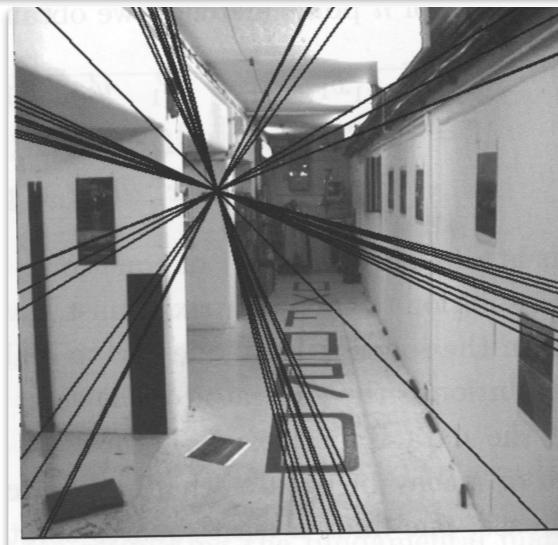
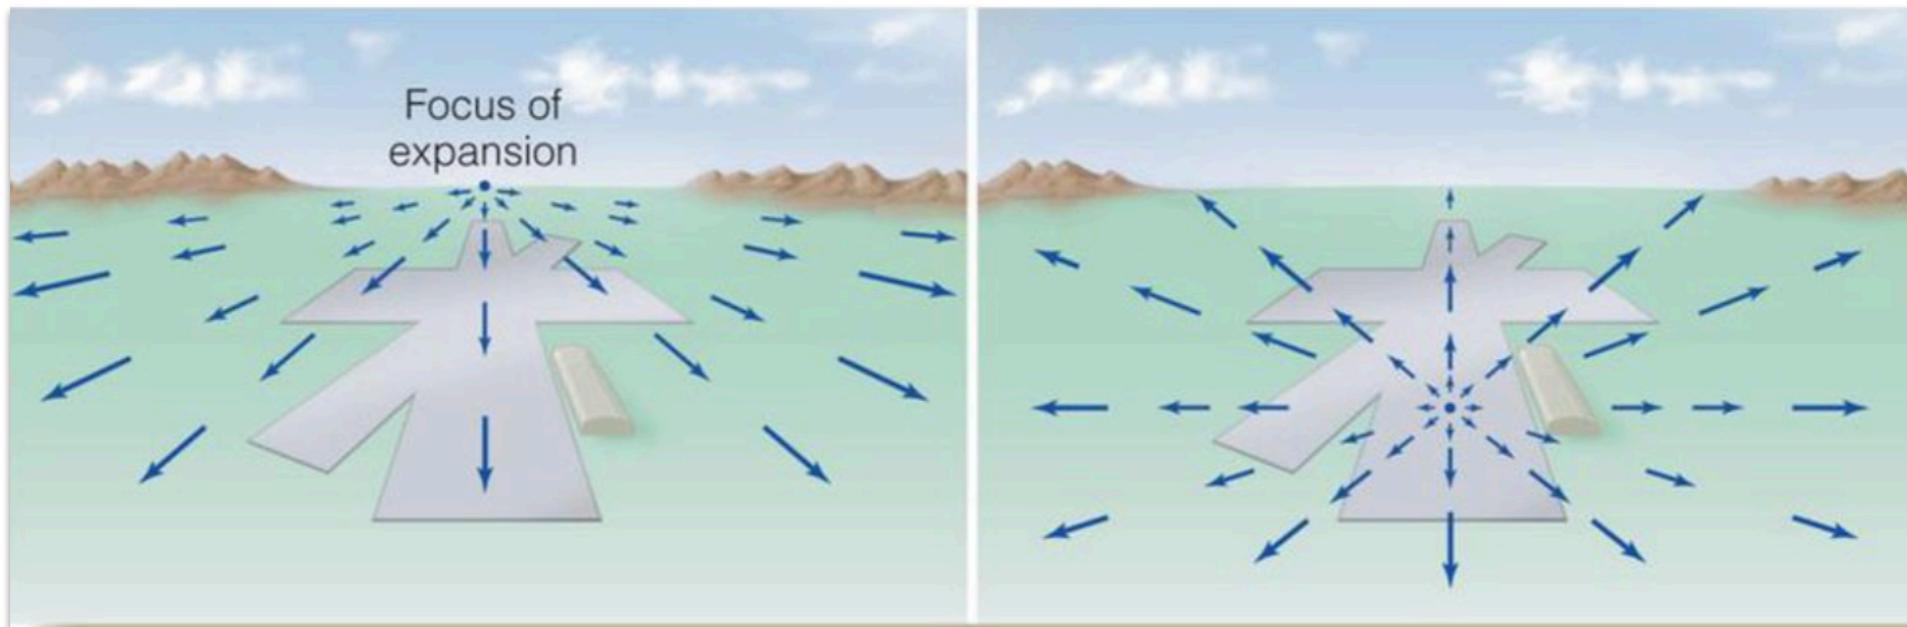
$$\text{Let } t_z = \lambda \qquad \text{Motion} = (\lambda x, \lambda y, \lambda) = \lambda(x, y, 1)$$

Thus, I know the direction that I am moving!

Focus of expansion

reveals the projection of camera translation velocity vector into image

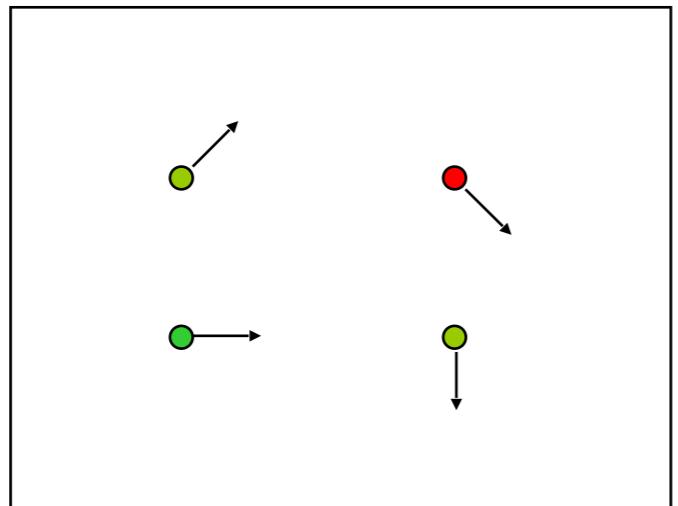
$$\begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} \Rightarrow \begin{bmatrix} \frac{t_x}{t_z} & \frac{t_y}{t_z} \end{bmatrix}$$



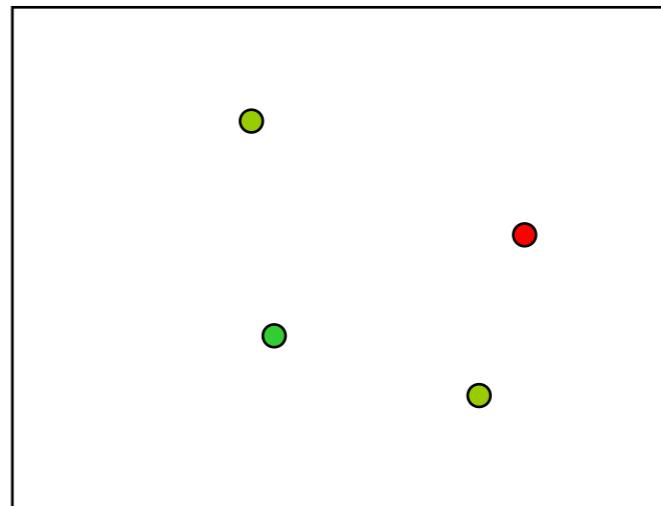
Summary

- If we can compute the 2D optical flow, then we can:
 - compute the time to contact
 - estimate the rotation of the camera
 - estimate the translation of the camera (if scene depth is known)
 - estimate the scene depth if camera translation is known

Problem Definition: Optical Flow



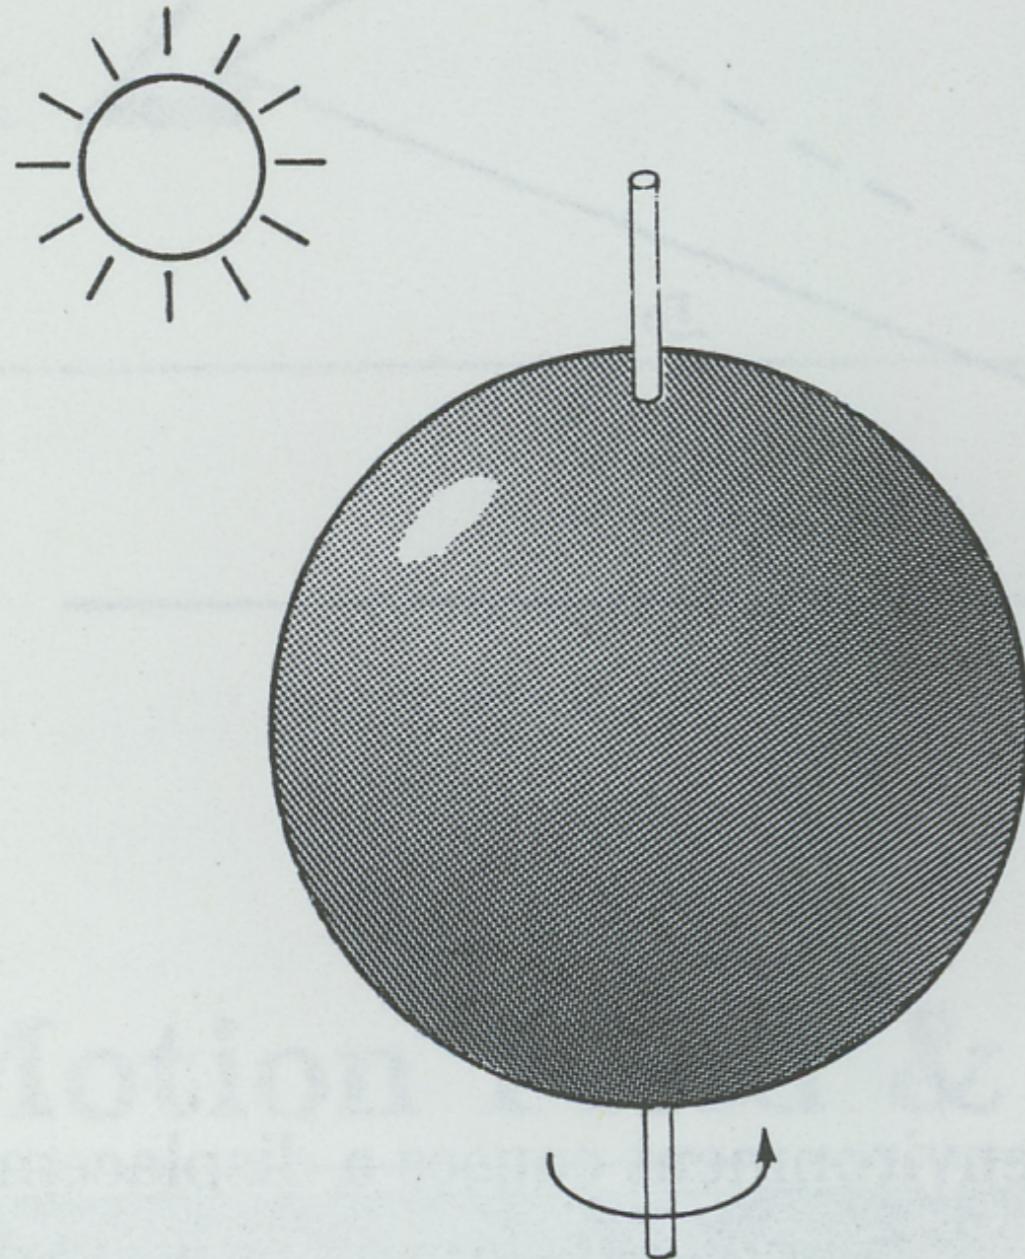
$H(x, y)$



$I(x, y)$

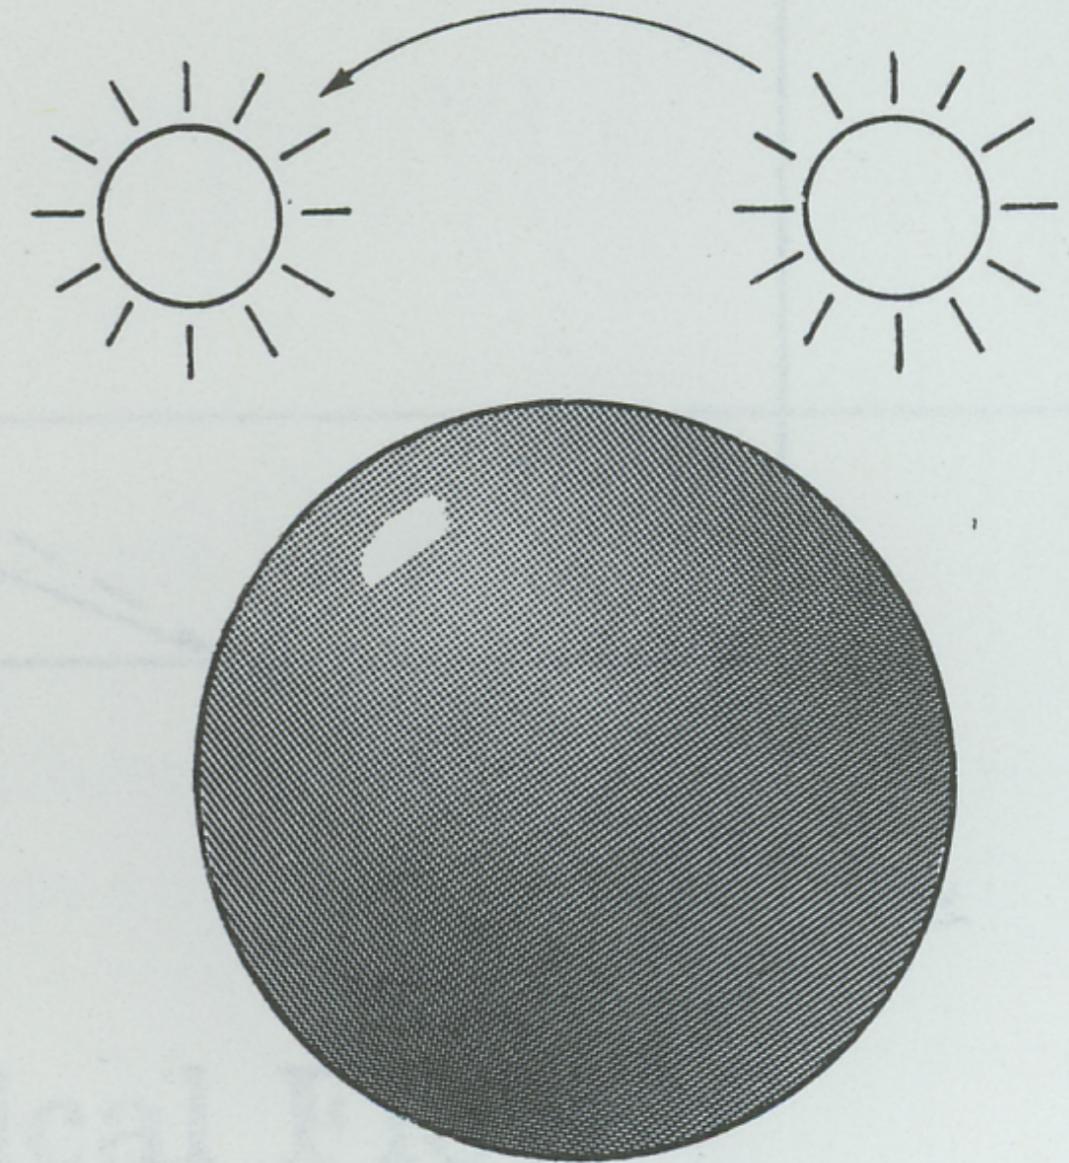
- How to estimate pixel motion from image H to image I ?
 - Find pixel correspondences
 - Given a pixel in H , look for nearby pixels of the same color in I
- Key assumption
 - **color constancy**: a point in H looks “the same” in image I
 - For grayscale images, this is **brightness constancy**

Caution: 2D measured optical flow \neq 3D scene flow



Motion field exists but no optical flow

(a)



No motion field but shading changes

(b)

General motion fields

Why do we want to compute motion fields beyond ego=motion?



Simple evidence that “single image” analysis won’t always suffice