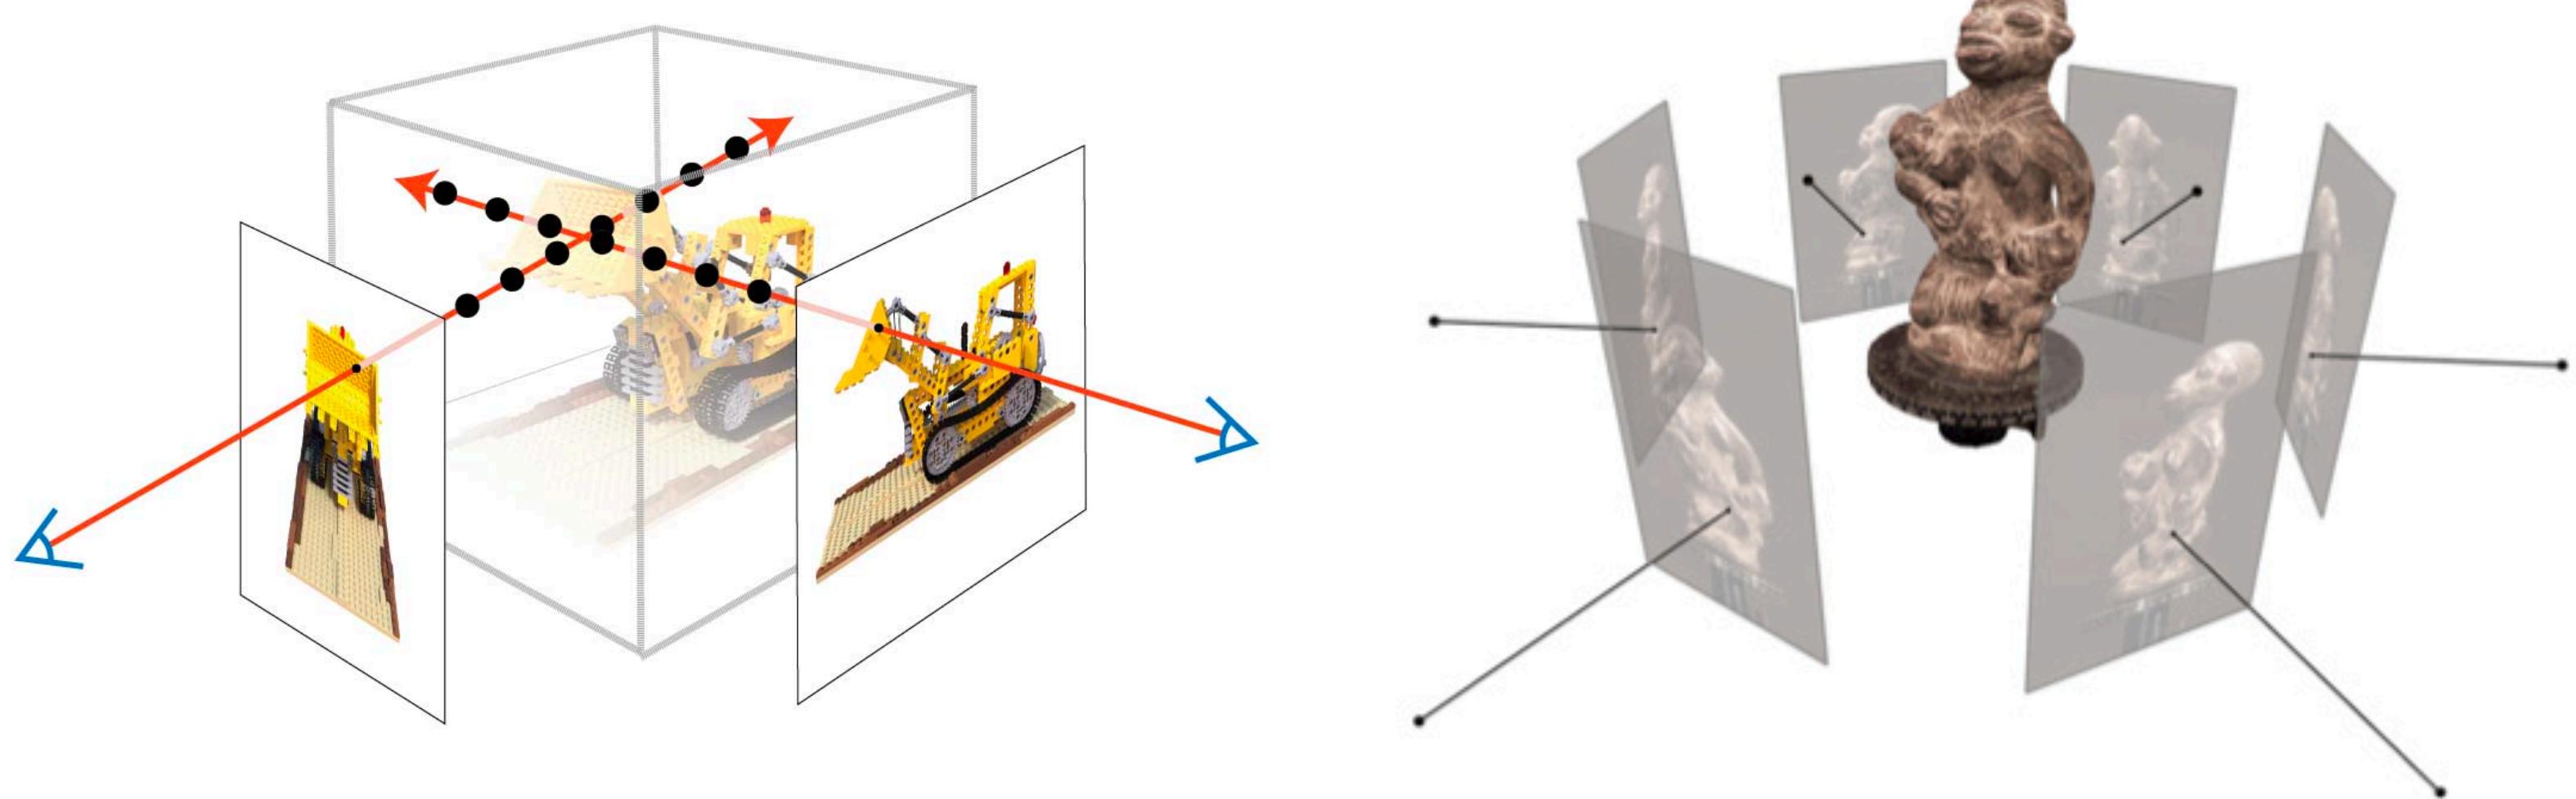


3D from Multi-view



Instructor: Shubham Tulsiani

A reminder to ask questions!



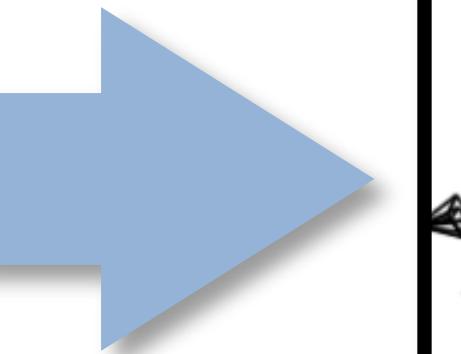
3D from Multi-view



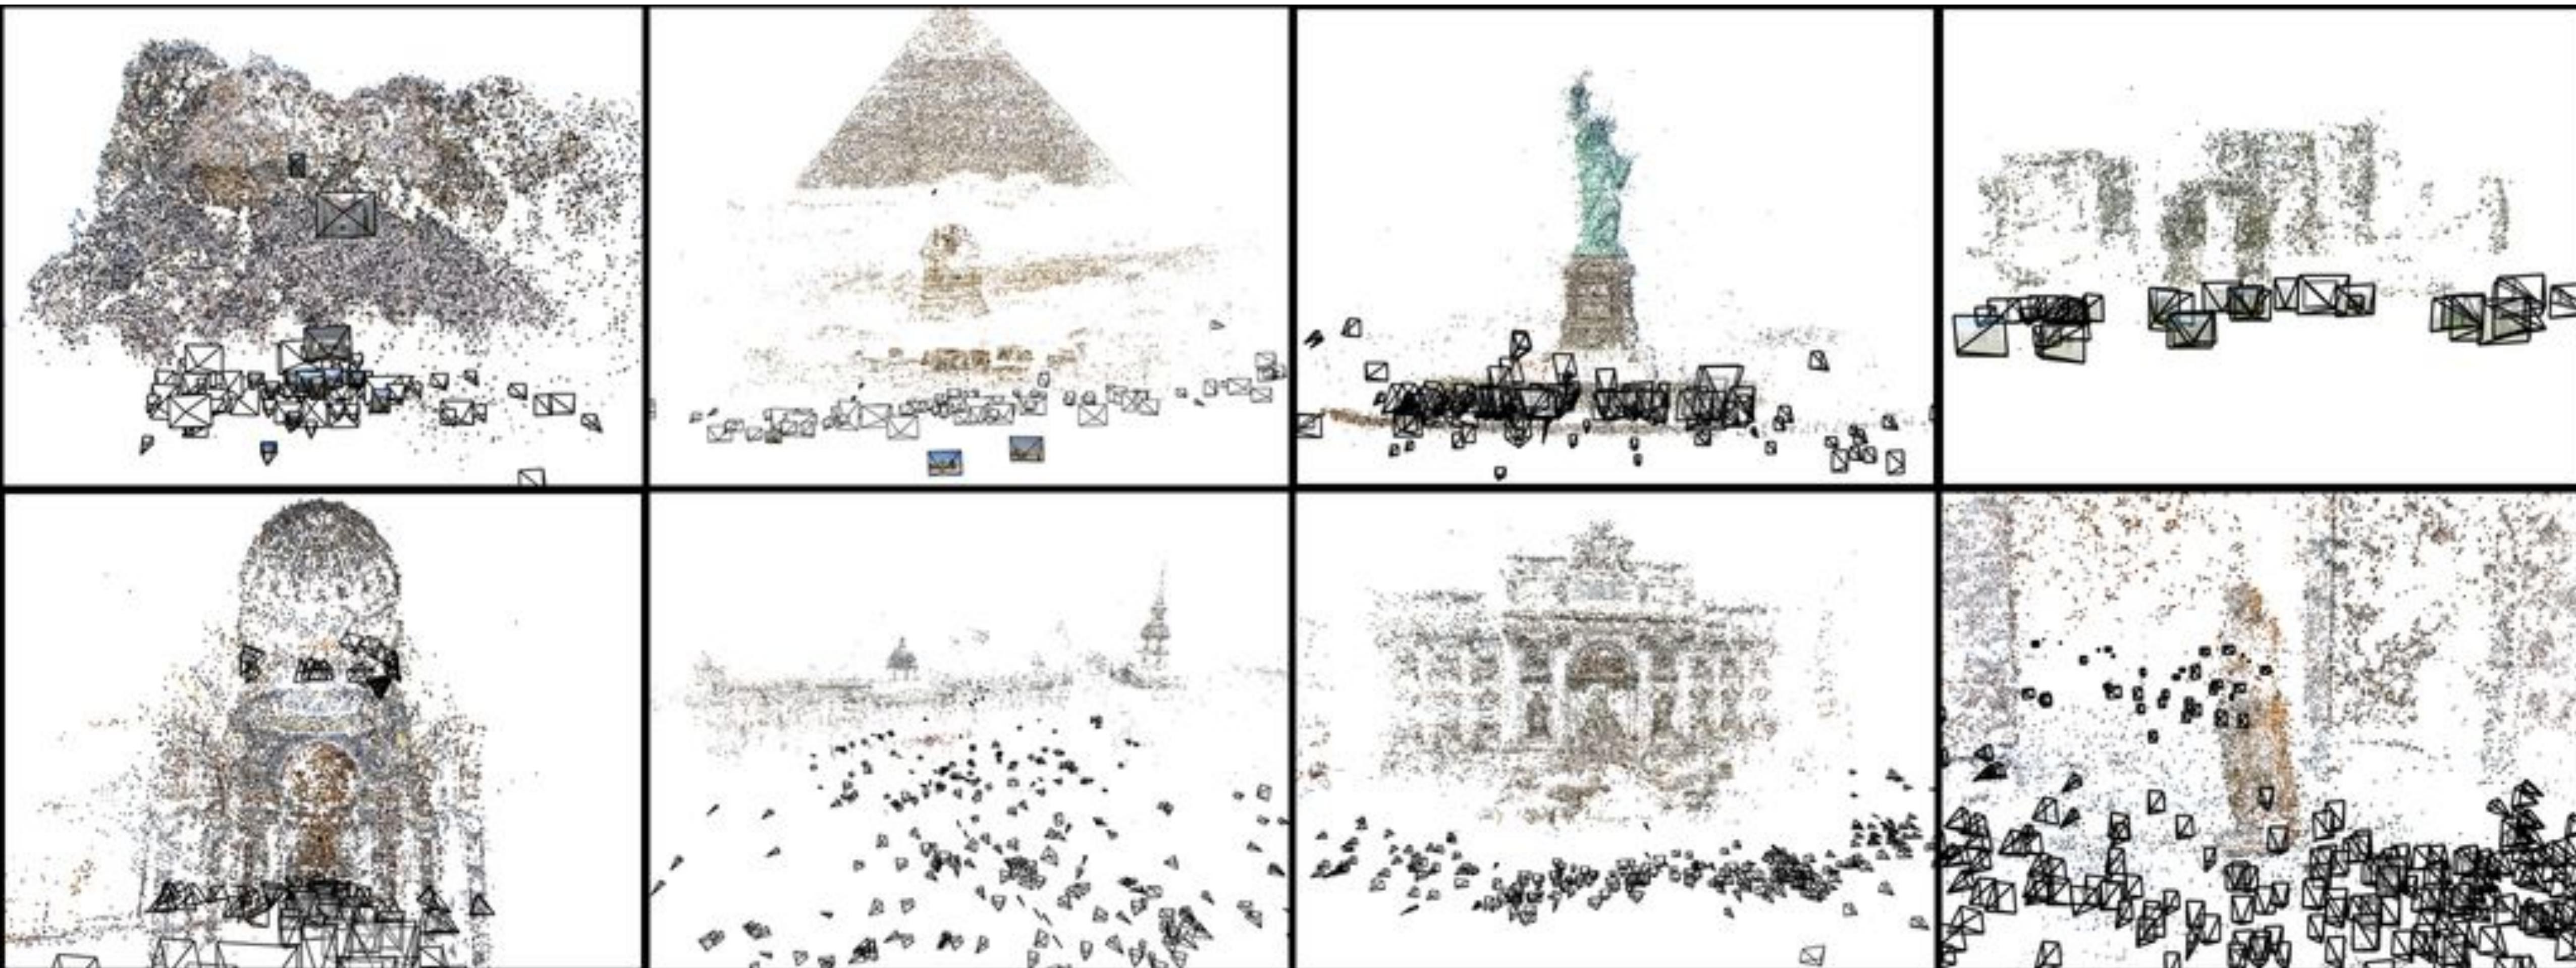
Images + Cameras

3D Representation

3D from Multi-view



3D from Multi-view



3D from Multi-view

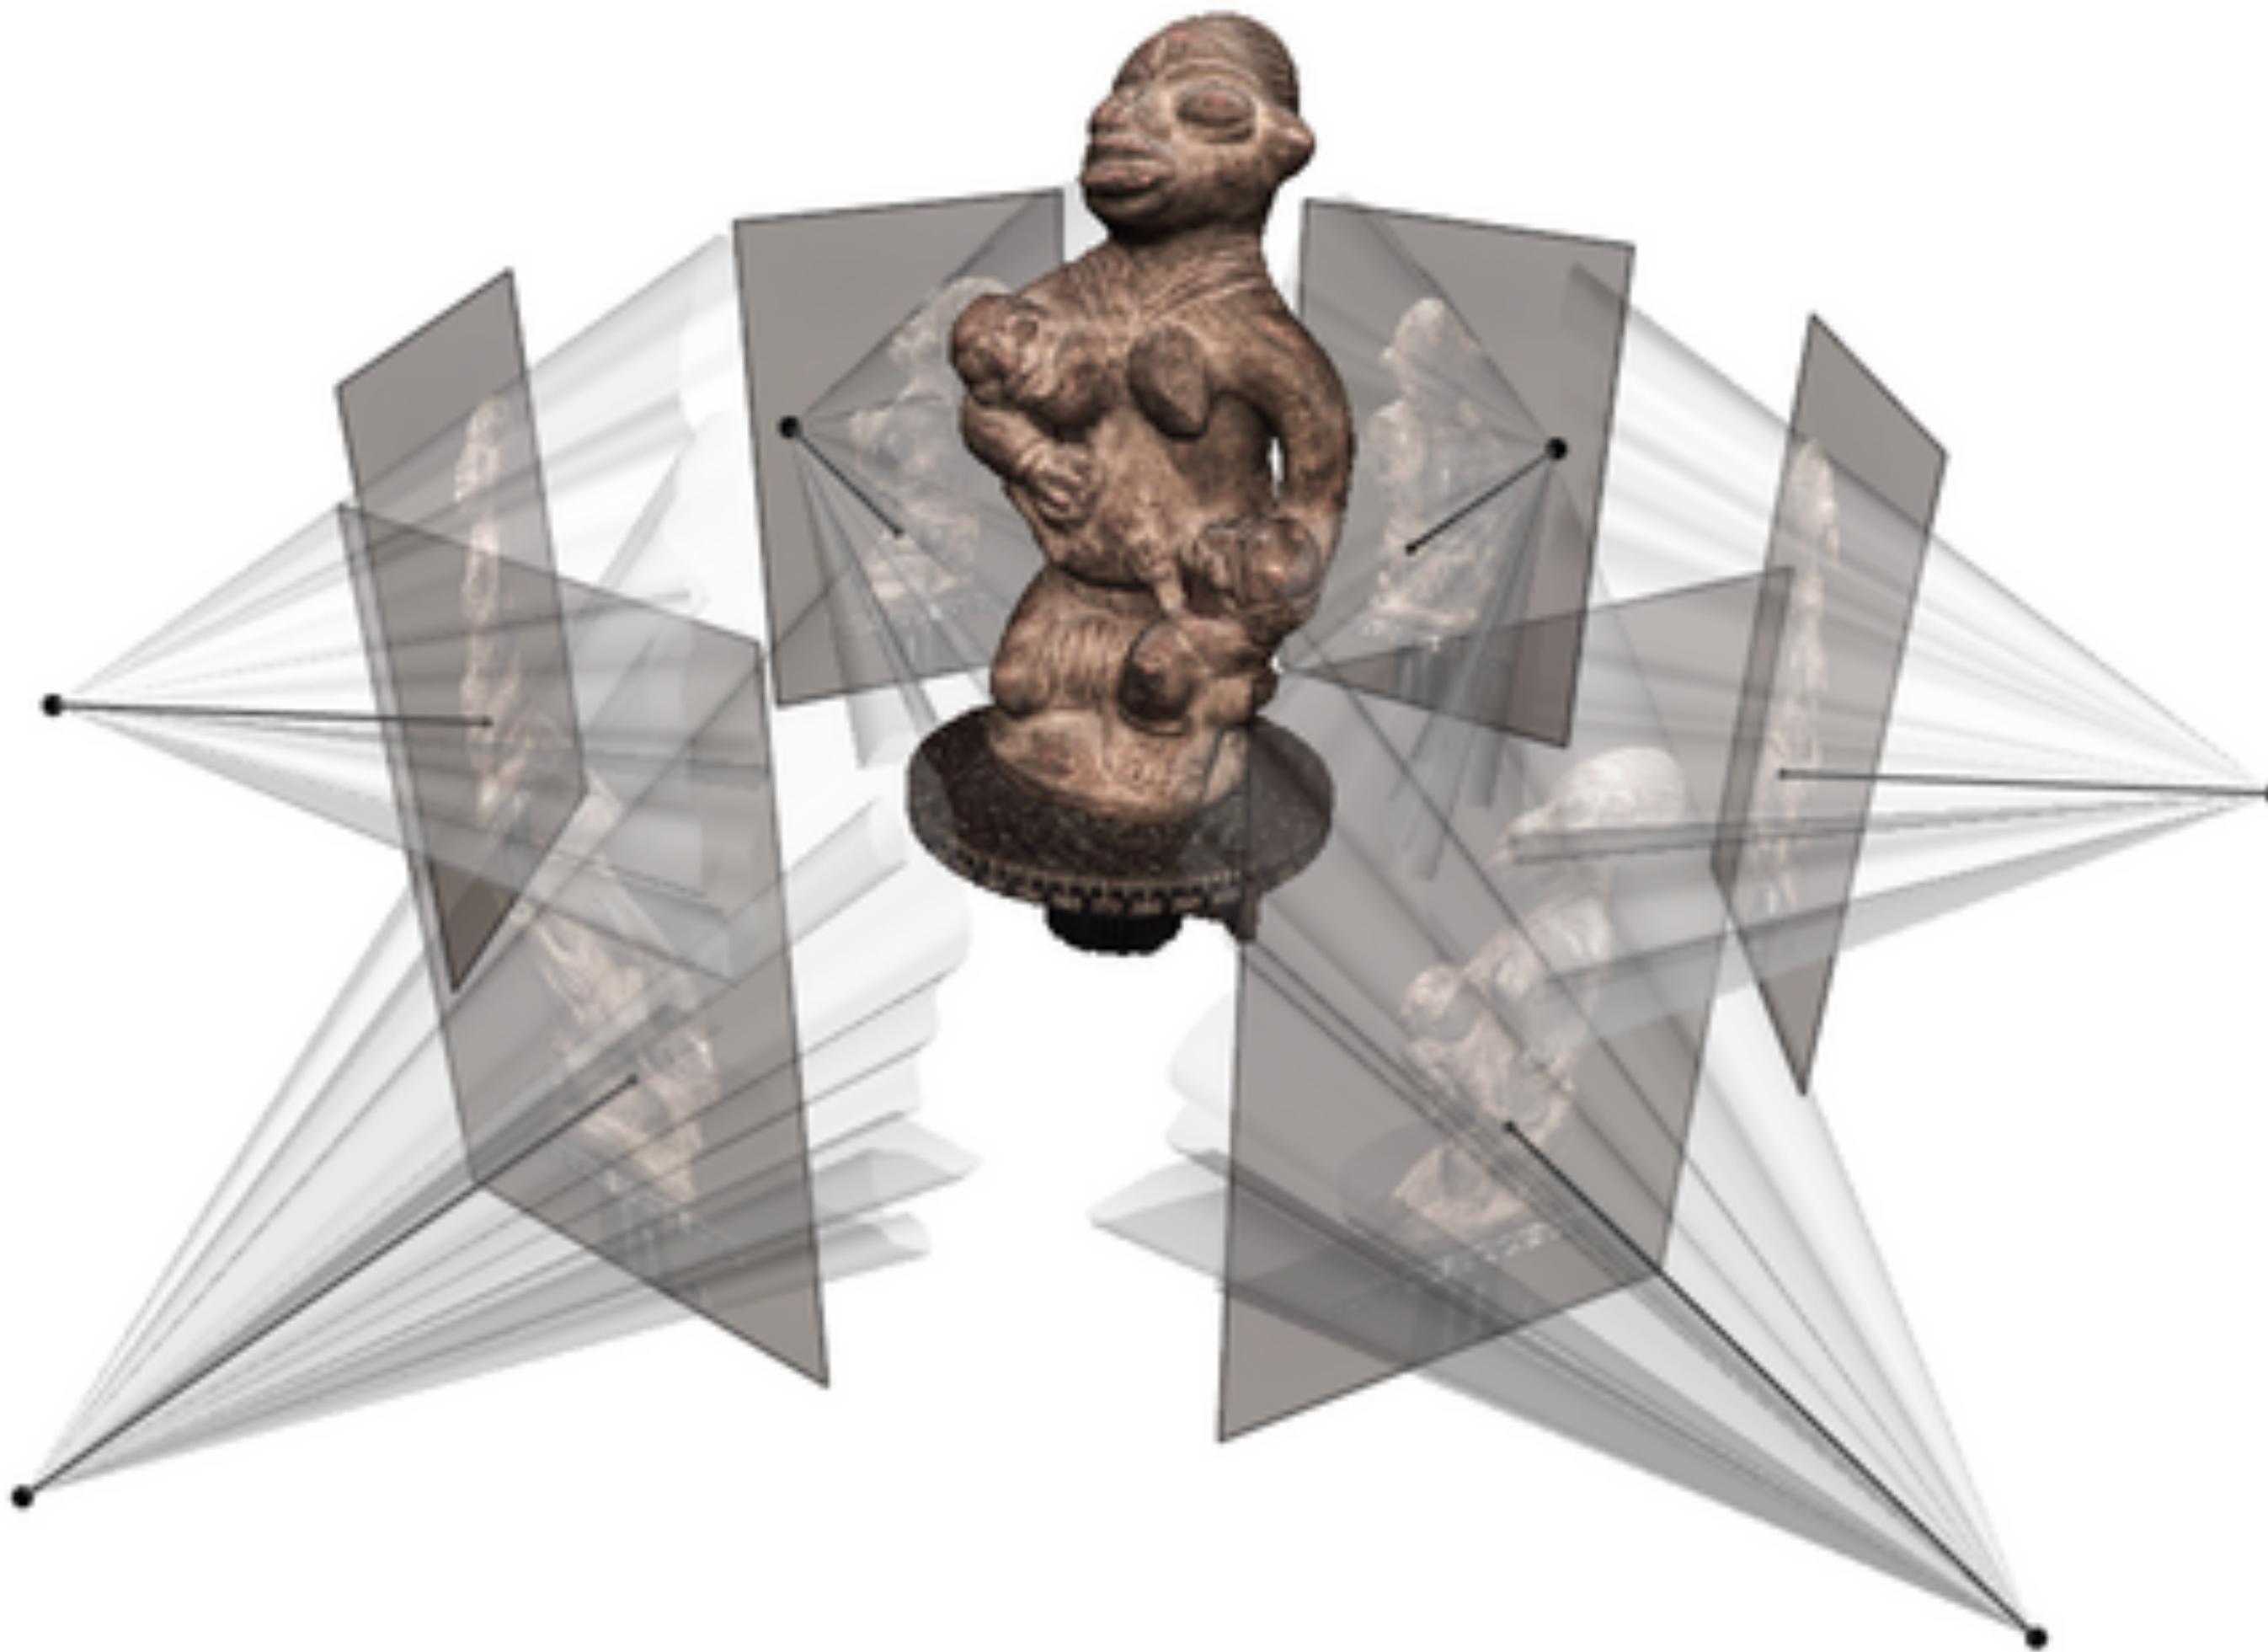
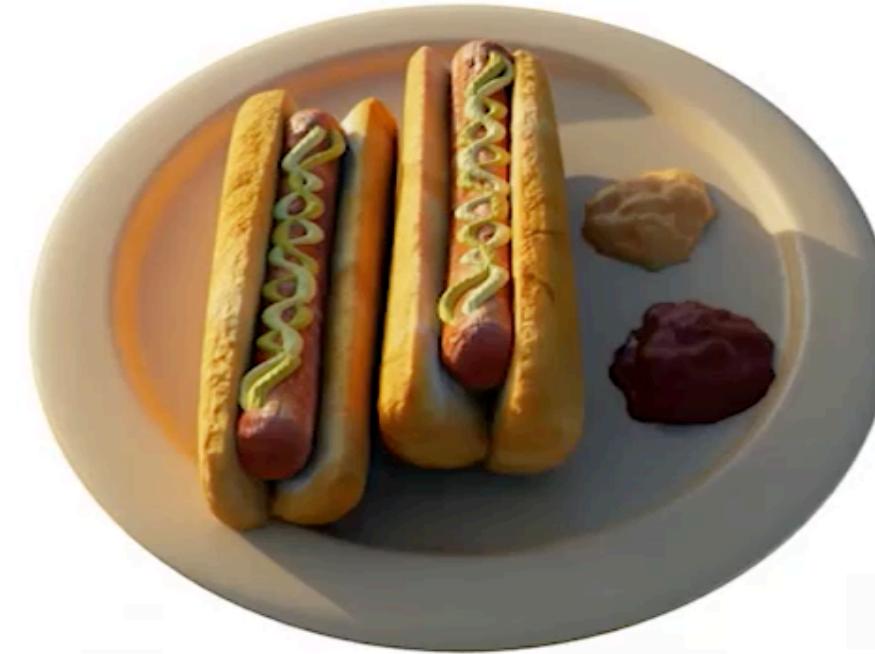
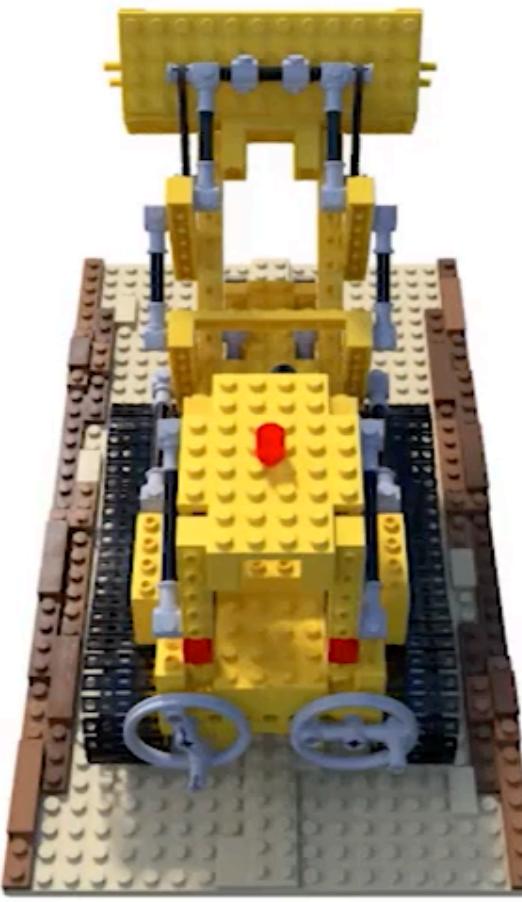


Image credits: Carlos Hernandez

3D from Multi-view



3D from Multi-view



Image credits: Paul Debevec

3D from Multi-view

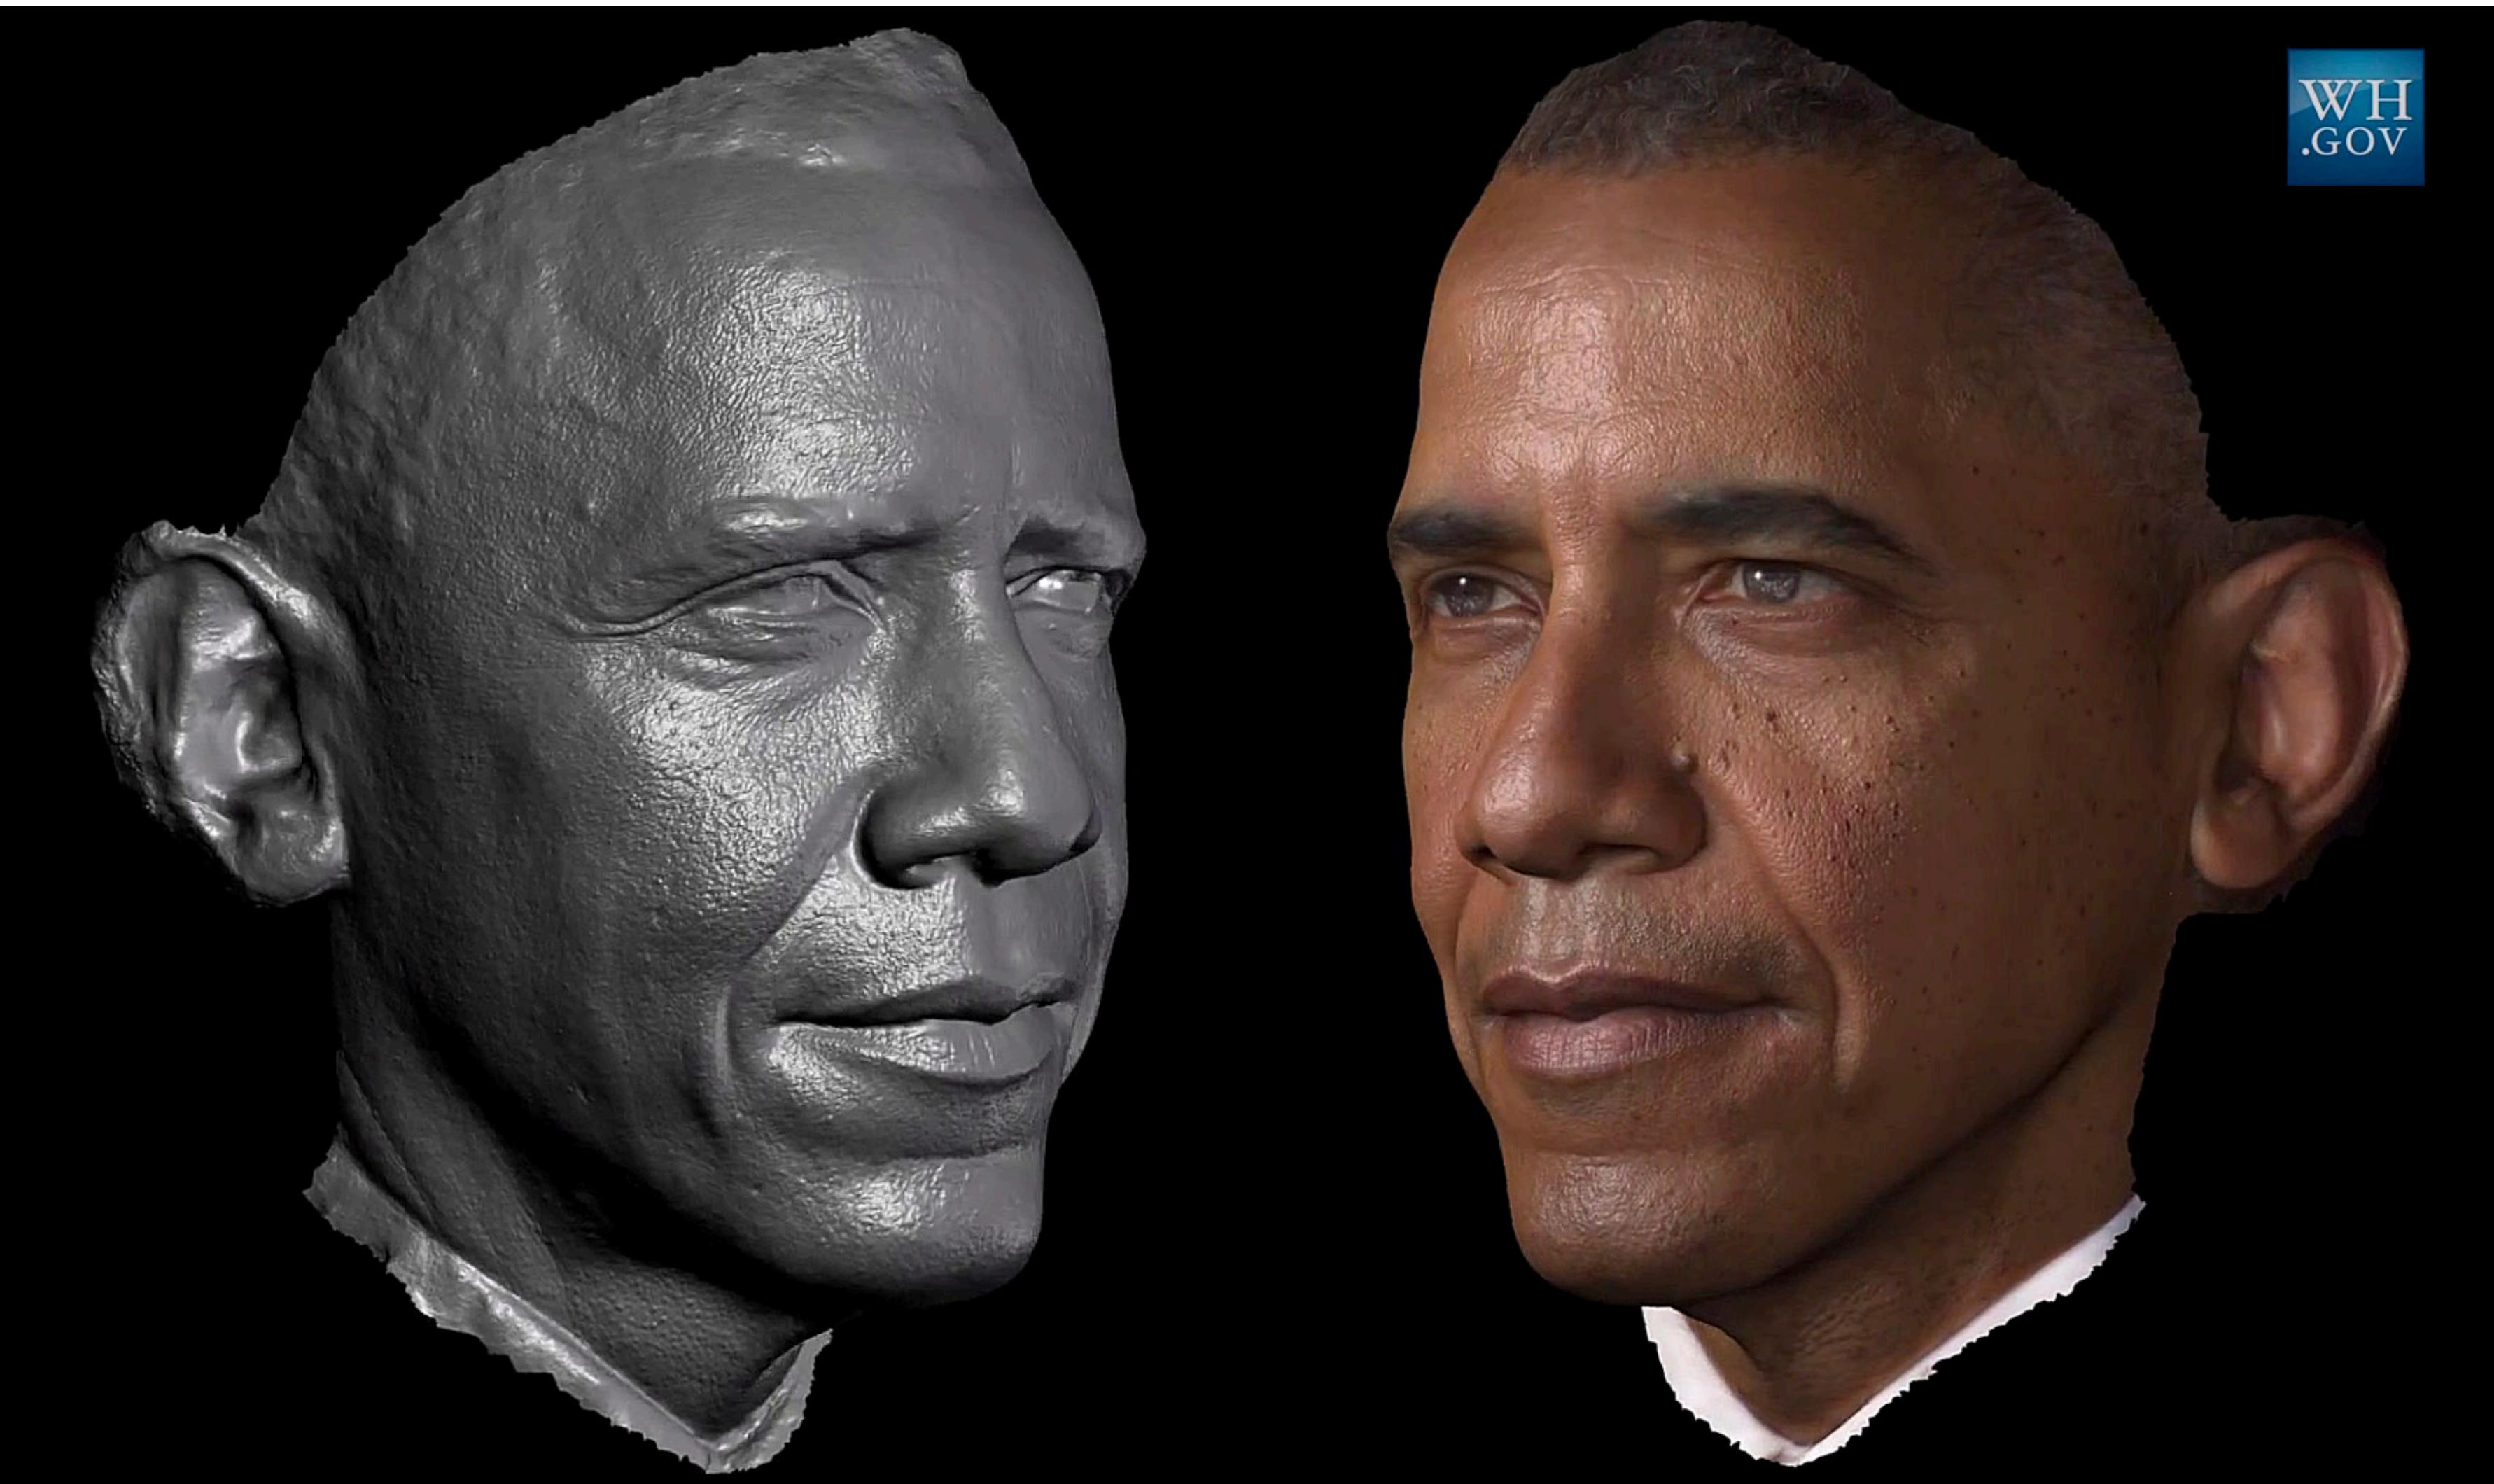


Image credits: Paul Debevec

3D from Multi-view



Images + Cameras

3D Representation

3D from Multi-view



Images + Cameras



3D Representation

What representation?

How can we infer it from
posed images?

3D from Multi-view



Images + Cameras



3D Representation

What representation?

**How can we infer it from
(posed) images?**

Key idea: Find a 3D representation (\mathbf{x}) whose renderings are consistent with all observed views

$$\min_{\mathbf{x}} \sum_i \| \text{Render}(\mathbf{x}, \pi_i) - I_i \|^2$$

3D from Multi-view



Images + Cameras



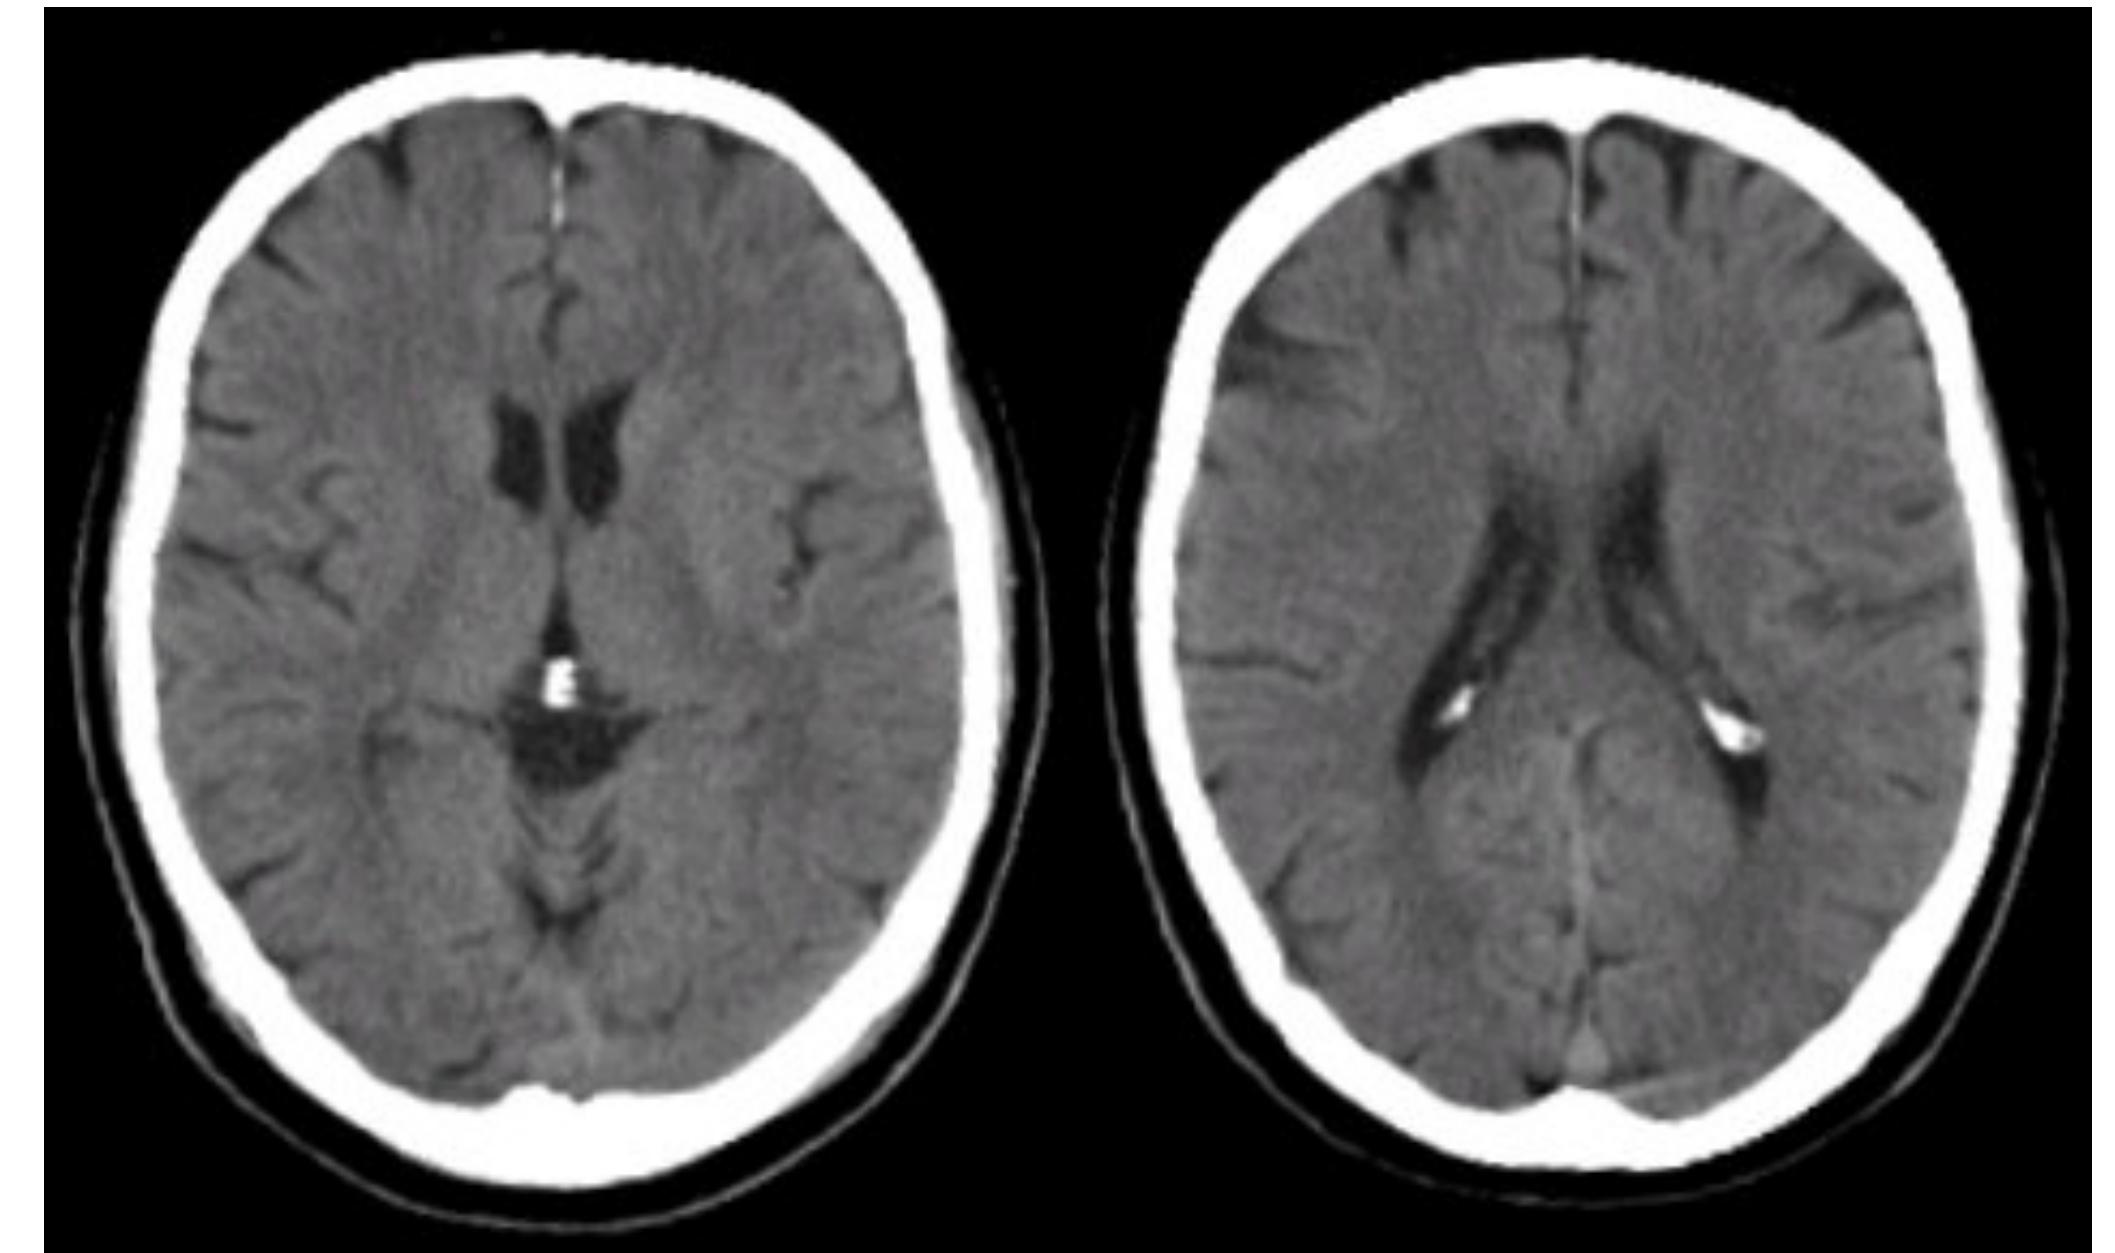
3D Representation

What representation?
(and how to render it)

How can we infer it from
posed images?

Today's goal: Understand two popular choices and corresponding rendering mechanisms

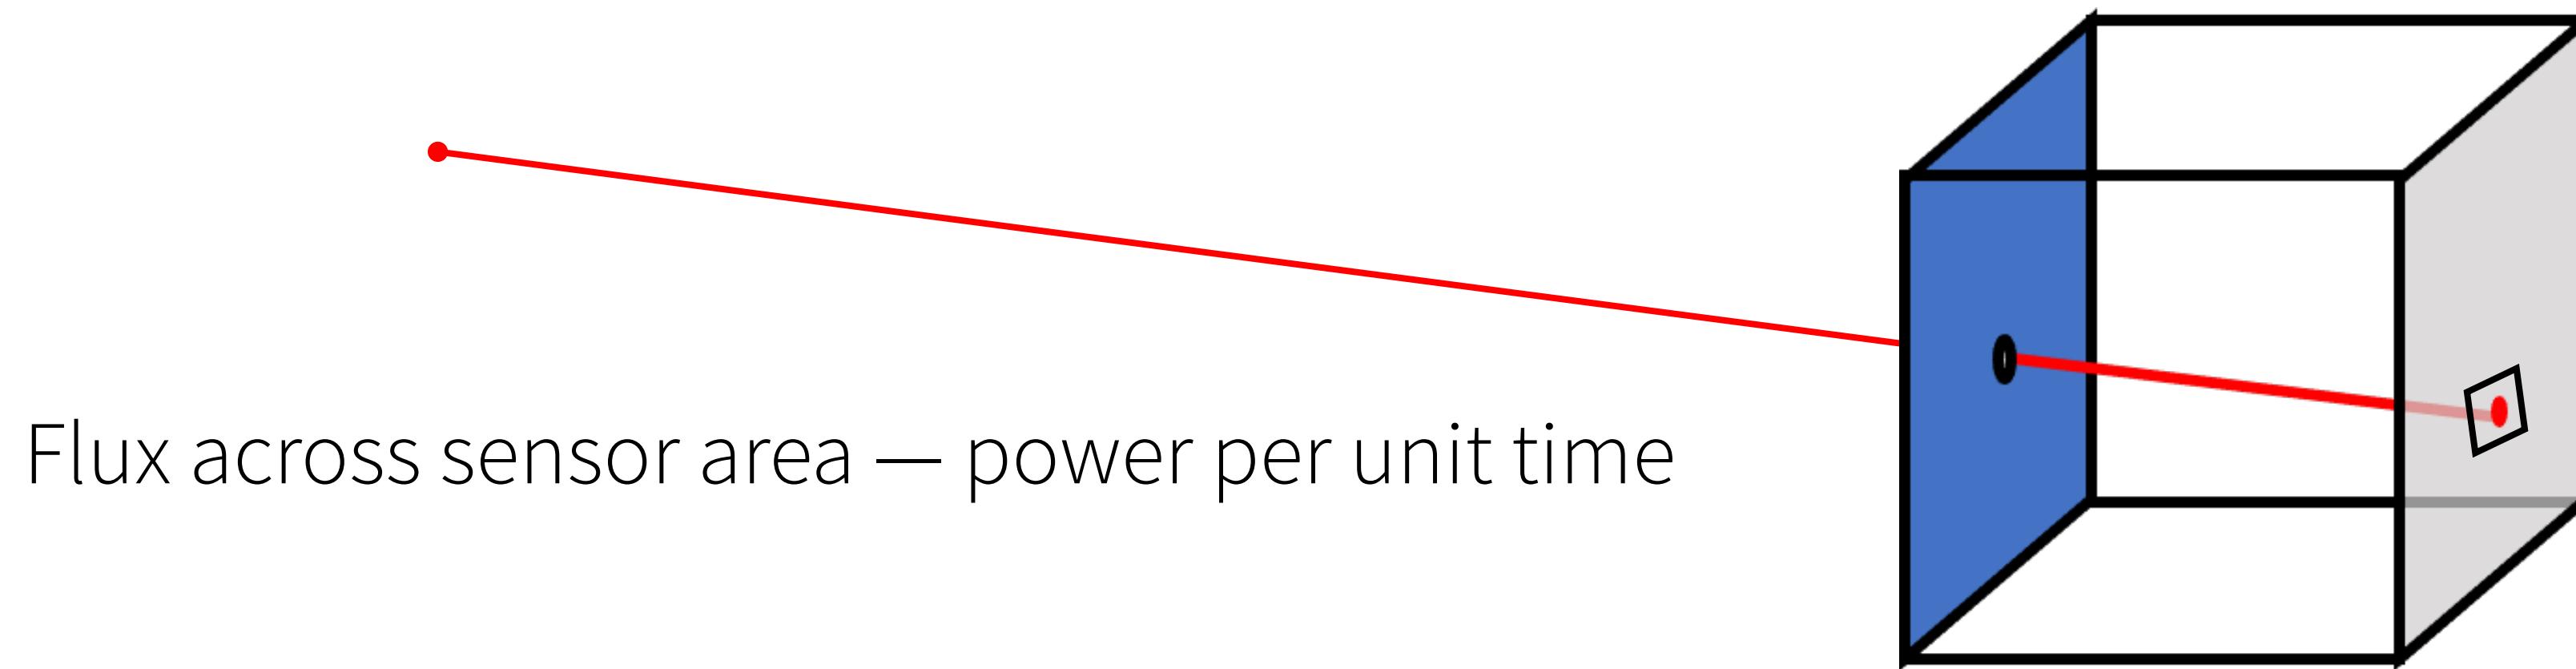
(Neural/Grid-based) Density Fields



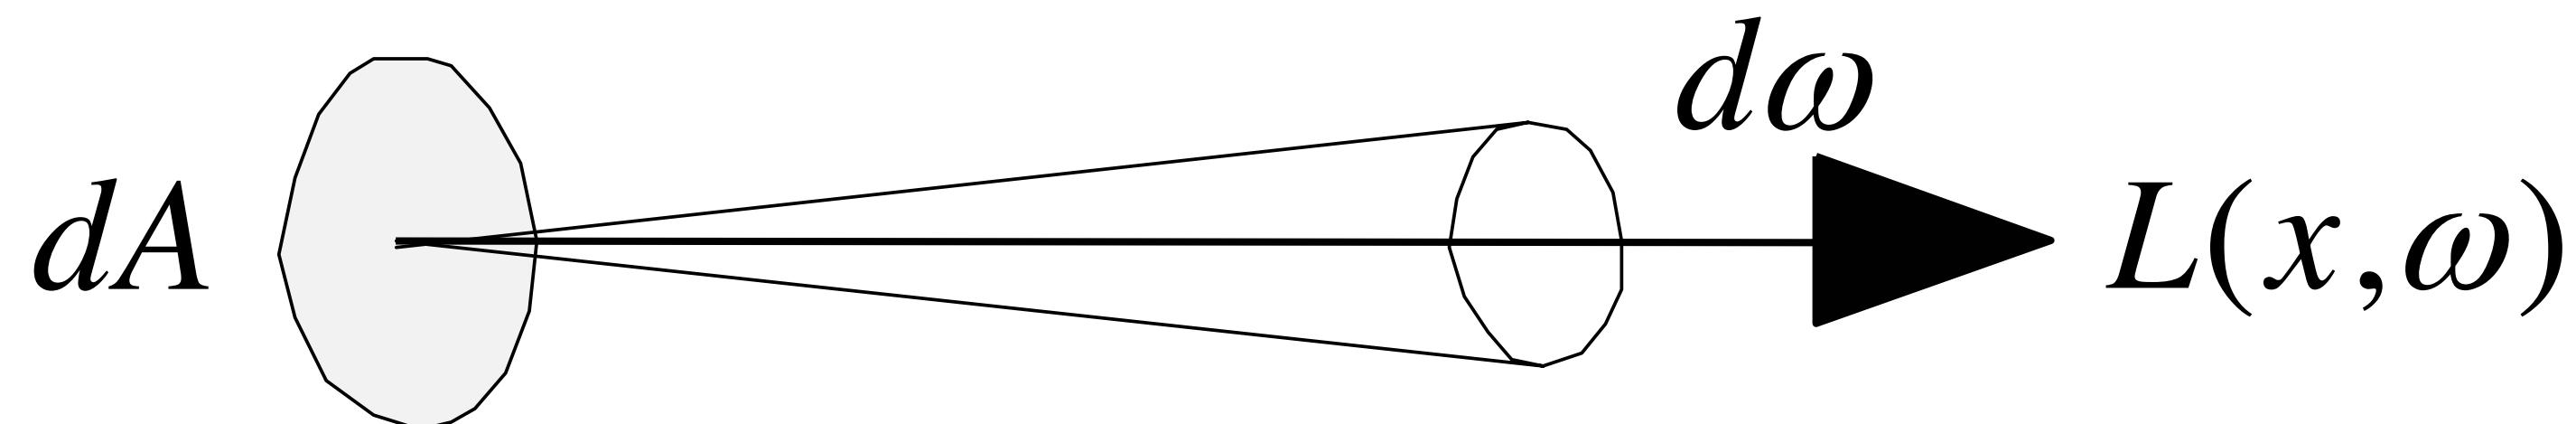
$$V[x, y, z] \in \mathbb{R}^+ \quad f(\mathbf{p}) \in \mathbb{R}^+$$

Q: How to render such volumetric representations?

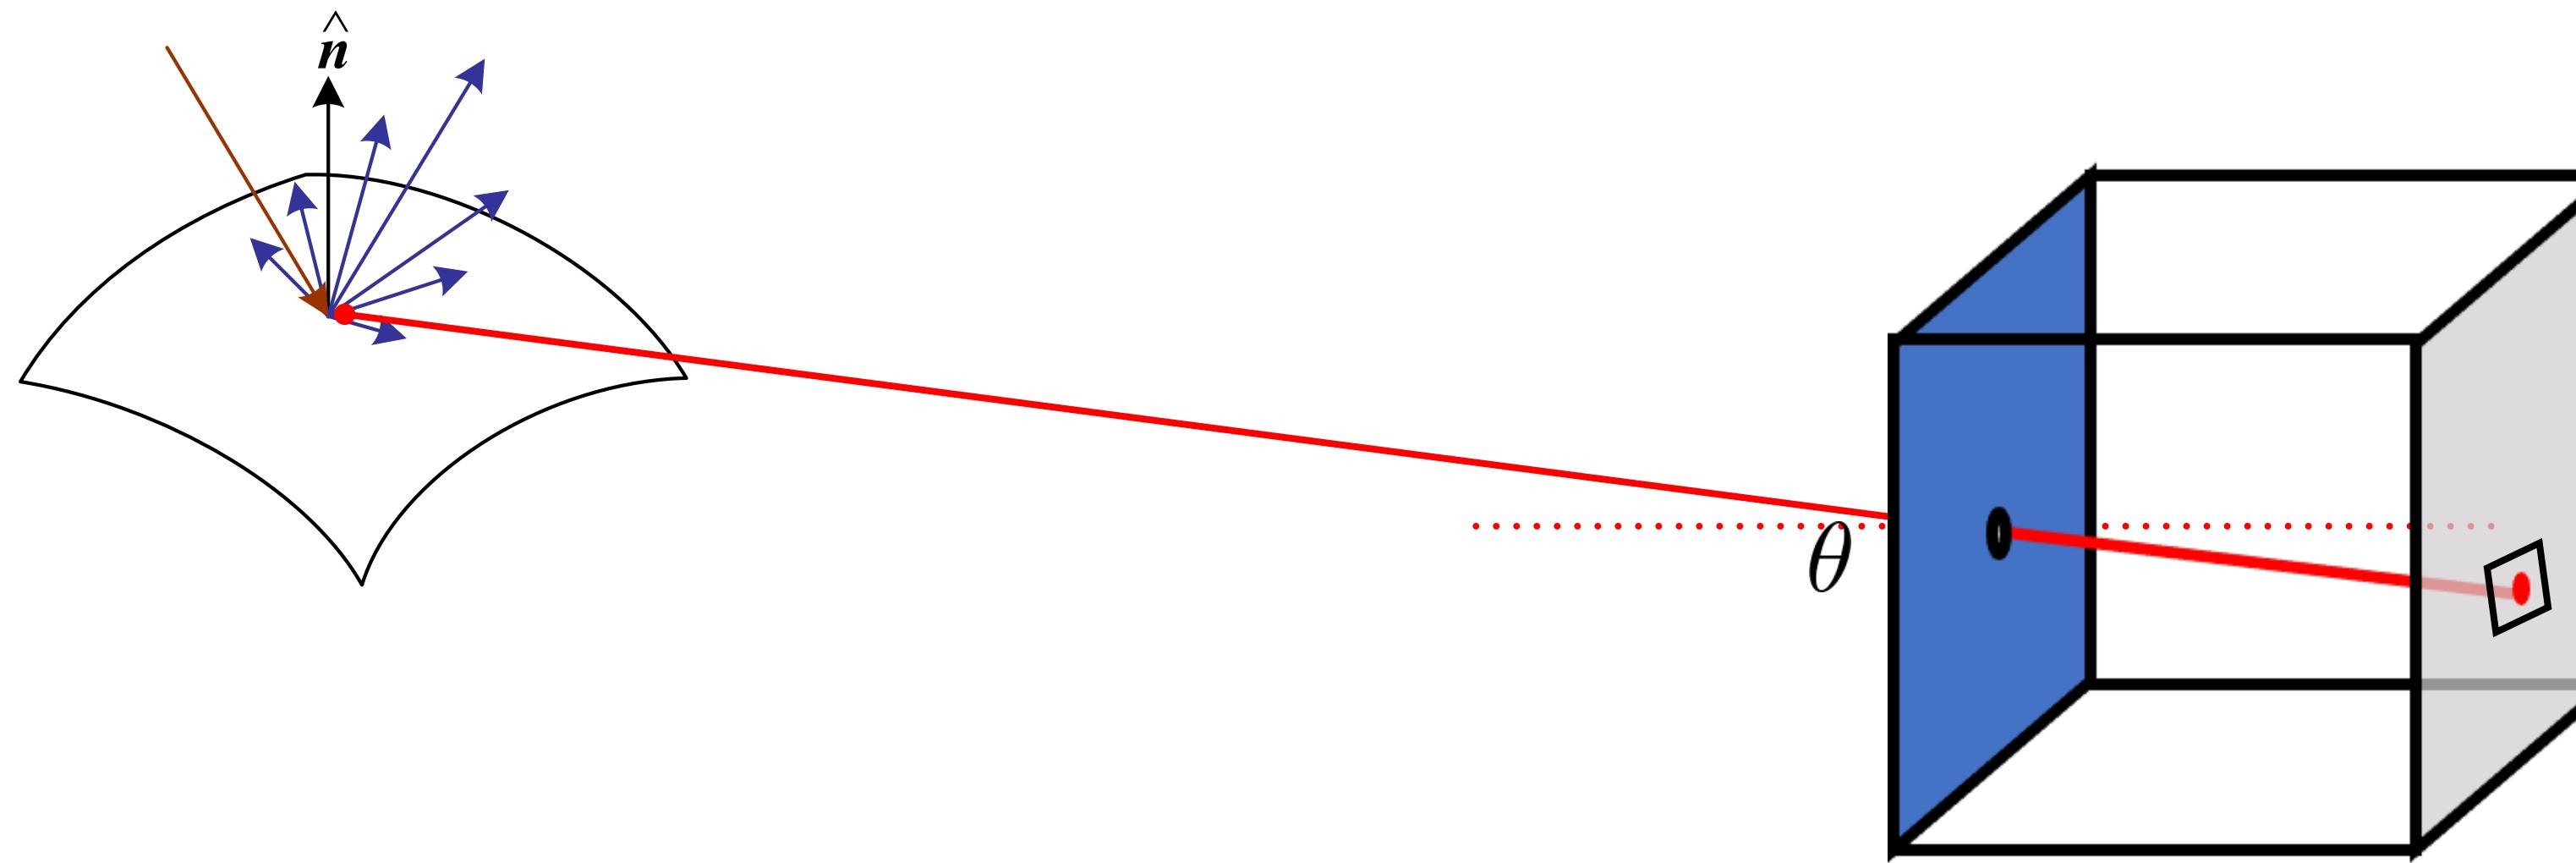
What does a pixel measure?



Definition: The field **radiance (luminance)** at a point in space in a given direction is the power per unit solid angle per unit area perpendicular to the direction



Surface Rendering



$$\propto L(\mathbf{x}^*, \omega)$$

radiance for: \mathbf{x}^* = optical centre, w = direction from \mathbf{x}^* to pixel sensor

$$= L(\mathbf{x}^* - \lambda\omega, \omega)$$

Pixel value → Emitted radiance at a **single** (surface) point

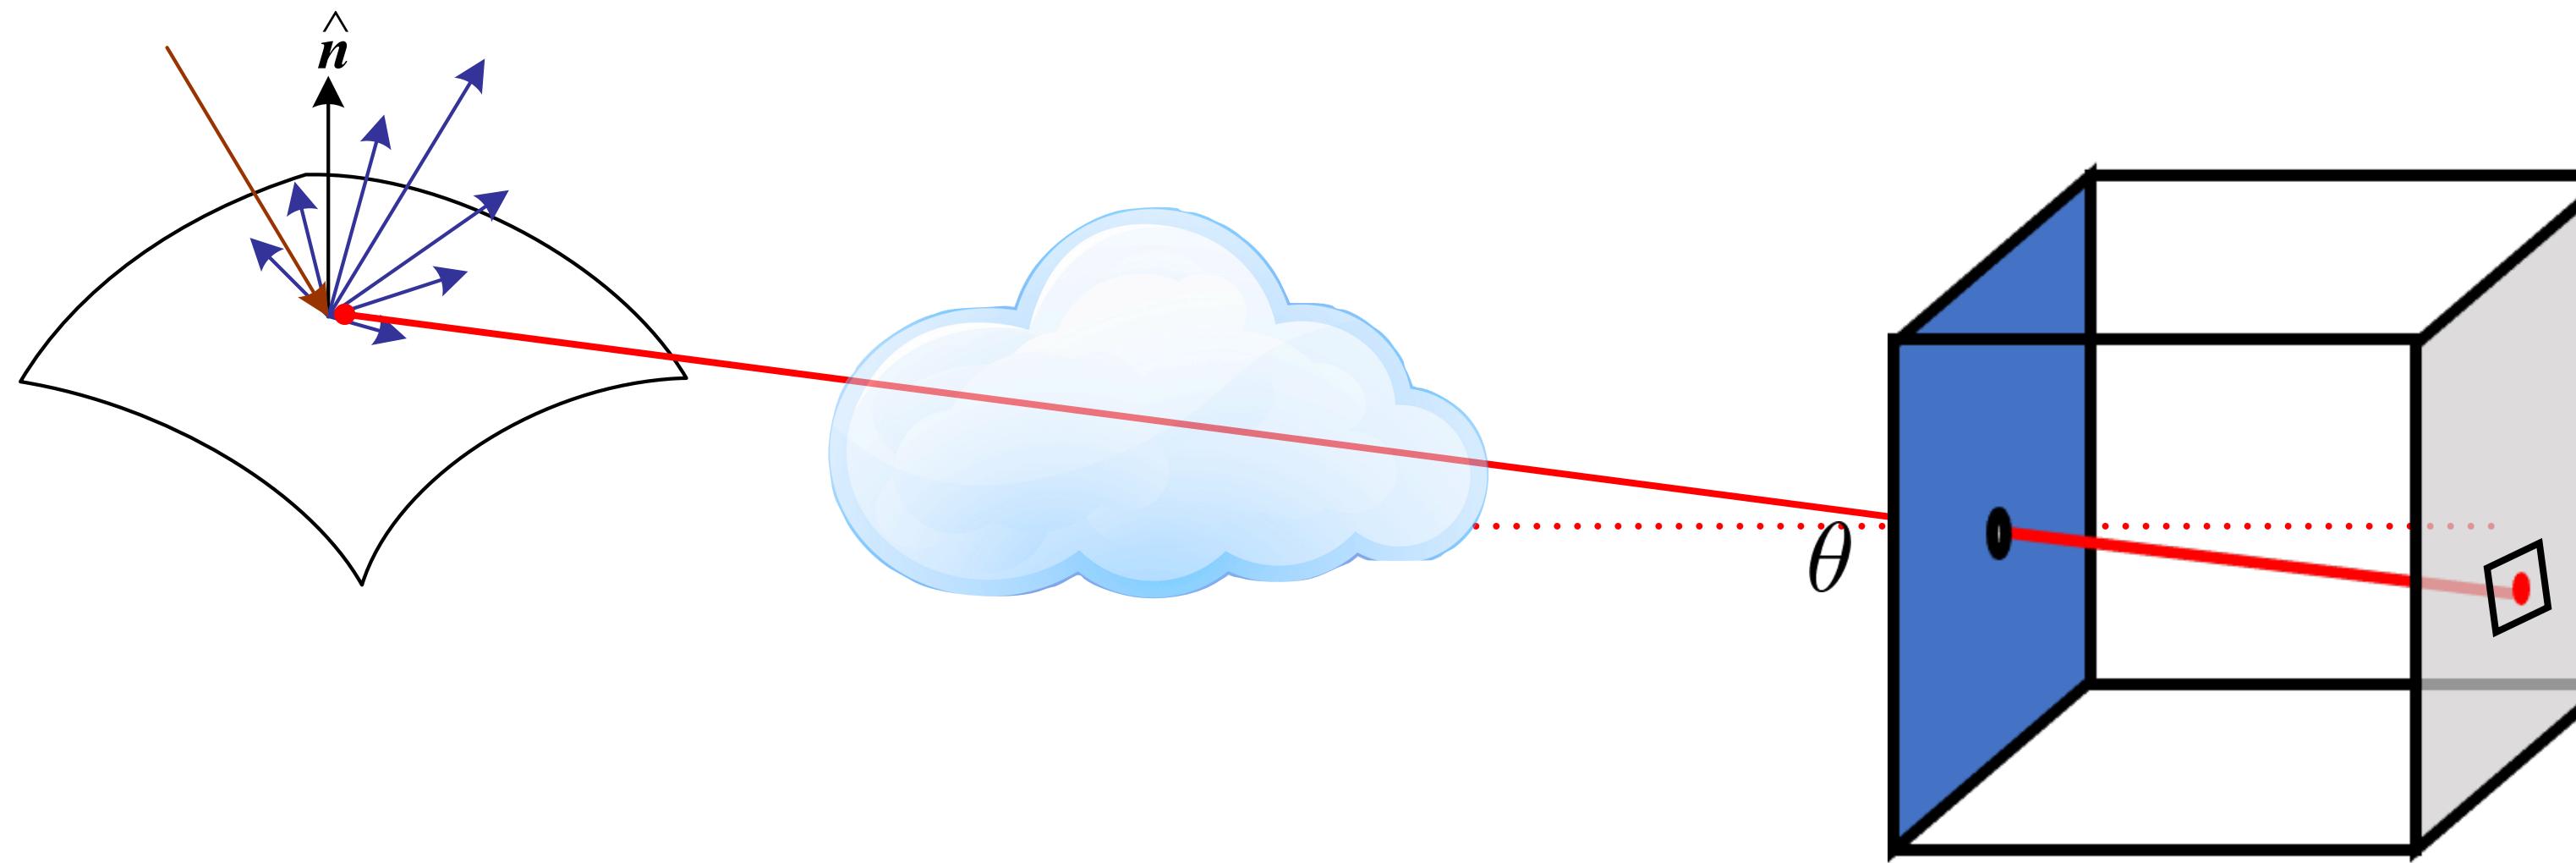
Beyond Surfaces



Not a valid assumption for fog, clouds, scattering, liquids etc.



Beyond Surfaces



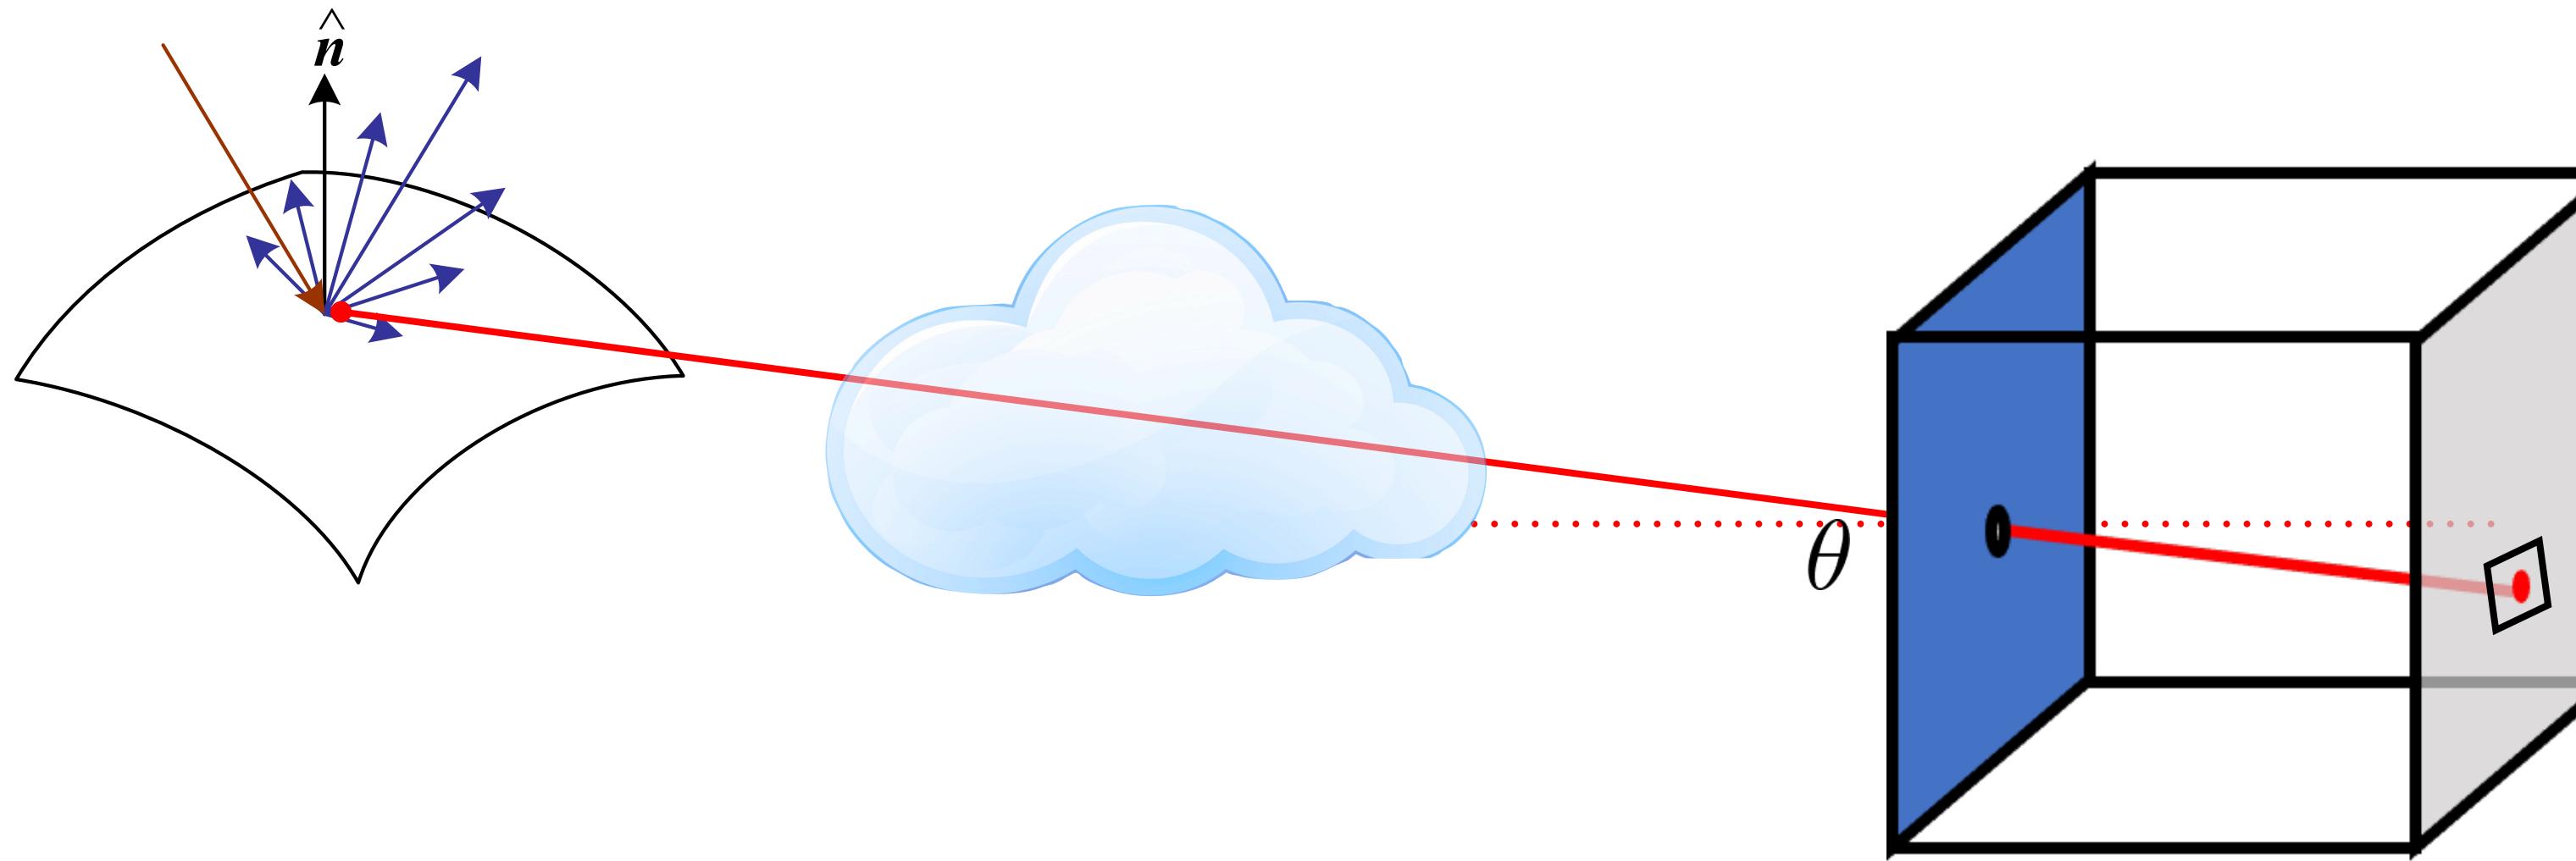
$$\propto L(\mathbf{x}^*, \omega)$$

radiance for: \mathbf{x}^* = pixel sensor centre, ω = direction from x to optical centre

$$\neq L(\mathbf{x}^* - \lambda\omega, \omega)$$

Pixel value **cannot** be reduced to emitted radiance by a single surface point

Volume Rendering



$$\propto L(\mathbf{x}^*, \omega)$$

radiance for: \mathbf{x}^* = pixel sensor centre, w = direction from x to optical centre

How does $L(x, w)$ vary along a ray?

Absorption

Light traveling in a medium
can get absorbed by it



<http://anordinarymom.wordpress.com>

Absorption

Light traveling in a medium
can get absorbed by it



Absorption



Thus, if one is to be five times as distant, make it five times bluer.

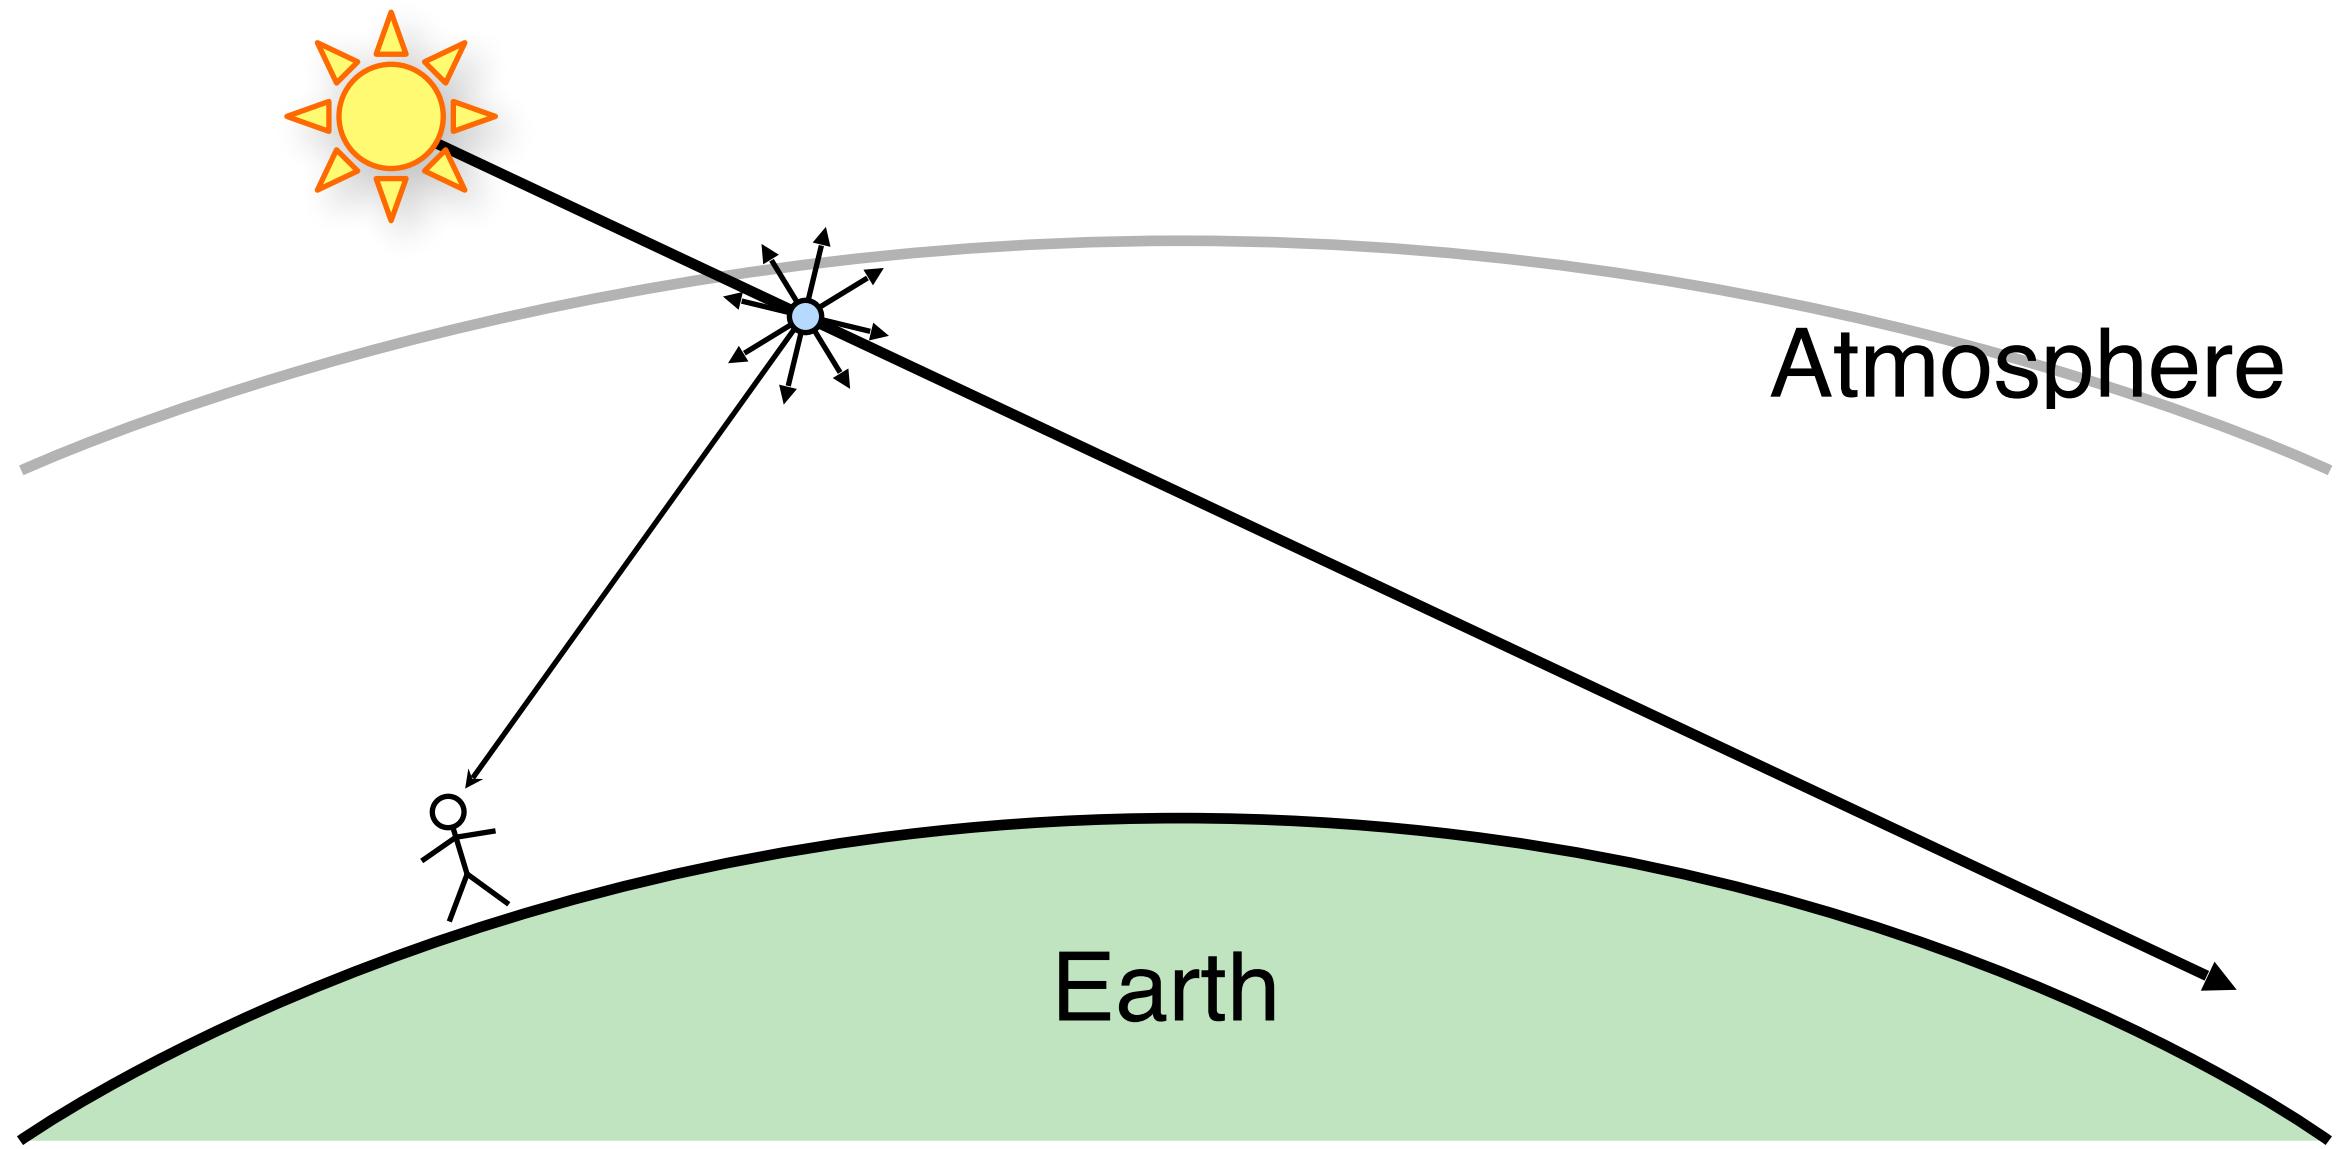
—Treatise on Painting, Leonardo Da Vinci, pp 295, circa 1480.

Emission

The medium can emit light on its own



Scattering

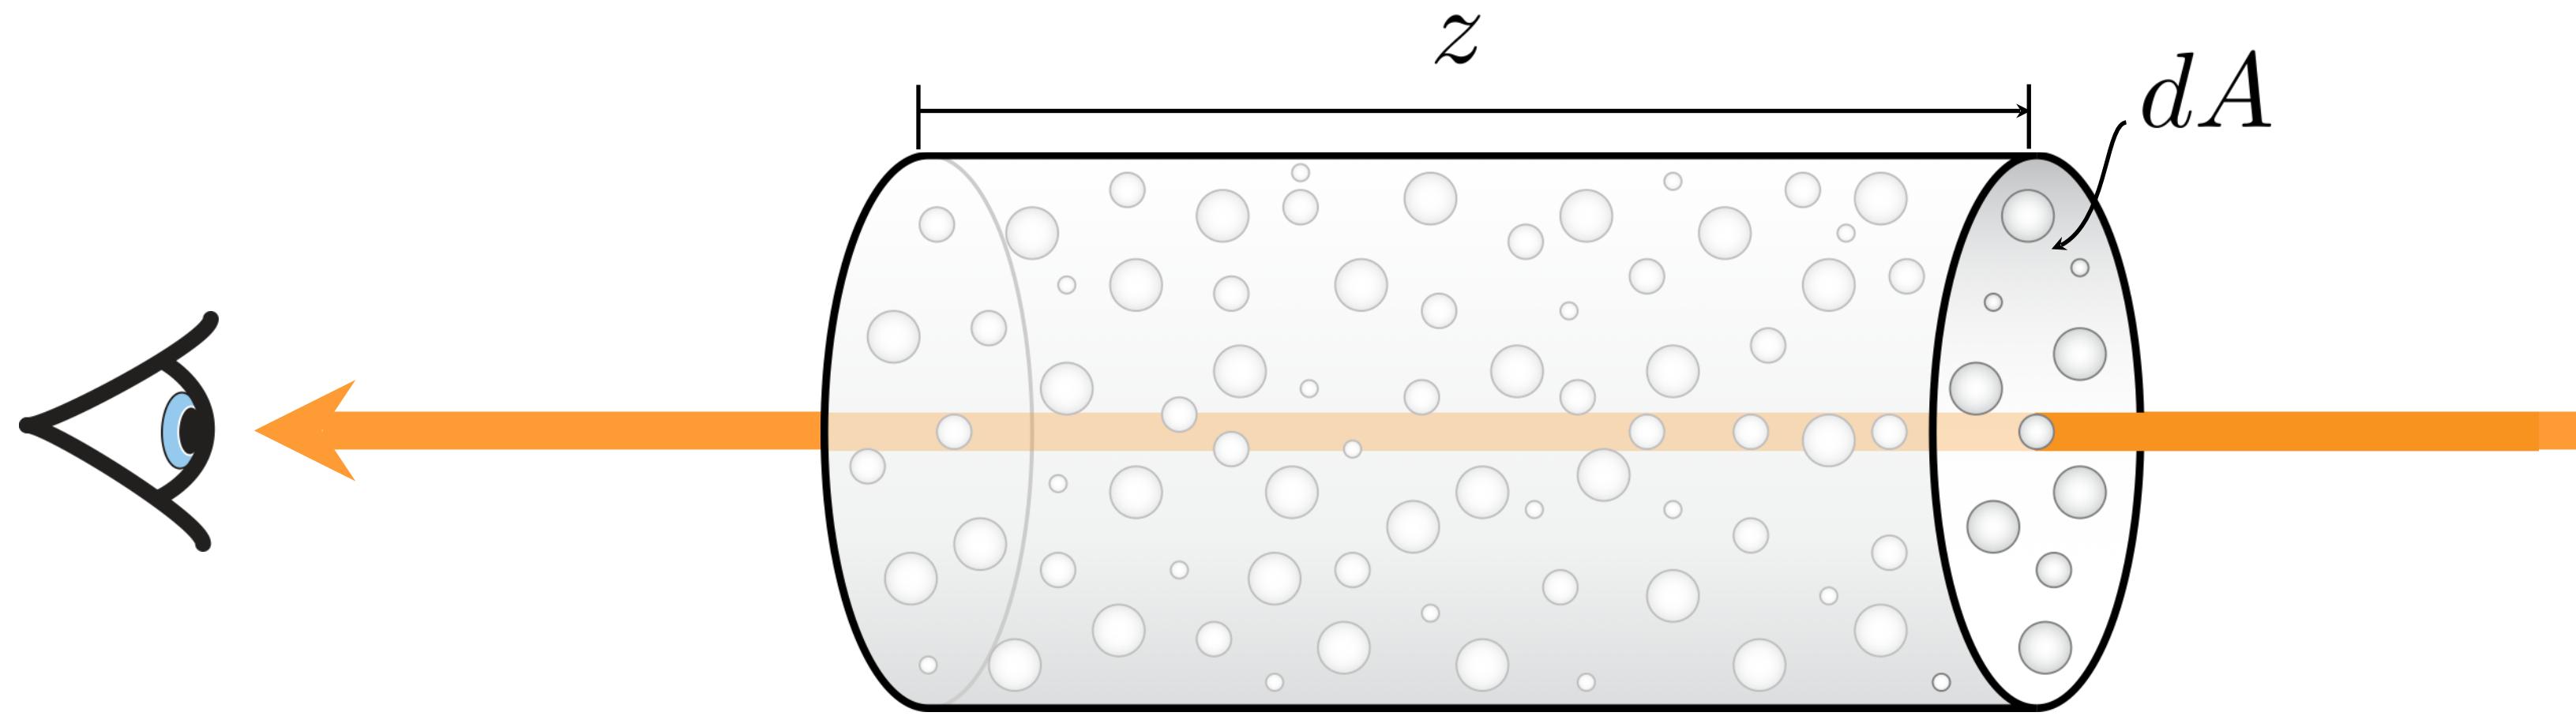


Light along a ray can depend on light
in other directions



(we will **not** model scattering effects in today's lecture)

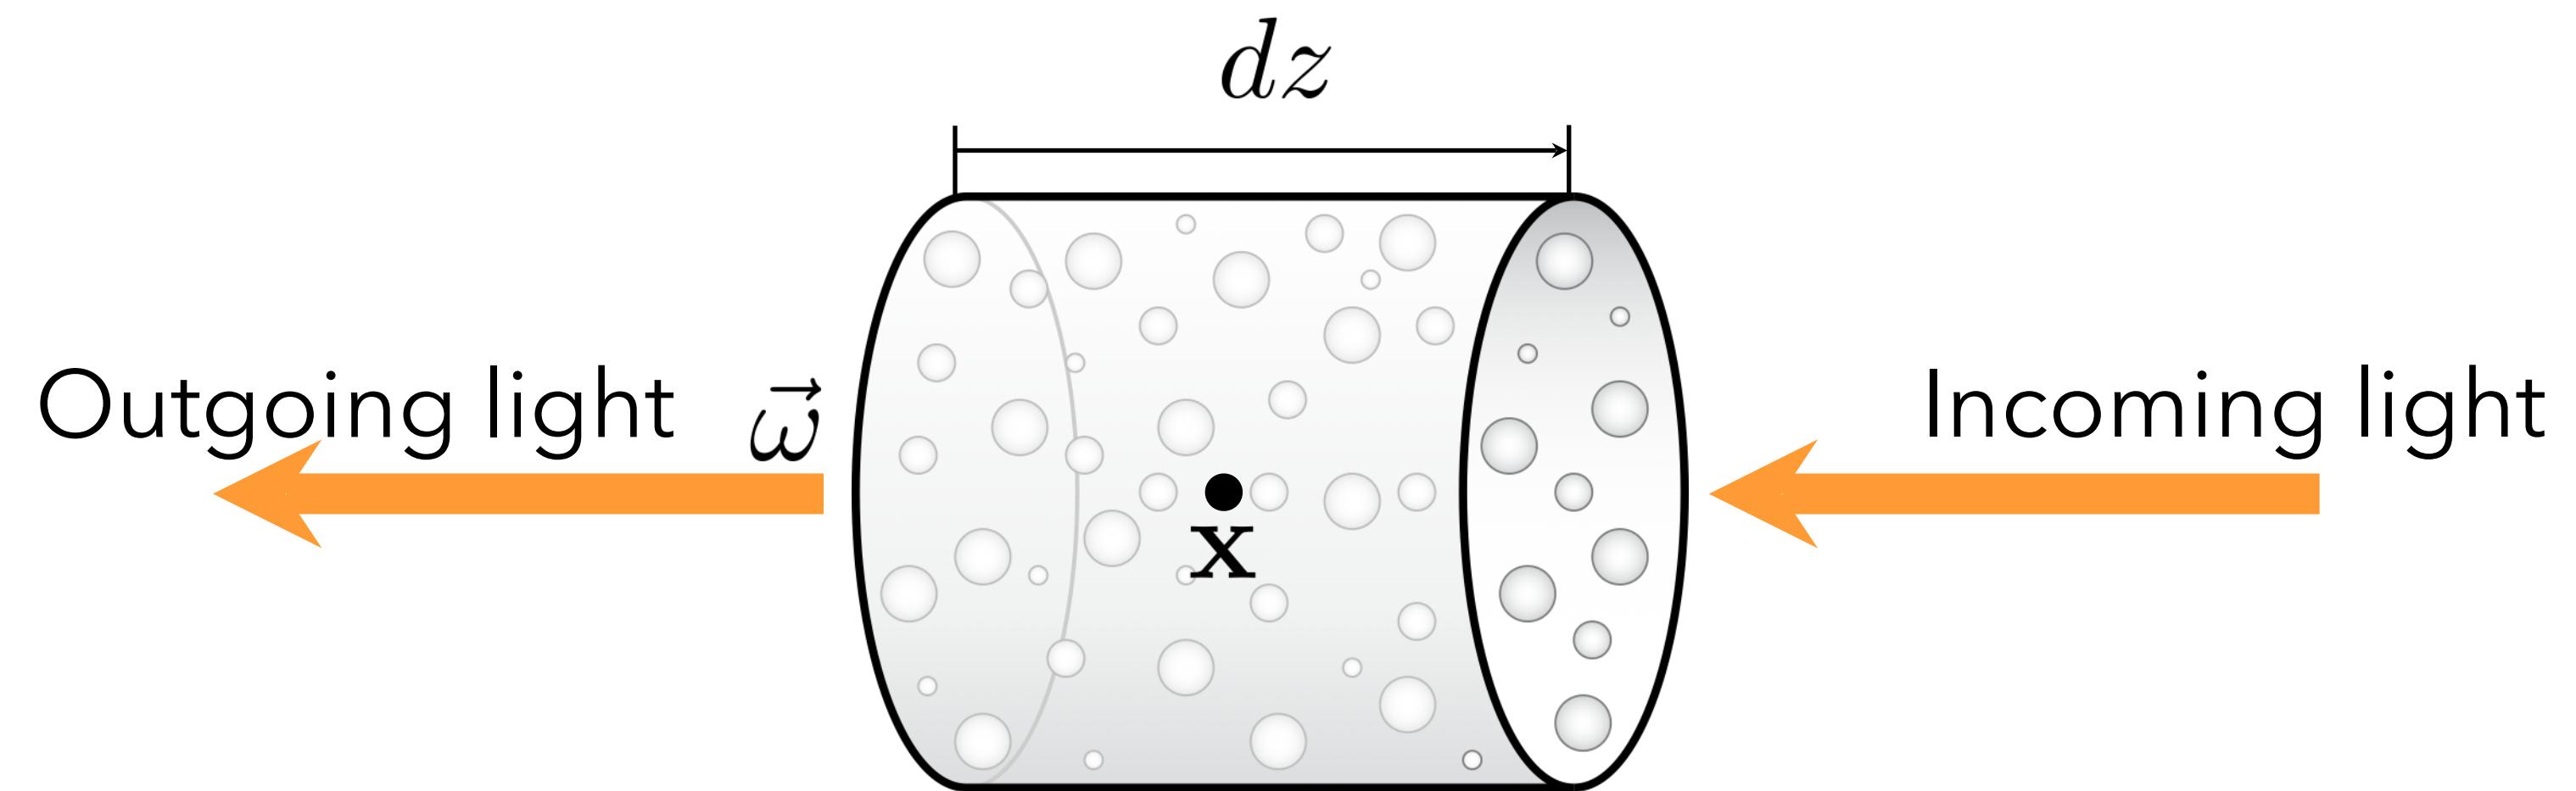
Volume Rendering a Ray



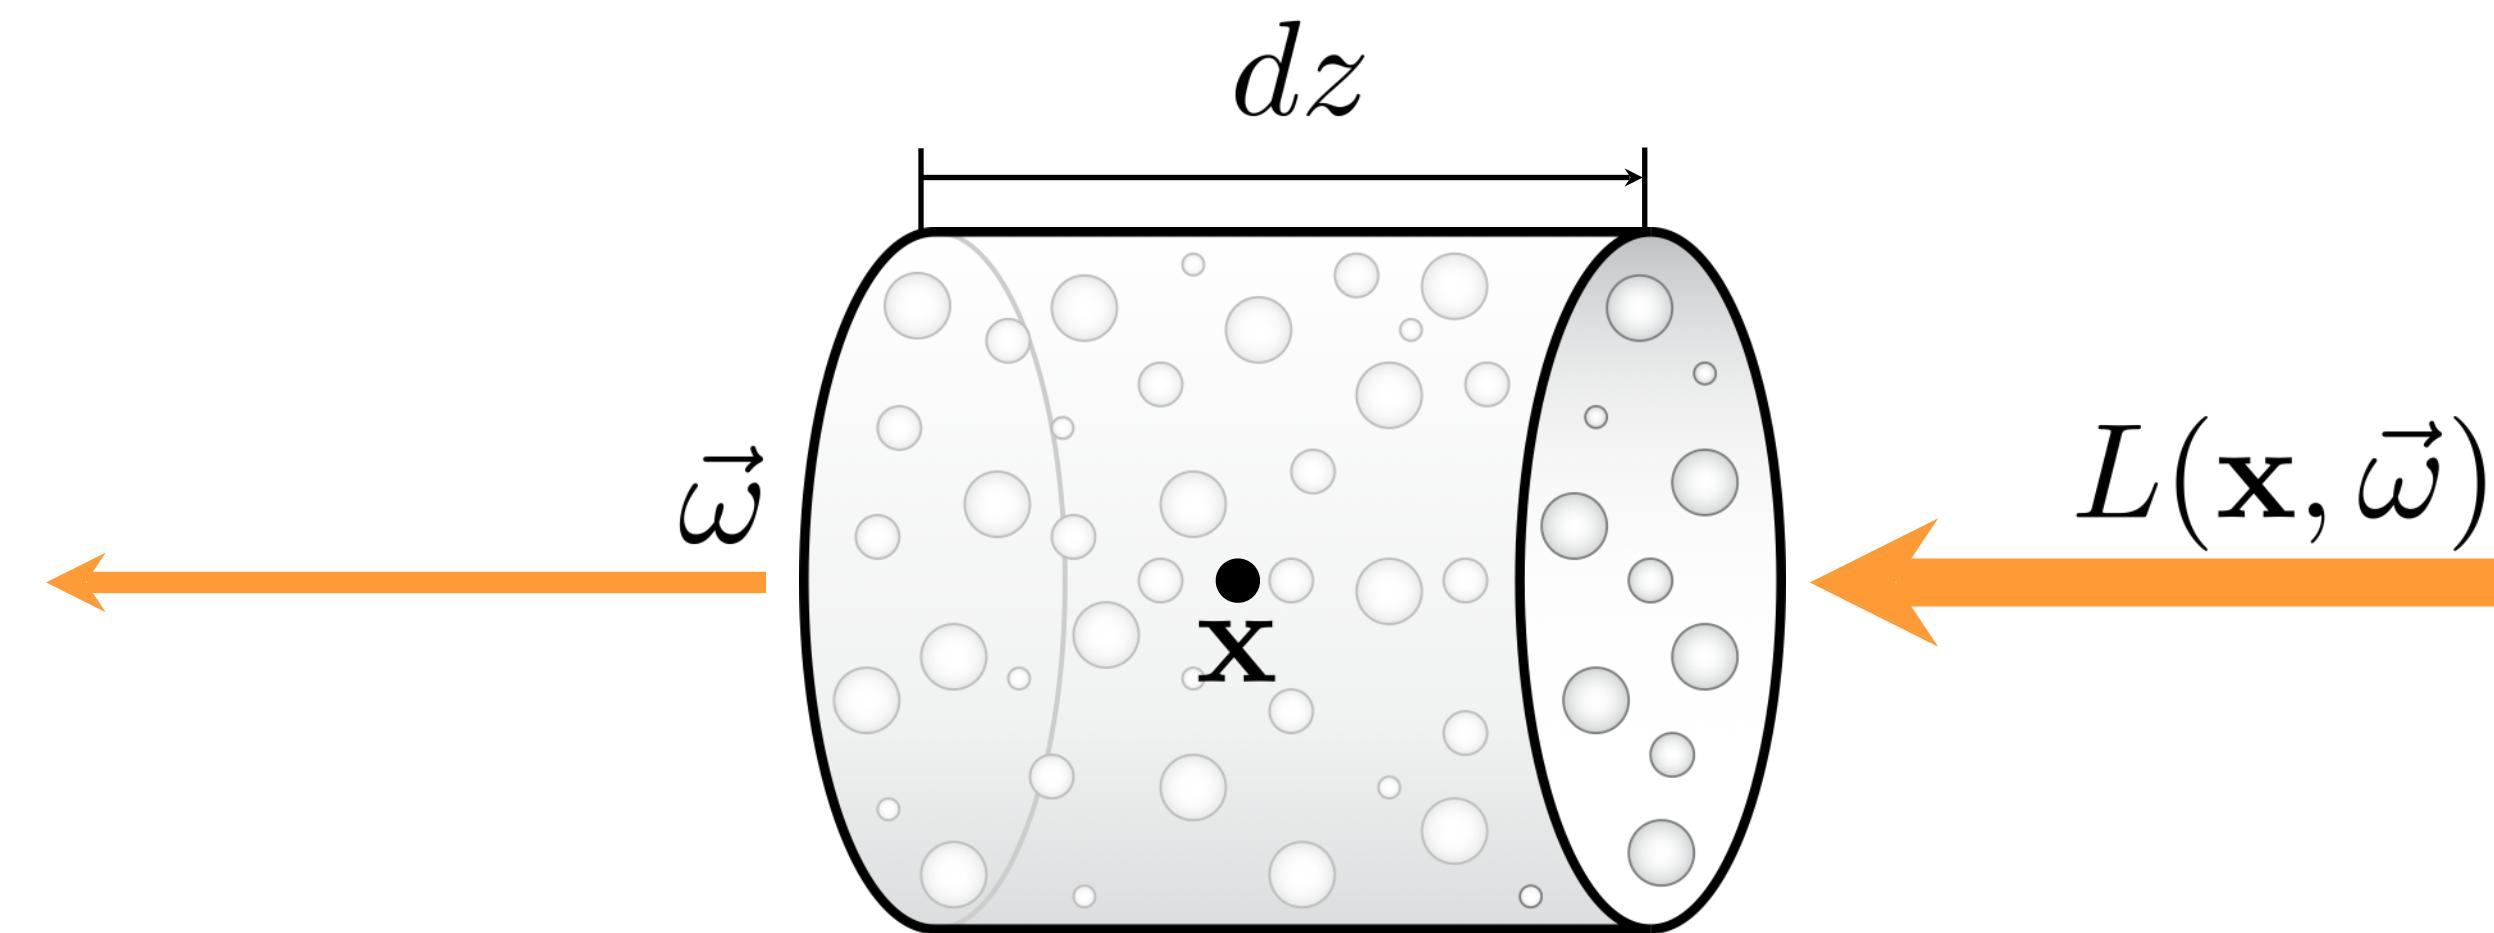
How much light is *lost/gained* along the differential beam due to interactions of light with the medium?

How does $L(x, w)$ vary along a ray?

Volume Rendering a Ray



Volume Rendering a Ray: Absorption

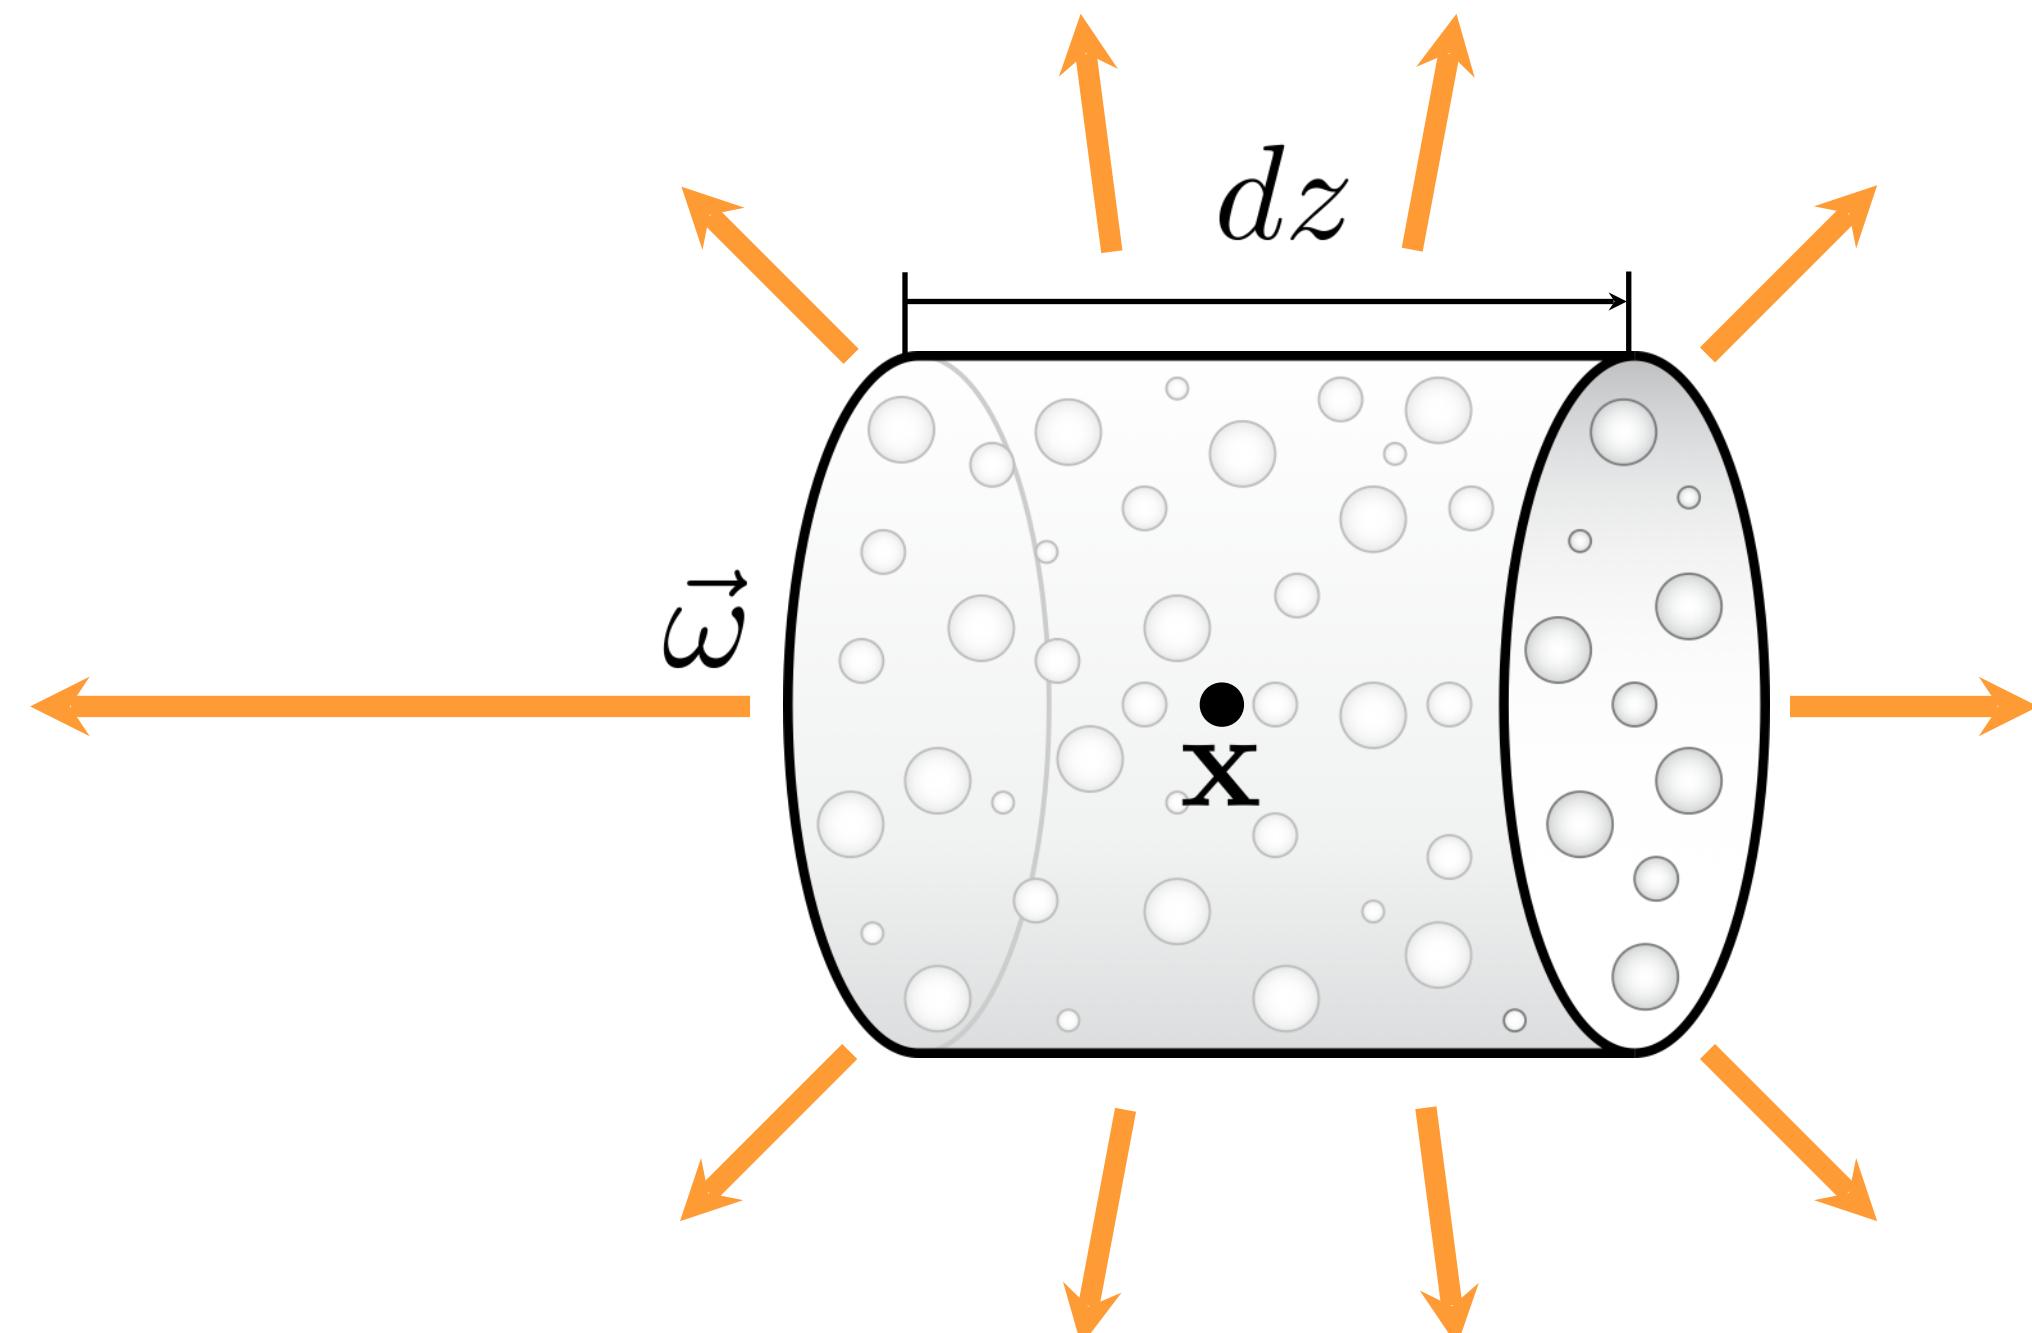


$$dL(\mathbf{x}, \vec{\omega}) = -\sigma_a(\mathbf{x})L(\mathbf{x}, \vec{\omega})dz$$

$\sigma_a(\mathbf{x})$: absorption coefficient

Q: Why should dL be **proportional** to L (incoming radiance)?

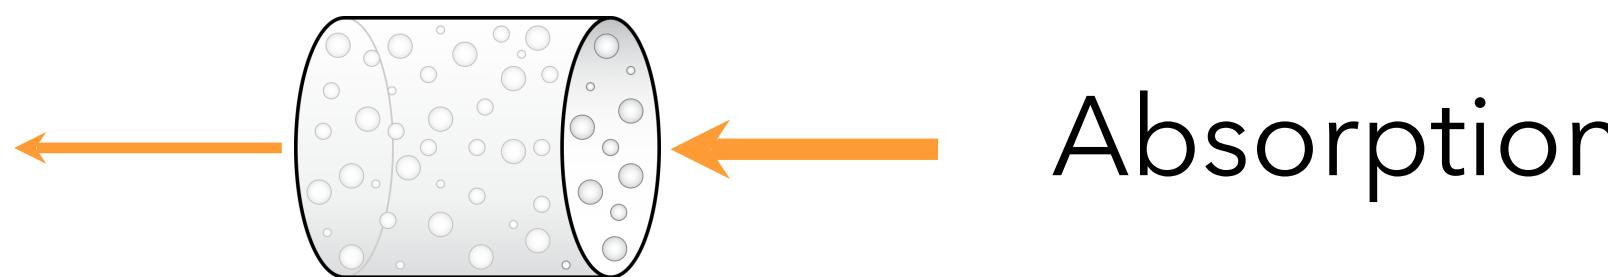
Volume Rendering a Ray: Emission



$$dL(\mathbf{x}, \vec{\omega}) = \sigma_a(\mathbf{x}) L_e(\mathbf{x}, \vec{\omega}) dz$$

$L_e(\mathbf{x}, \vec{\omega})$: emitted radiance

Volume Rendering: Emission-Absorption Model

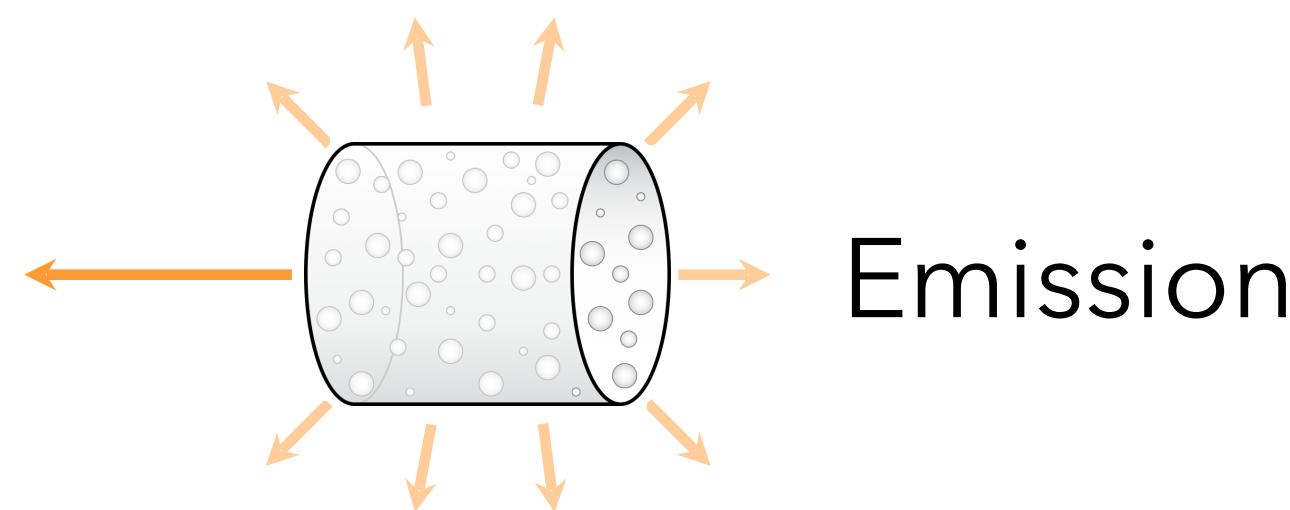


$$dL(\mathbf{x}, \vec{\omega}) = \boxed{-\sigma_a(\mathbf{x})L(\mathbf{x}, \vec{\omega})dz}$$

Losses

$$\boxed{+\sigma_a(\mathbf{x})L_e(\mathbf{x}, \vec{\omega})dz}$$

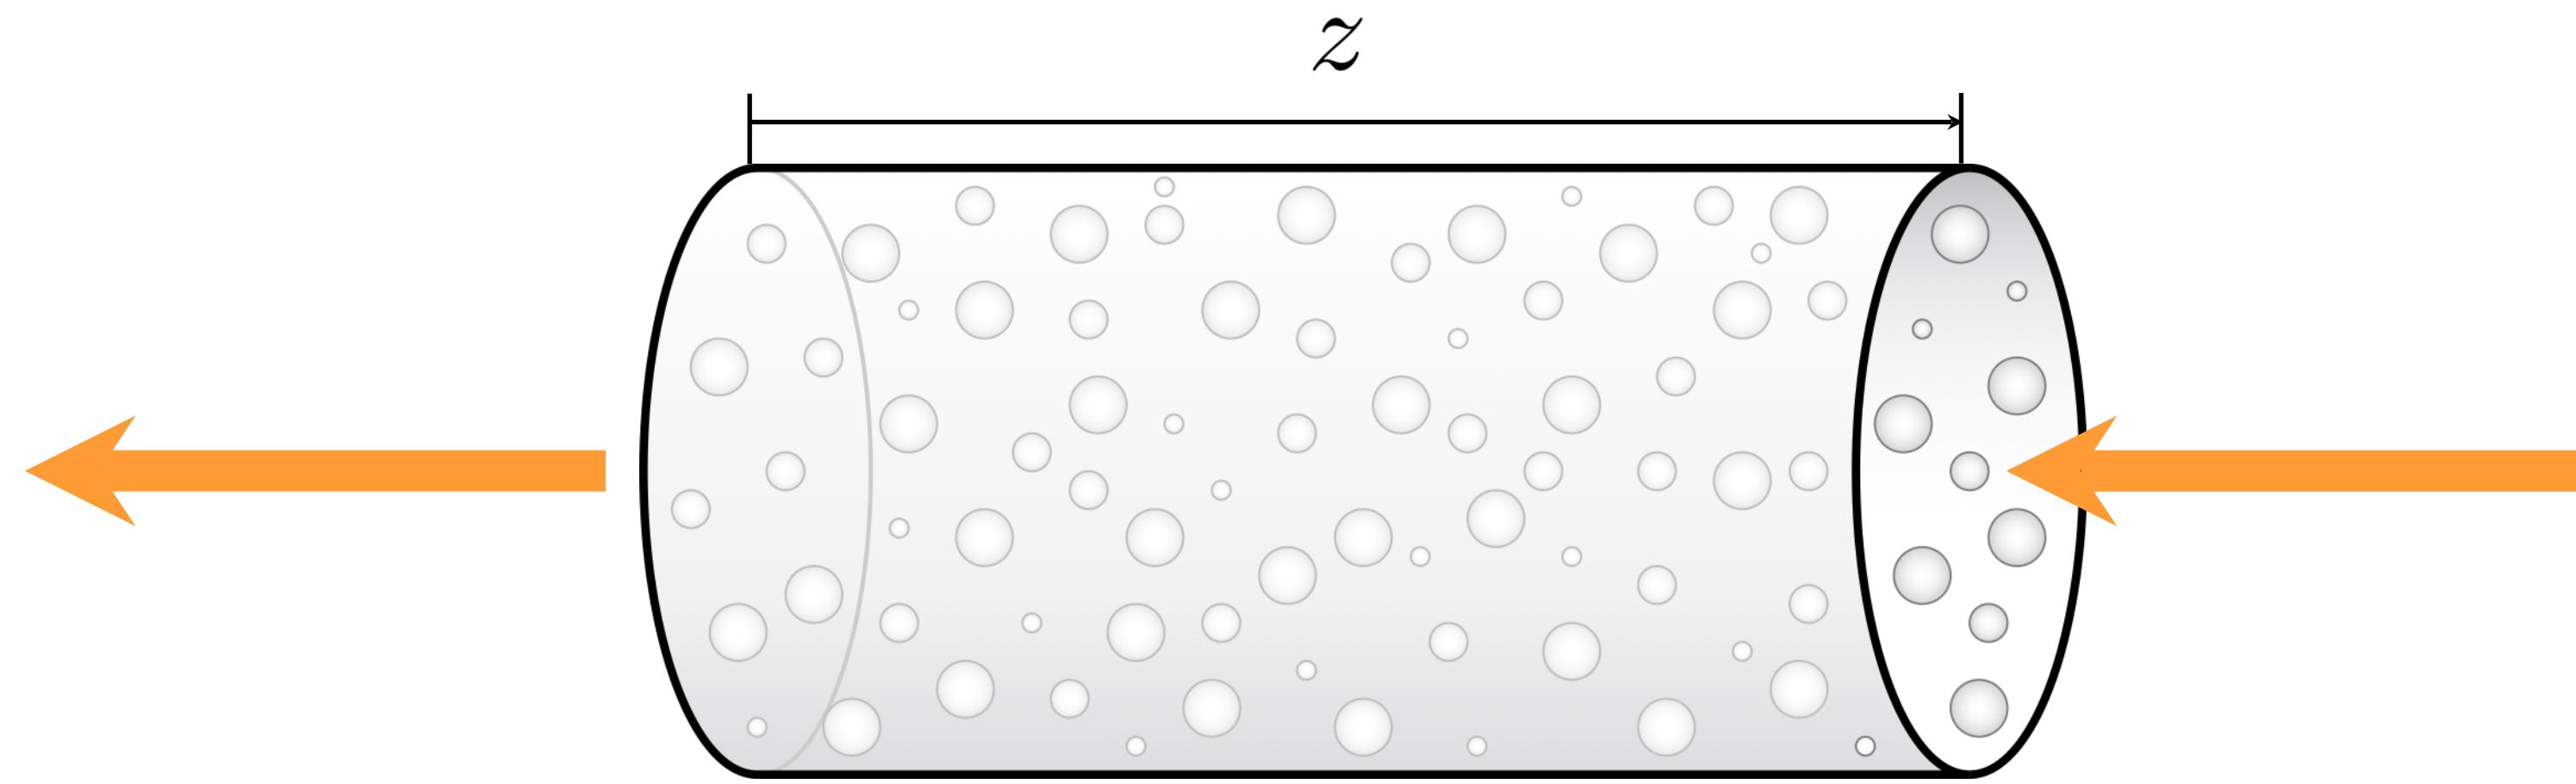
Gains



Will drop the subscript moving forward

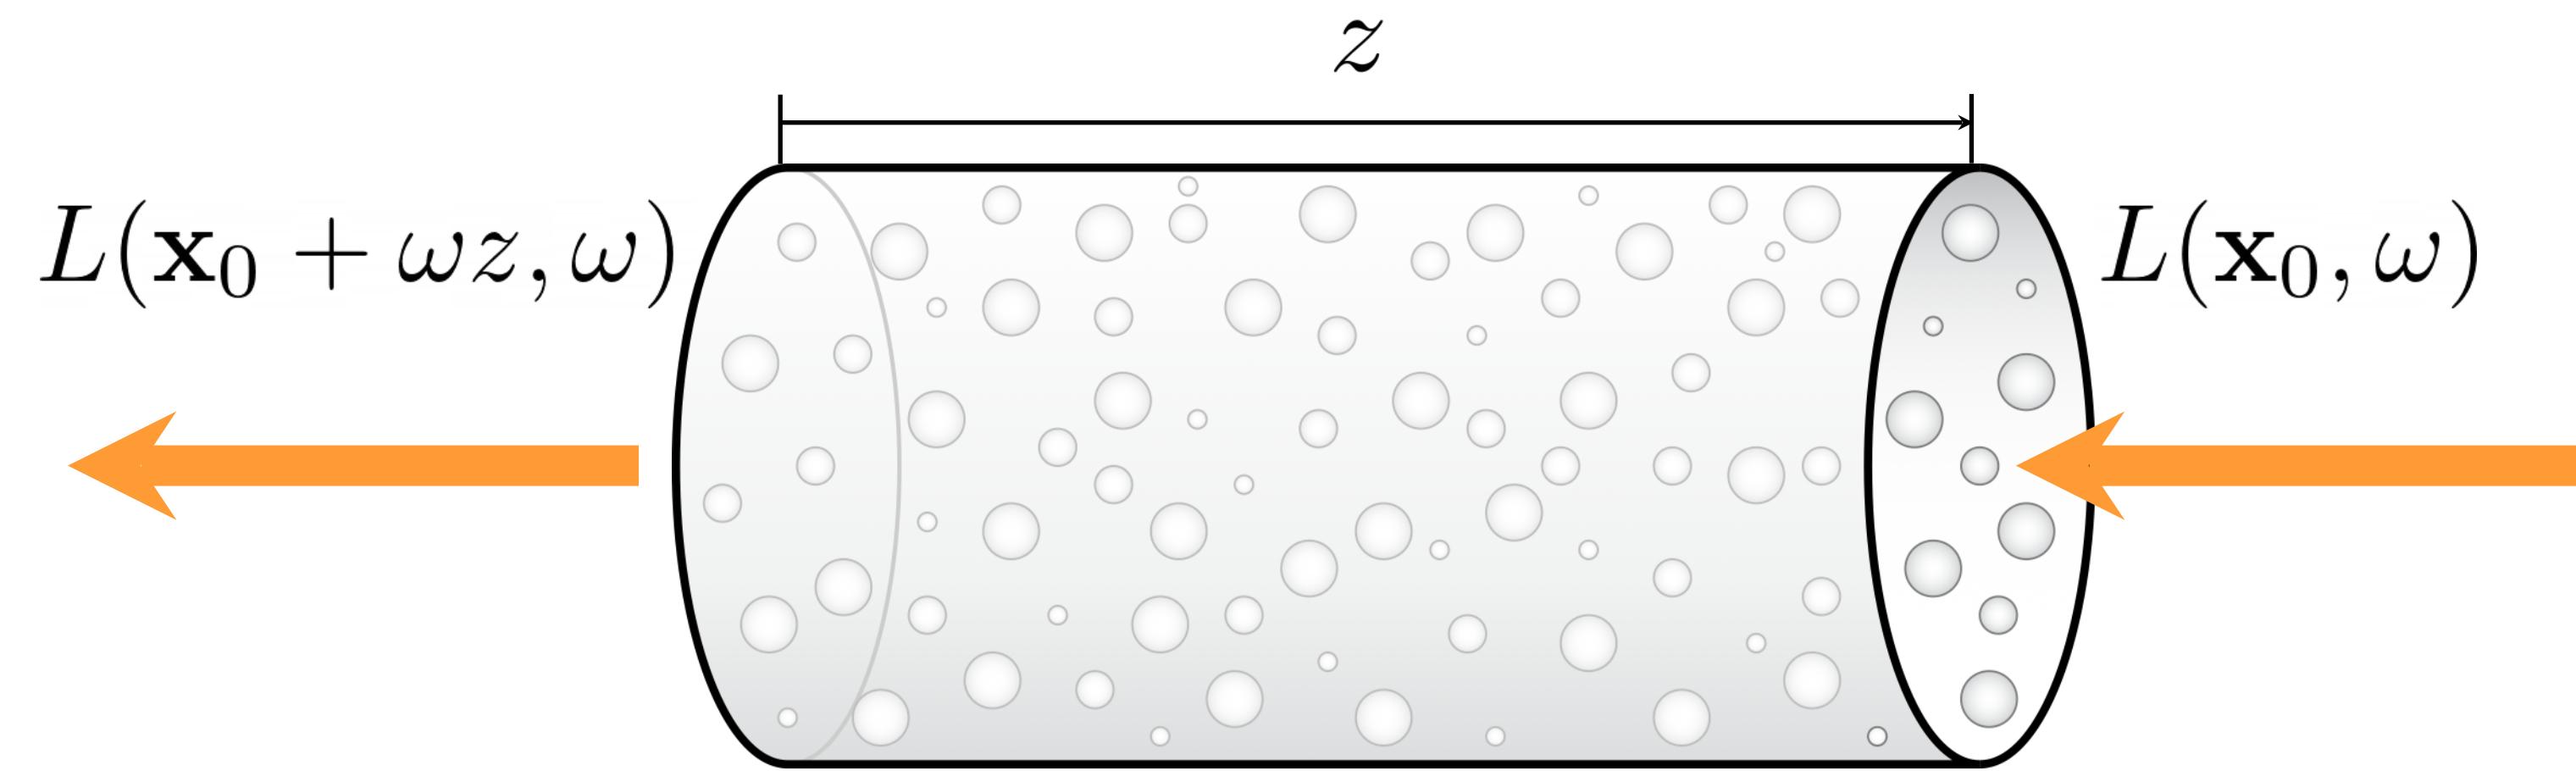
$$\sigma_a \equiv \sigma$$

From infinitesimal to finite lengths



- 1. Model absorption-only case**
- 2. Allow emission**

Absorption-only Volume Rendering

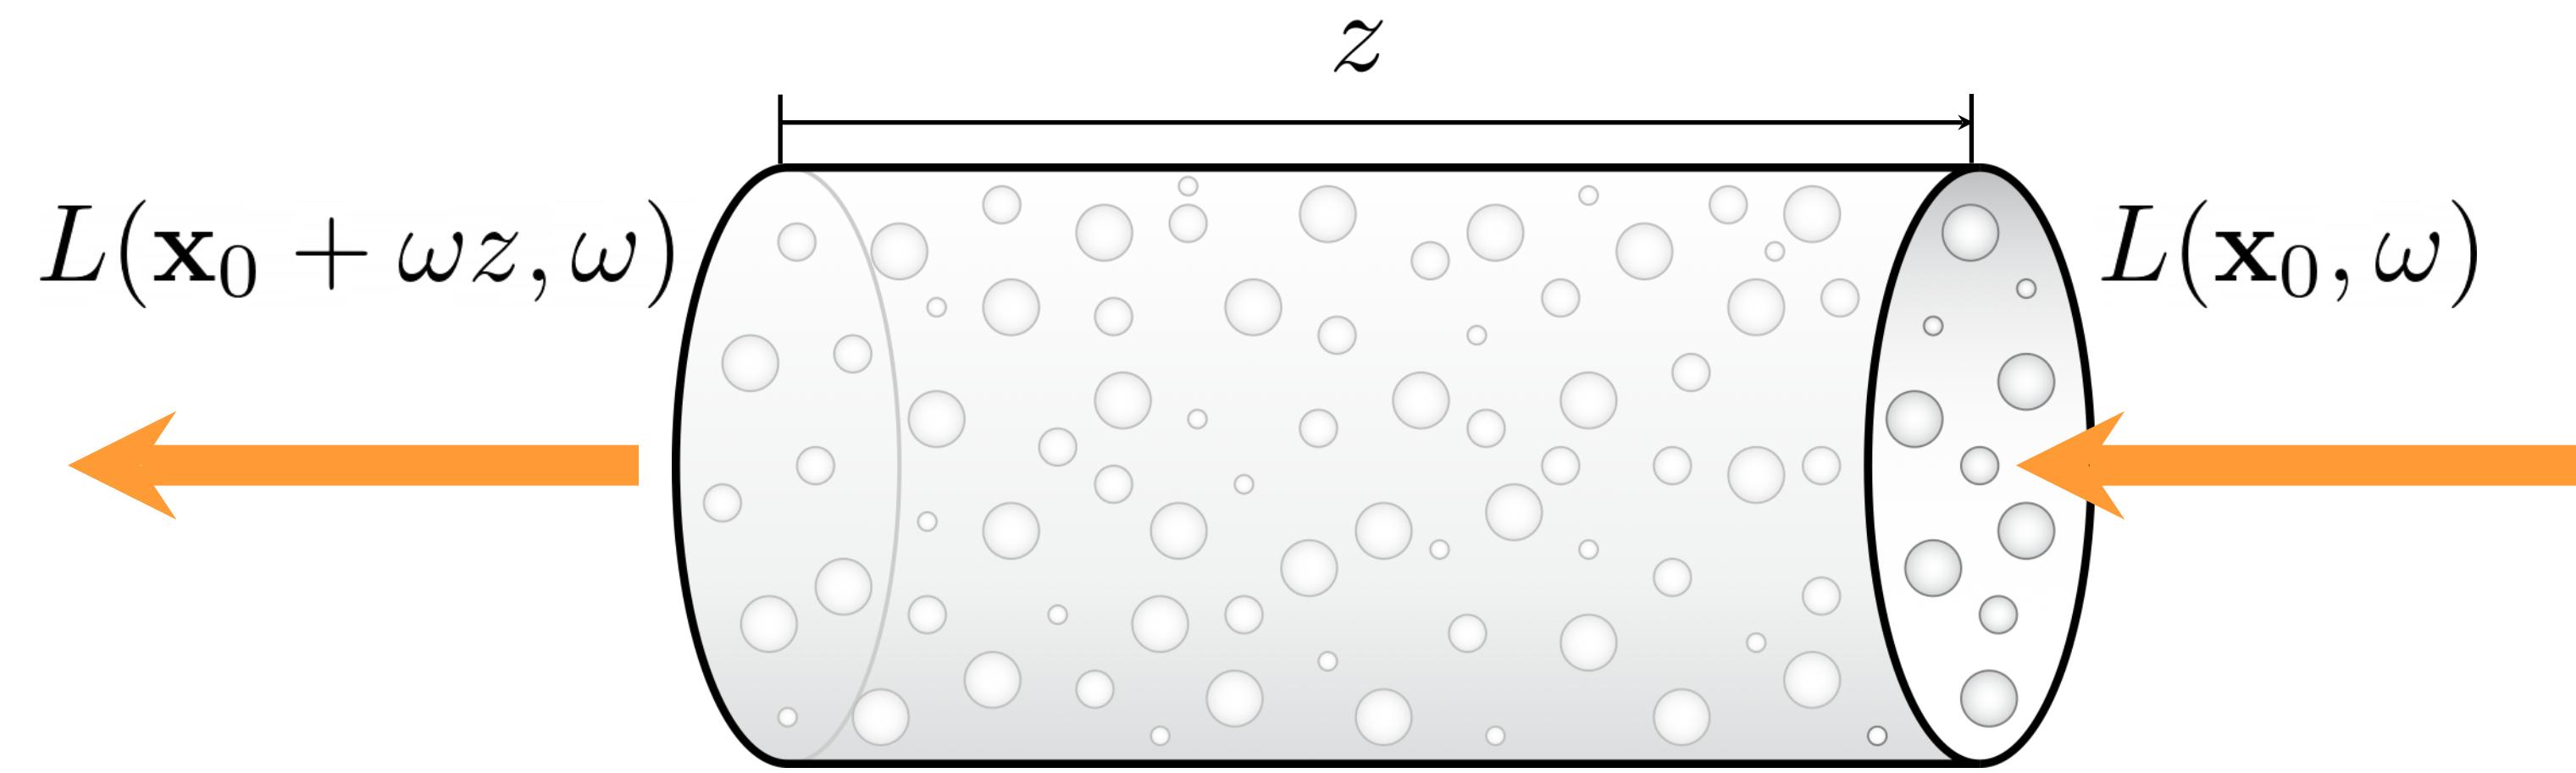


Q: What if we have a homogenous medium? (uniform coefficient)

$$dL(\mathbf{x}, \omega) = -\sigma L(\mathbf{x}, \omega) dz$$

$$L(\mathbf{x}_0 + \omega z, \omega) = e^{-\sigma z} L(\mathbf{x}_0, \omega)$$

Absorption-only Volume Rendering



Q: What if we have a non-homogenous medium?

$$dL(\mathbf{x}, \omega) = -\sigma(\mathbf{x})L(\mathbf{x}, \omega)dz$$

$$L(\mathbf{x}_0 + \omega z, \omega) = e^{-\int_{t=0}^z \sigma(\mathbf{x} + \omega \mathbf{t}) dt} L(\mathbf{x}_0, \omega)$$

Transmittance

Transmittance

$$T(\mathbf{x}, \mathbf{y})$$

What fraction of radiance at \mathbf{x} in direction of \mathbf{y} , reaches \mathbf{y} ?
(along a straight line under absorption-only model)

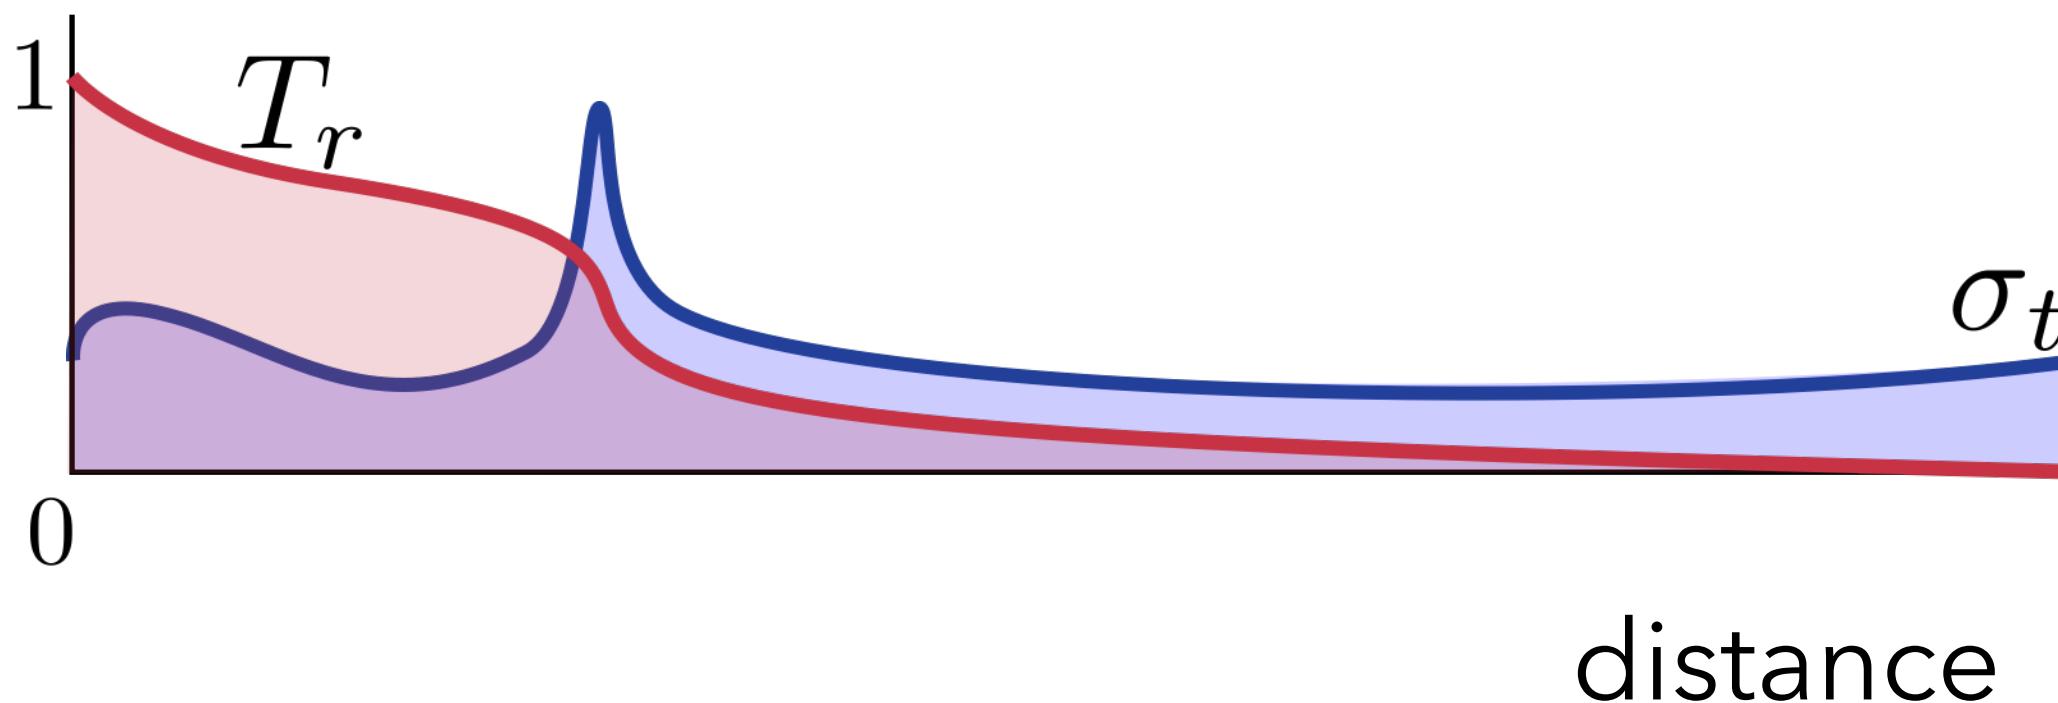
Q: should this be symmetric?

Homogenous Medium:

$$e^{-\sigma \|\mathbf{x}-\mathbf{y}\|}$$

Non-Homogenous Medium:

$$e^{-\int_{t=0}^{\|\mathbf{x}-\mathbf{y}\|} \sigma(\mathbf{x}+\omega t) dt}$$



Transmittance

$$T(\mathbf{x}, \mathbf{y})$$

What fraction of radiance at \mathbf{x} in direction of \mathbf{y} , reaches \mathbf{y} ?
(along a straight line under absorption-only model)

Homogenous Medium:

$$e^{-\sigma \|\mathbf{x}-\mathbf{y}\|}$$

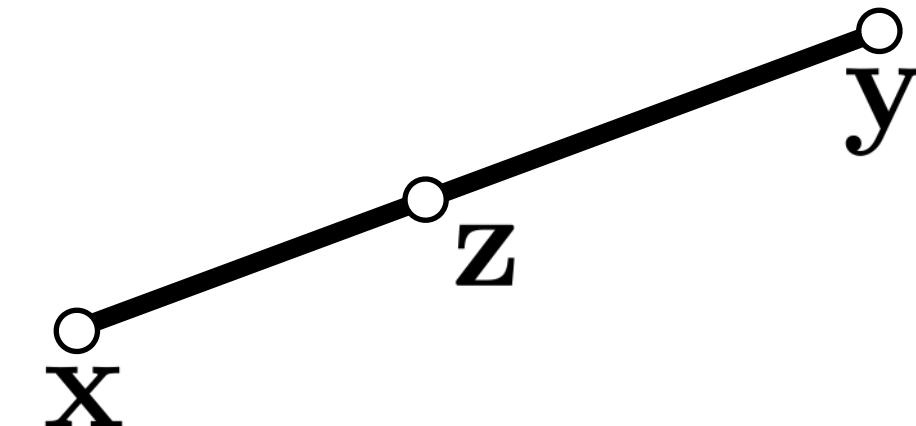
Non-Homogenous Medium:

$$e^{-\int_{t=0}^{\|\mathbf{x}-\mathbf{y}\|} \sigma(\mathbf{x}+\omega\mathbf{t}) dt}$$

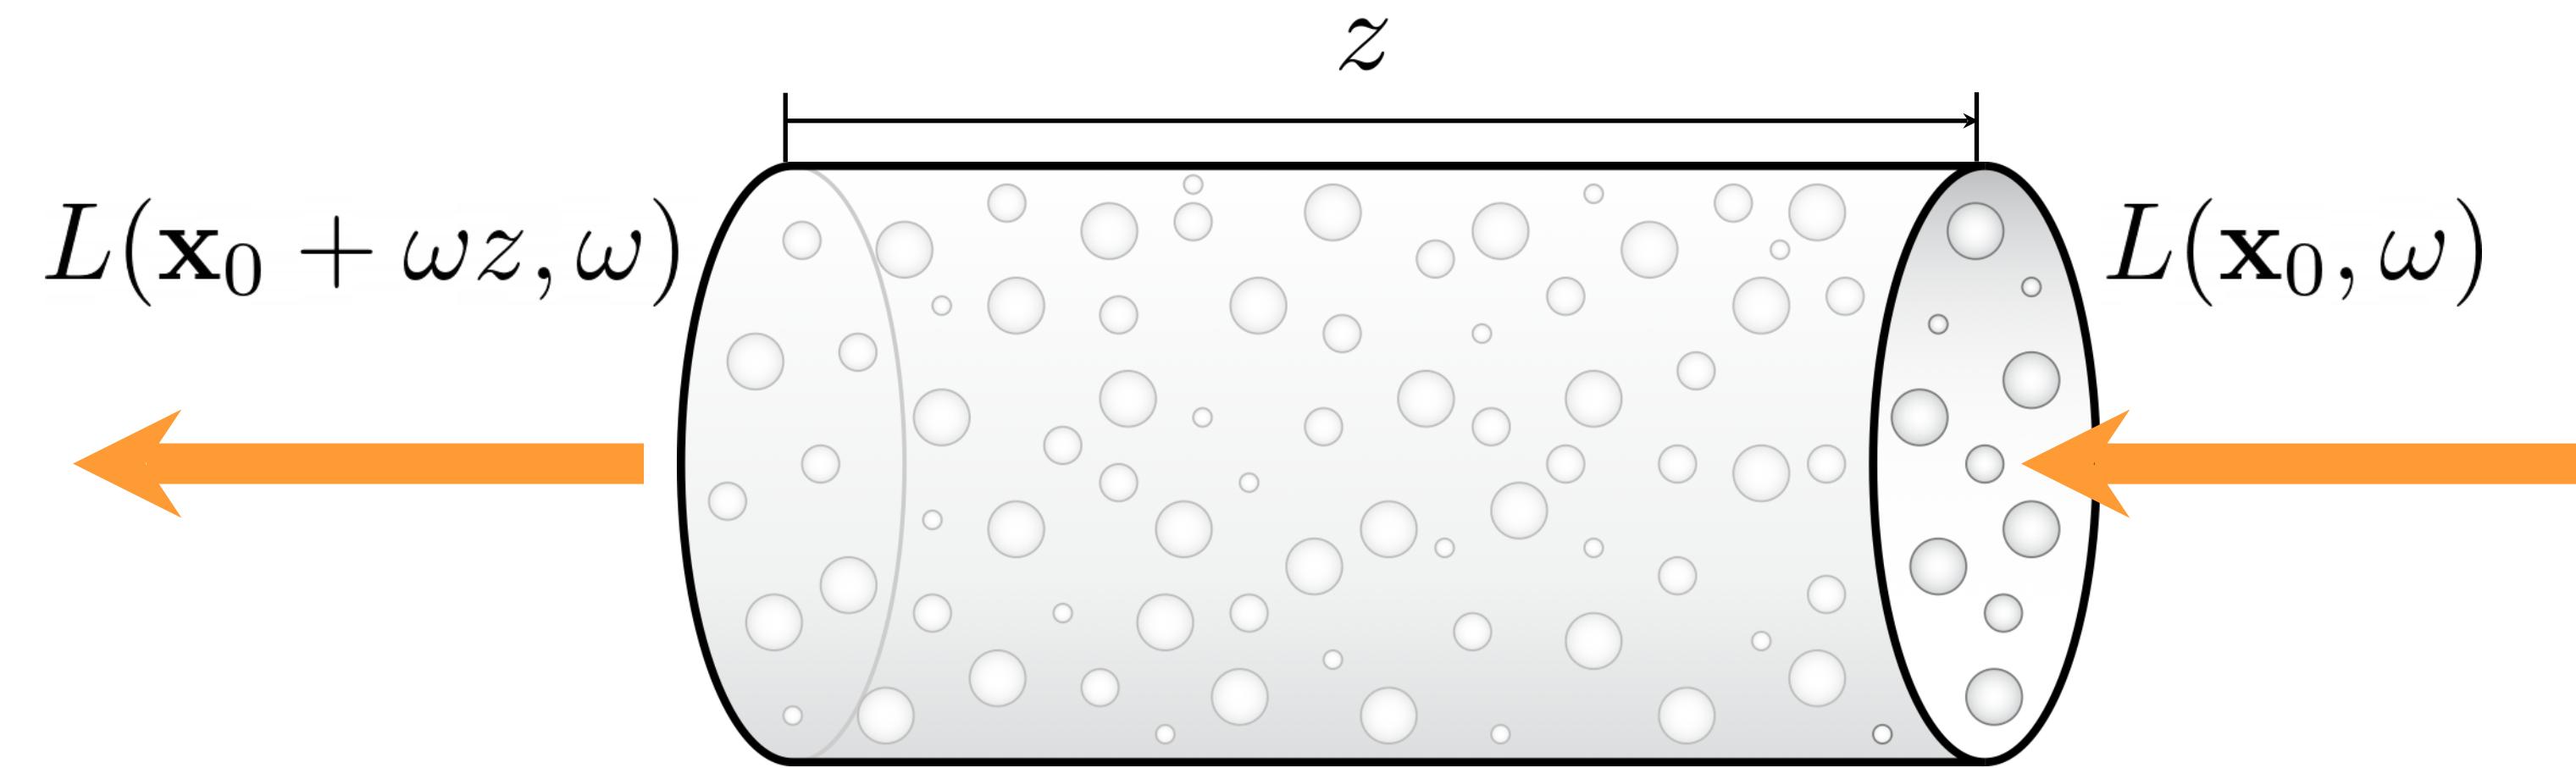
Multiplicativity:

$$T(\mathbf{x}, \mathbf{y}) = T(\mathbf{x}, \mathbf{z})T(\mathbf{z}, \mathbf{y})$$

Why?



Absorption-only Volume Rendering

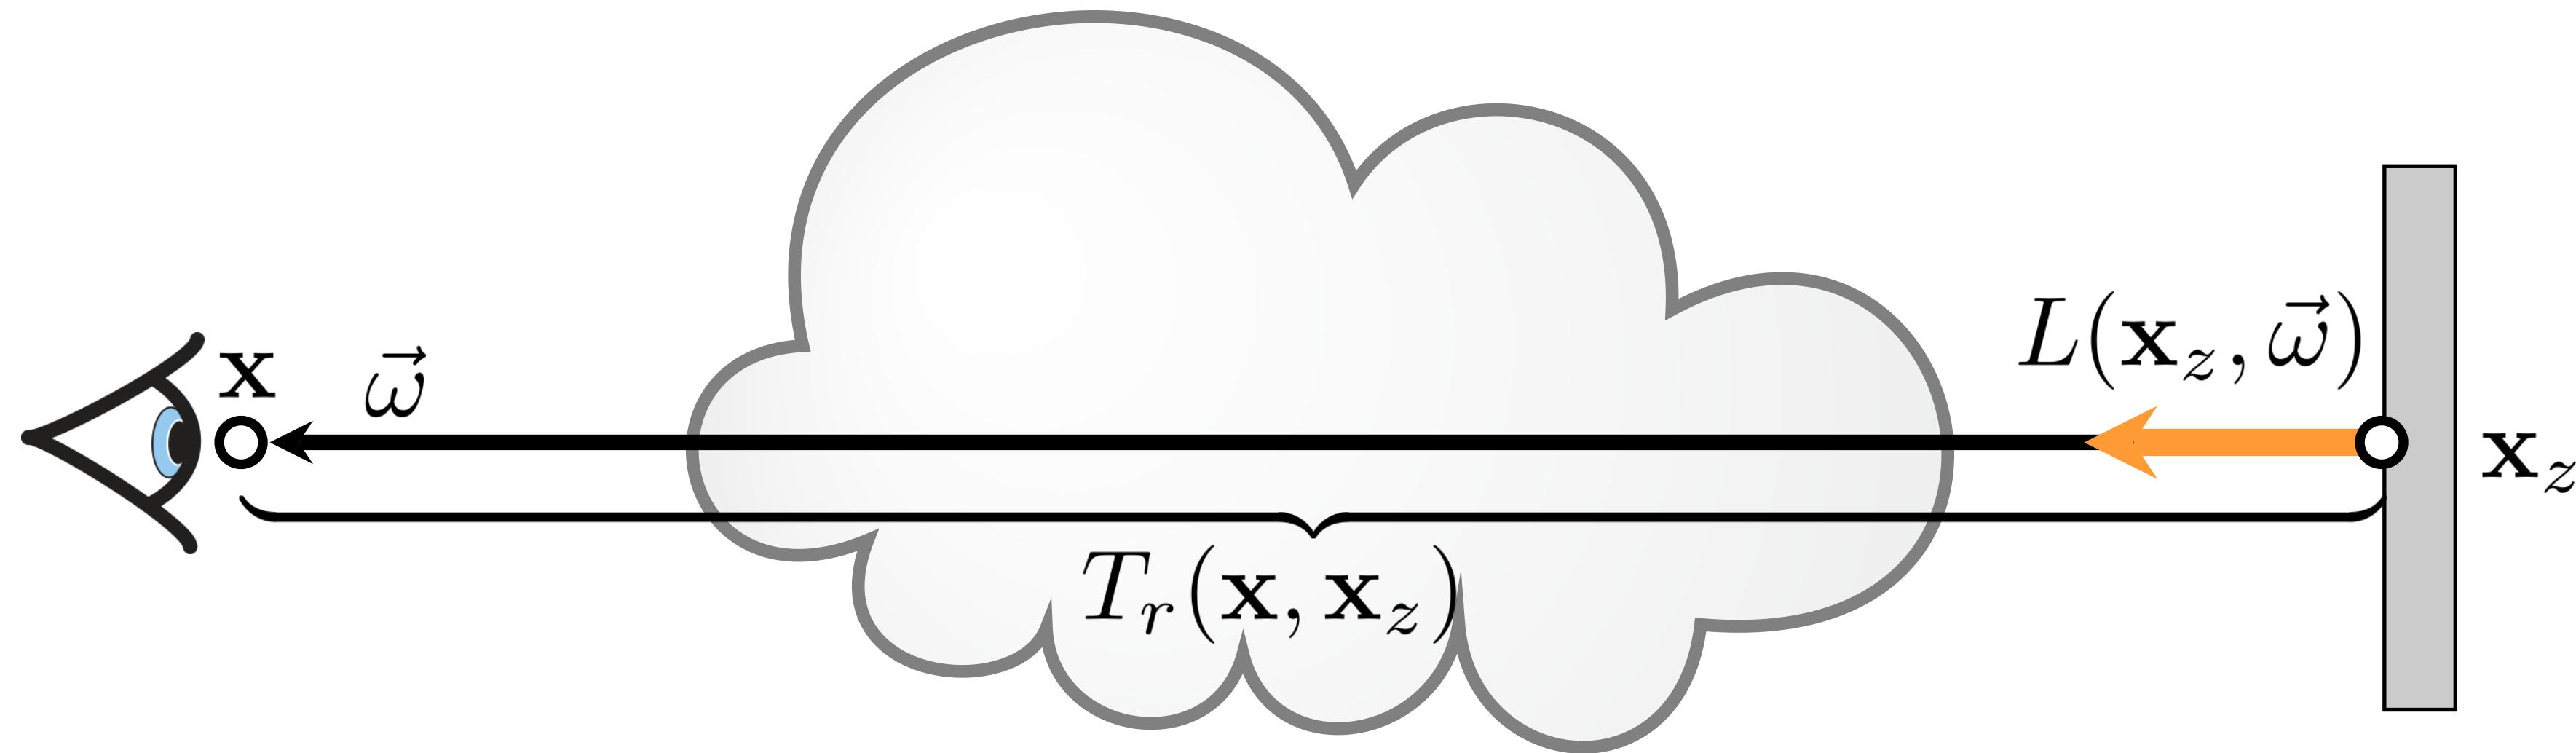


$$L(\mathbf{x}_0 + \omega z, \omega) = T(\mathbf{x}_0, \mathbf{x}_0 + \omega z)L(\mathbf{x}_0, \omega)$$

Absorption-only Volume Rendering



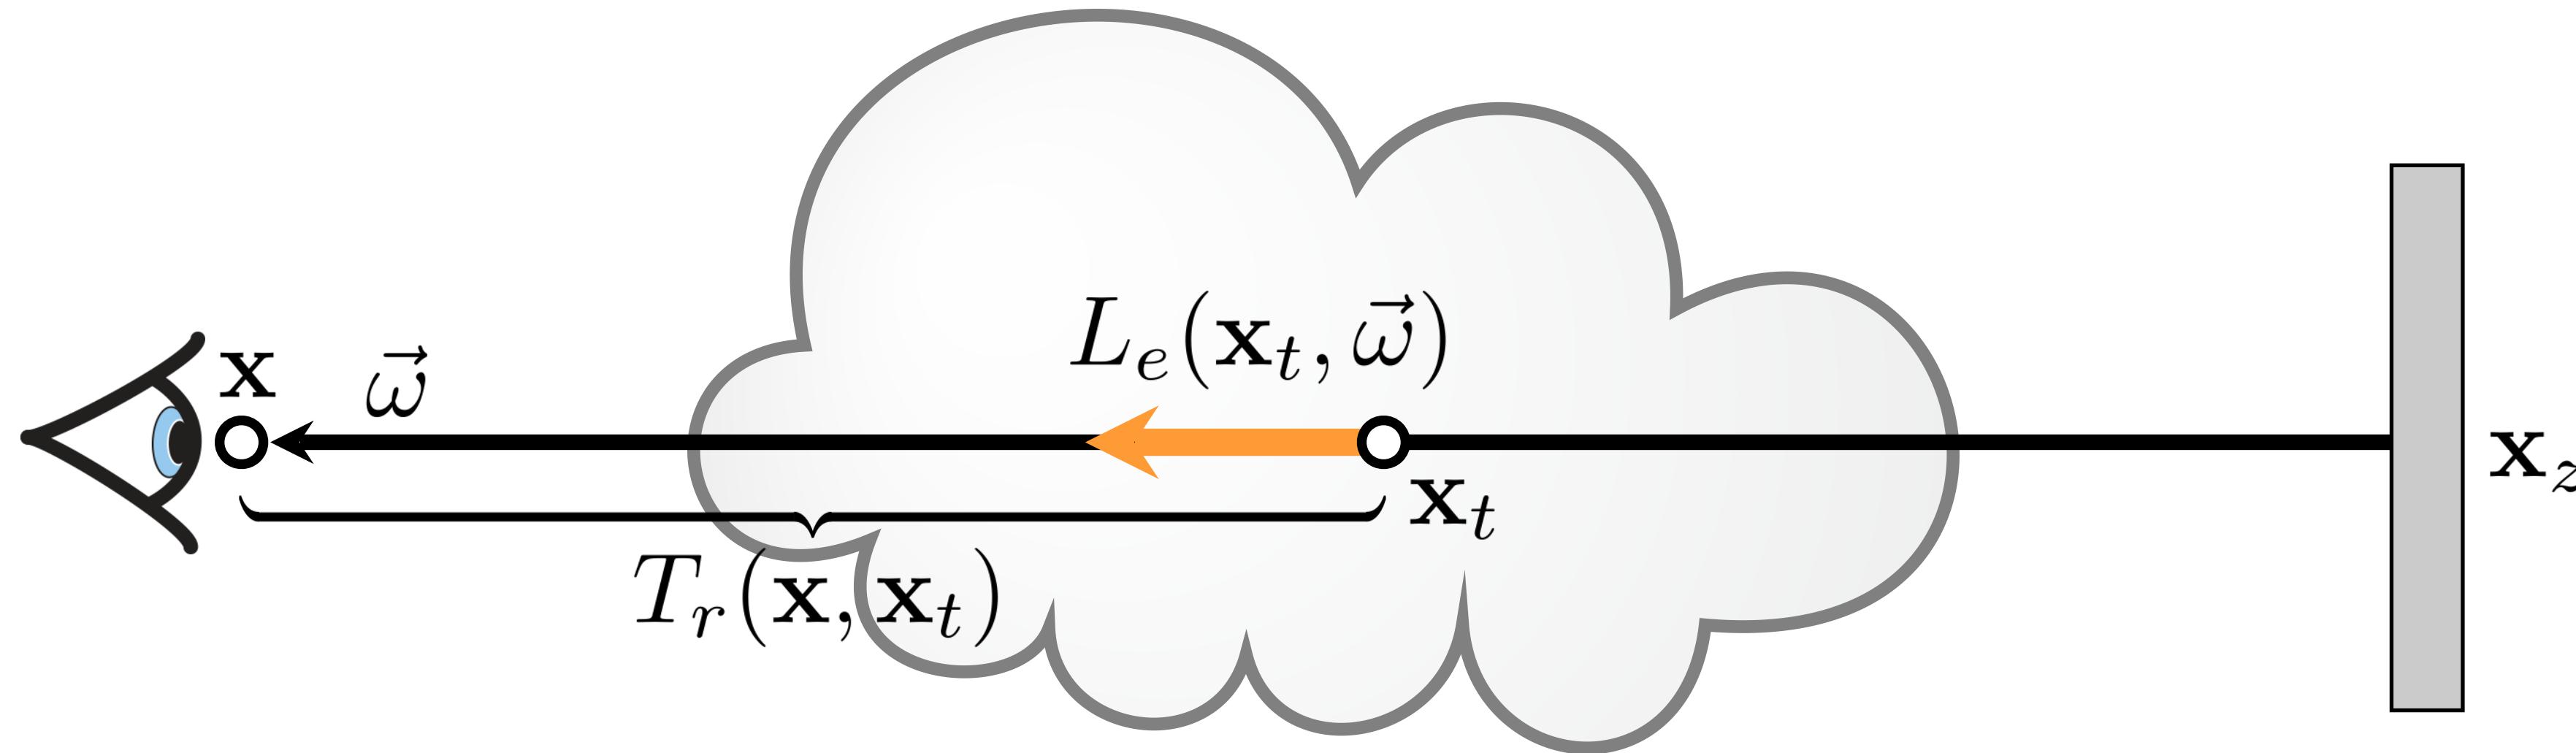
Absorption-only Volume Rendering



$$L(\mathbf{x}, \omega) = T(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \omega)$$

Radiance from ‘outside’ the medium

Emission-Absorption Volume Rendering

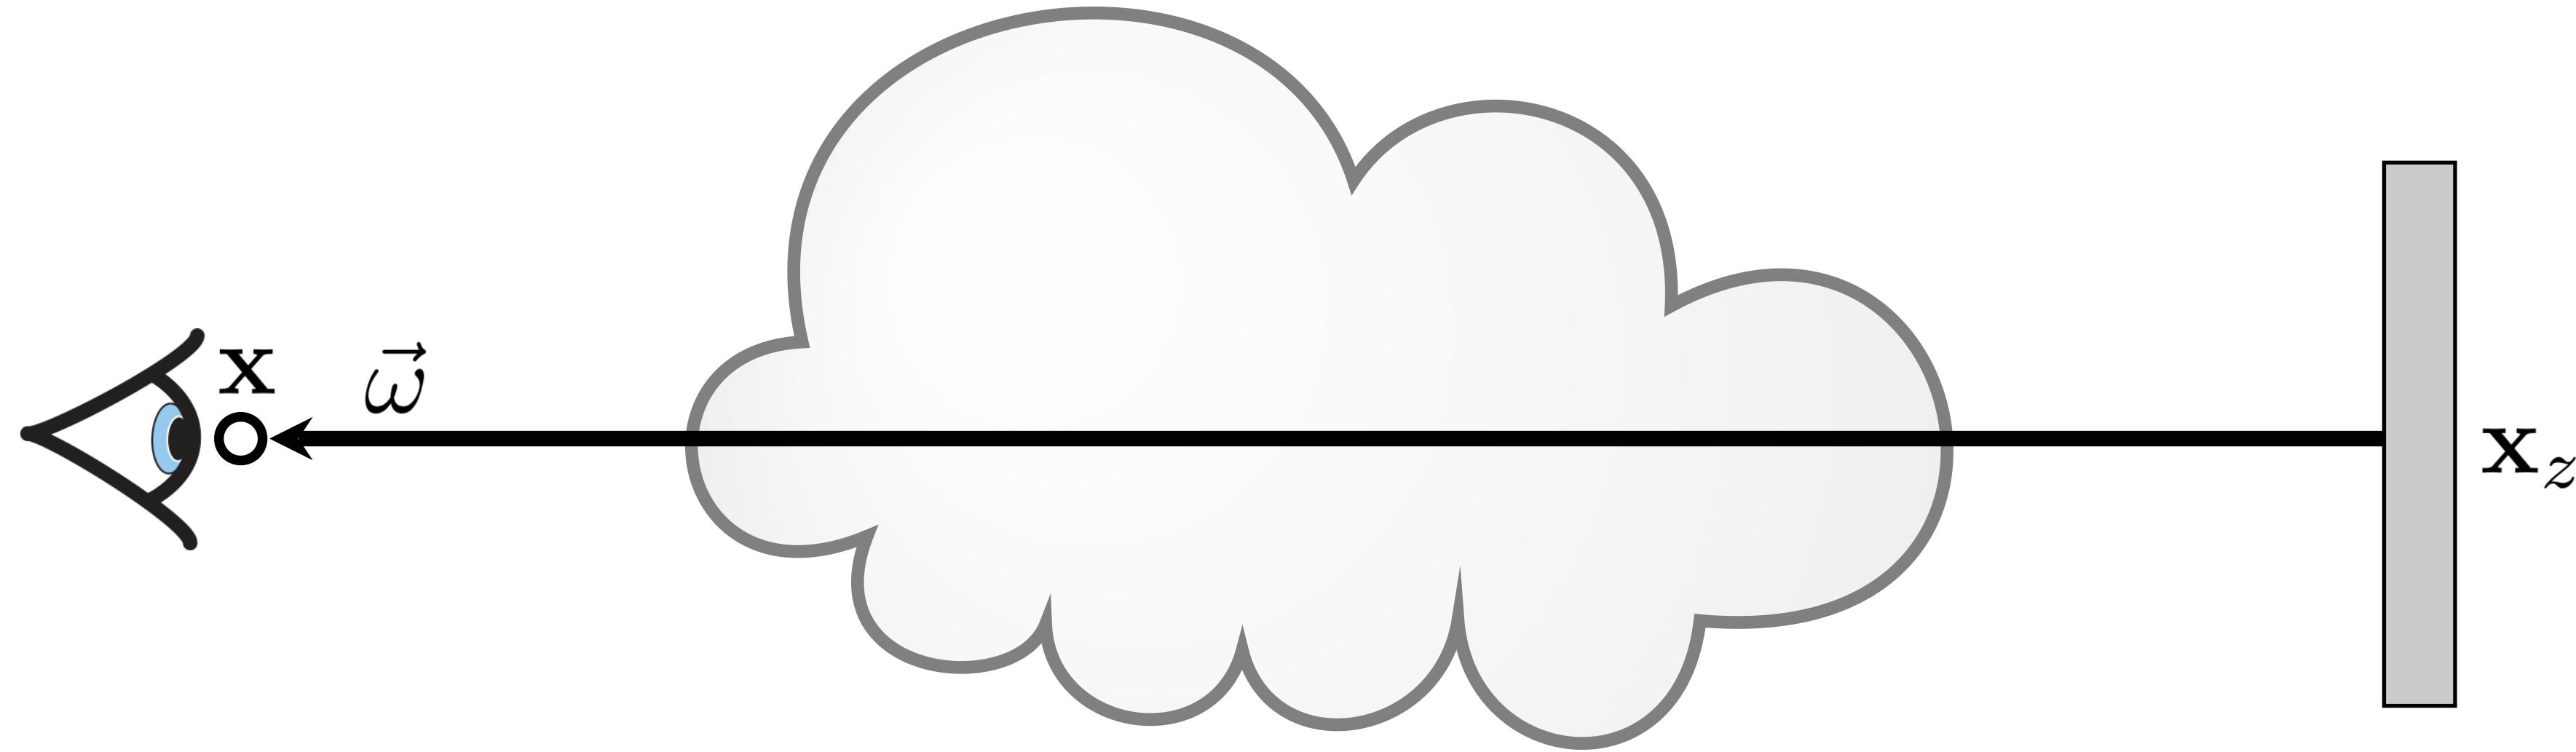


$$L(\mathbf{x}, \omega) = T(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \omega)$$

$$+ \int_0^z T(\mathbf{x}, \mathbf{x}_t) \sigma(\mathbf{x}_t) \underline{L_e(\mathbf{x}_t, \omega)} dt$$

Accumulated Emitted Radiance from inside

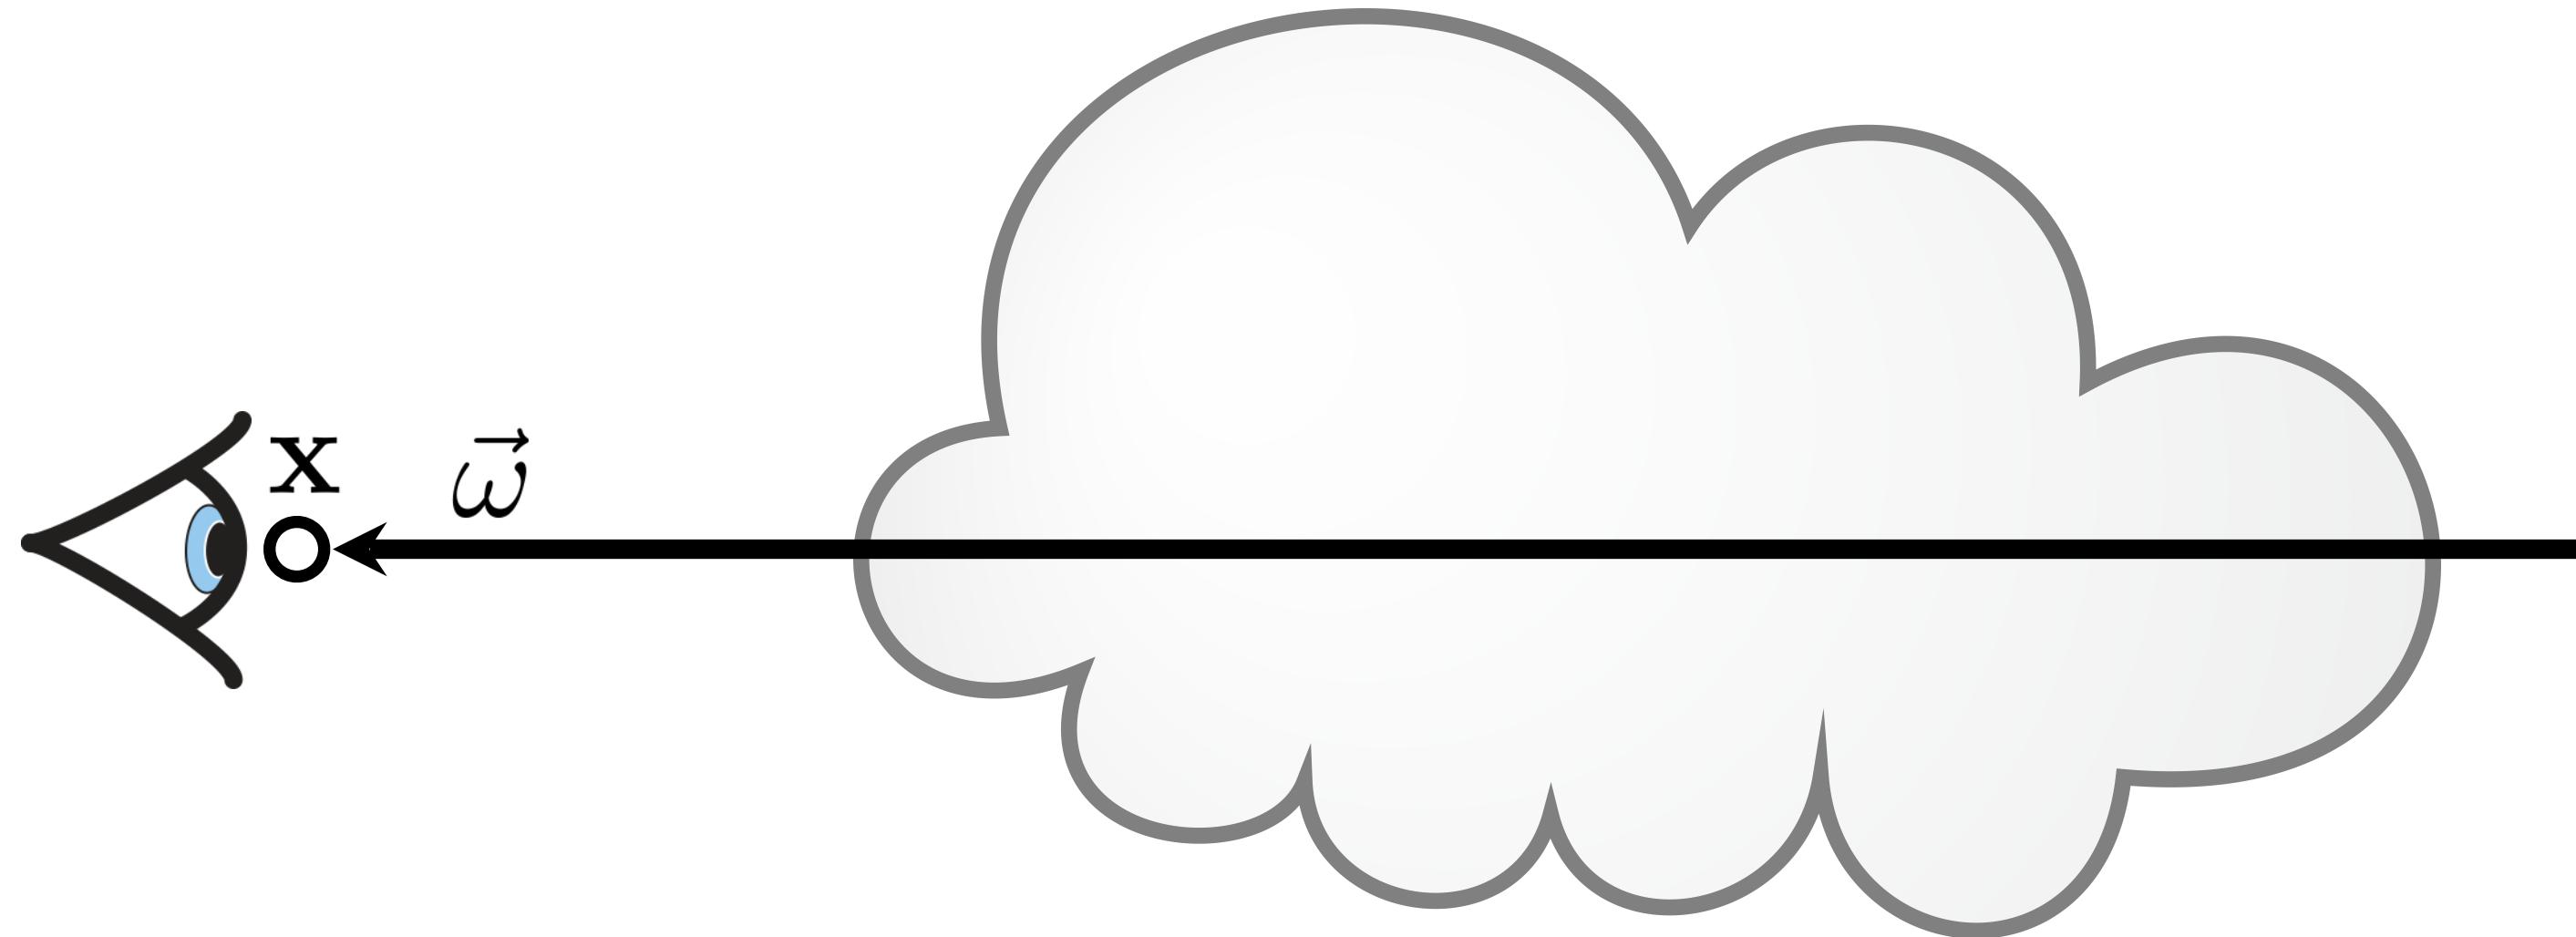
Emission-Absorption Volume Rendering



$$L(\mathbf{x}, \omega) = T(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \omega) + \int_0^z T(\mathbf{x}, \mathbf{x}_t)\sigma(\mathbf{x}_t)L_e(\mathbf{x}_t, \omega)dt$$

Can we compute this analytically? Approximately?

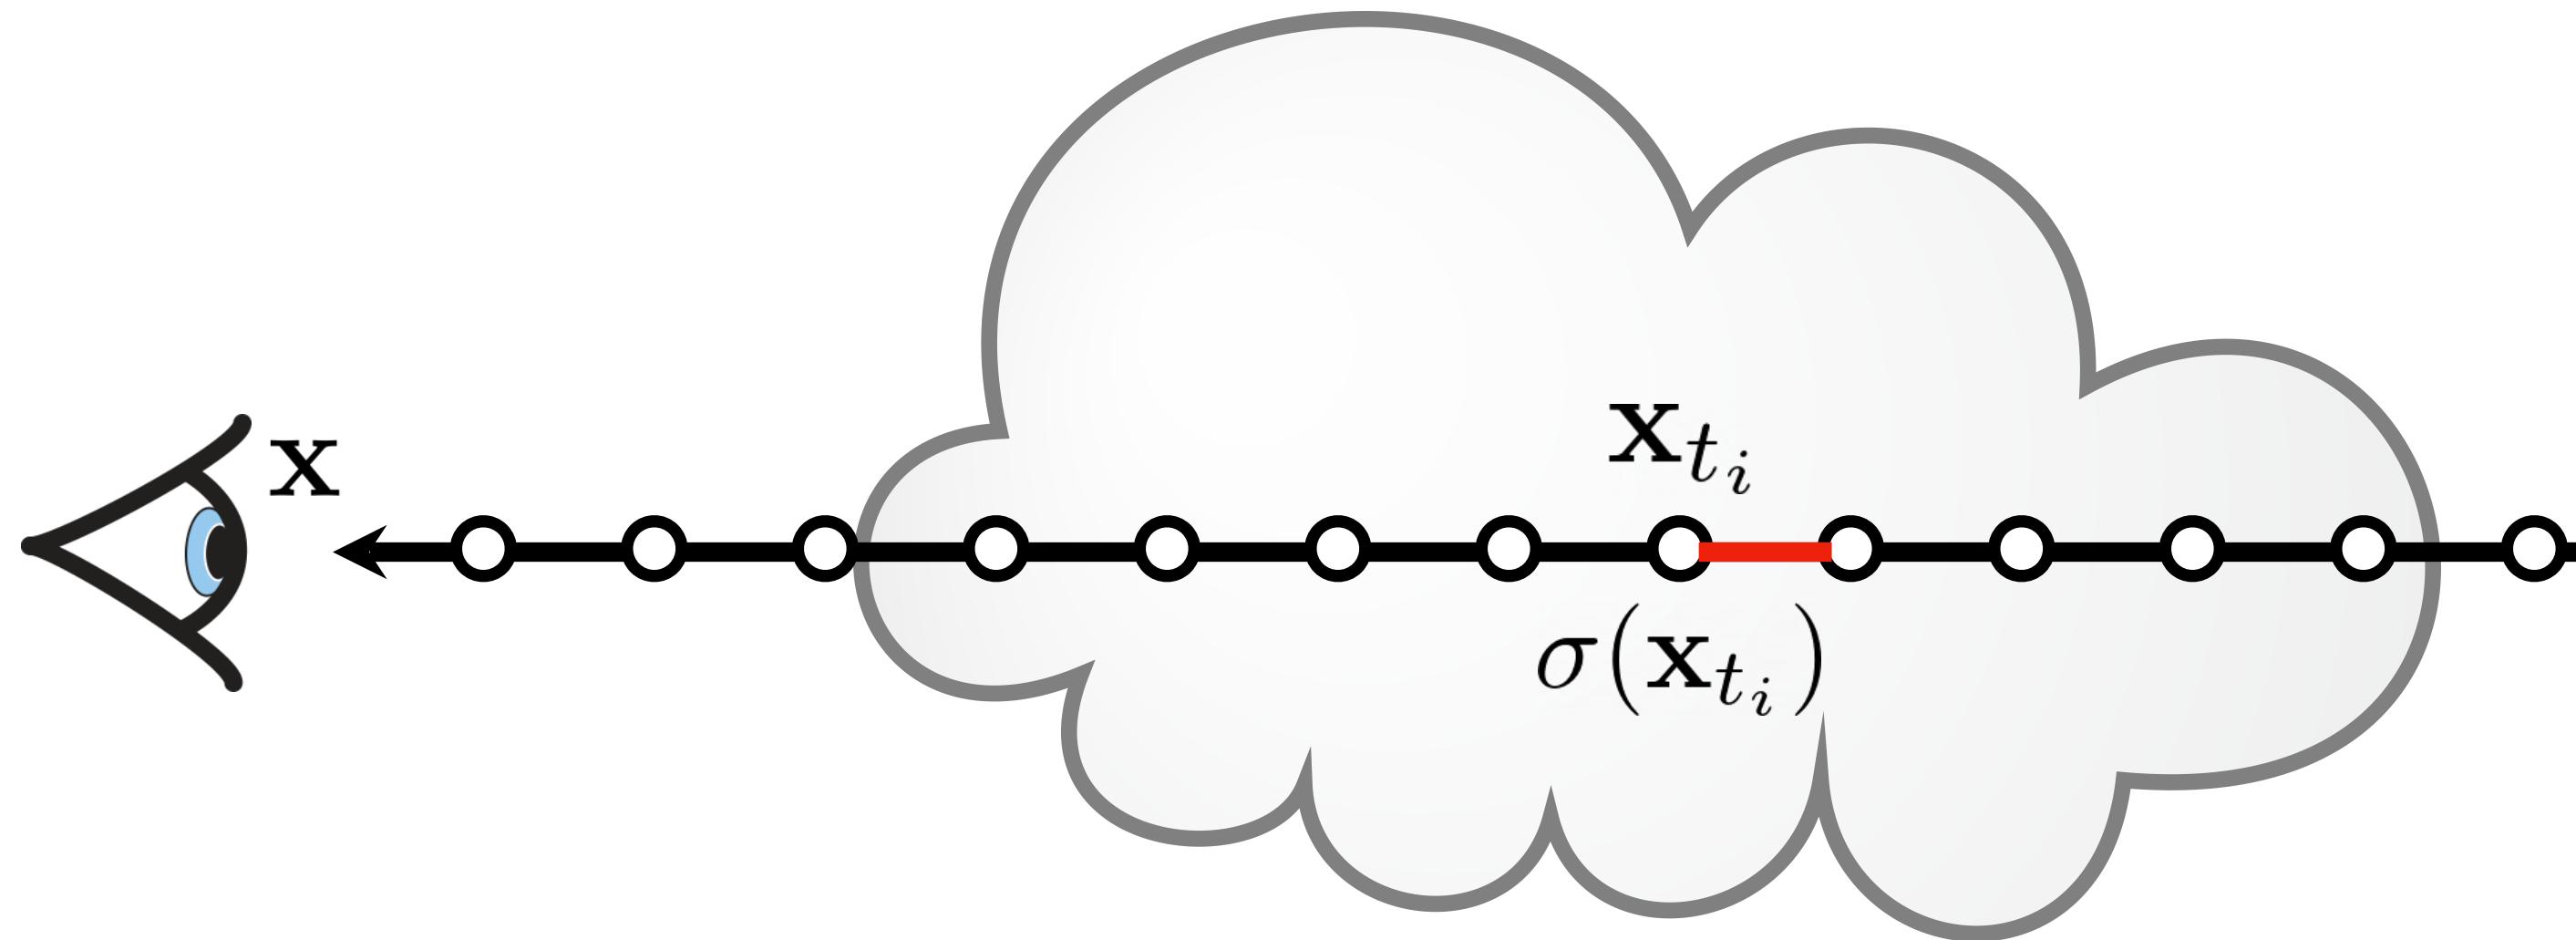
Computational Volume Rendering



$$L(\mathbf{x}, \omega) = \int_0^z T(\mathbf{x}, \mathbf{x}_t) \sigma(\mathbf{x}_t) L_e(\mathbf{x}_t, \omega) dt$$

Let's ignore light from outside medium (for now)

Computational Volume Rendering: Ray Marching



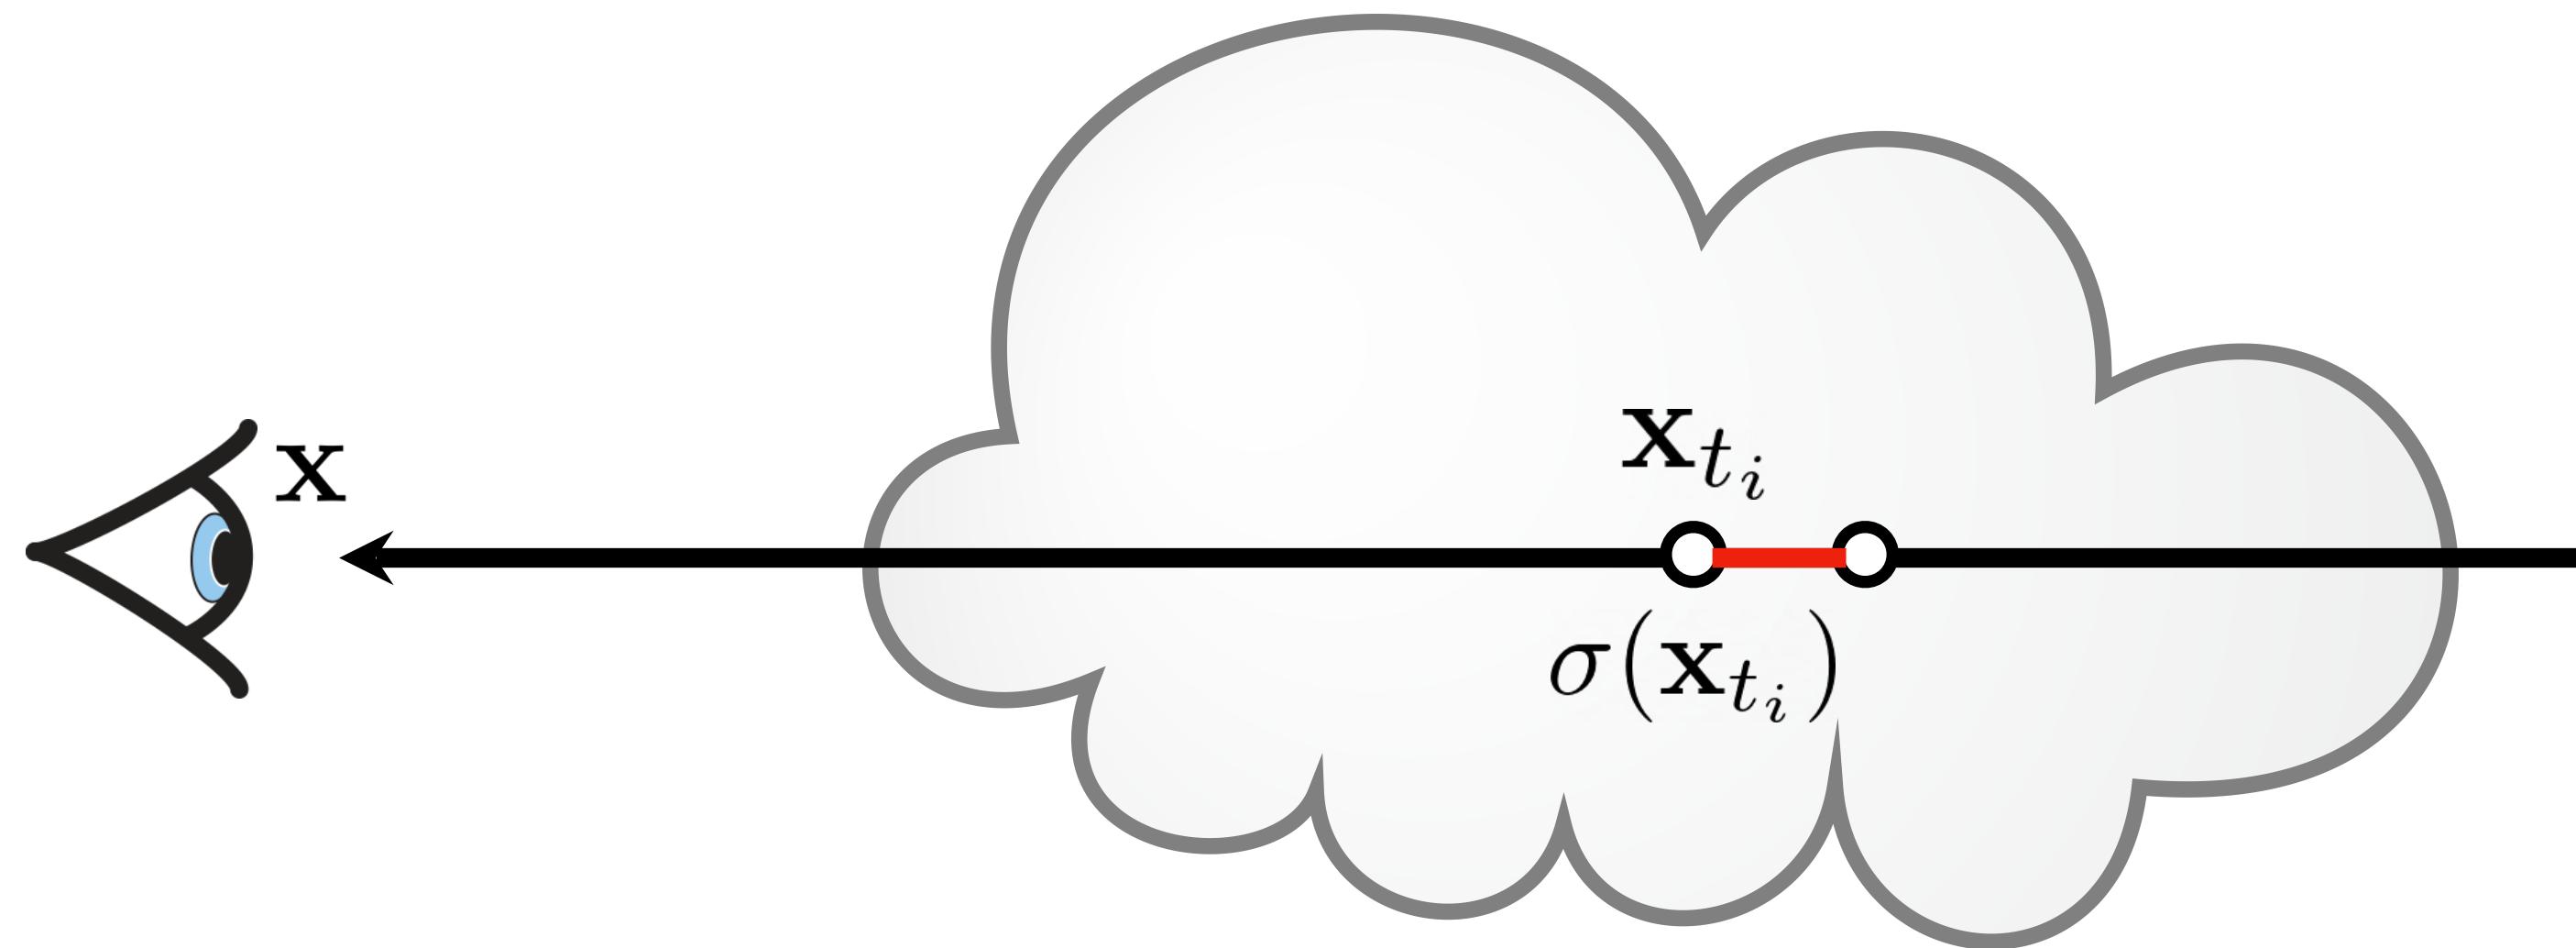
$$L(\mathbf{x}, \omega) = \sum_{i=1}^N \text{(contribution from } i^{\text{th}} \text{ segment)}$$

Approximate with a discrete sum

\mathbf{x}_{t_i} : i^{th} sample along ray at depth t_i

Δt : distance between successive samples

Computational Volume Rendering: Ray Marching

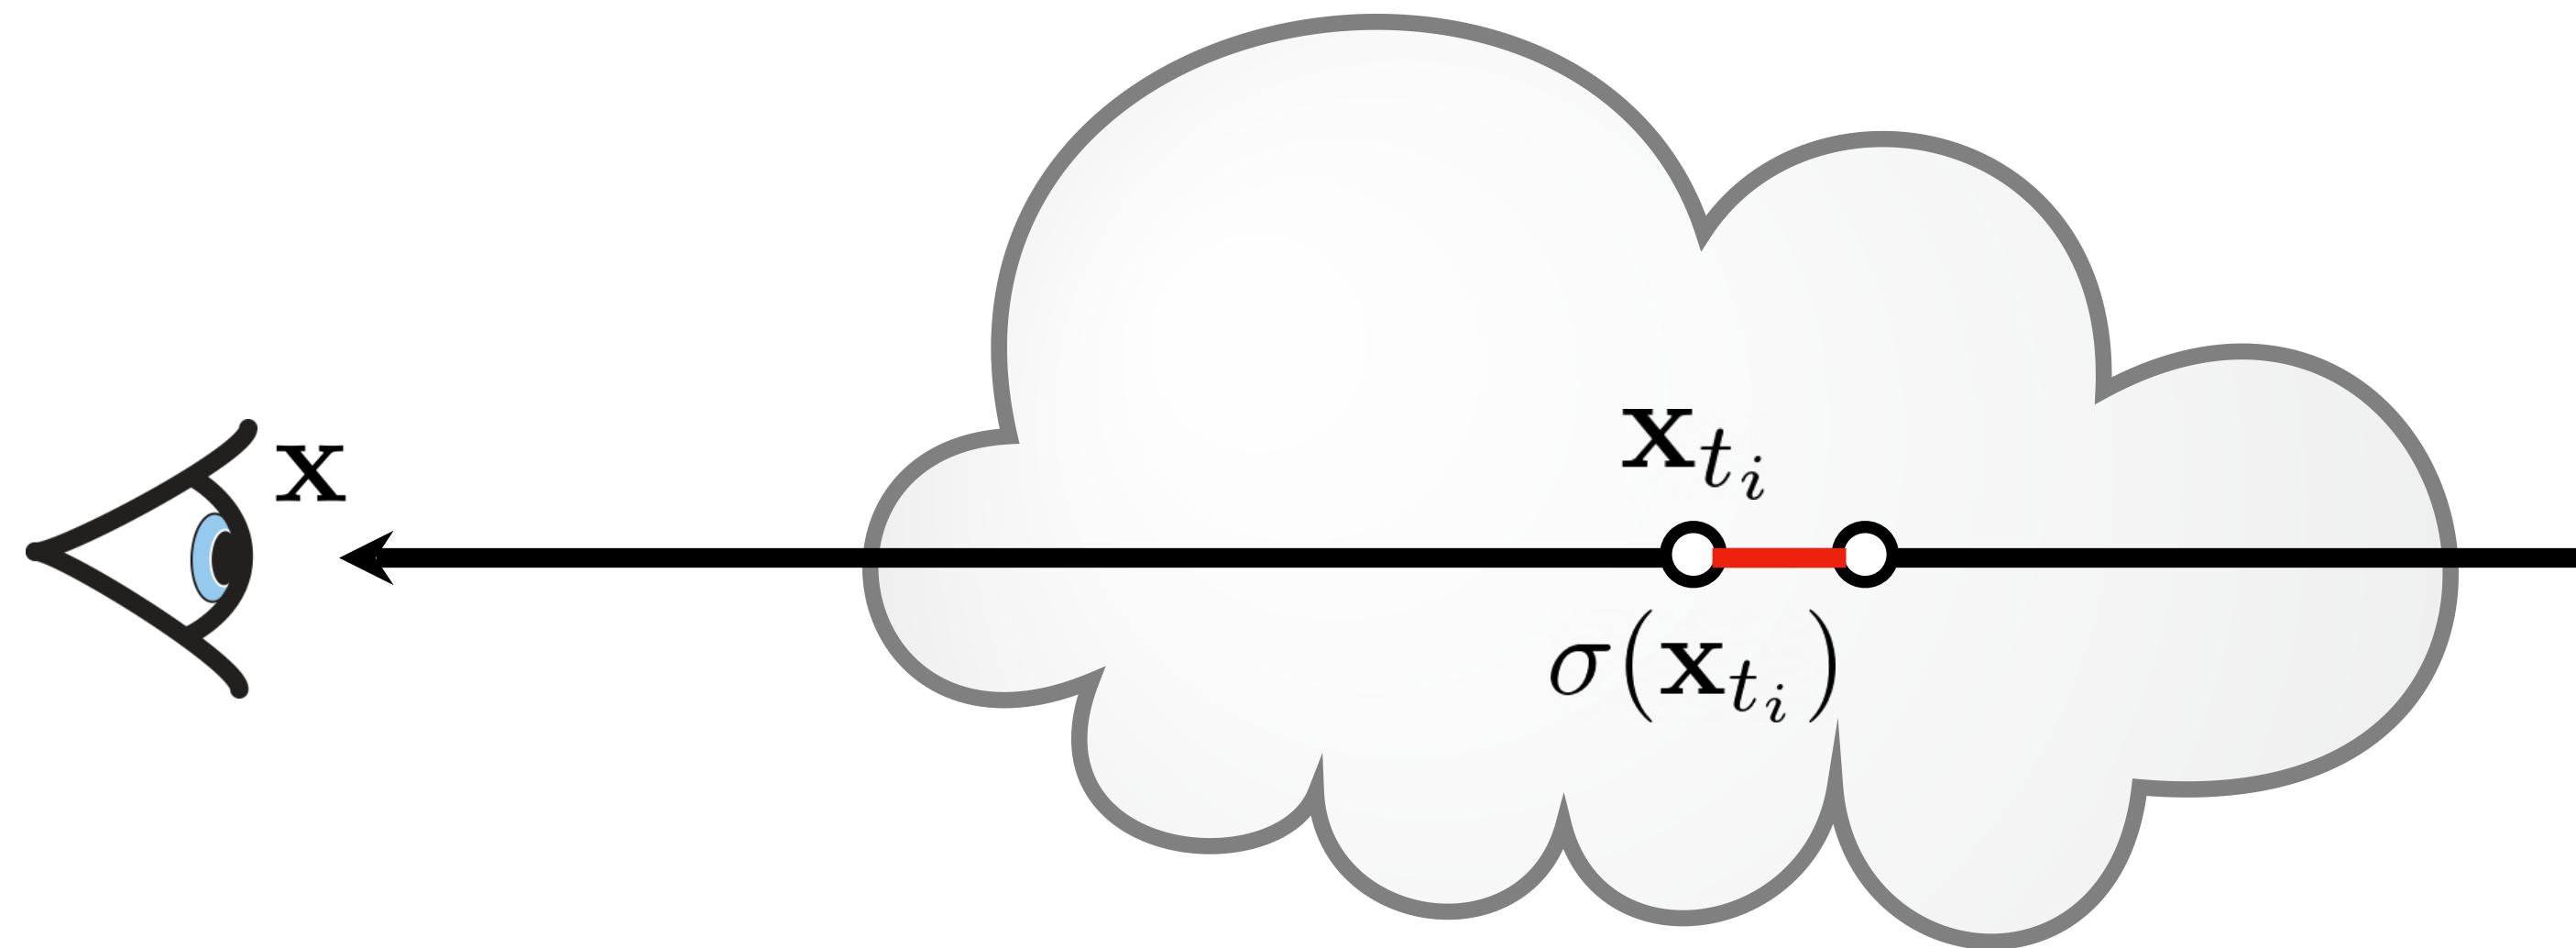


$$L(\mathbf{x}, \omega) = \sum_{i=1}^N \text{(contribution from } i^{\text{th}} \text{ segment)}$$

\mathbf{x}_{t_i} : i^{th} sample along ray at depth t_i

Δt : distance between successive samples

Computational Volume Rendering: Ray Marching

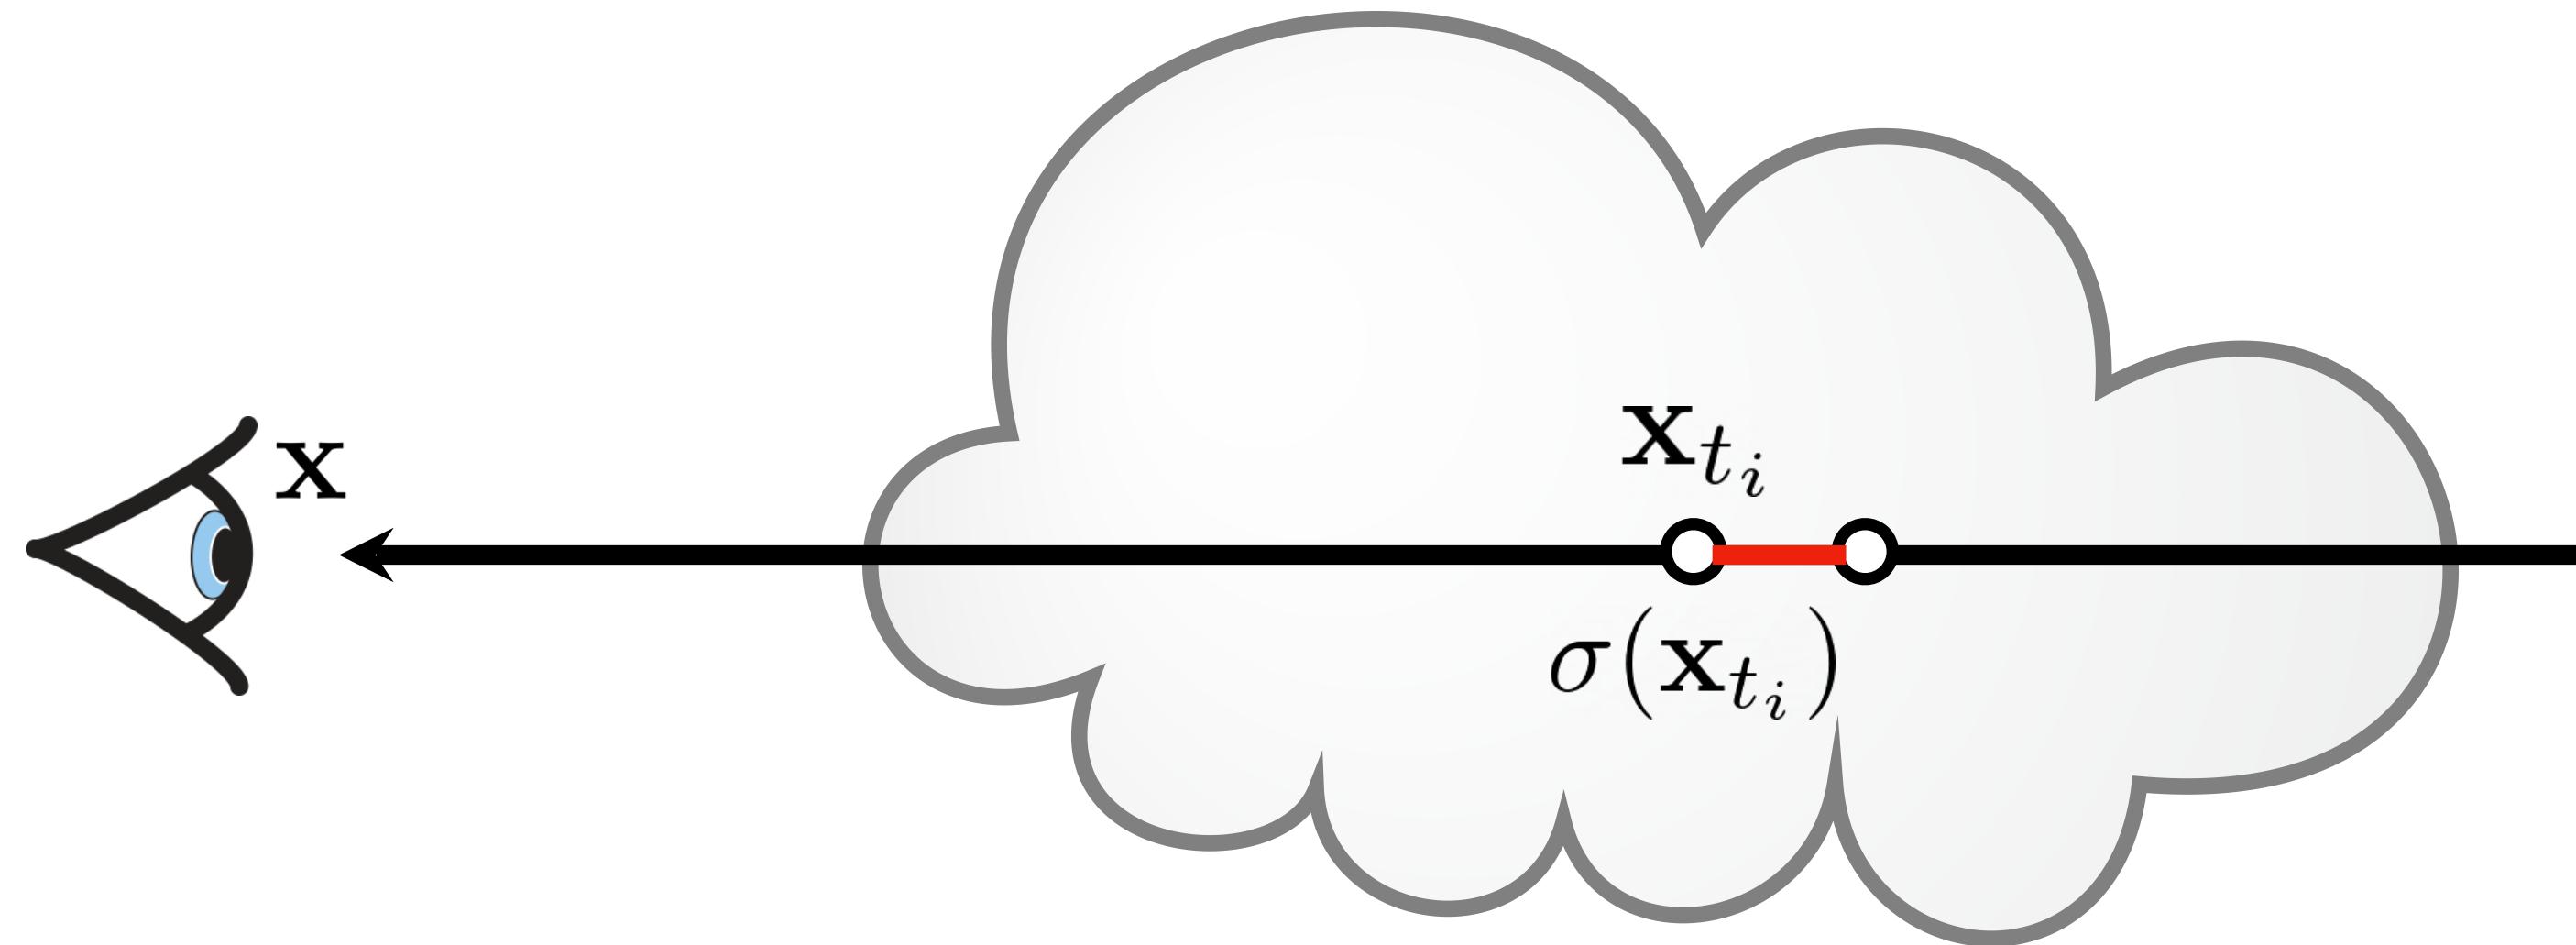


$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) \cdot x \text{ (emission from } i^{\text{th}} \text{ segment)}$$

\mathbf{x}_{t_i} : i^{th} sample along ray at depth t_i

Δt : distance between successive samples

Computational Volume Rendering: Ray Marching



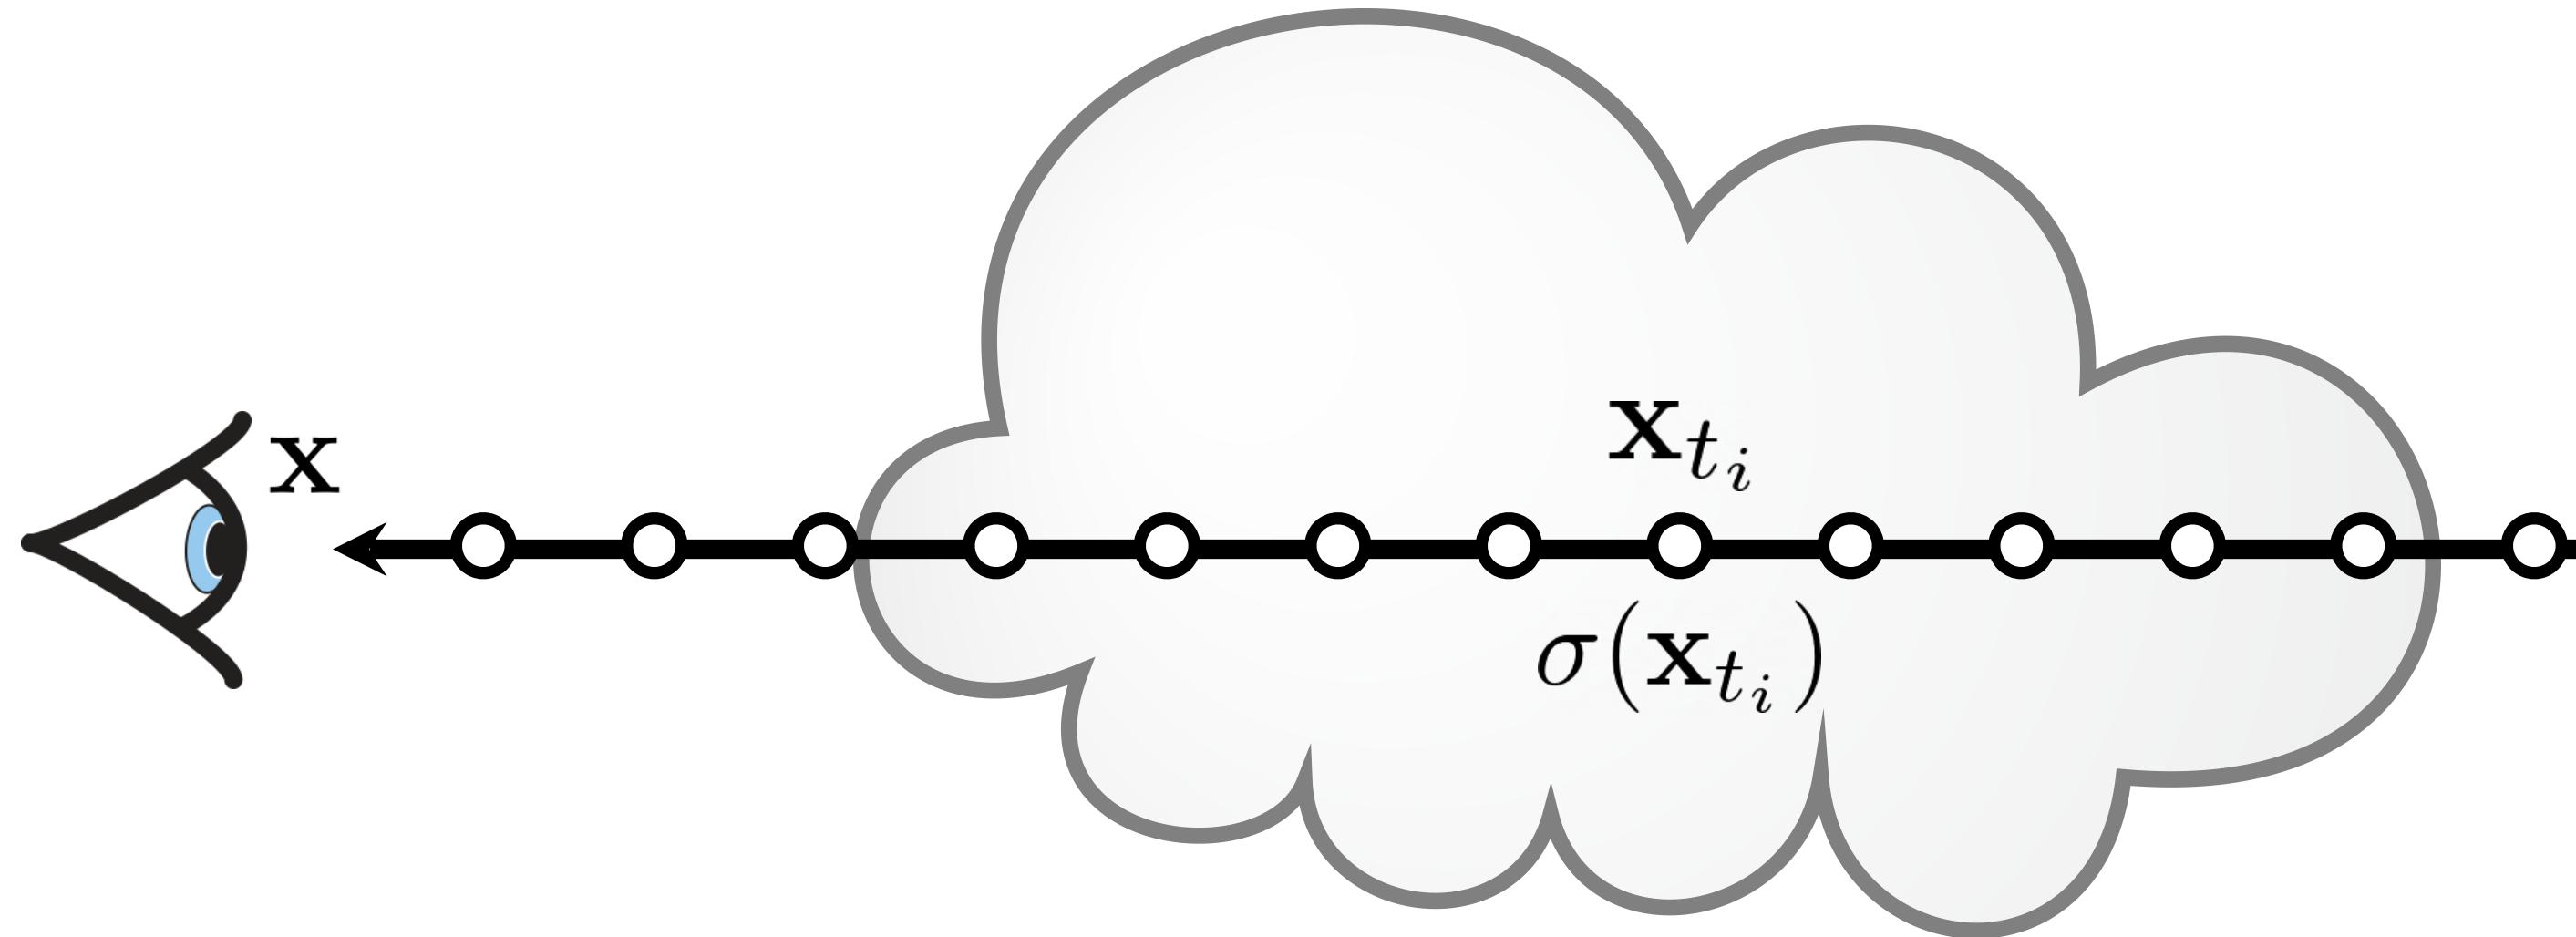
$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) \cdot (1 - e^{-\sigma_{t_i} \Delta t}) L_e(\mathbf{x}_{t_i}, \omega)$$

Requires a longer derivation which we'll skip

\mathbf{x}_{t_i} : i^{th} sample along ray at depth t_i

Δt : distance between successive samples

Computational Volume Rendering: Ray Marching



$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) \cdot (1 - e^{-\sigma_{t_i} \Delta t}) L_e(\mathbf{x}_{t_i}, \omega)$$

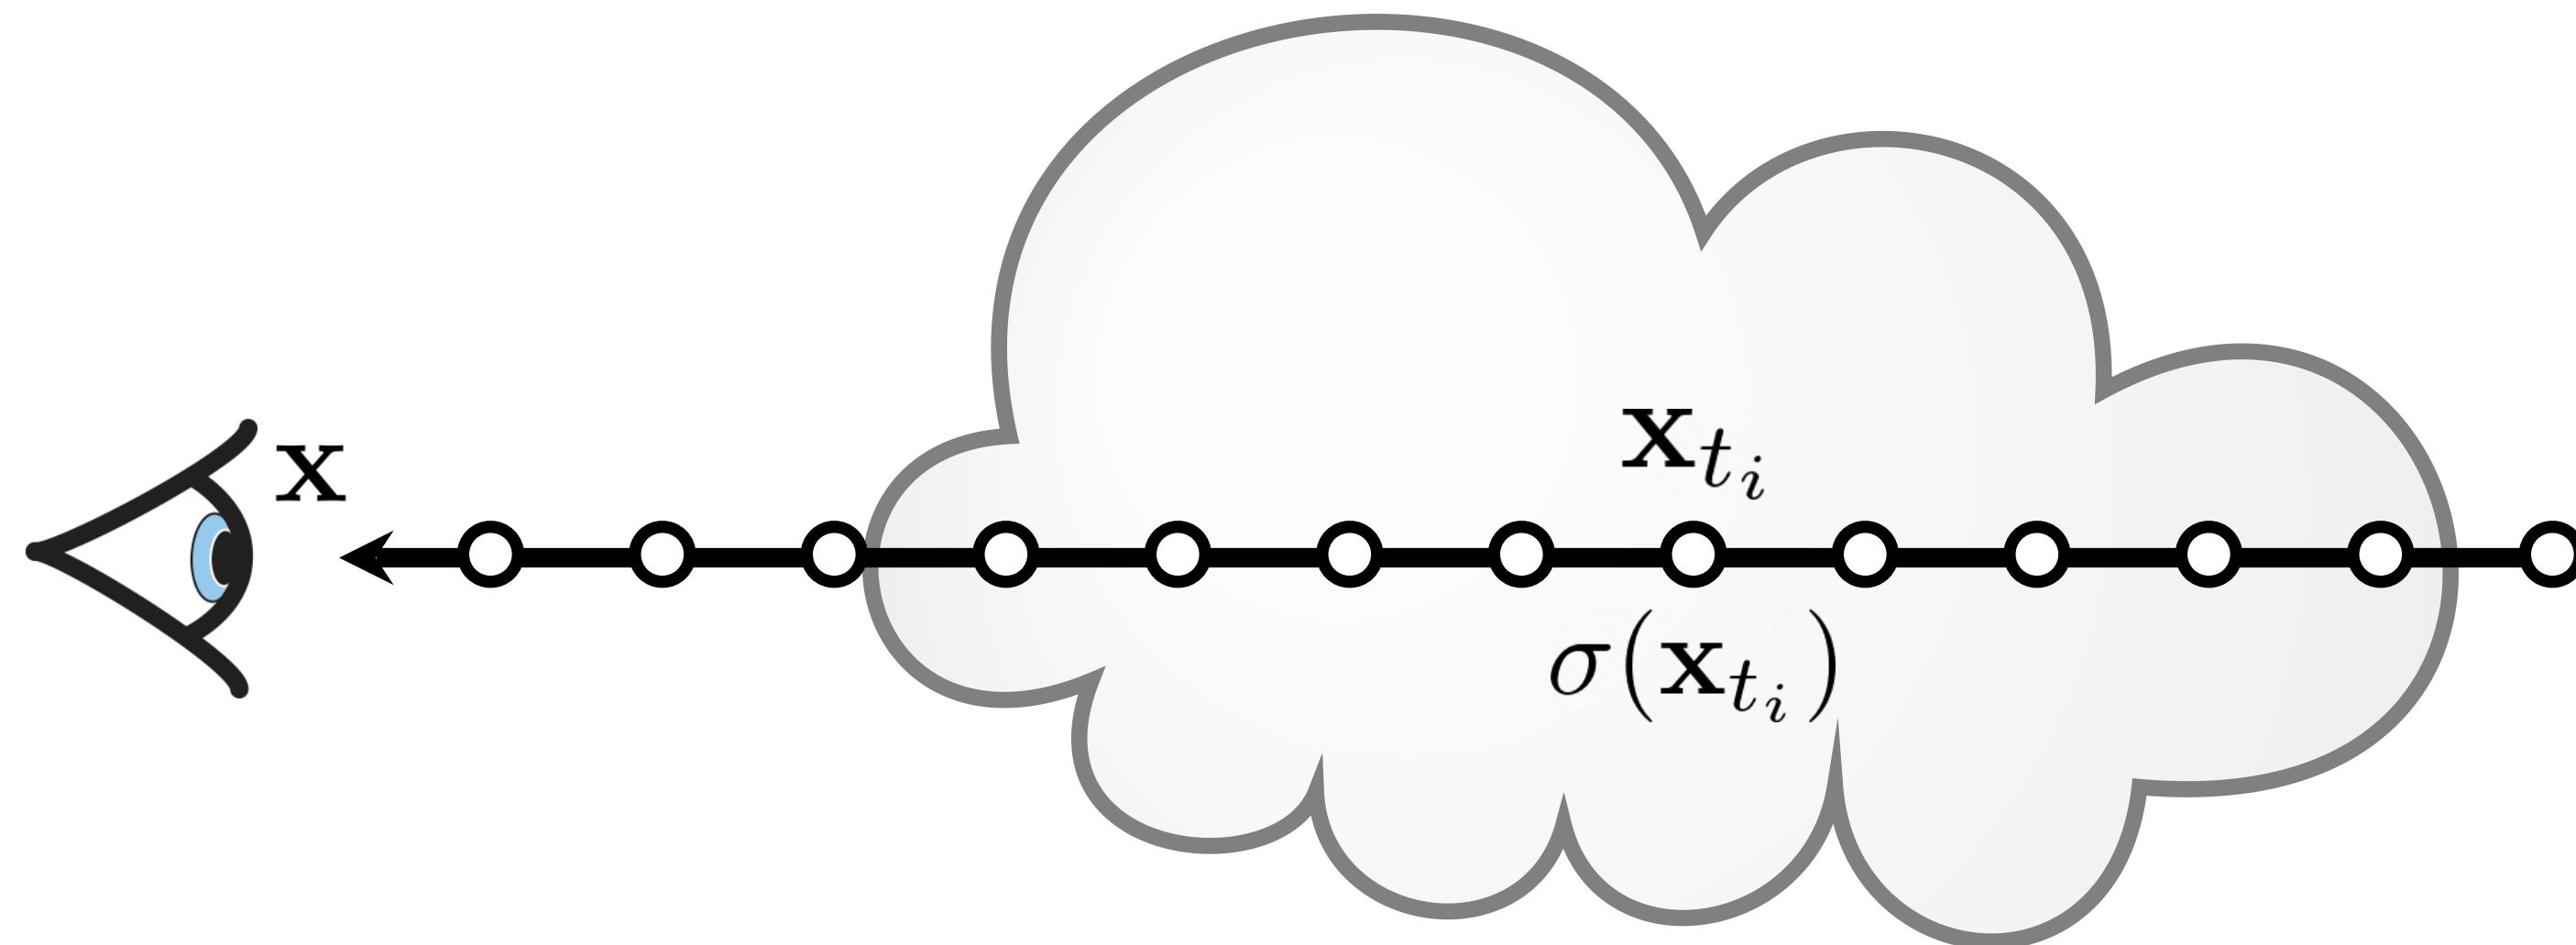
Multiplicativity

Base case: $T(\mathbf{x}, \mathbf{x}_{t_1}) = 1$

$$T(\mathbf{x}, \mathbf{x}_{t_i}) = T(\mathbf{x}, \mathbf{x}_{t_{i-1}}) e^{-\sigma_{t_{i-1}} \Delta t}$$

**Assume constant coefficient
between samples**

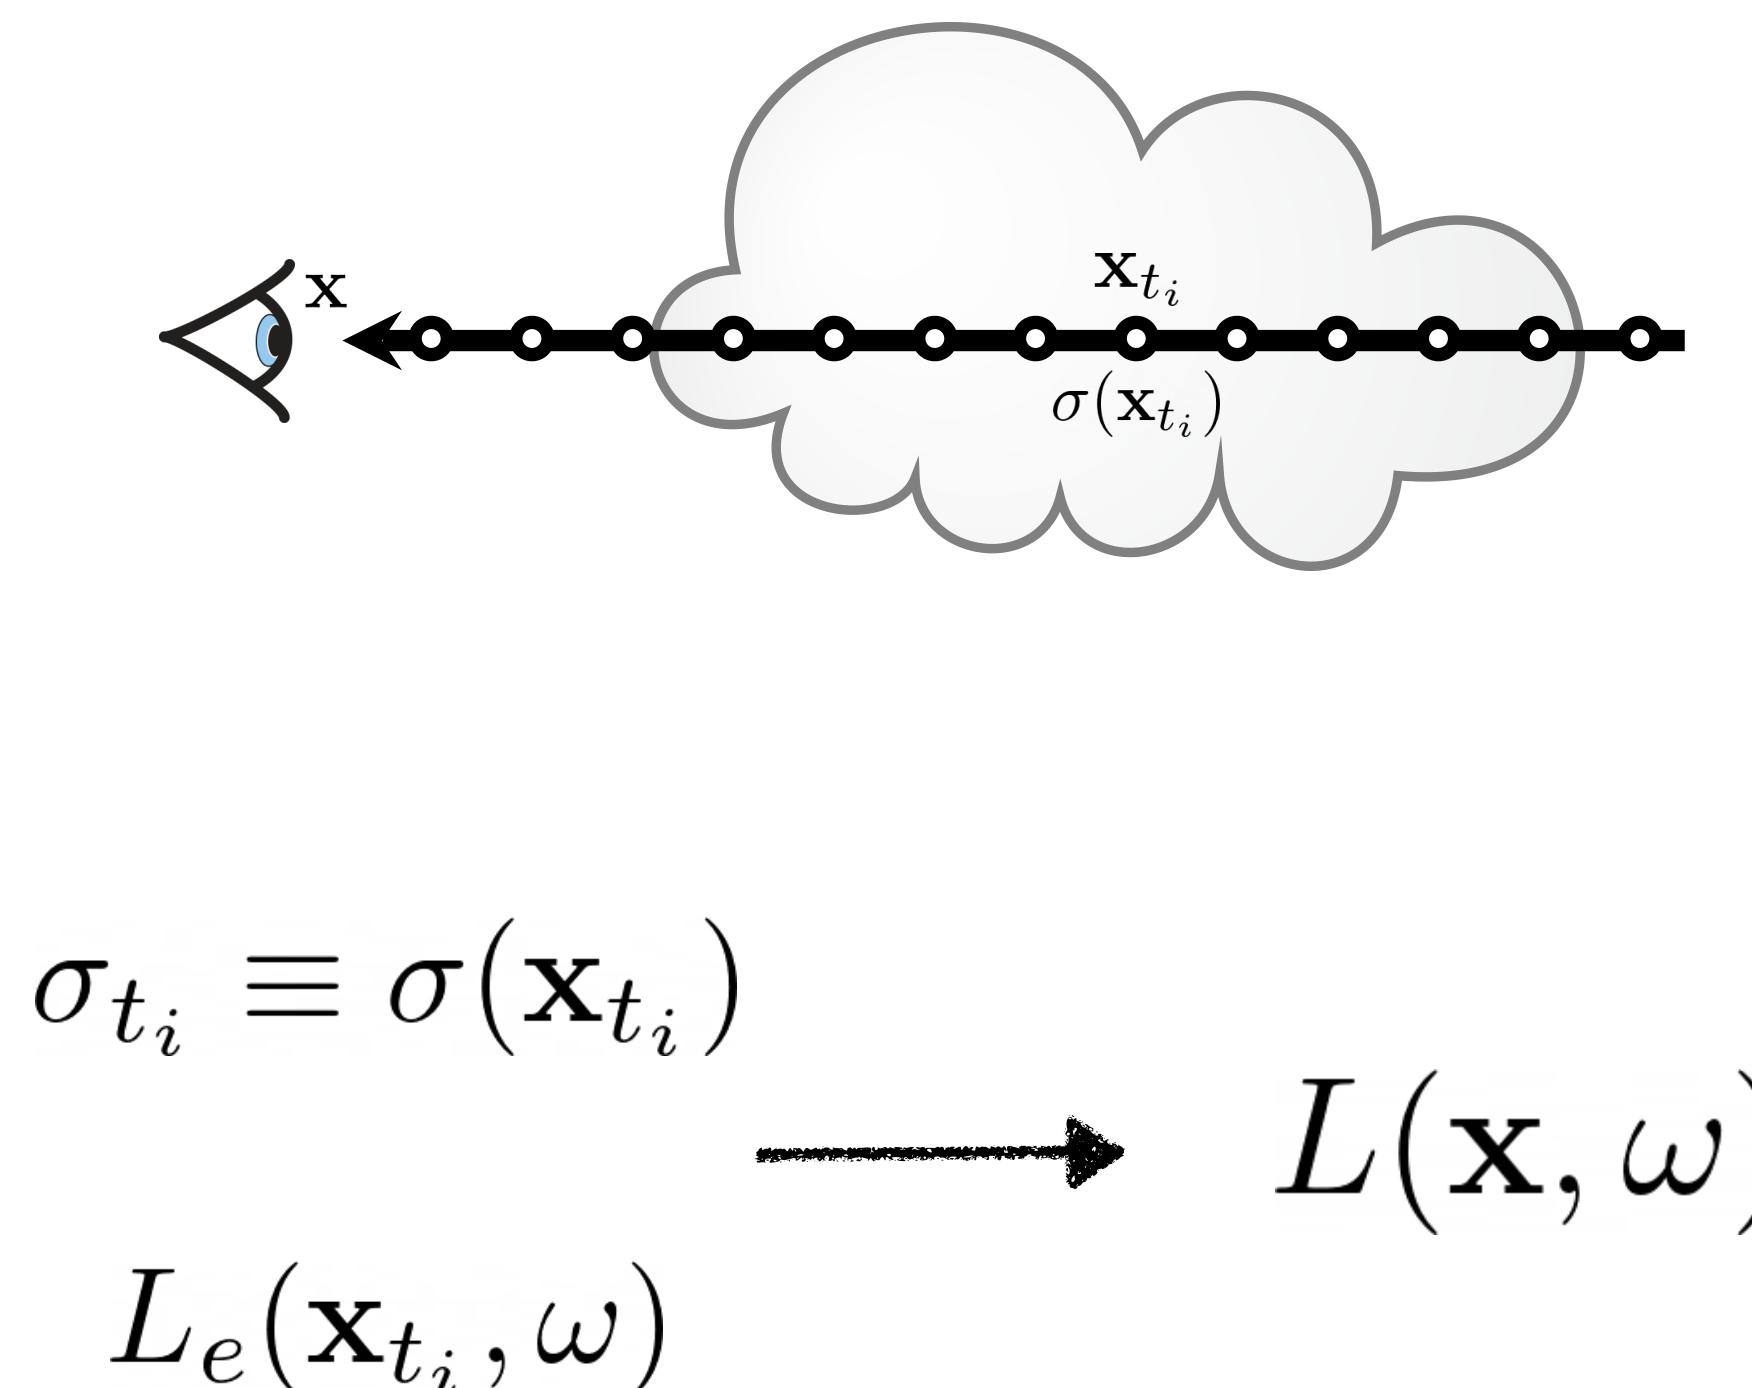
Computational Volume Rendering: Ray Marching



1. Draw uniform samples along a ray (N segments, or N+1 points)
2. Compute transmittance between camera and each sample
3. Aggregate contributions across segments to get overall radiance (color)

(minor modifications can allow non-uniform samples)

Computational Volume Rendering: A summary



If we can compute:

- a) (per-point) density
- b) (per-point, direction) emitted light,

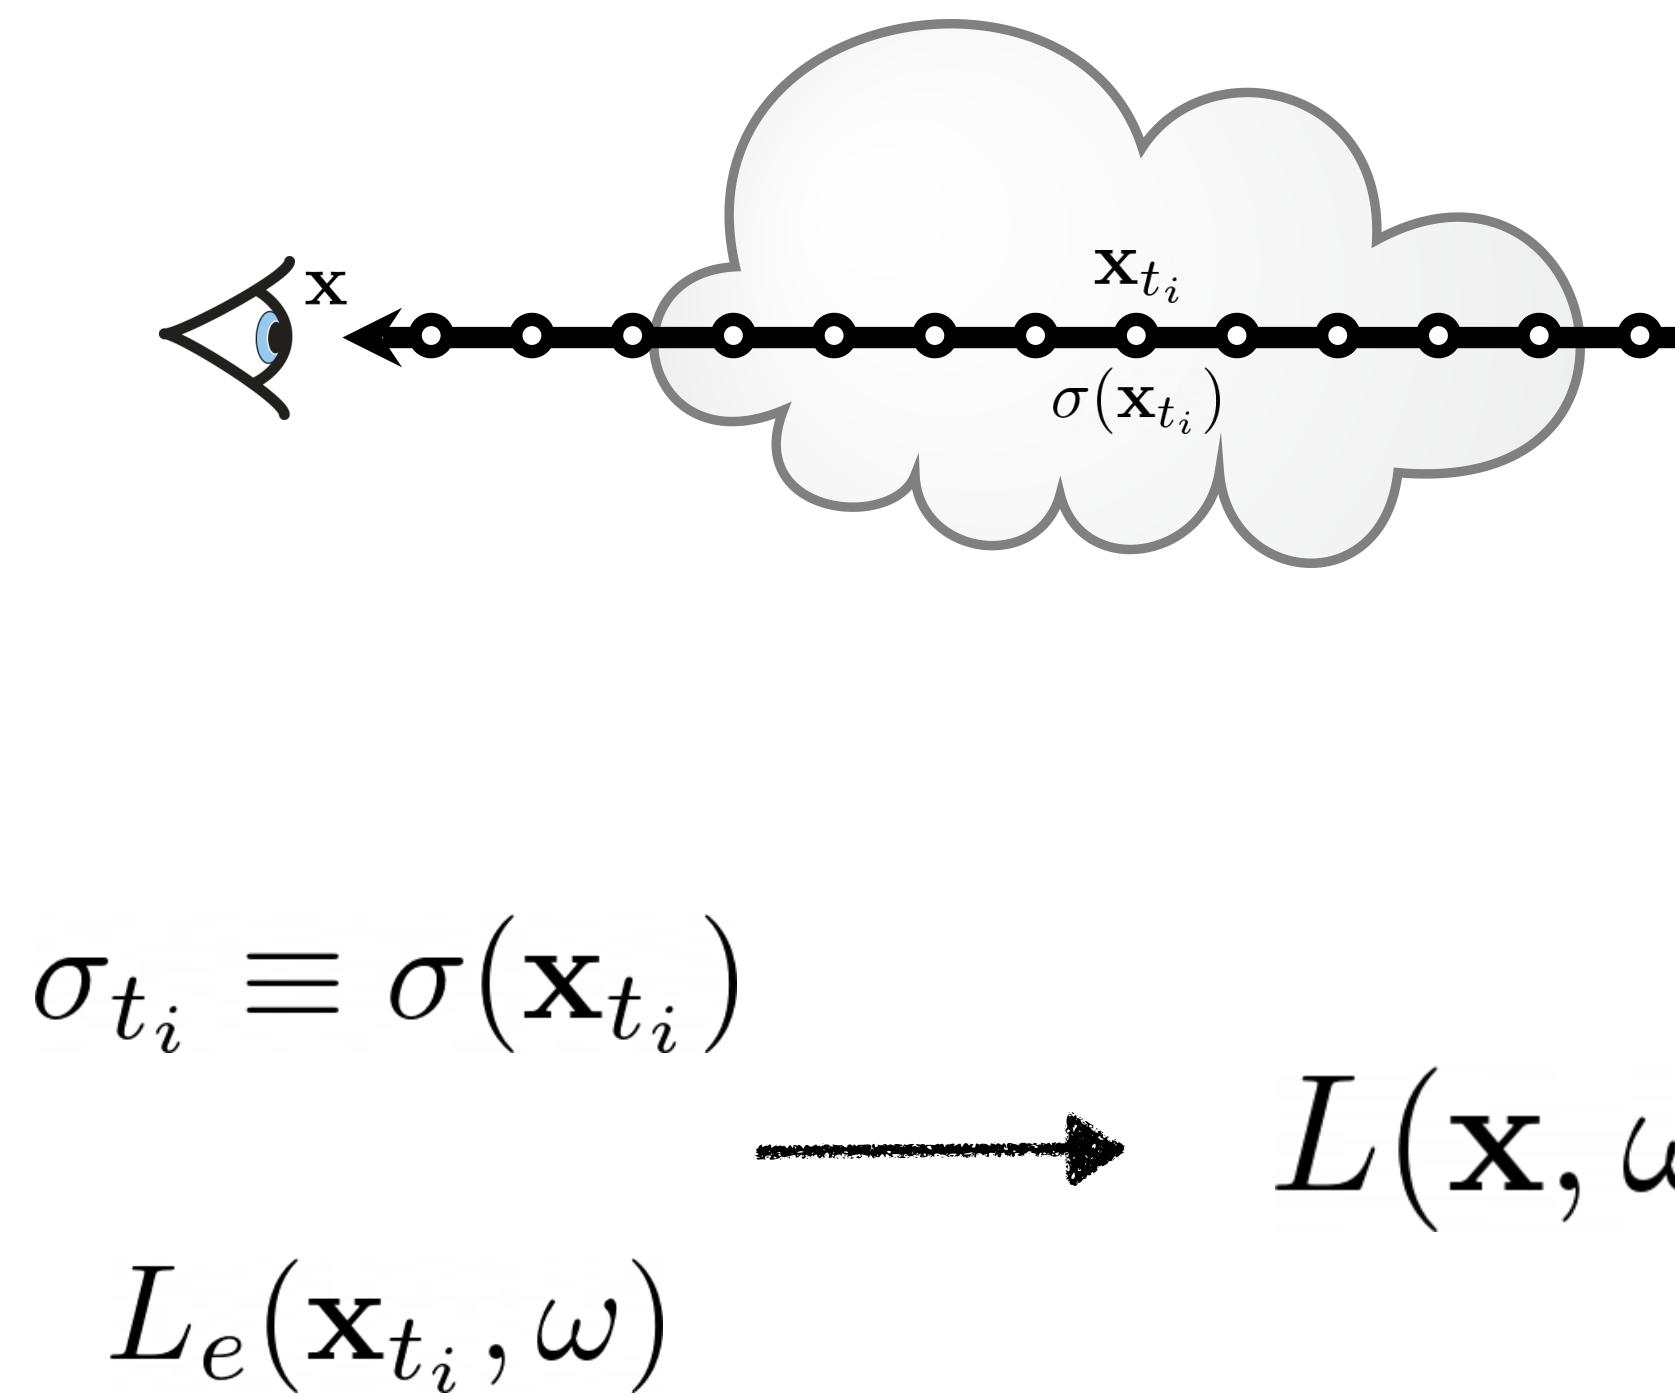
we can render **any** ray through the medium

Equivalently, we can render an image from any camera viewpoint
(using H^*W rays)

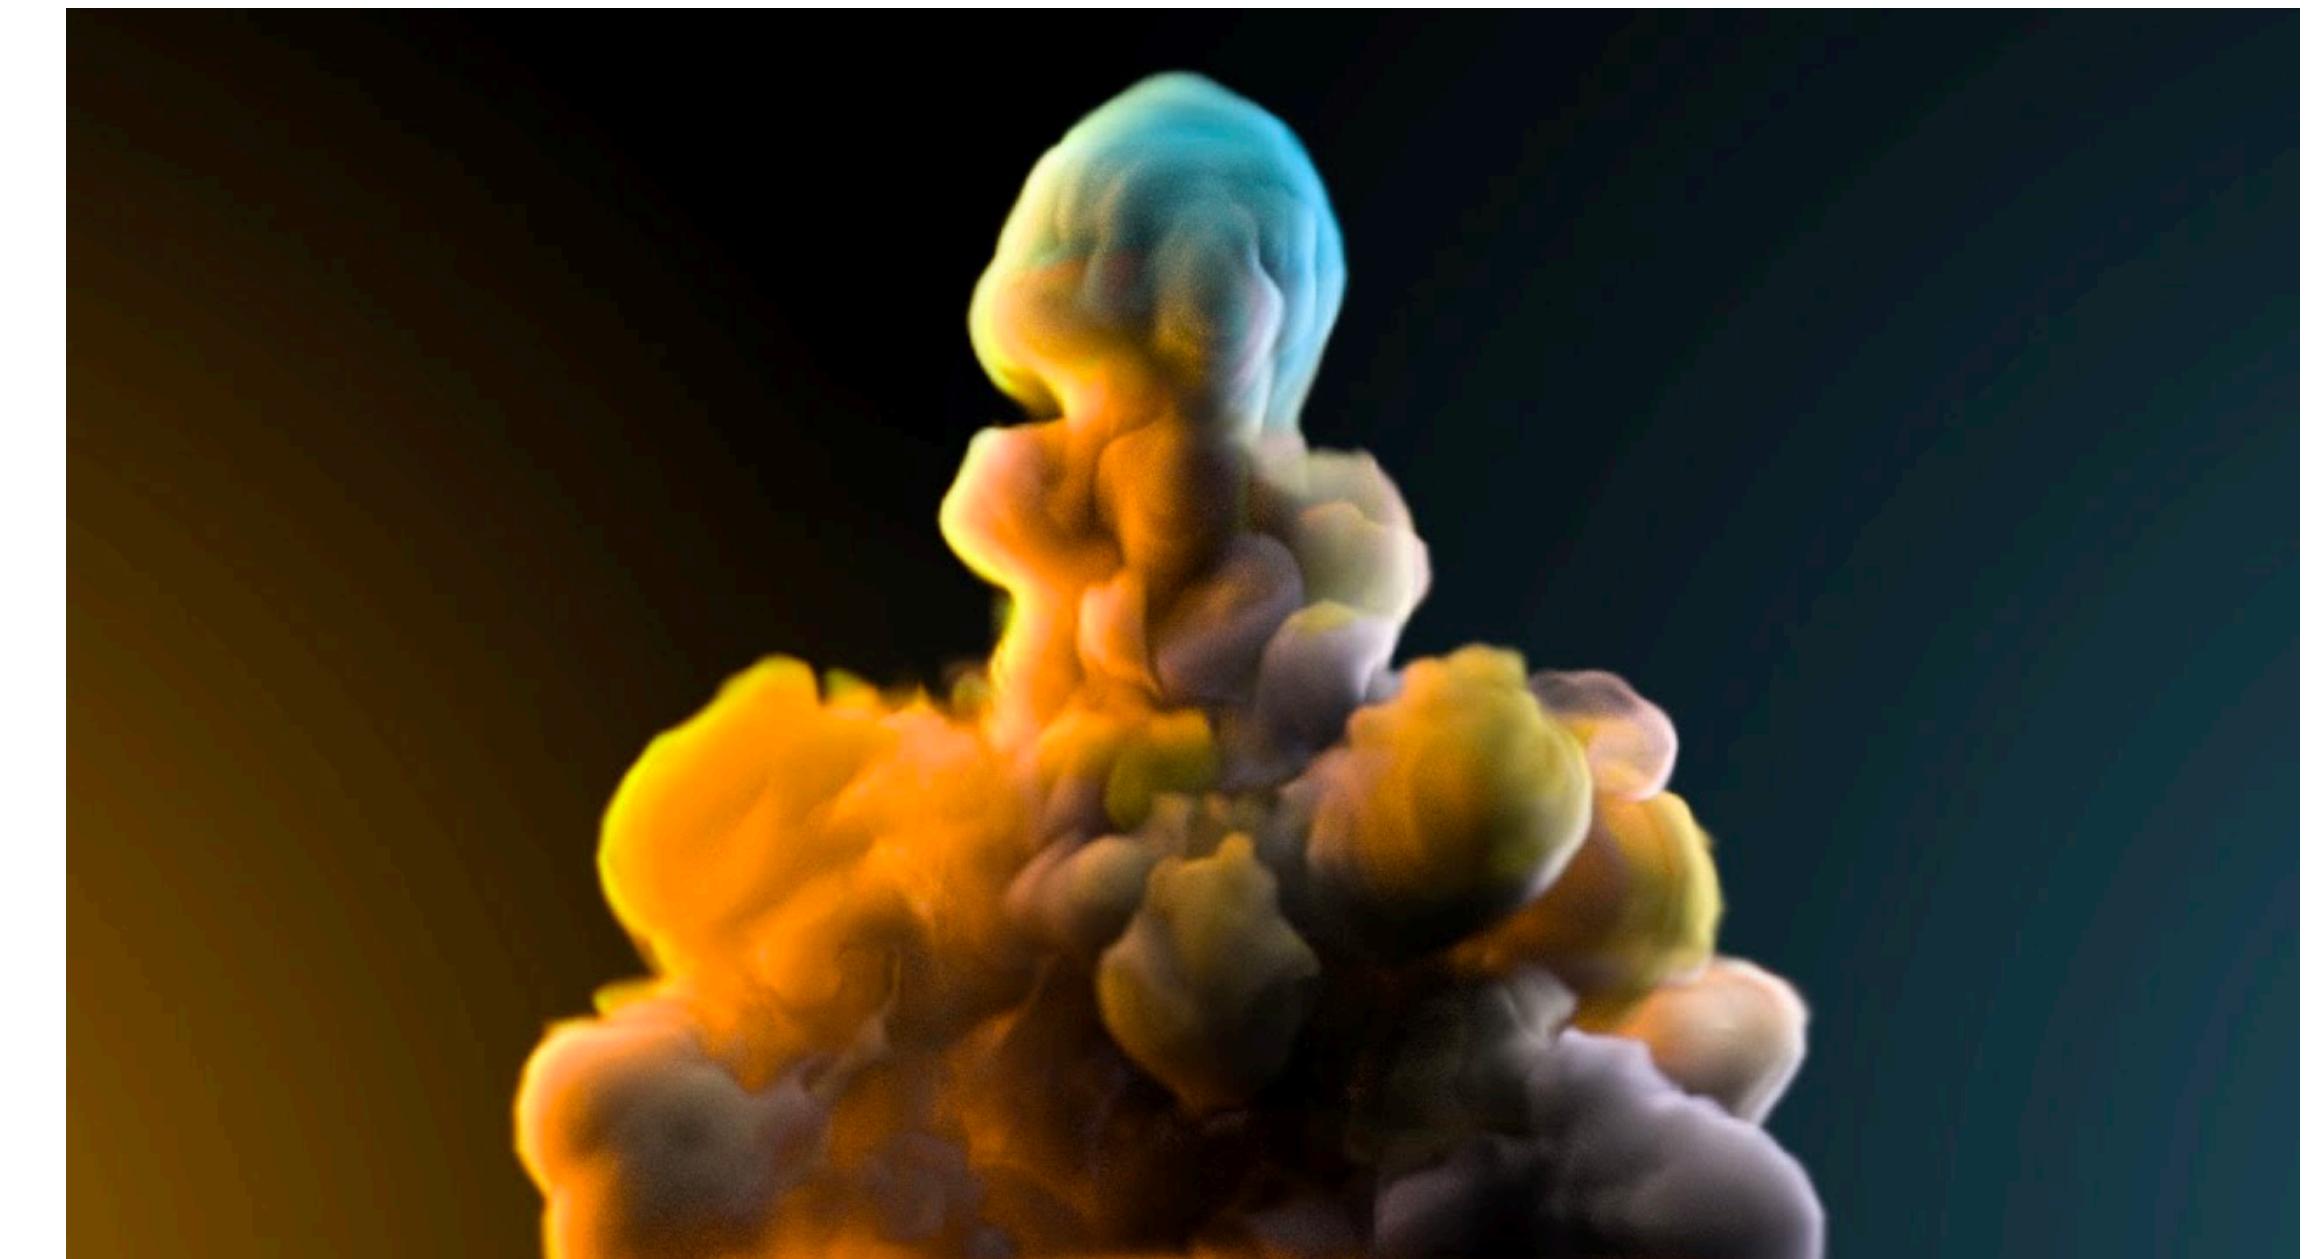
Note: **Differentiable** process w.r.t. the **density, emitted light**

and also camera parameters if density, emission are differentiable
functions of position, direction

Volumes: Rendering and Representation



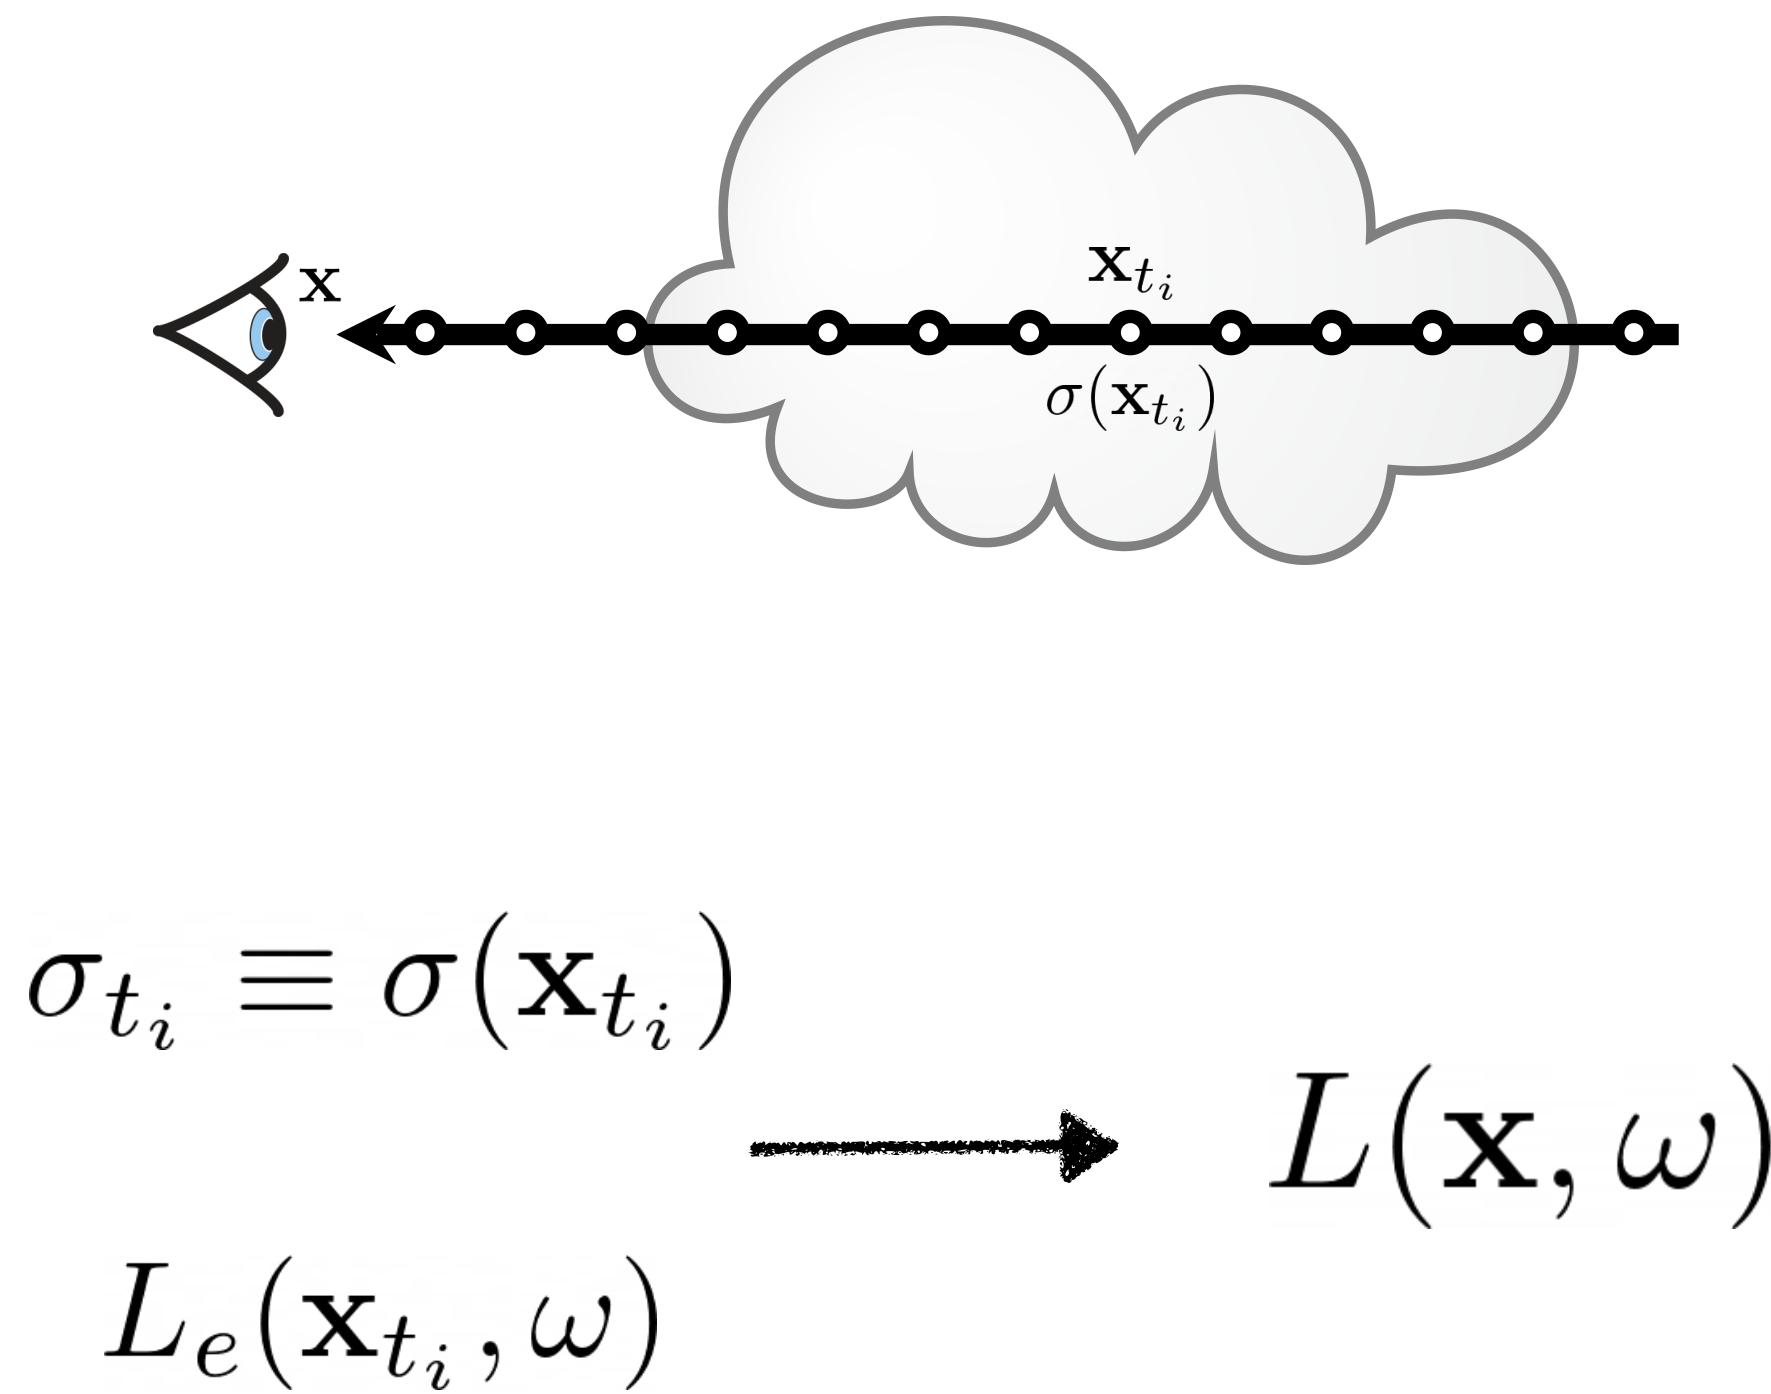
Rendering Algorithm



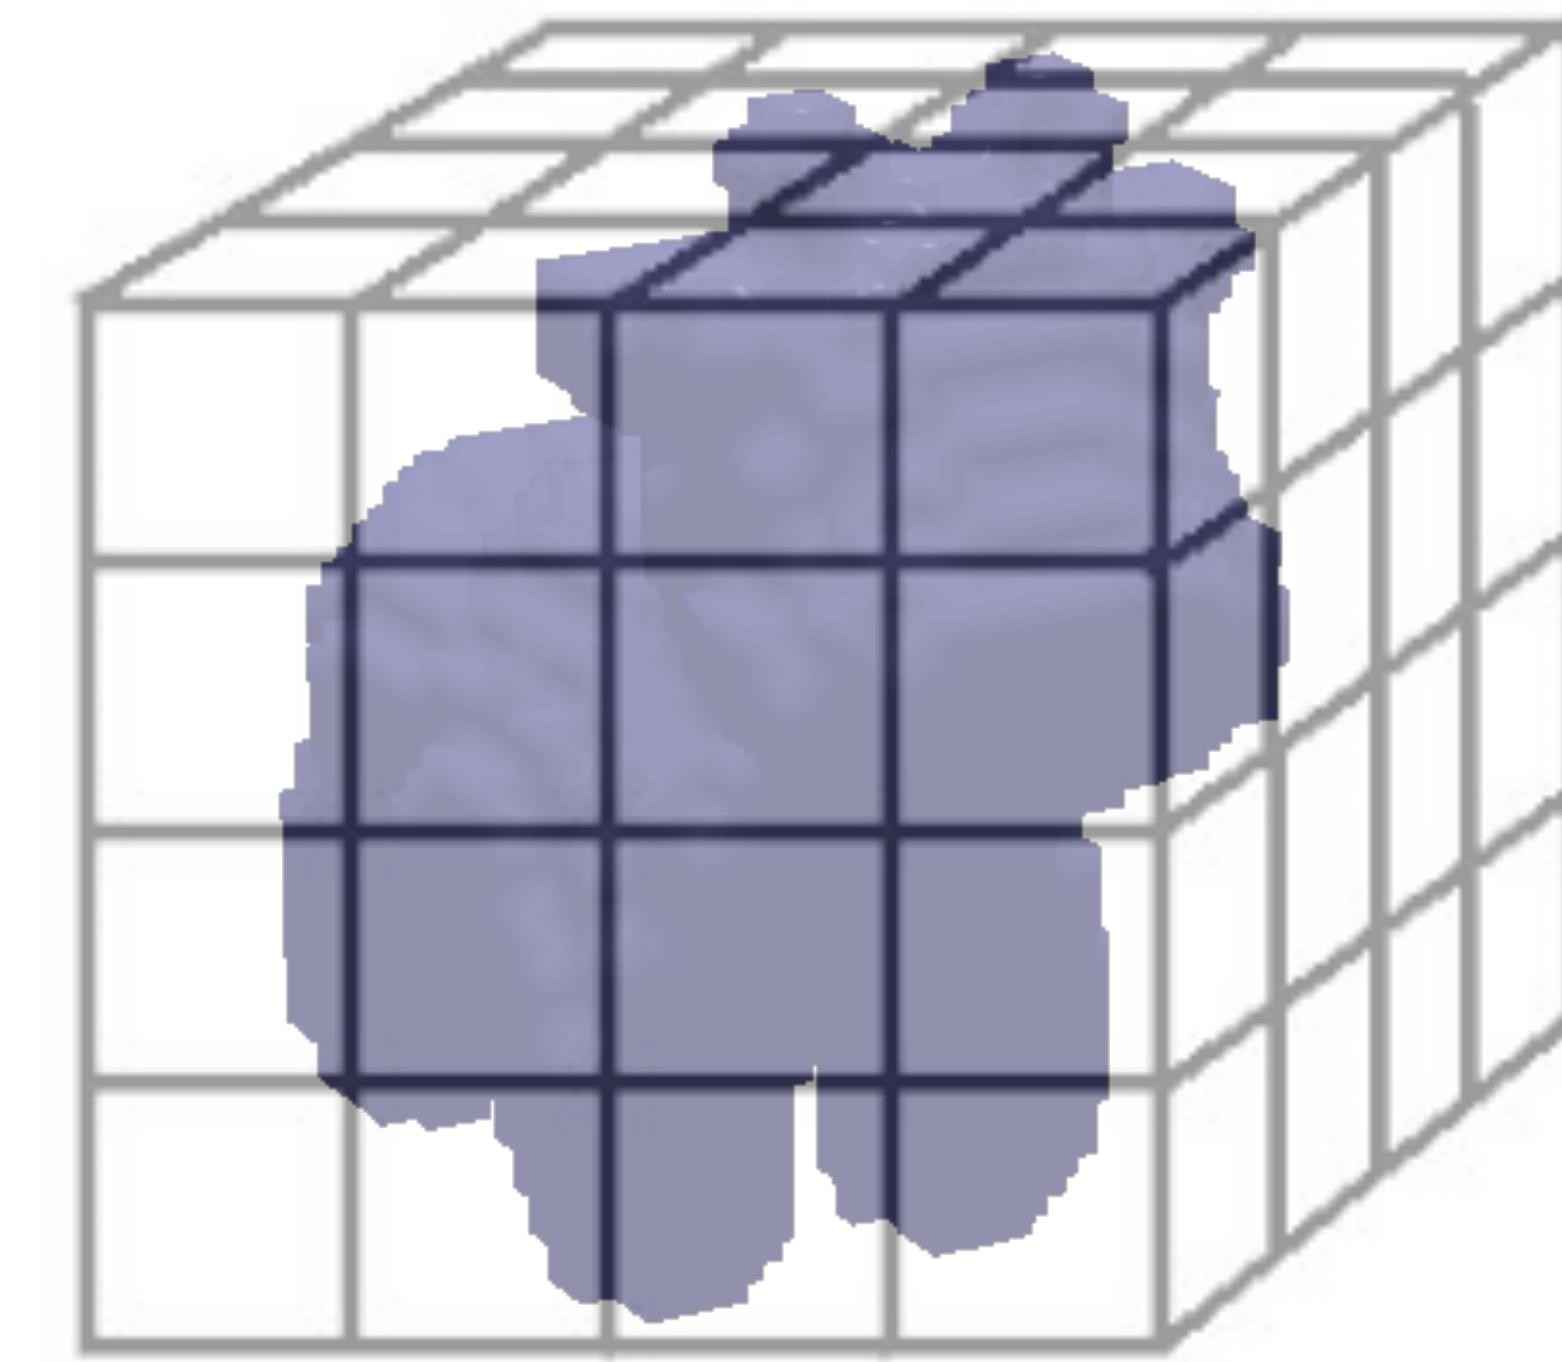
How to represent volumes?

(such that we can compute pointwise density and emitted light)

Volumes: Rendering and Representation



Rendering Algorithm

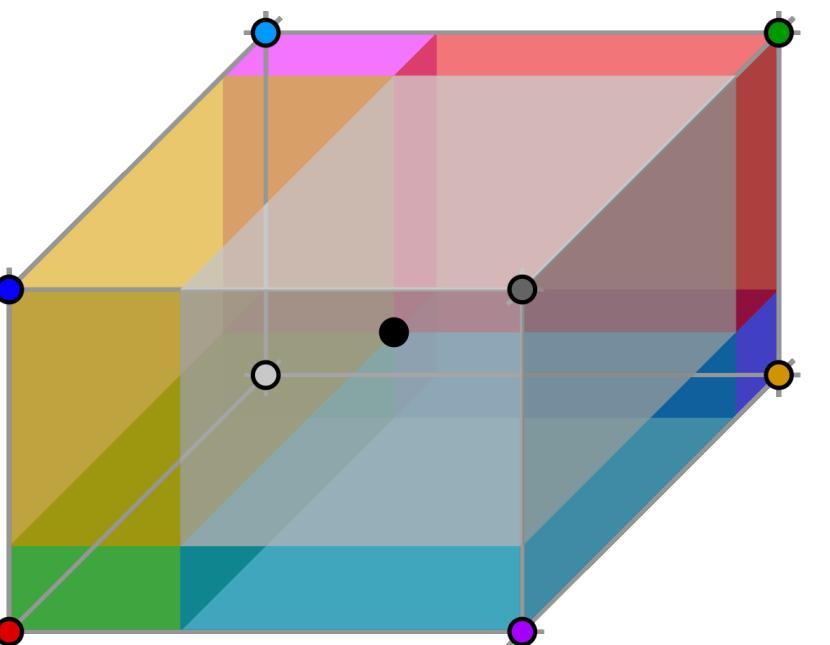


Option 1: A grid

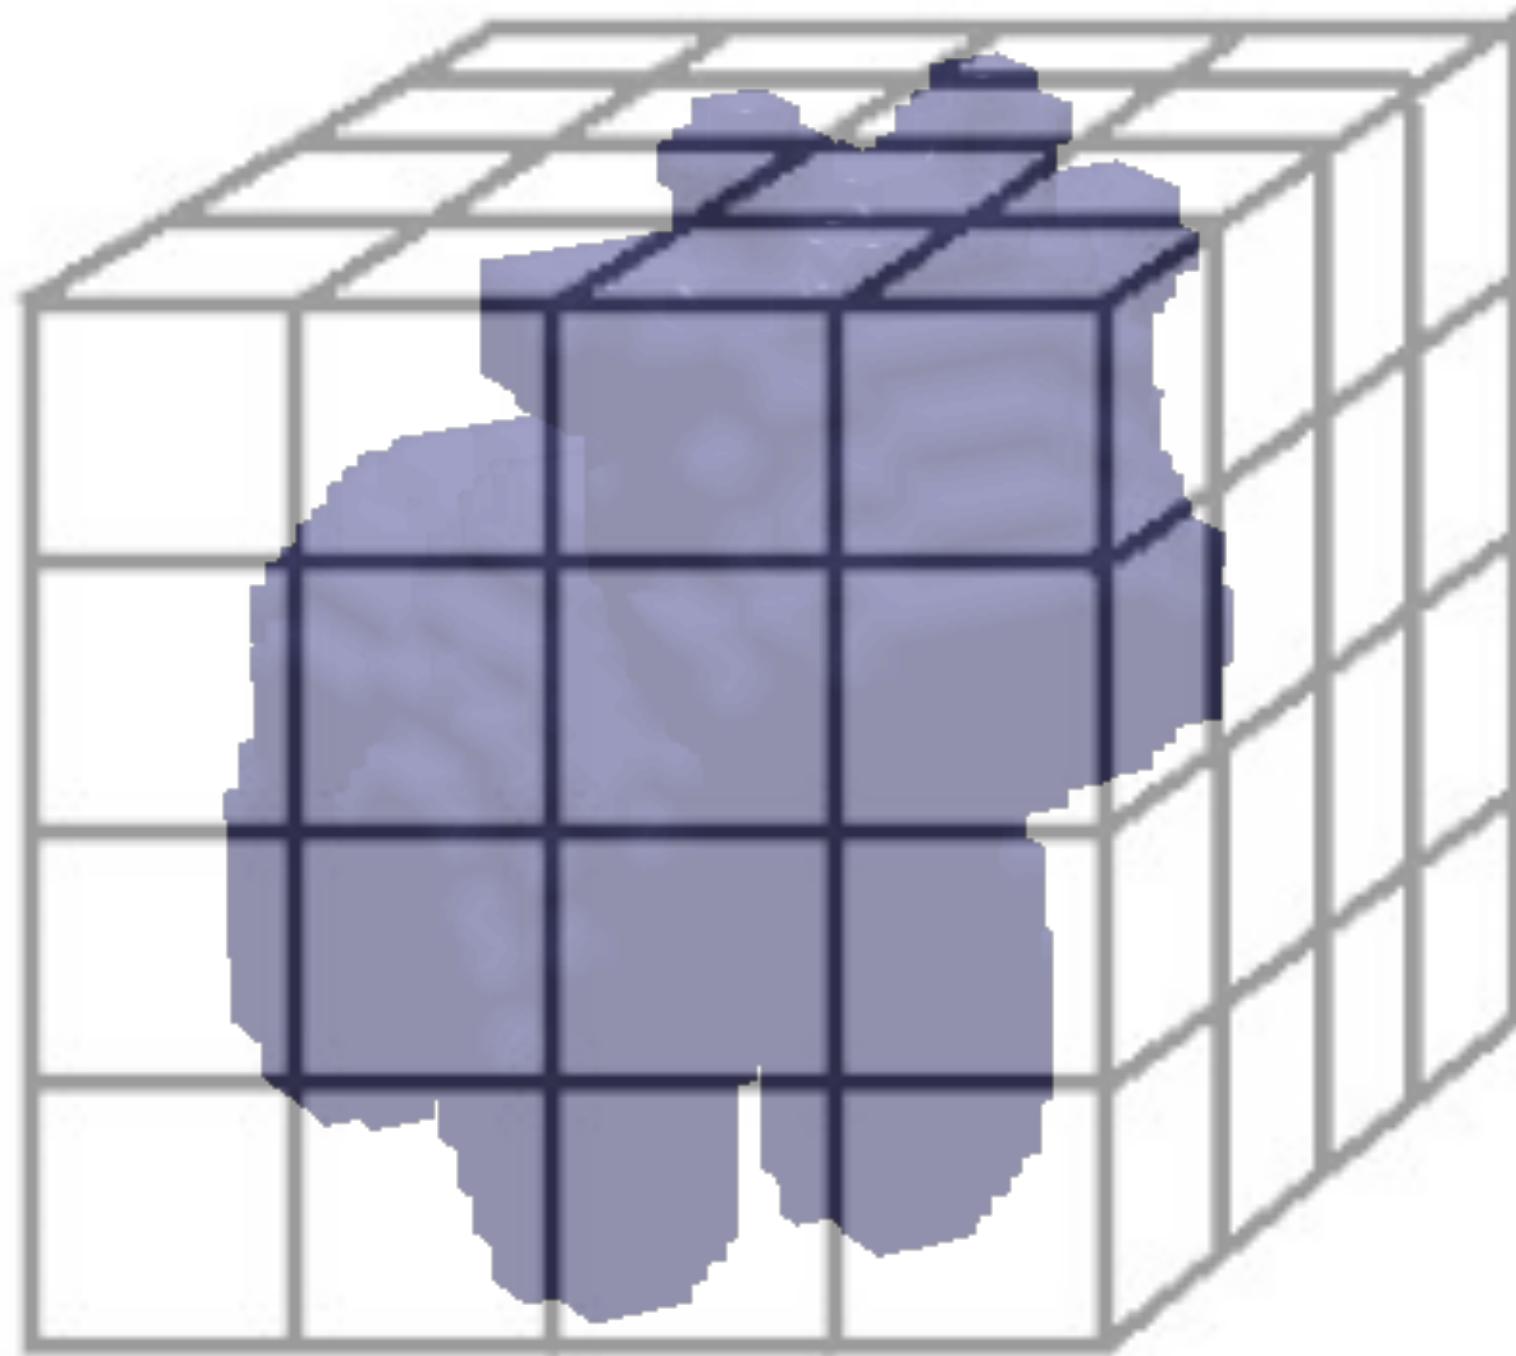
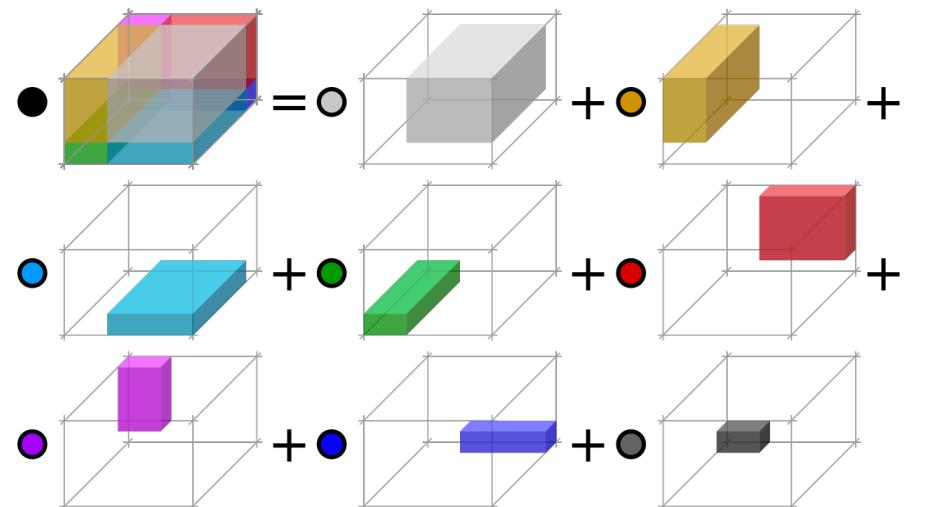
Grid Representation

$$\sigma[x, y, z] \in \mathbb{R}^+$$

Store a positive scalar in each voxel (representing density)



Trilinear
Interpolation



Grid Representation

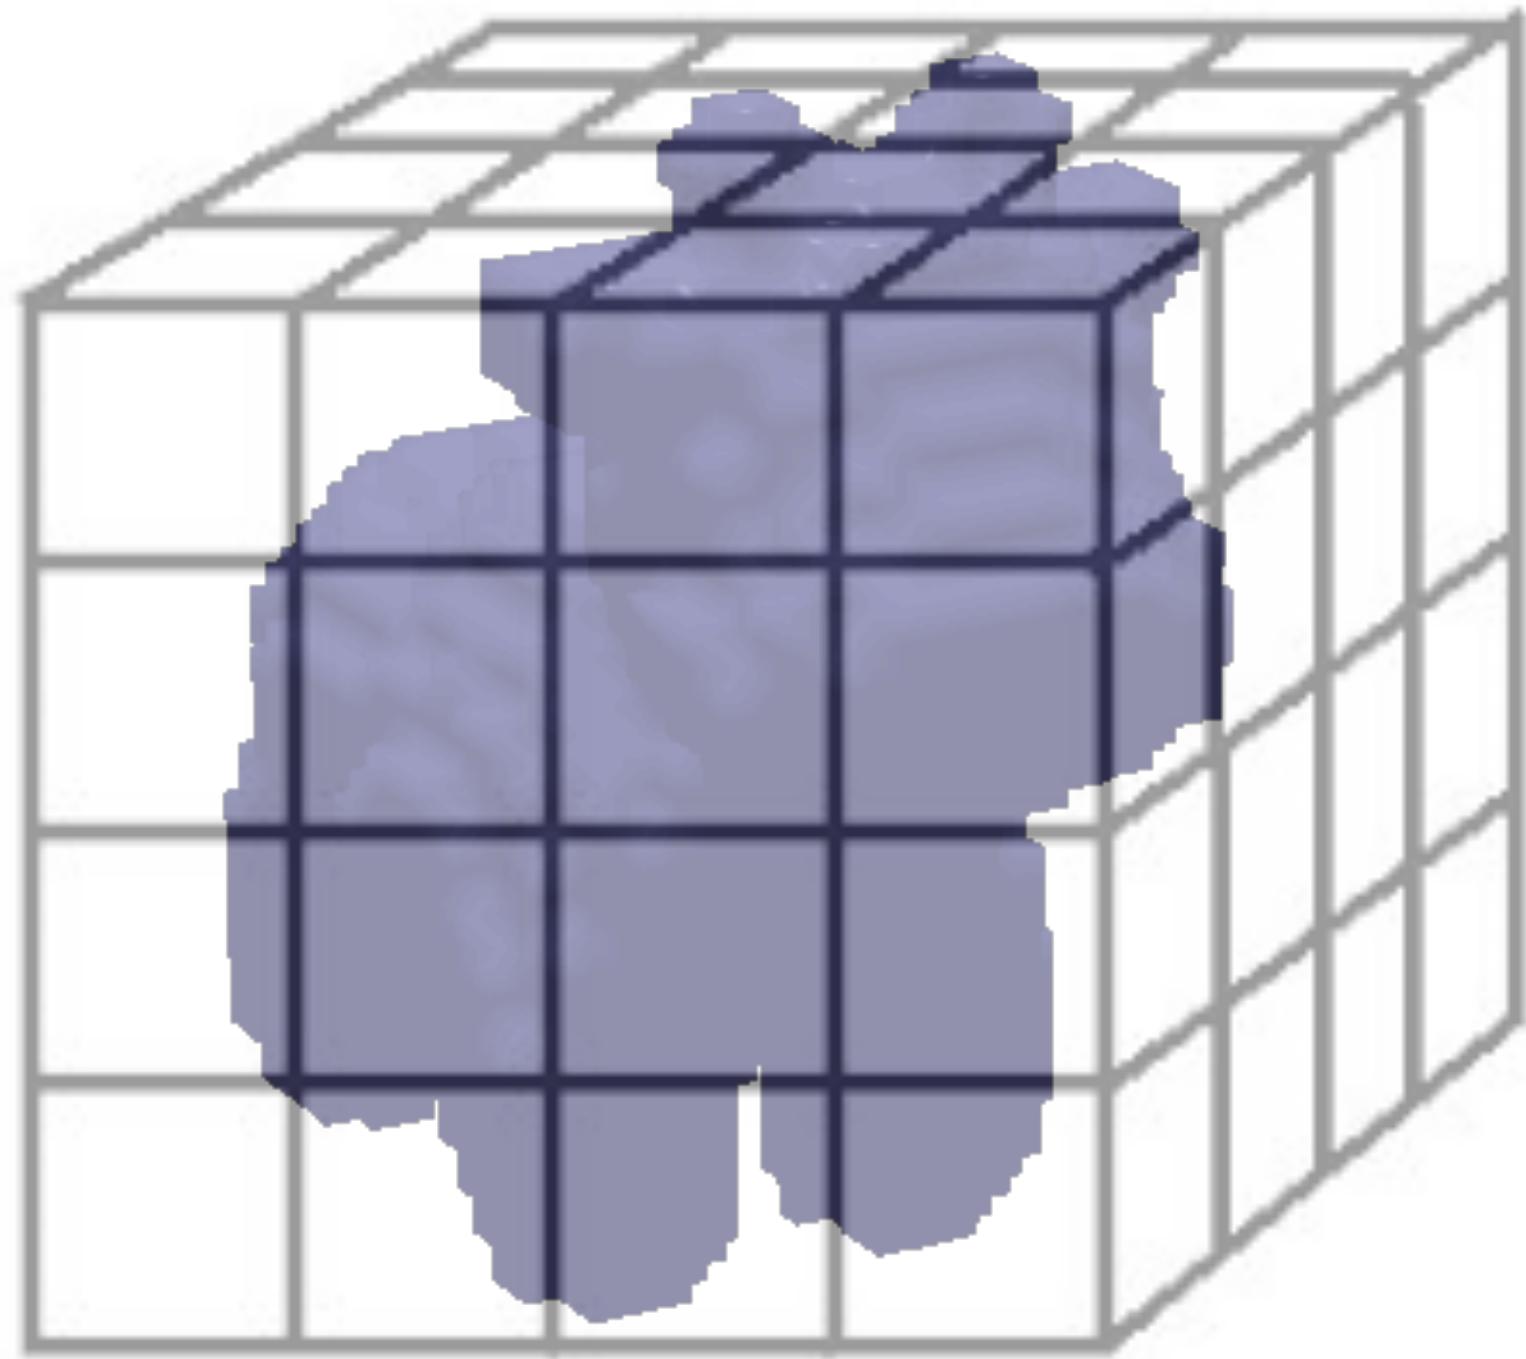
$$\sigma[x, y, z] \in \mathbb{R}^+$$

Store a positive scalar in each voxel (representing density)

Q: How to model view-dependent emission? $L_e(\mathbf{x}, \omega)$

A: Ignore view-dependence in emission

$$L_e(\mathbf{x}, \omega) \equiv c(\mathbf{x})$$



Grid Representation

$$\sigma[x, y, z] \in \mathbb{R}^+$$

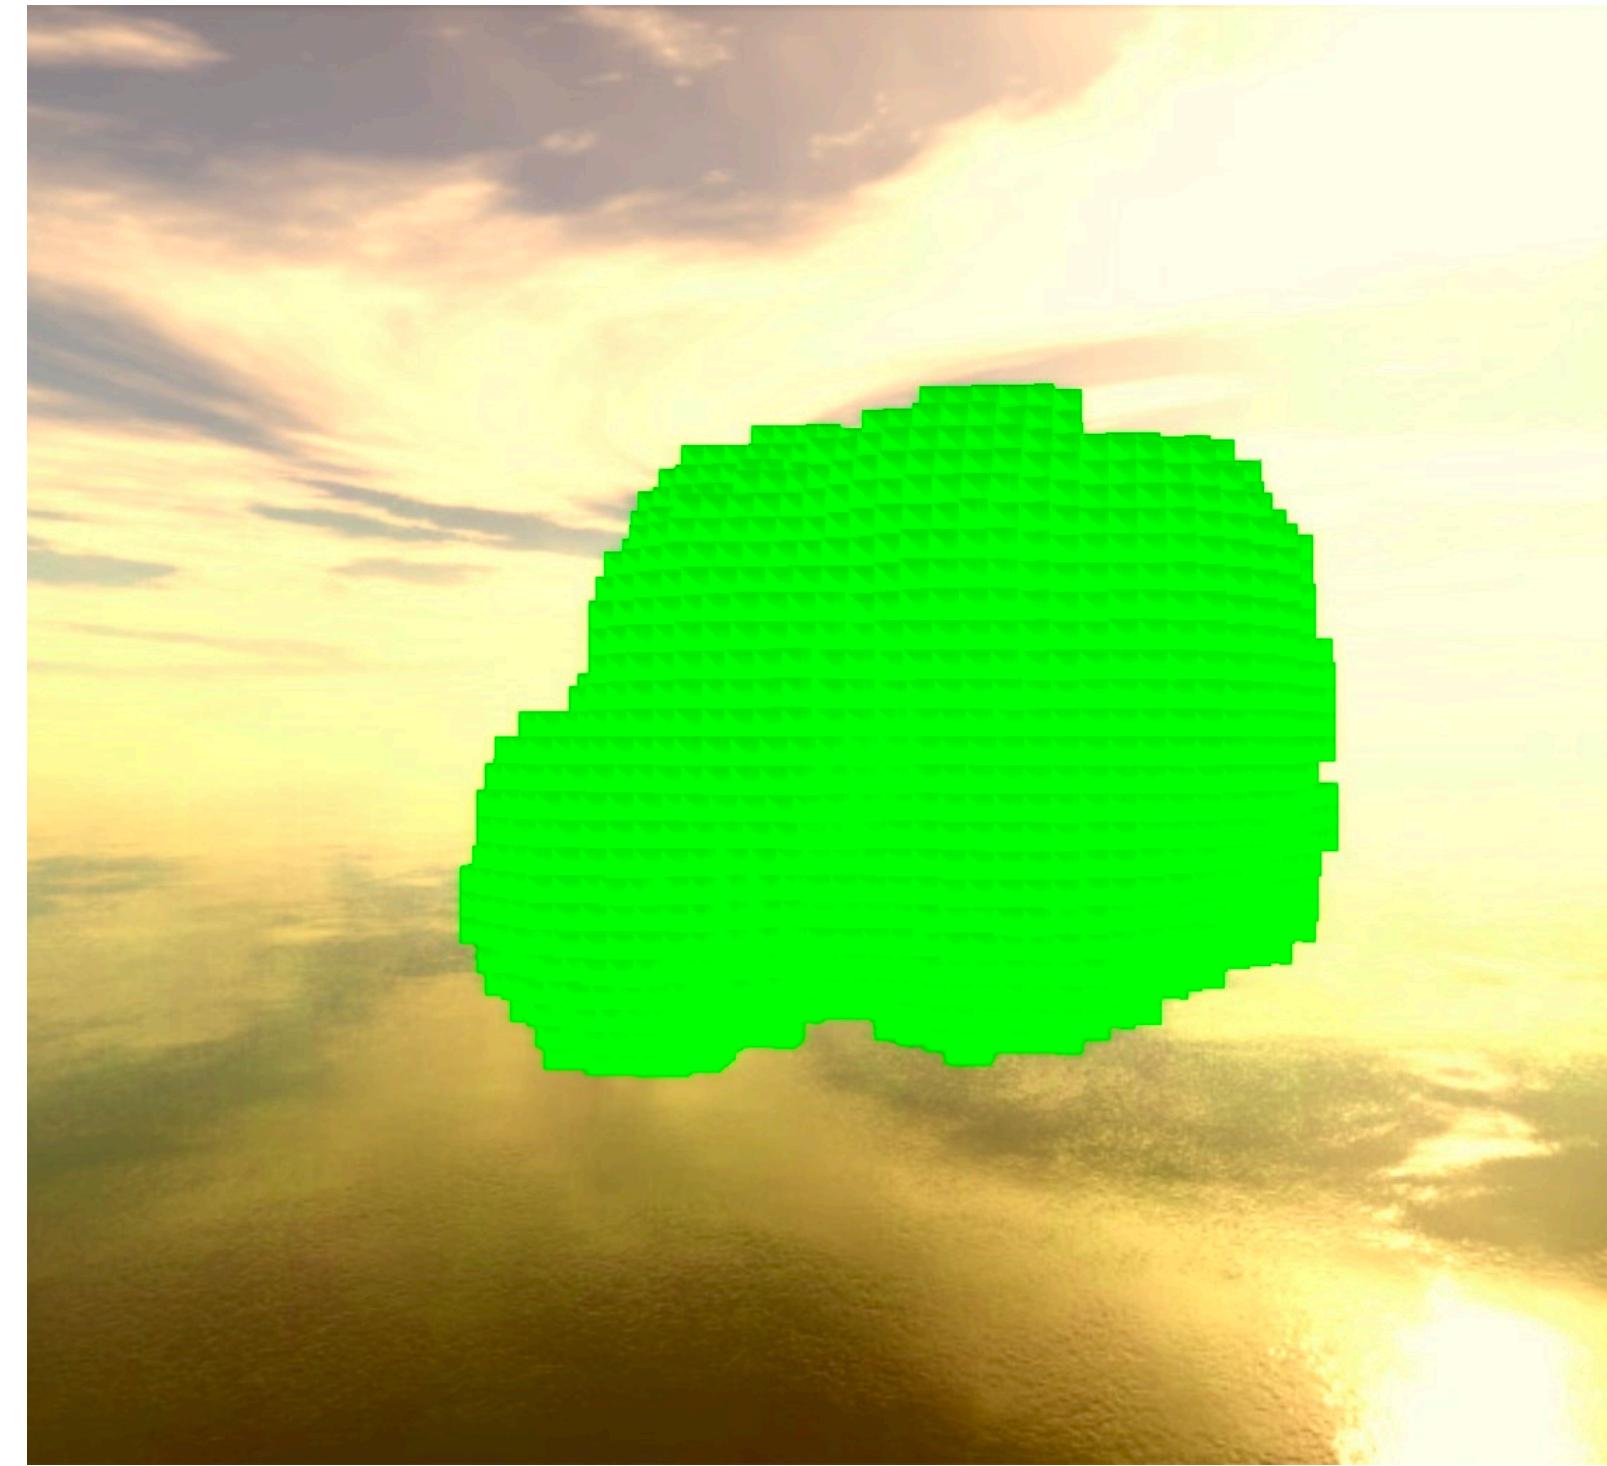
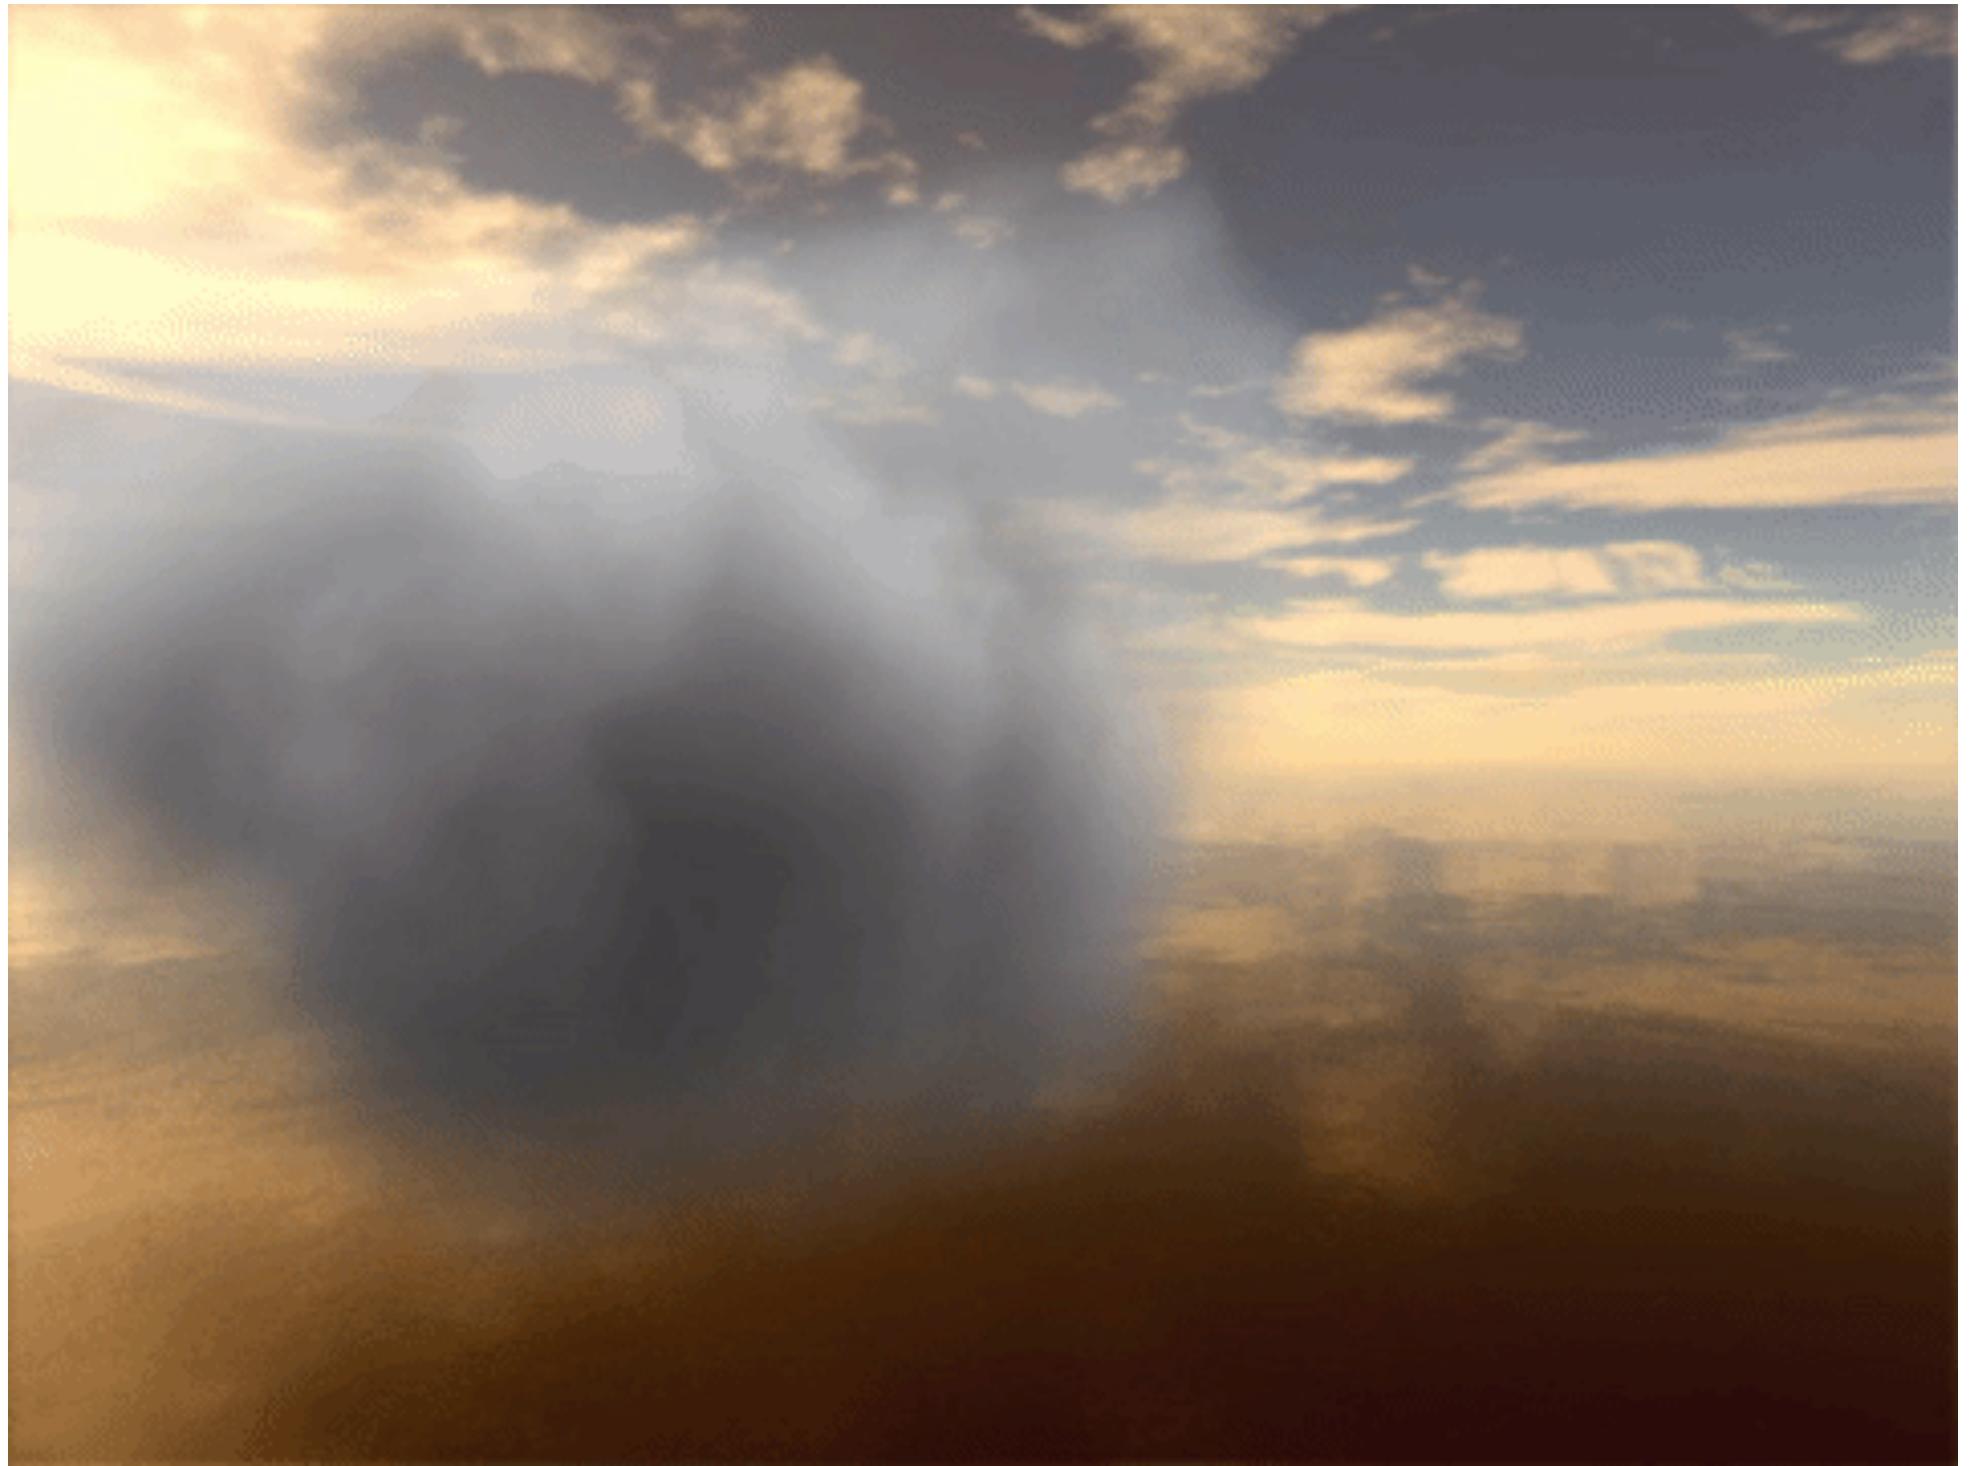
Store a positive scalar in each voxel (representing density)

$$c[x, y, z] \in [0, 1]^3$$

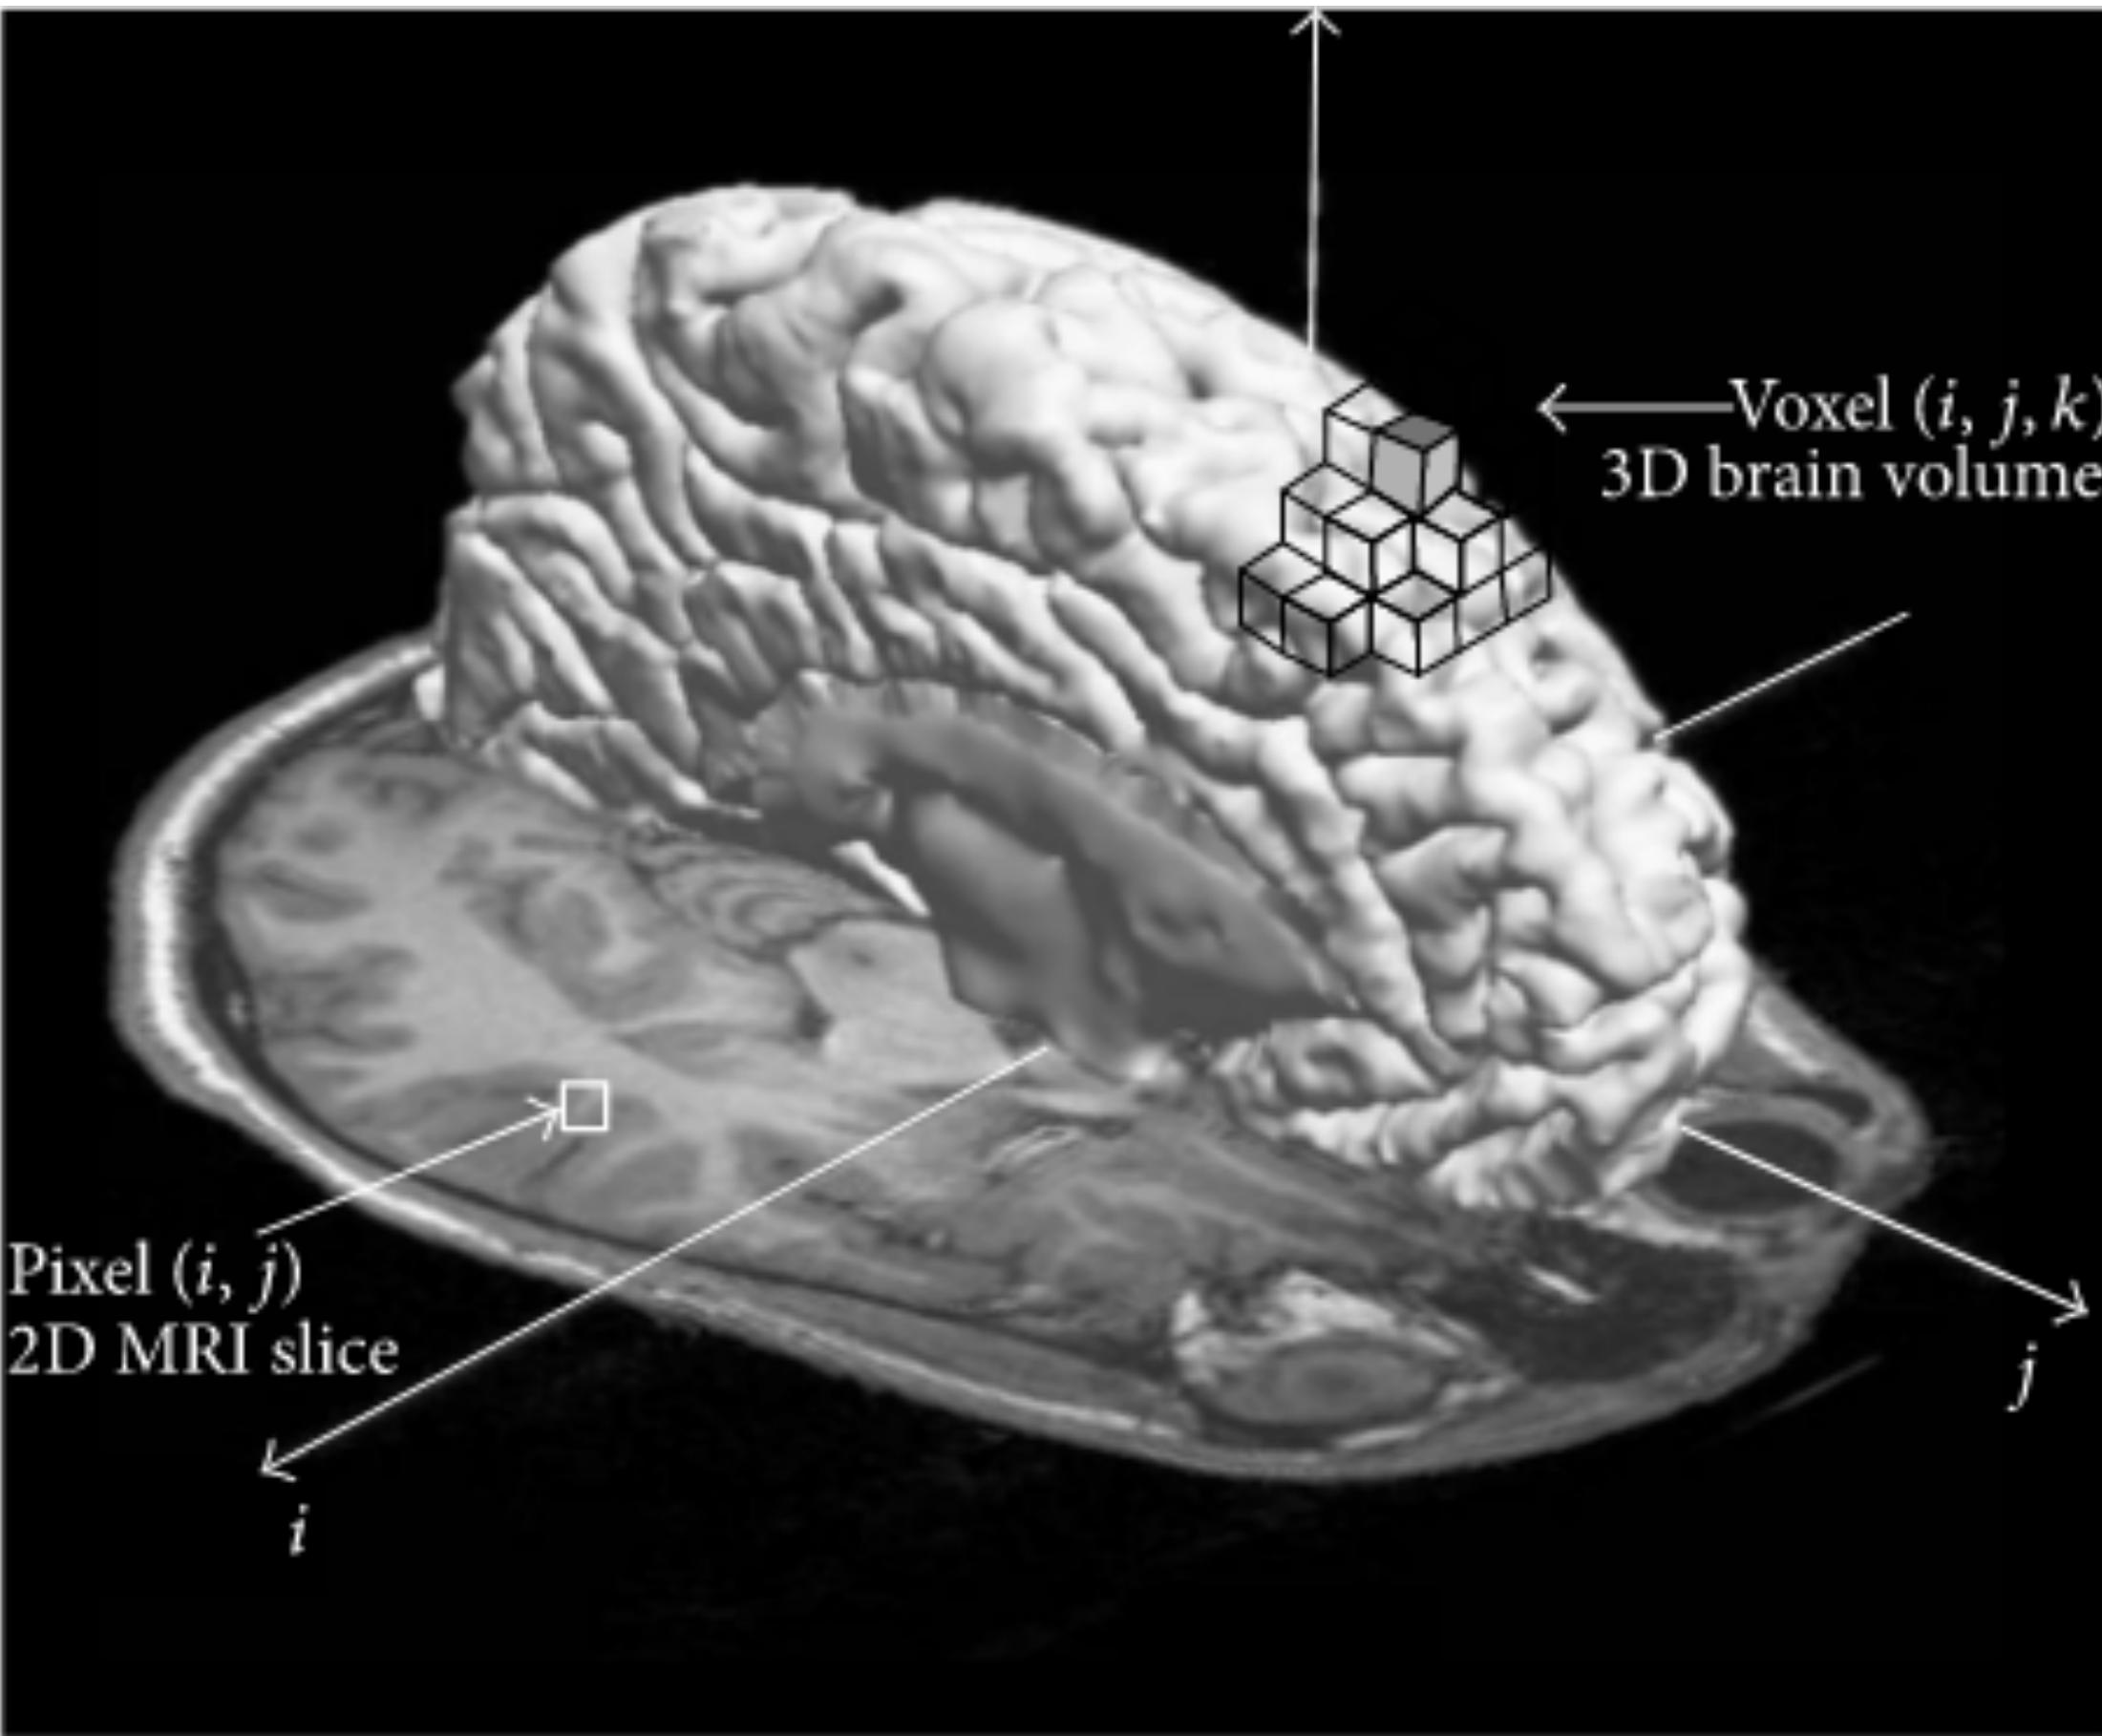
Store a color in each voxel (representing emitted light)



Grid Representation



Grid Representation



Grid Representation

$$\sigma[x, y, z] \in \mathbb{R}^+$$

Store a positive scalar in each voxel (representing density)

$$c[x, y, z] \in [0, 1]^3$$

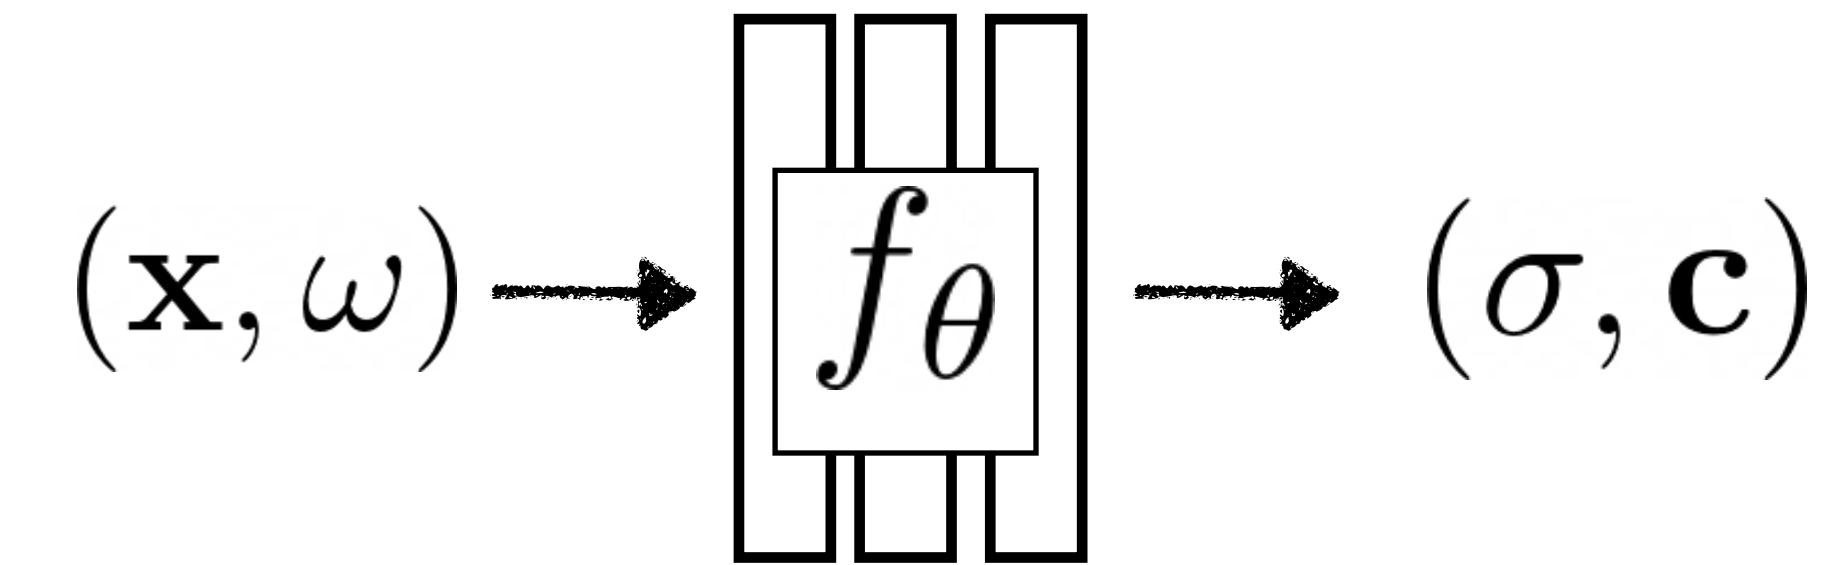
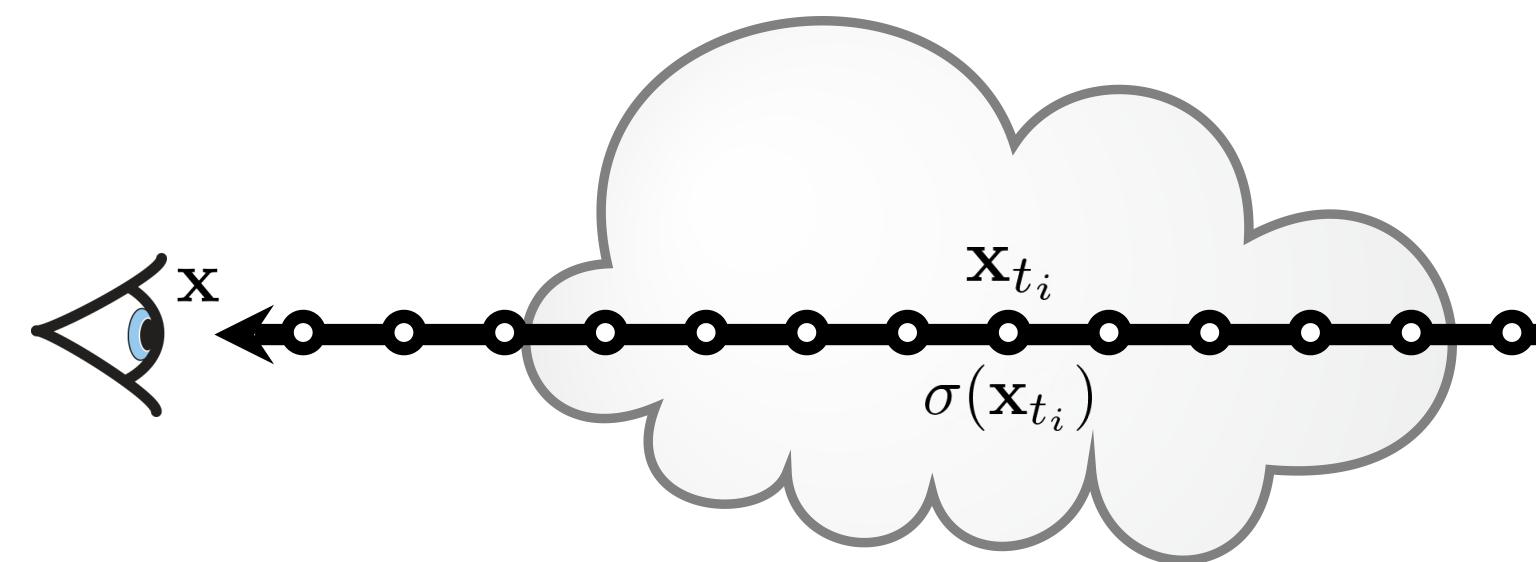
Store a color in each voxel (representing emitted light)



Cannot model view-dependent effects

Cubic growth of grid size with scene dimension

Volumes: Rendering and Representation



$$\sigma \in \mathbb{R}^+ \quad \mathbf{c} \in [0, 1]^3$$

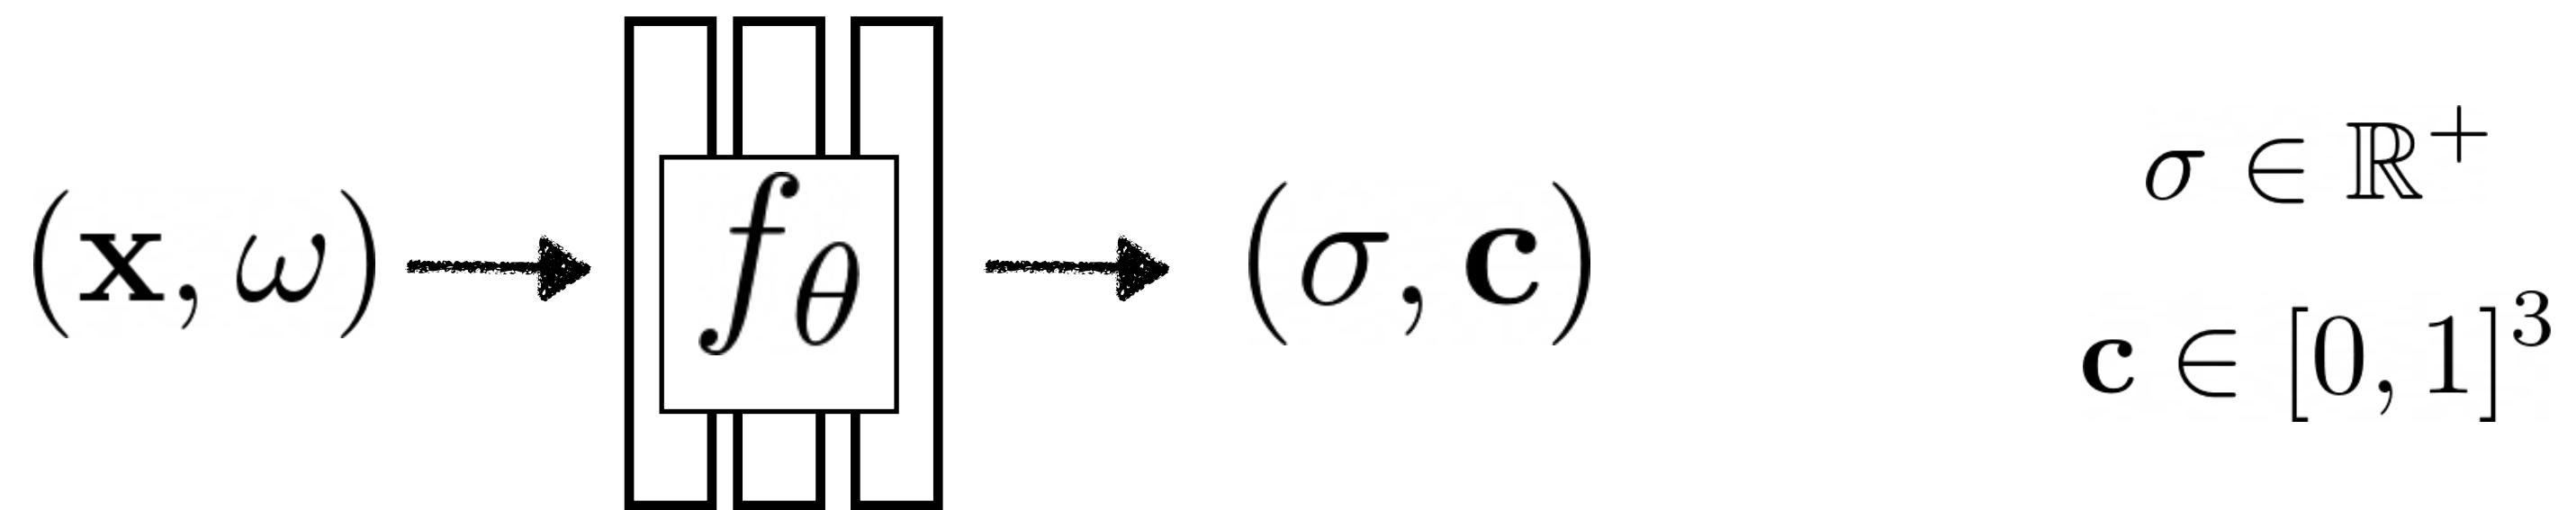
$$\begin{aligned}\sigma_{t_i} &\equiv \sigma(\mathbf{x}_{t_i}) \\ &\longrightarrow L(\mathbf{x}, \omega) \\ L_e(\mathbf{x}_{t_i}, \omega) &\end{aligned}$$

Ensure with designing MLP that density only depends on \mathbf{x}

Rendering Algorithm

Option 2: An MLP

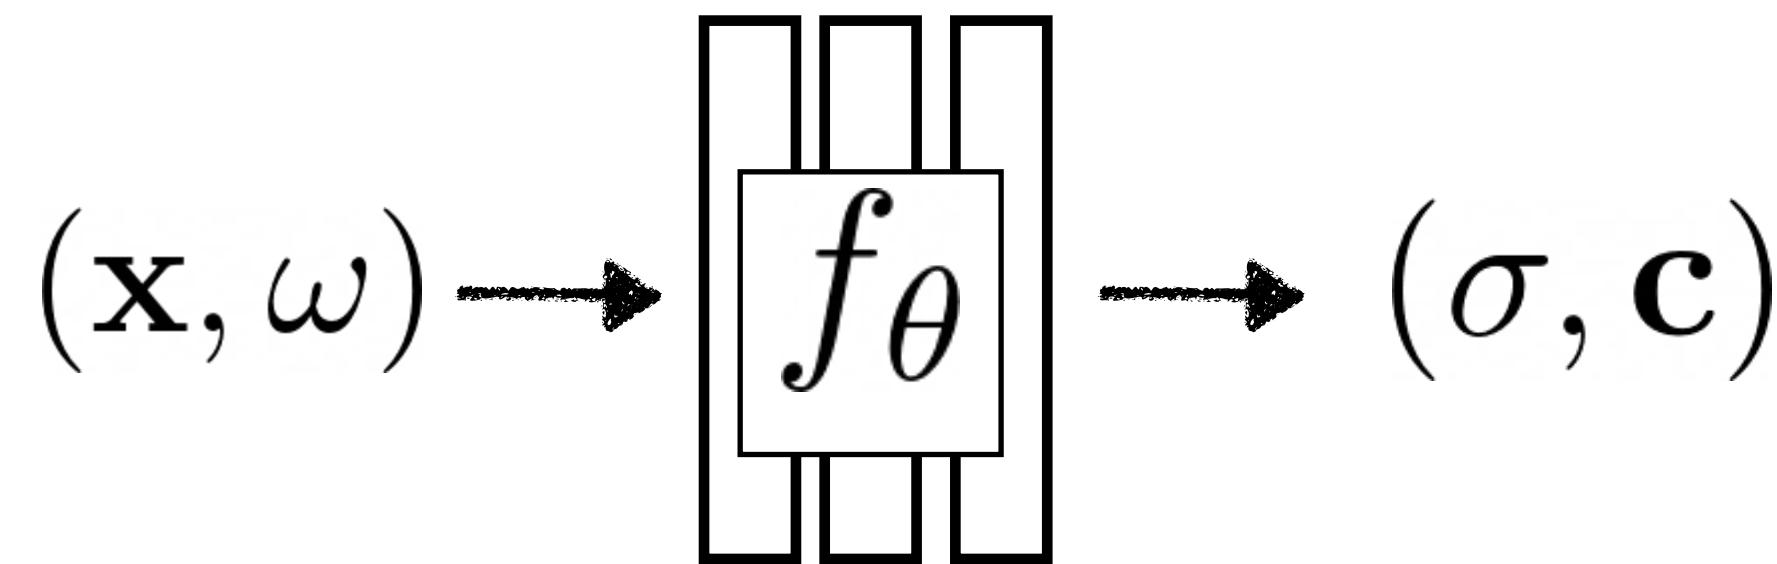
Neural Radiance Fields



A scene is represented as an MLP that:

- 1) Given an input point, predicts density
- 2) Given a point and viewing direction, predicts radiance (color)

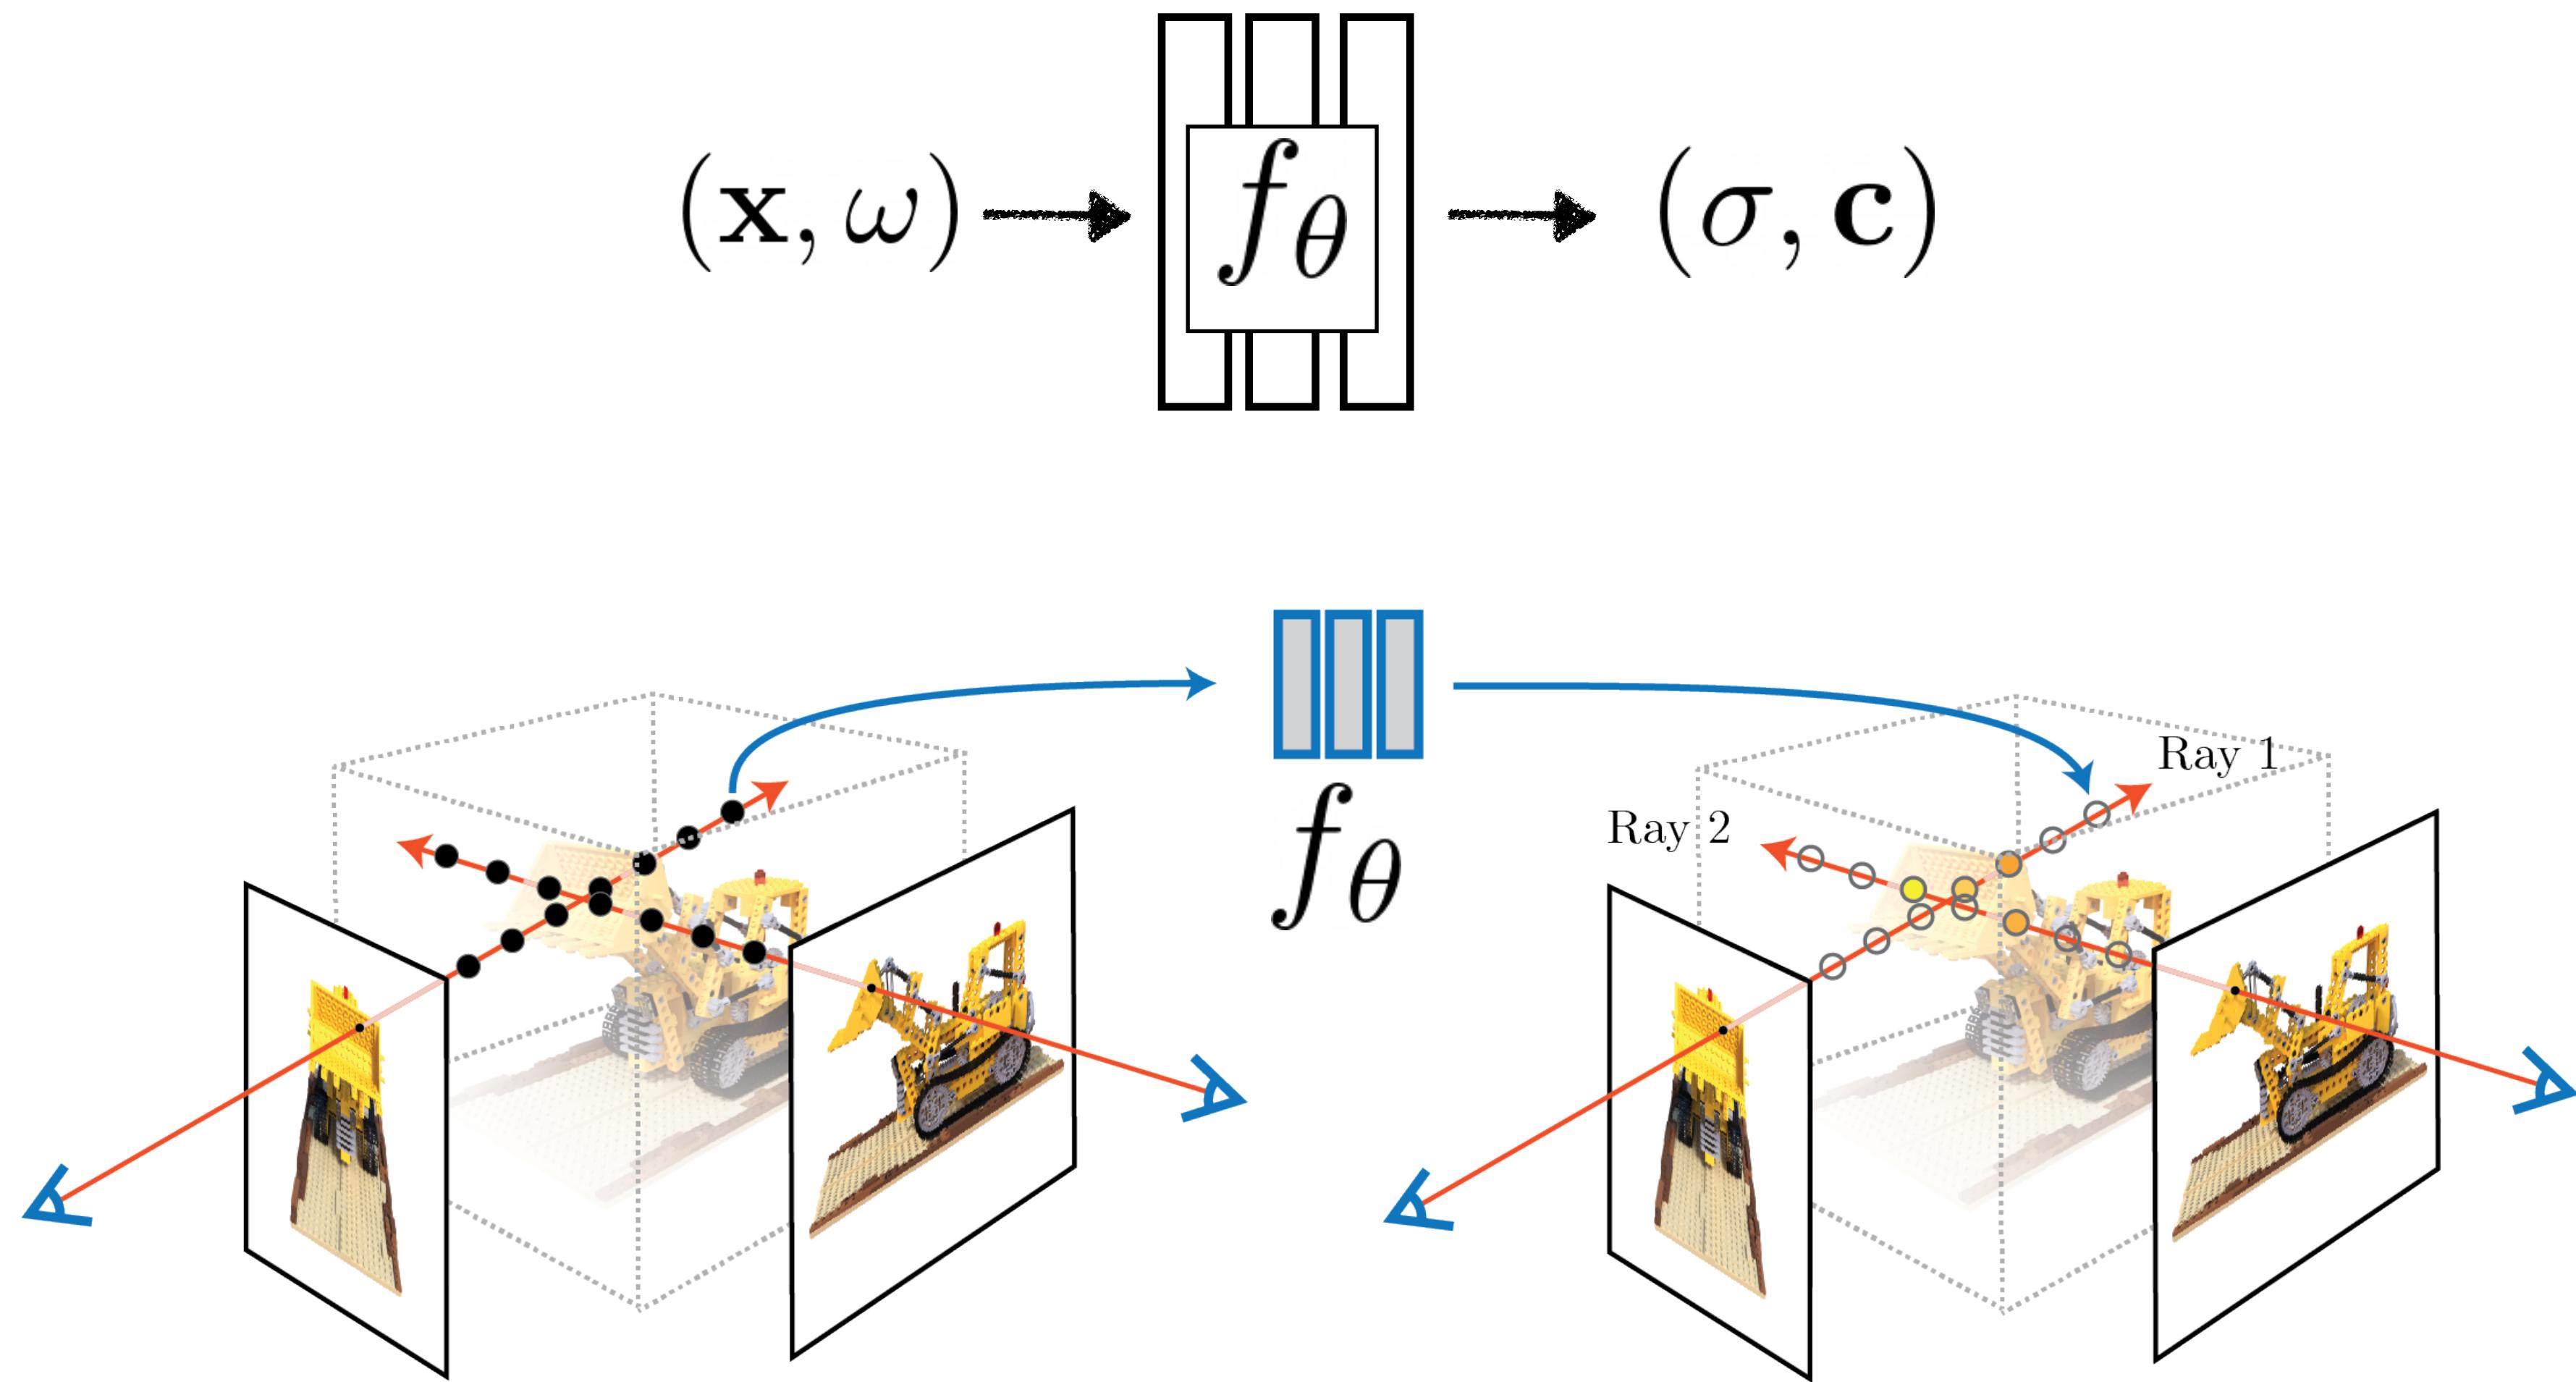
Neural Radiance Fields



How to train this MLP?

Key idea: Learn a 3D representation consistent with observed views

Learning Neural Radiance Fields

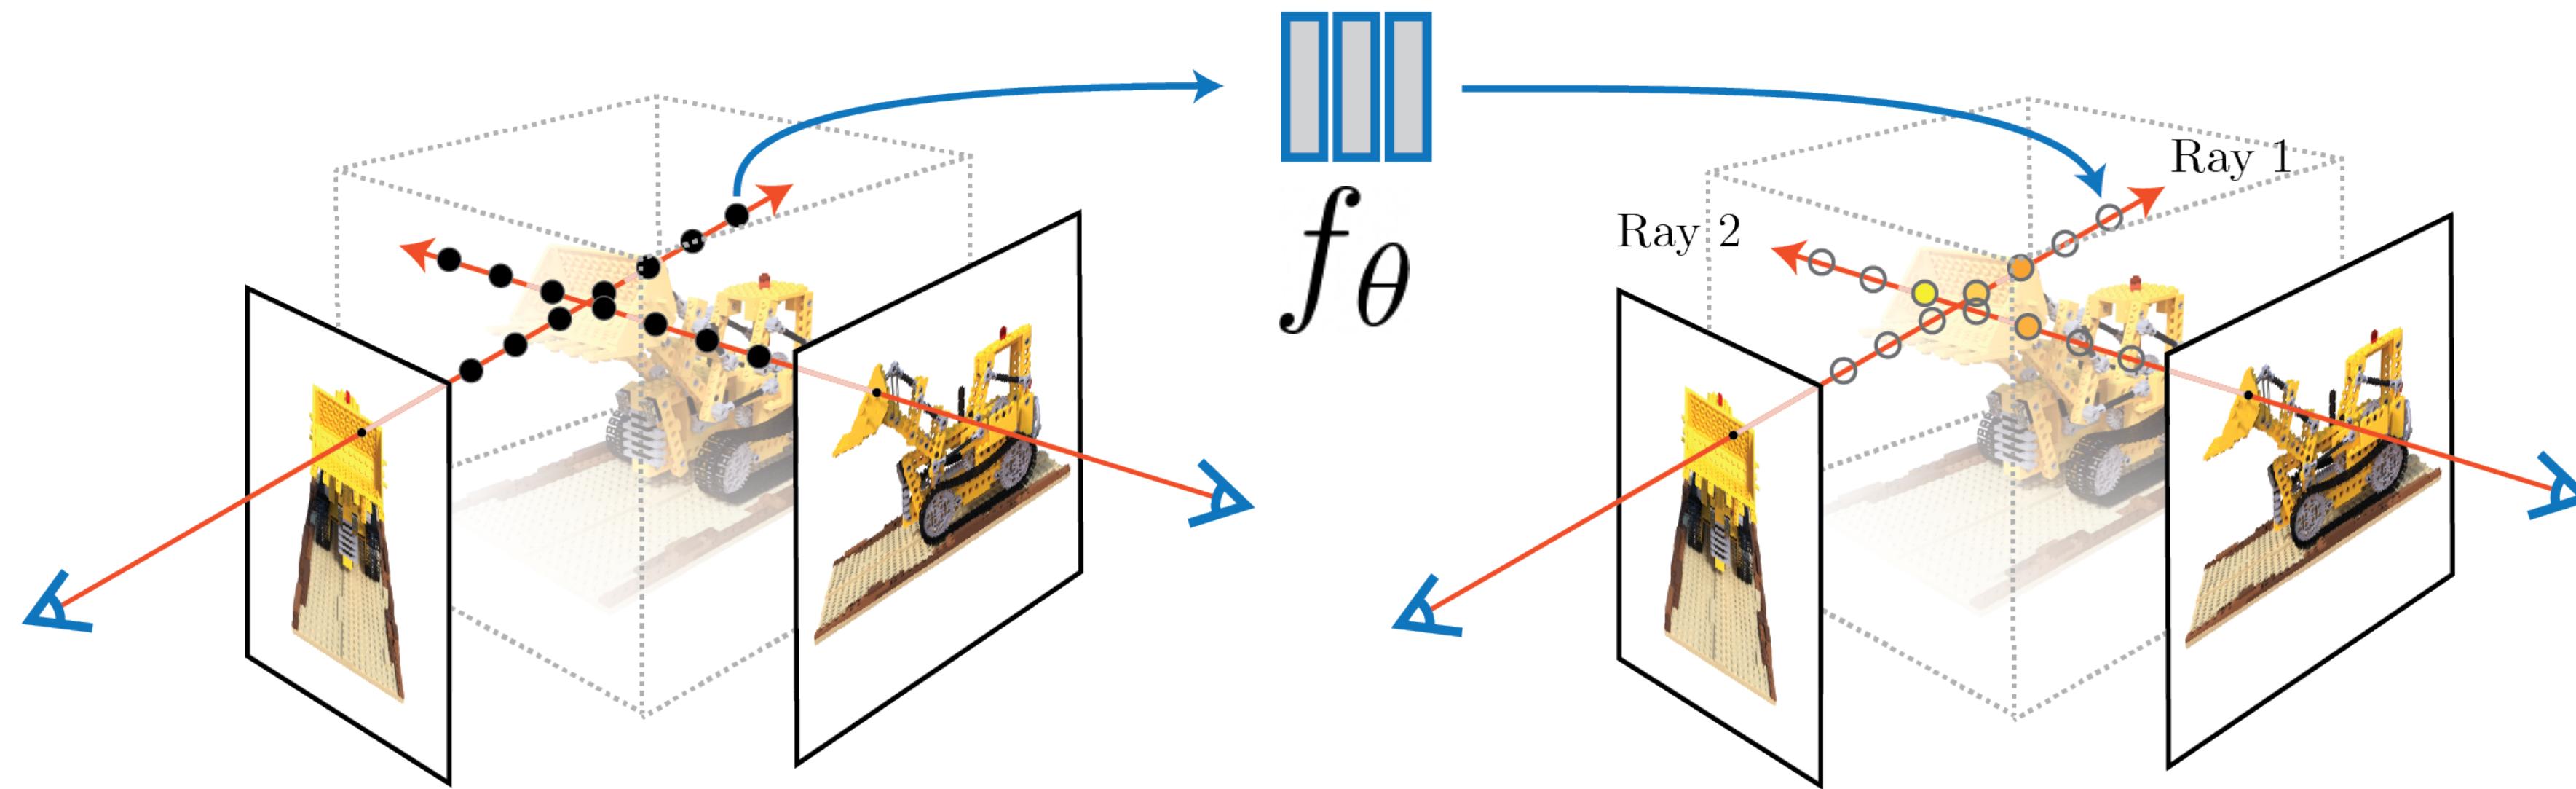


NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis.

Learning Neural Radiance Fields

$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \left\| \underbrace{\text{render}(\mathbf{p}, \pi; f_{\theta}) - I[\mathbf{p}] \|^2}_{\text{volume rendering}} \right\|$$

(pixel, camera) \rightarrow ray

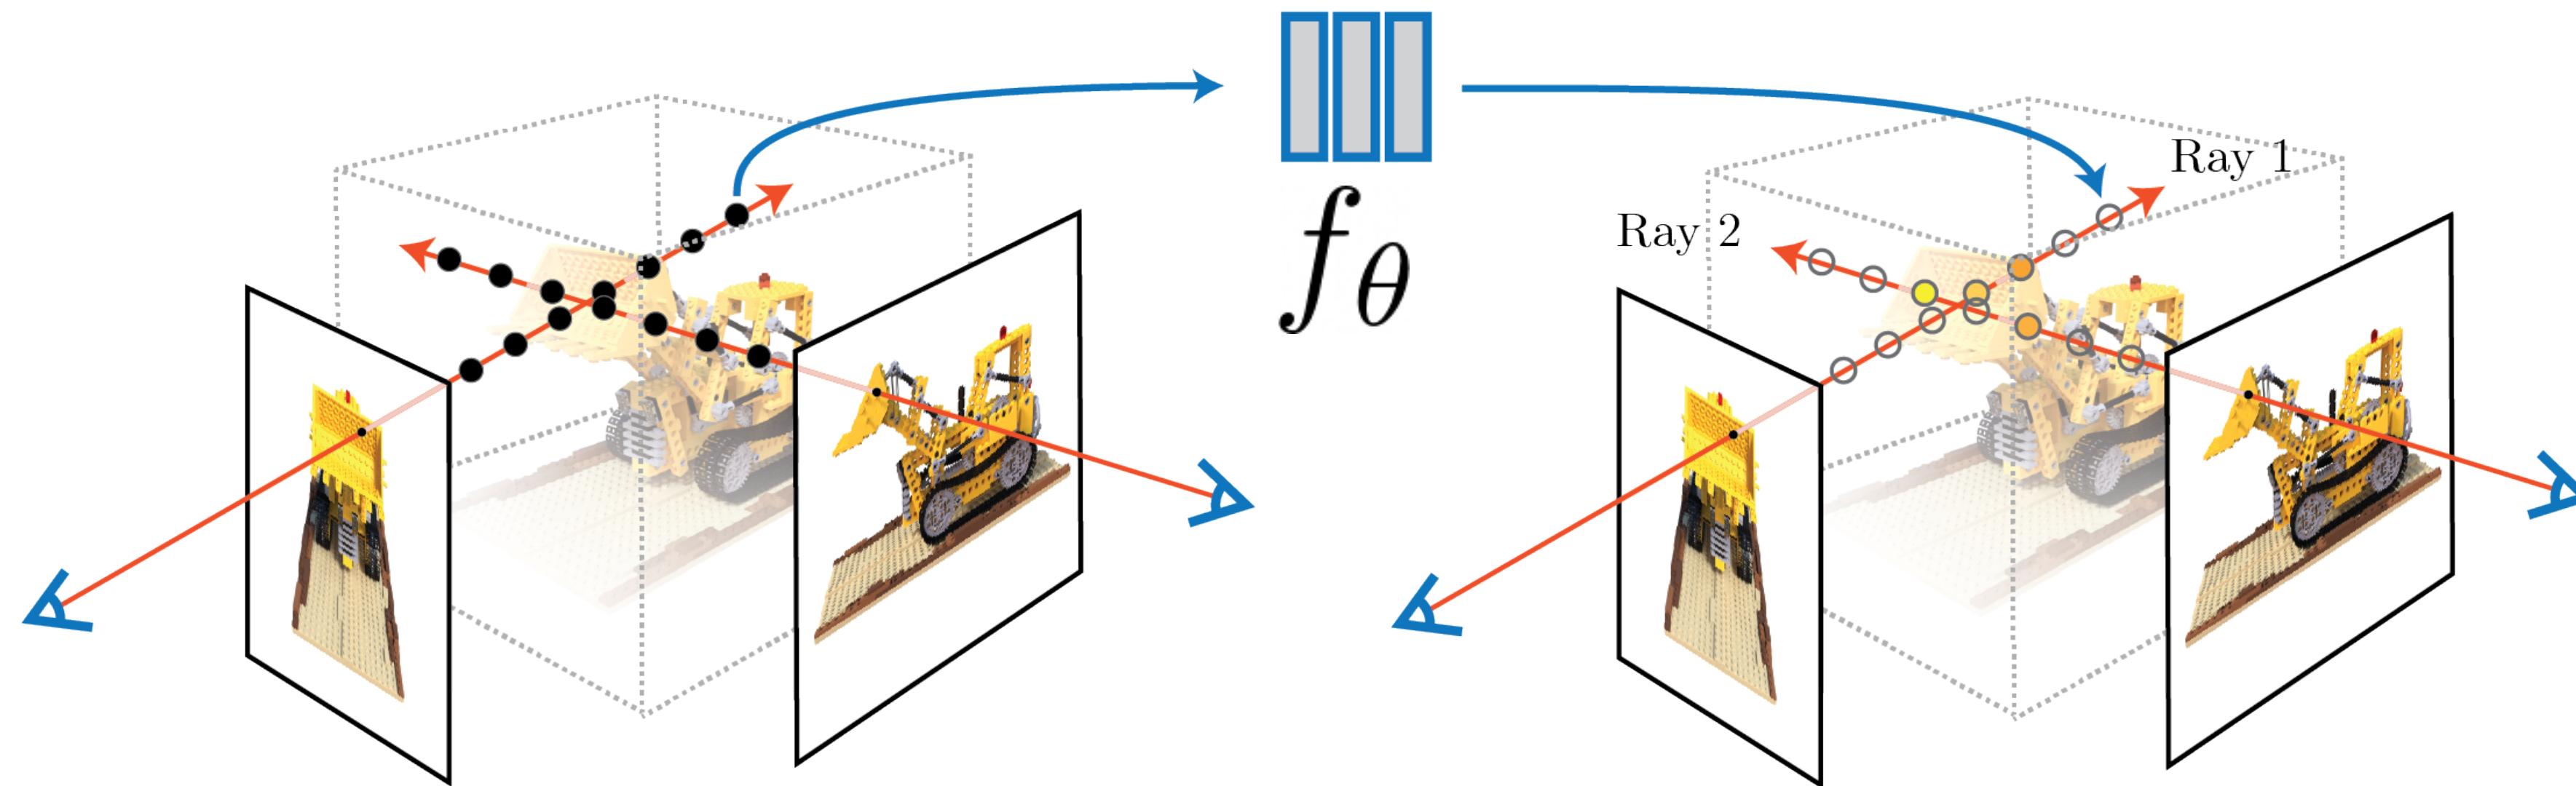


Learning Neural Radiance Fields

$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \|\text{render}(\mathbf{p}, \pi; f_{\theta}) - I[\mathbf{p}]\|^2$$

Typically require multiple images (> 50) with known camera parameters

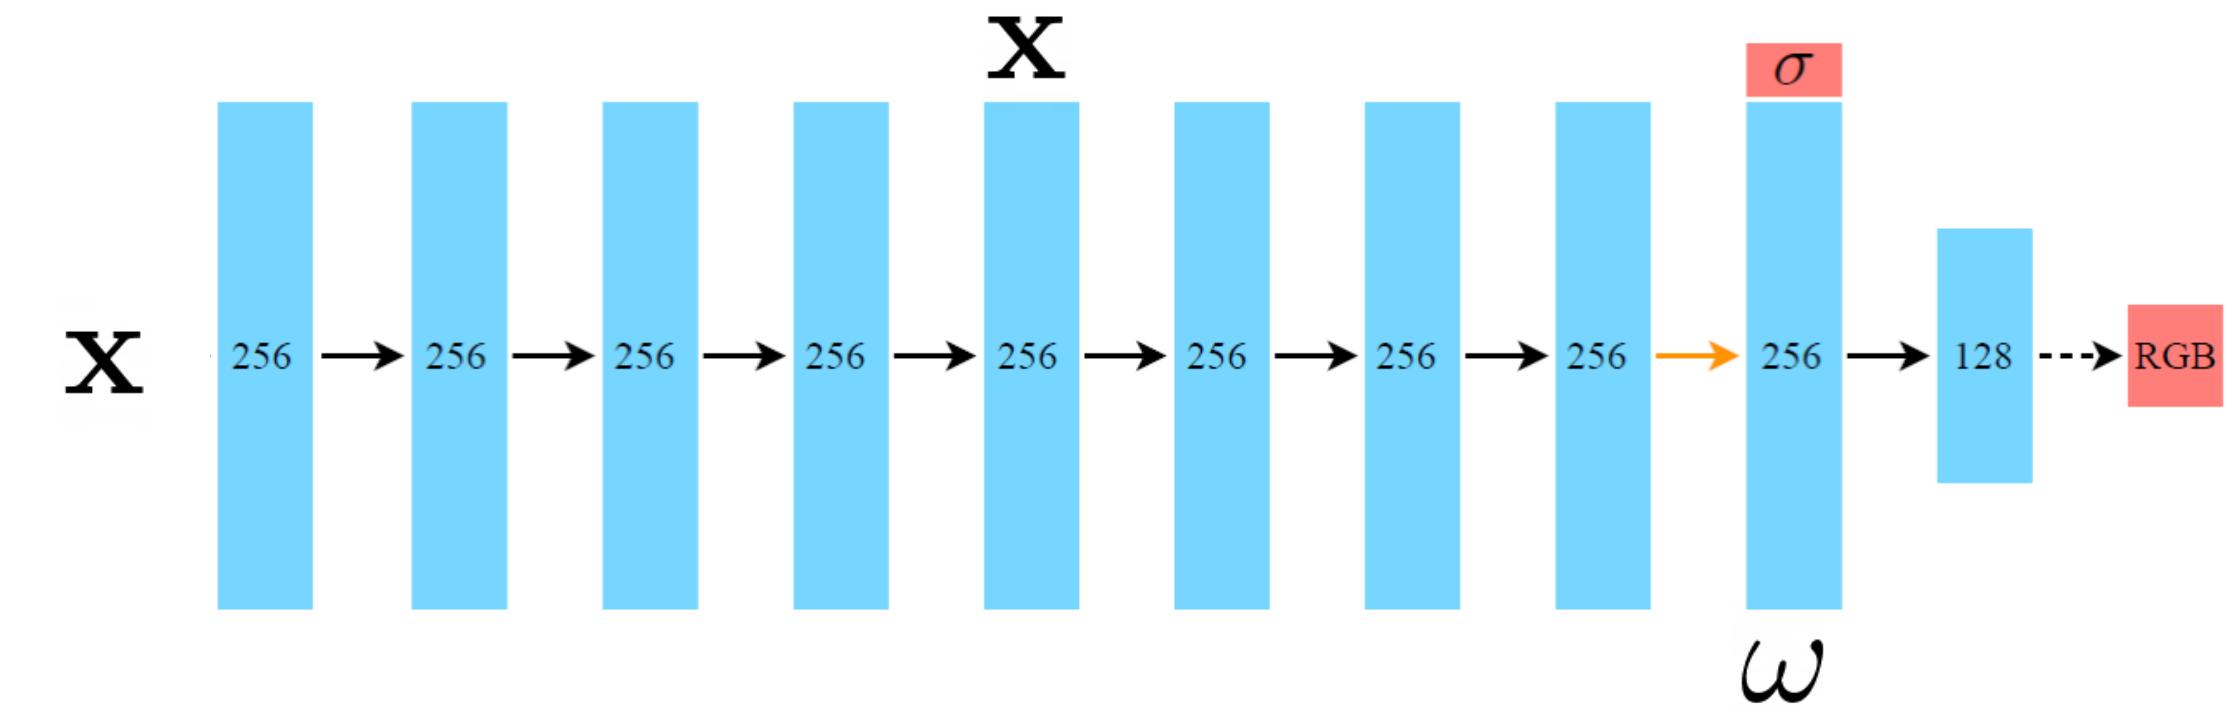
In practice, randomly sample a few pixels (rays) per optimization iteration



Learning Neural Radiance Fields



1) Acquire multiple images of a scene with associated camera viewpoints



2) Design a neural network $(\sigma, \mathbf{c}) = f_\theta(\mathbf{x}, \omega)$

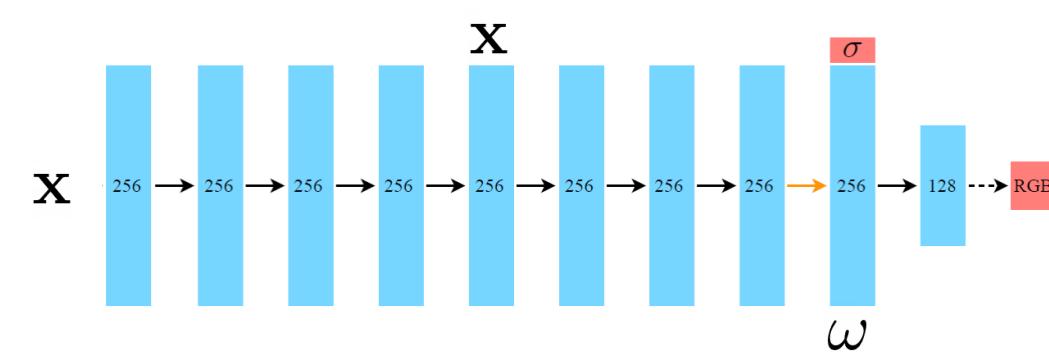
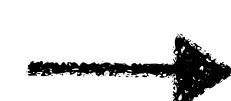
$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \|\text{render}(\mathbf{p}, \pi; f_\theta) - I[\mathbf{p}]\|^2$$

3) Train with a view-synthesis loss using volume rendering

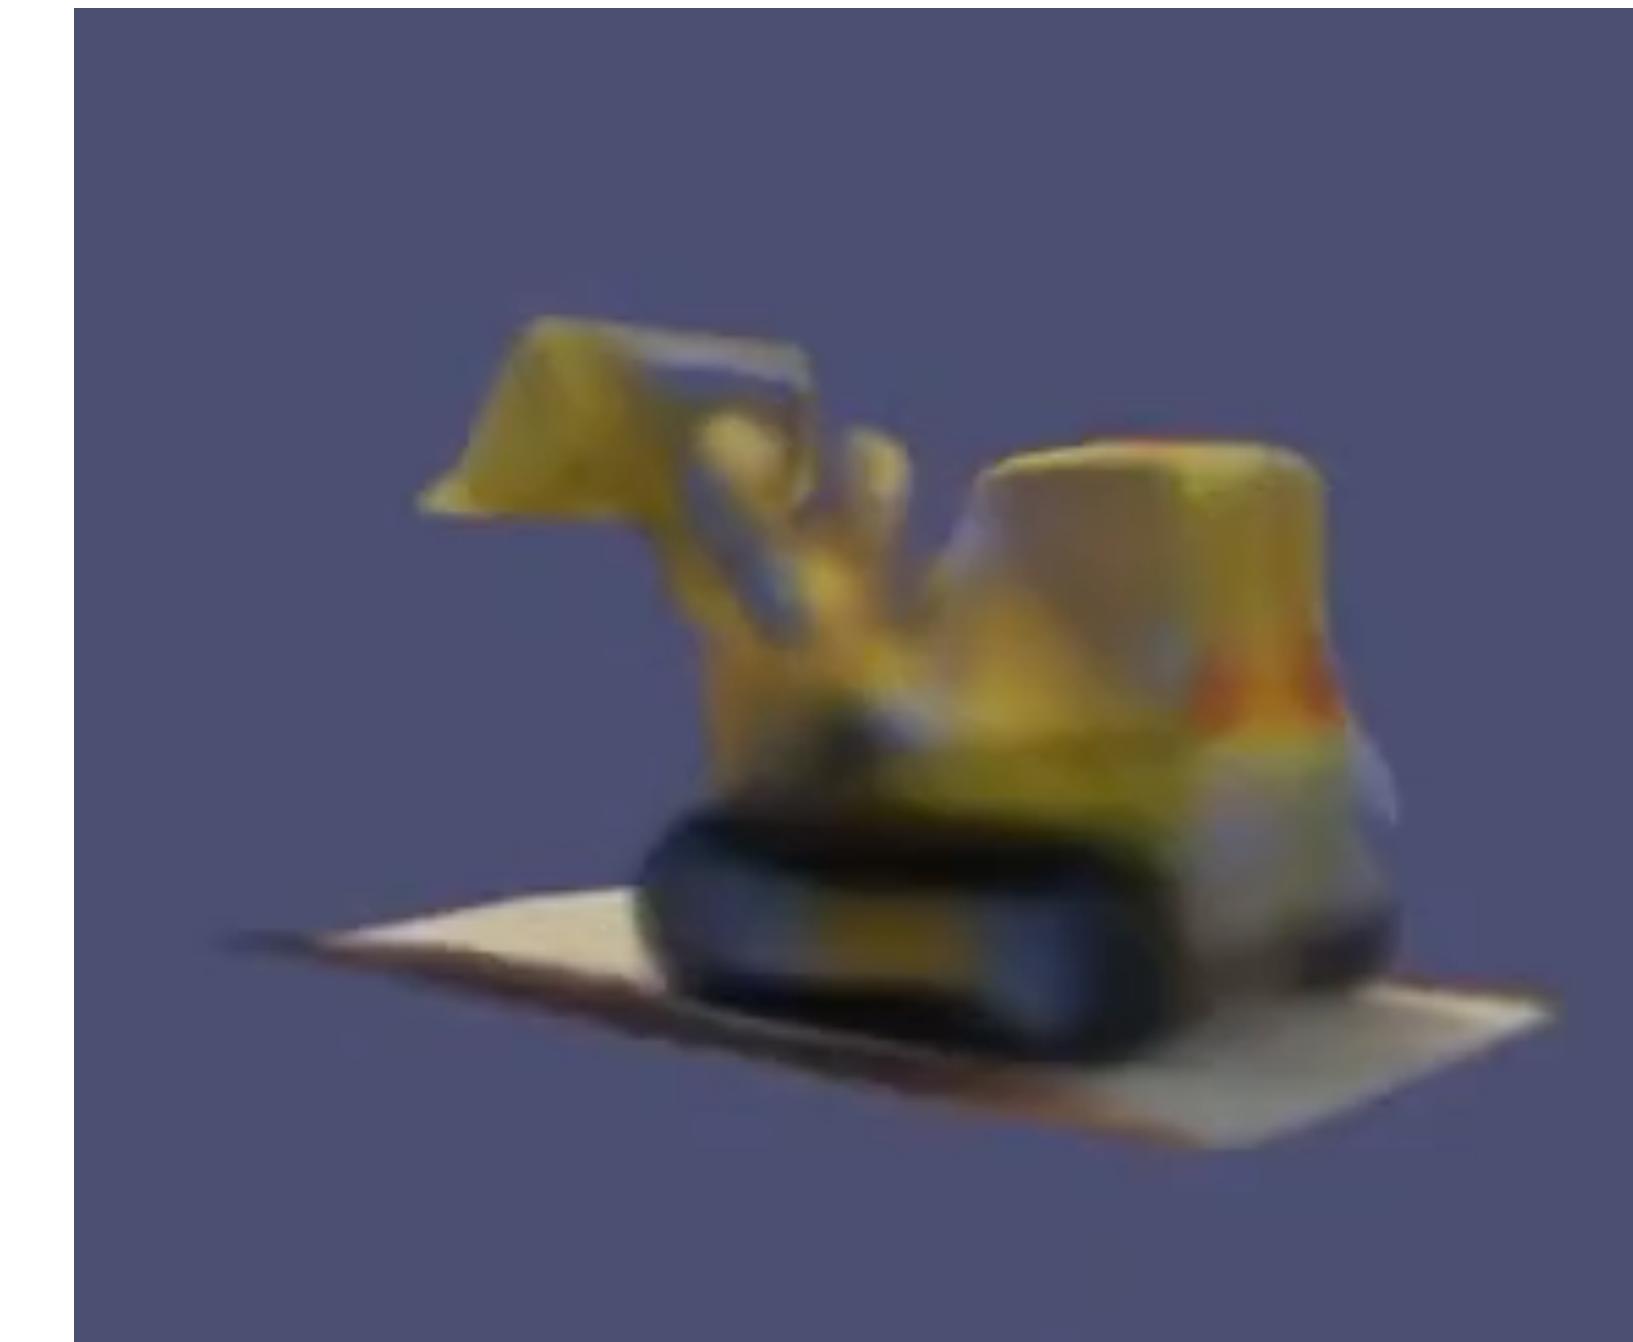
Learning Neural Radiance Fields: An Initial Attempt



100 training images



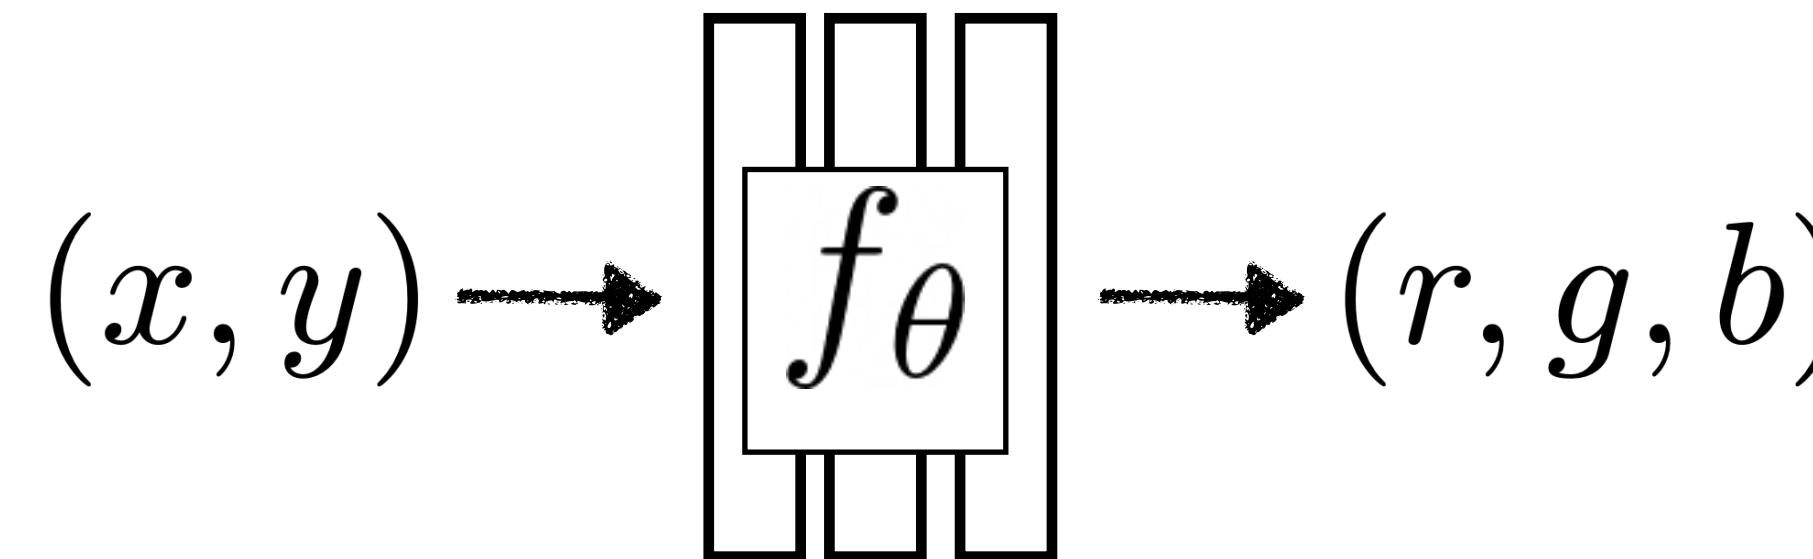
Optimized Neural Net



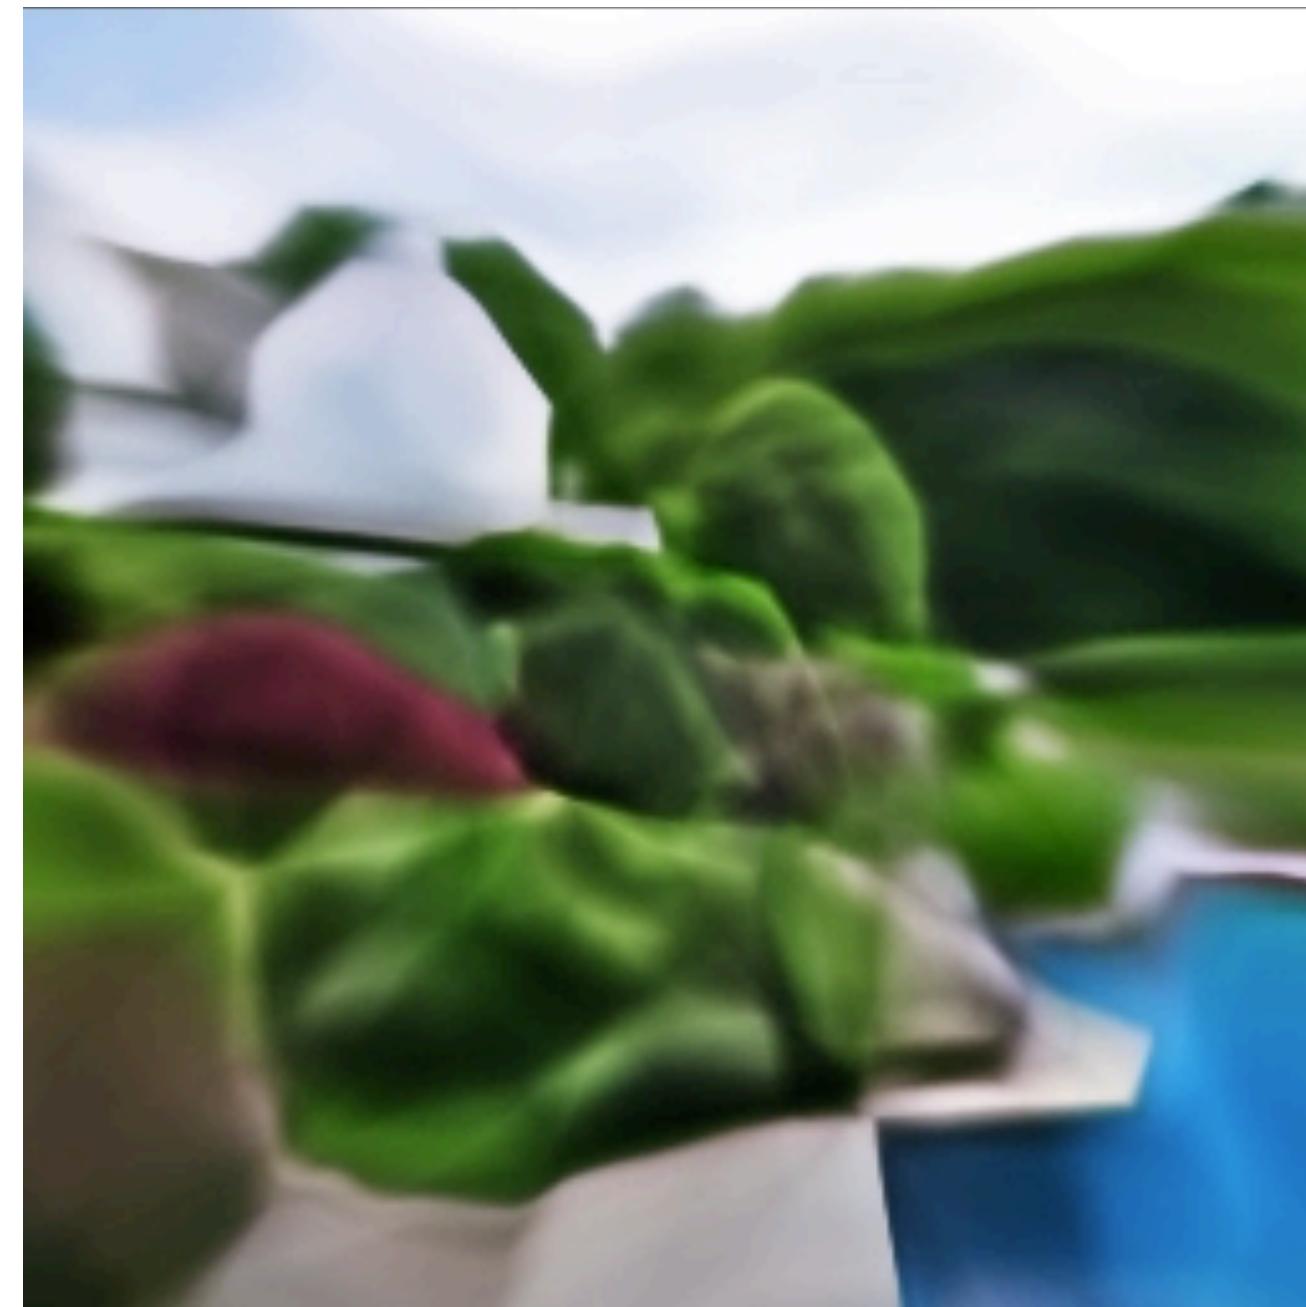
Novel-view Renderings

Cannot capture high-frequency details

Modeling High Frequencies: A Toy Problem



Input Image



Using a ‘standard’ MLP

v

$$\sin(\mathbf{v}), \cos(\mathbf{v})$$

$$\sin(2\mathbf{v}), \cos(2\mathbf{v})$$

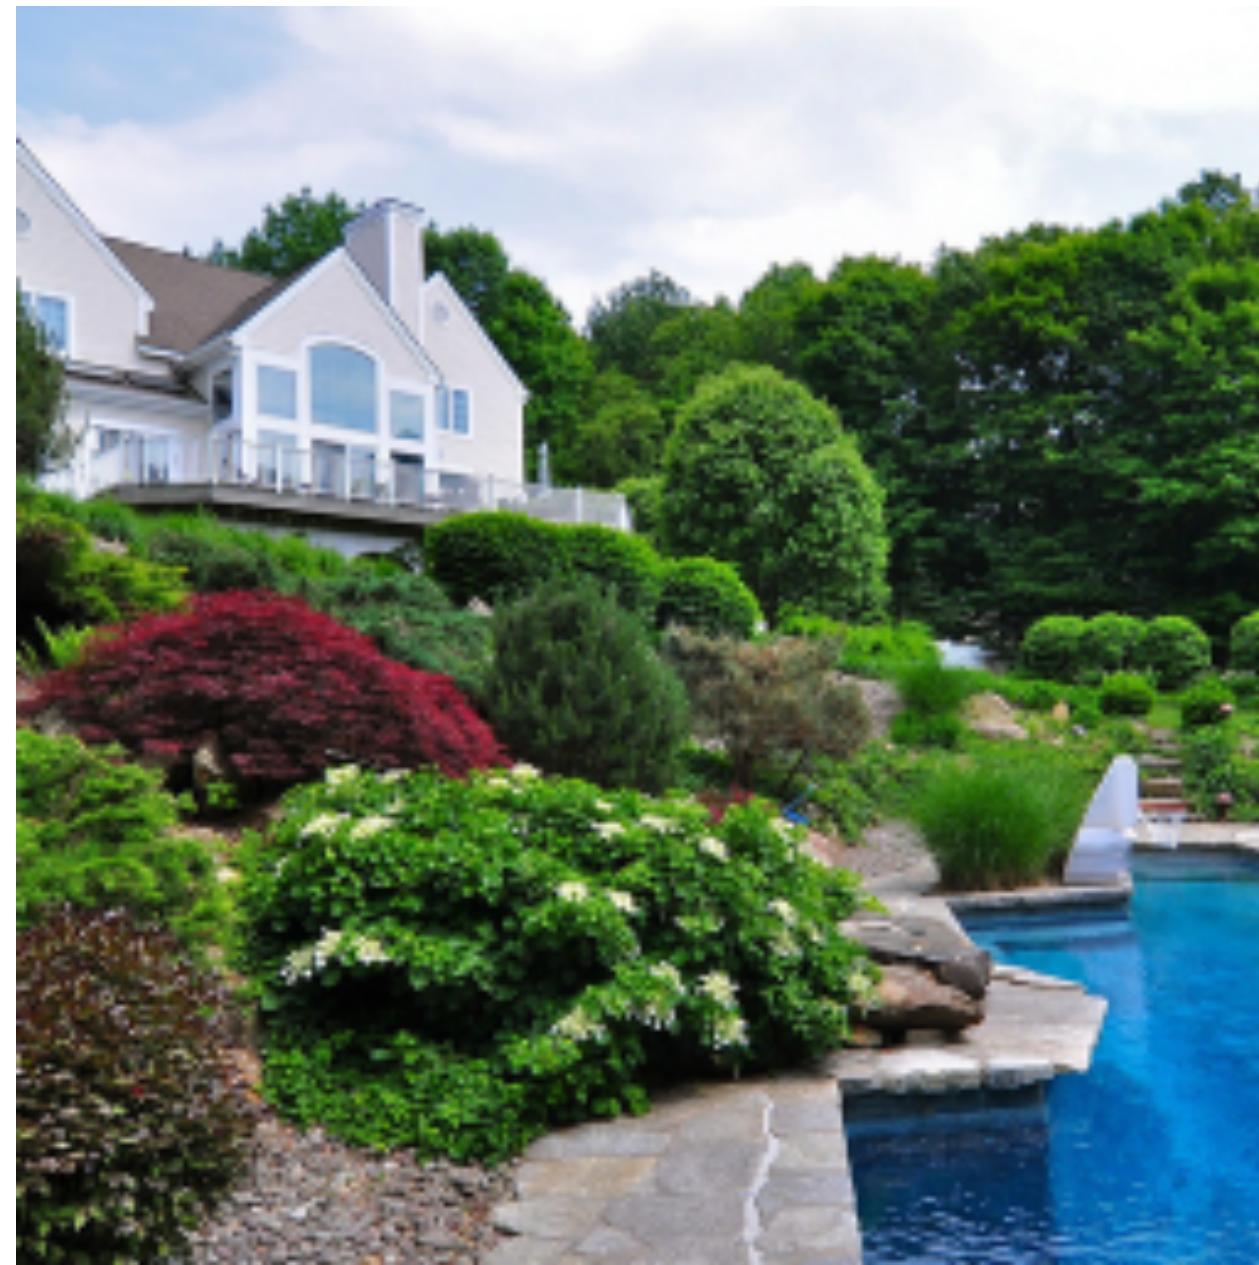
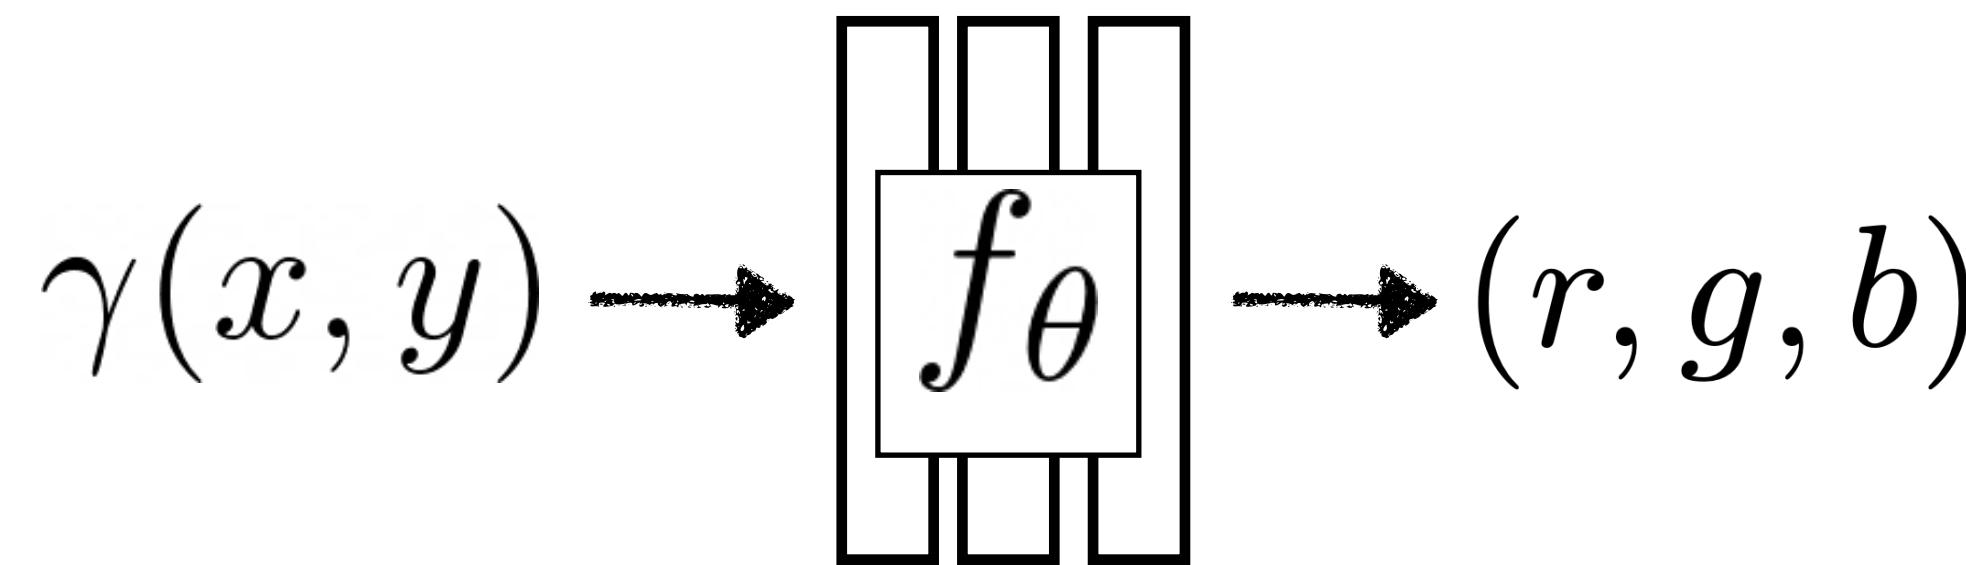
$$\sin(4\mathbf{v}), \cos(4\mathbf{v})$$

...

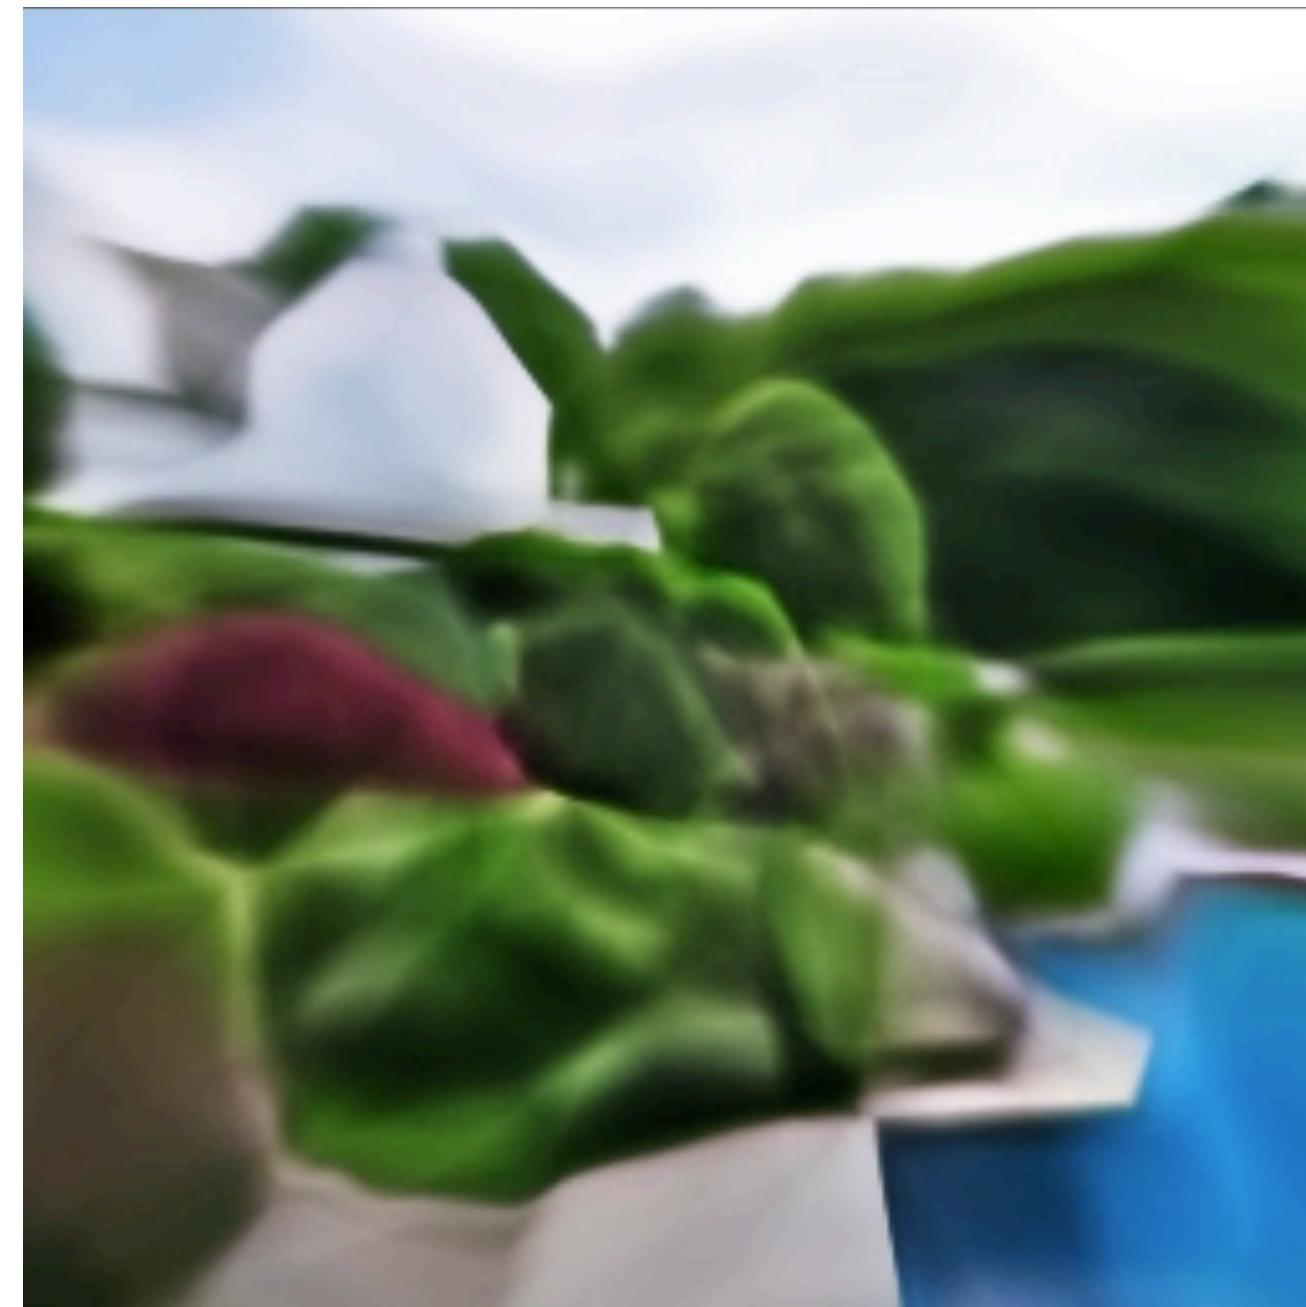
$$\sin(2^{L-1}\mathbf{v}), \cos(2^{L-1}\mathbf{v})$$

$$\gamma(\mathbf{v})$$

Modeling High Frequencies: A Toy Problem



Input Image



Using a 'standard' MLP

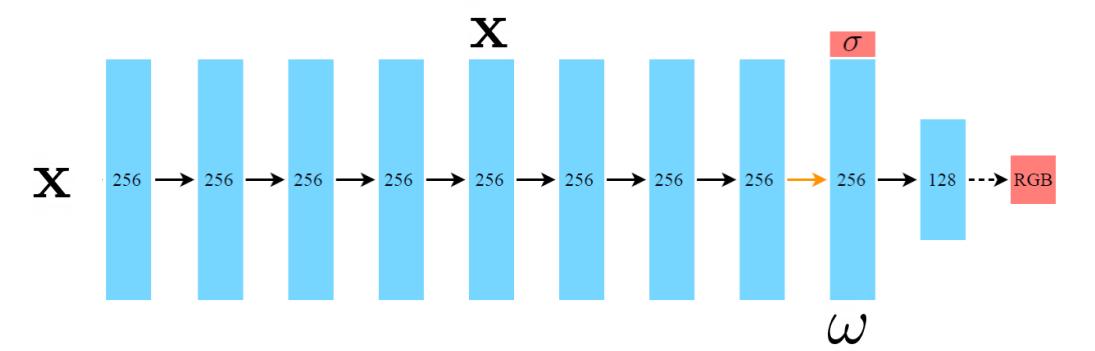


Using position encoding

Learning Neural Radiance Fields



100 training images



(using position encodings in input)

Optimized Neural Net



Novel-view Renderings

A great example that ‘execution matters’

Learning Neural Radiance Fields



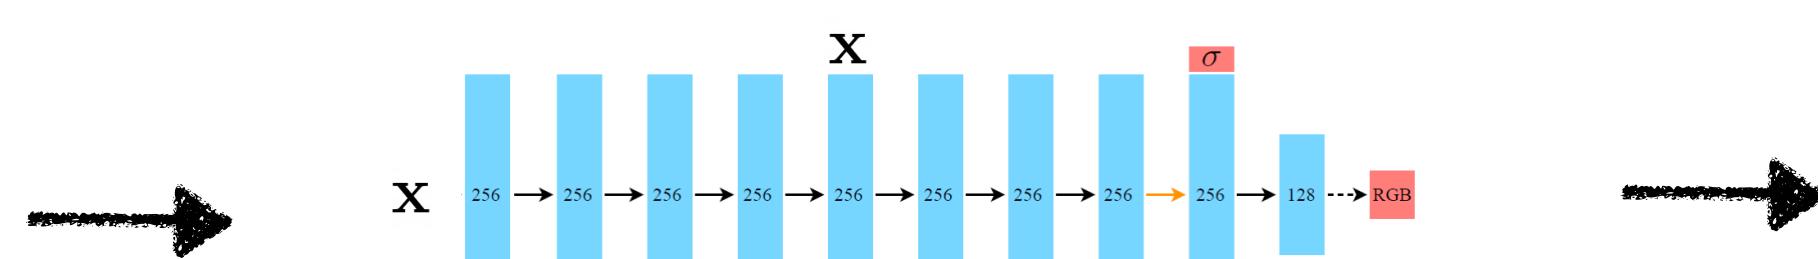
NeRF (Naive)



NeRF (with positional encoding)



Learning Neural Radiance Fields



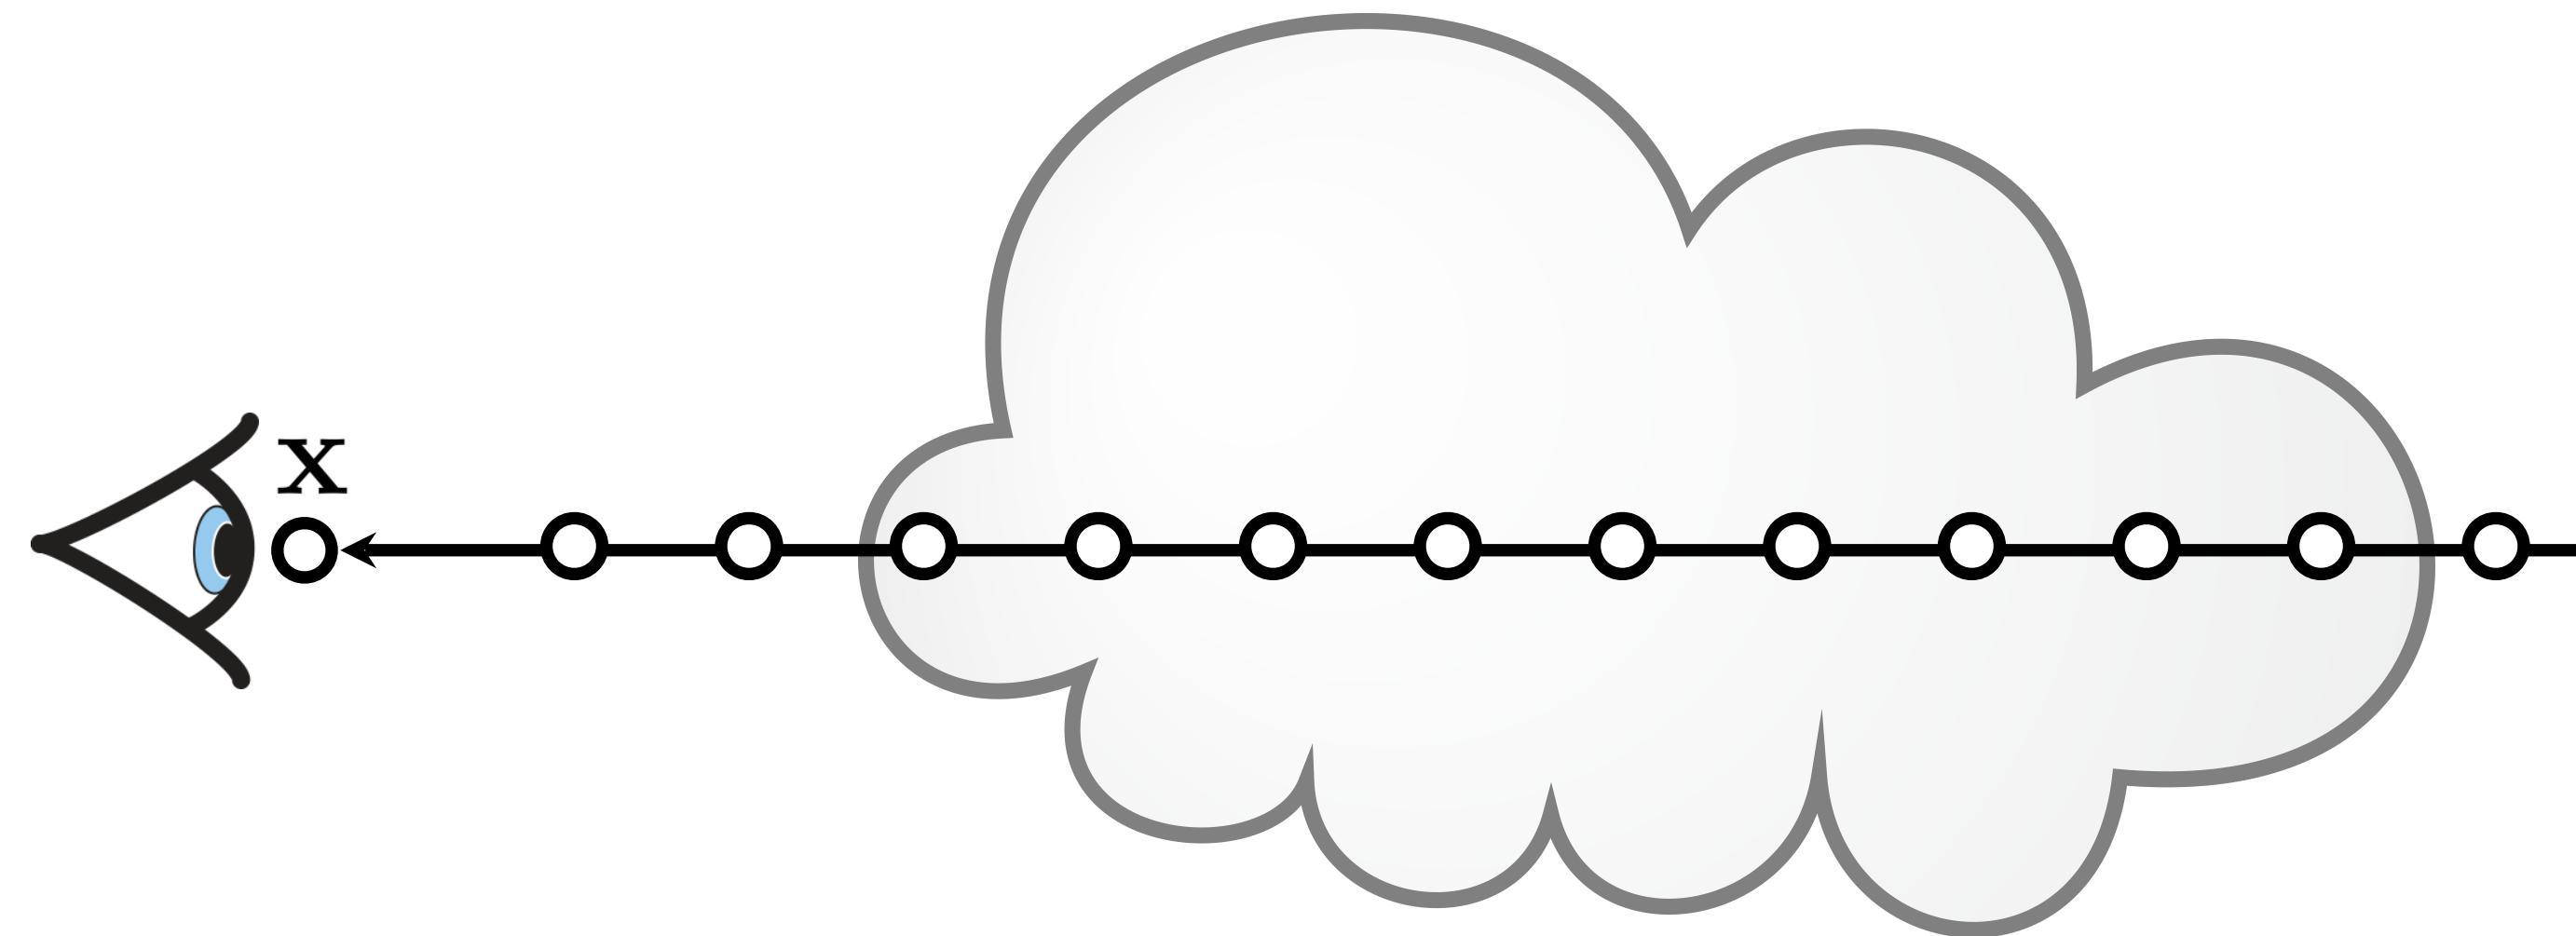
(using position encodings in input)



A very ‘simple’ representation and learning approach

Has inspired **many, many** extensions and followups

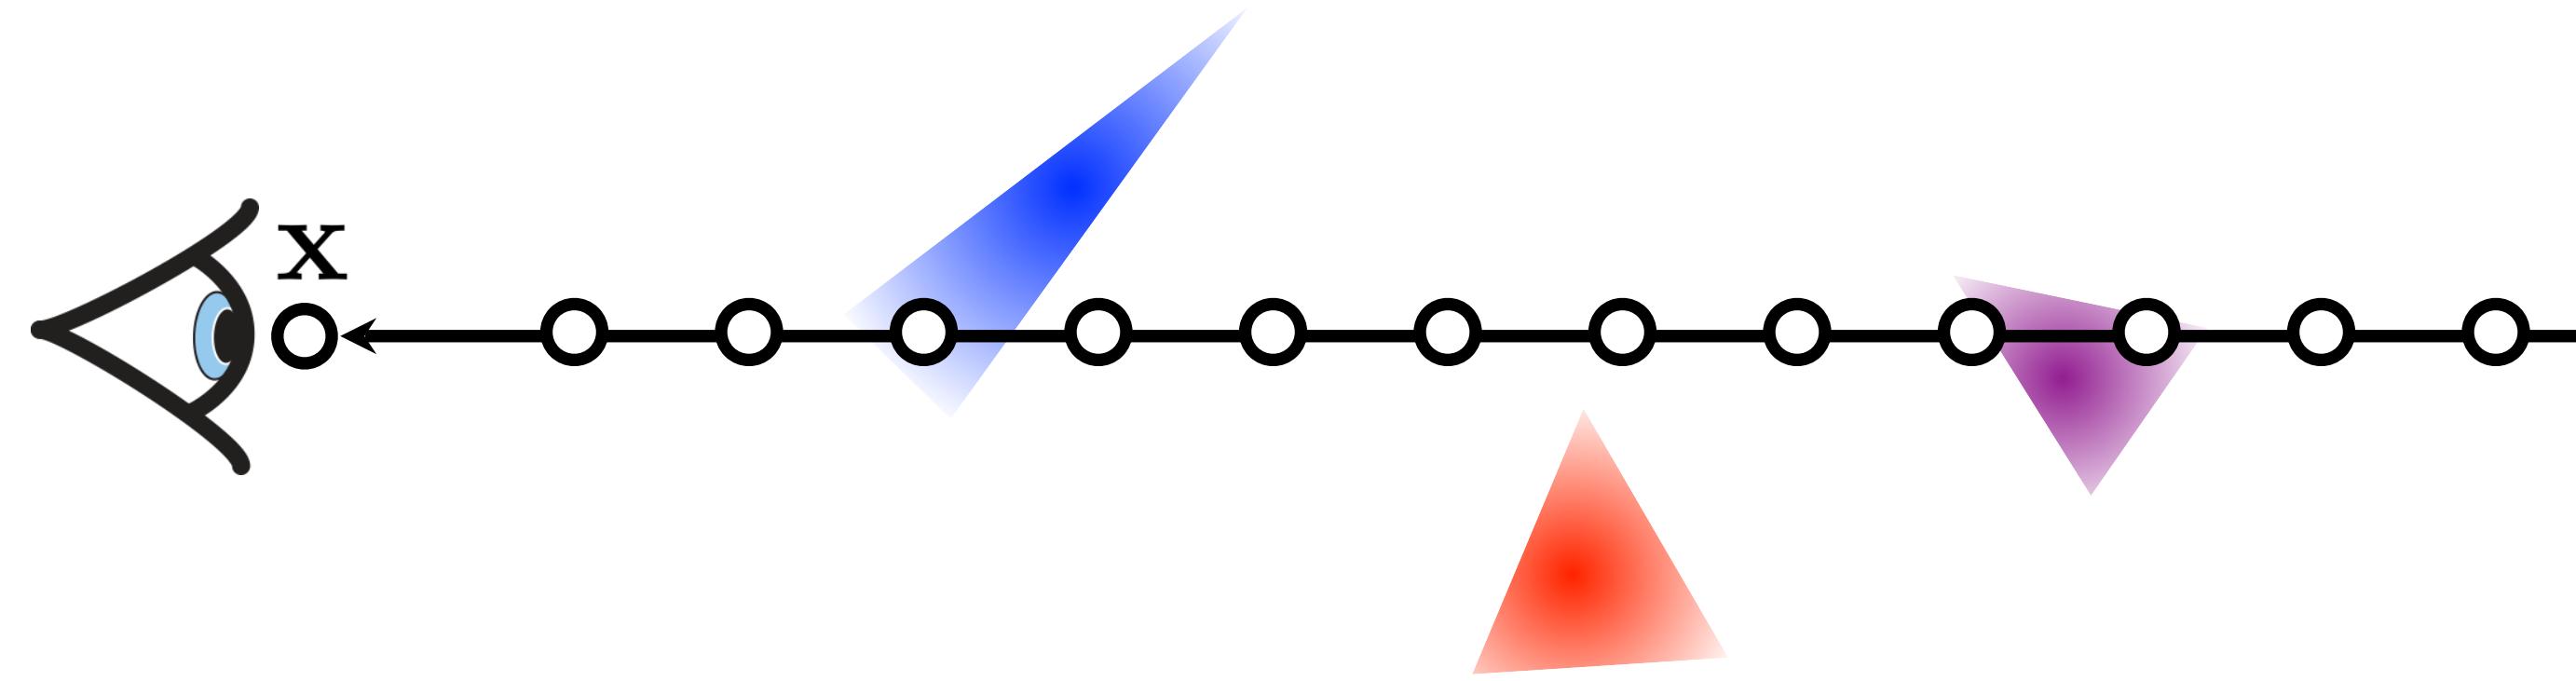
Recap: Rendering Volumes



1. Draw samples along the ray

2. Aggregate their contributions to render

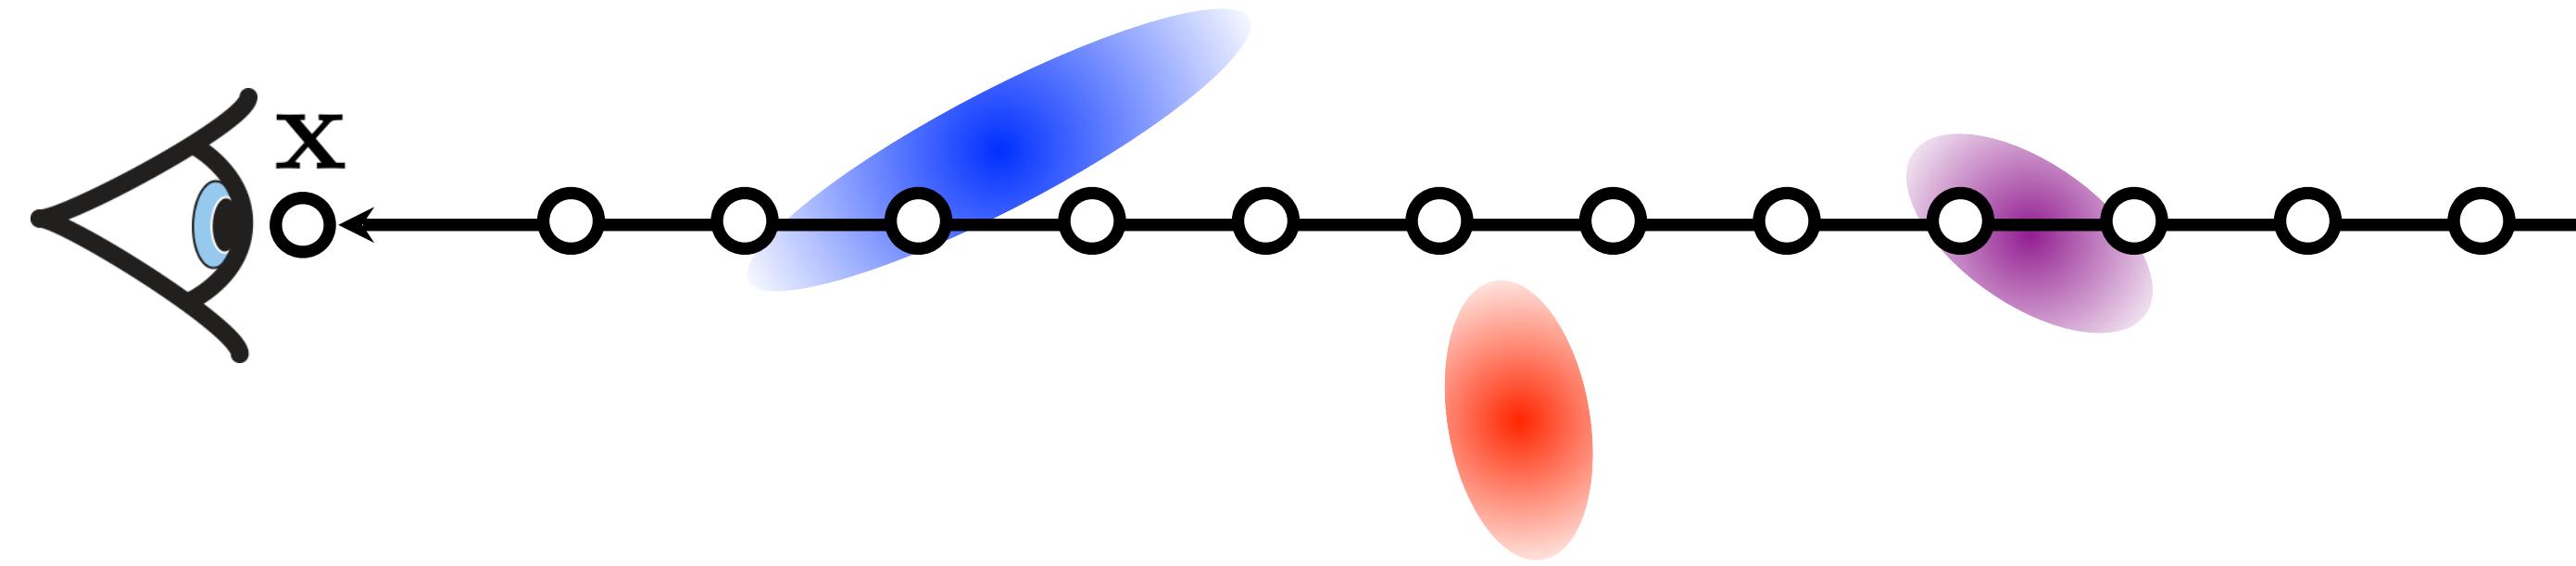
Rendering Primitives (e.g. Triangles/Meshes)



1. Draw samples along the ray

2. Aggregate their contributions to render

Rendering Primitives (e.g. Gaussians)

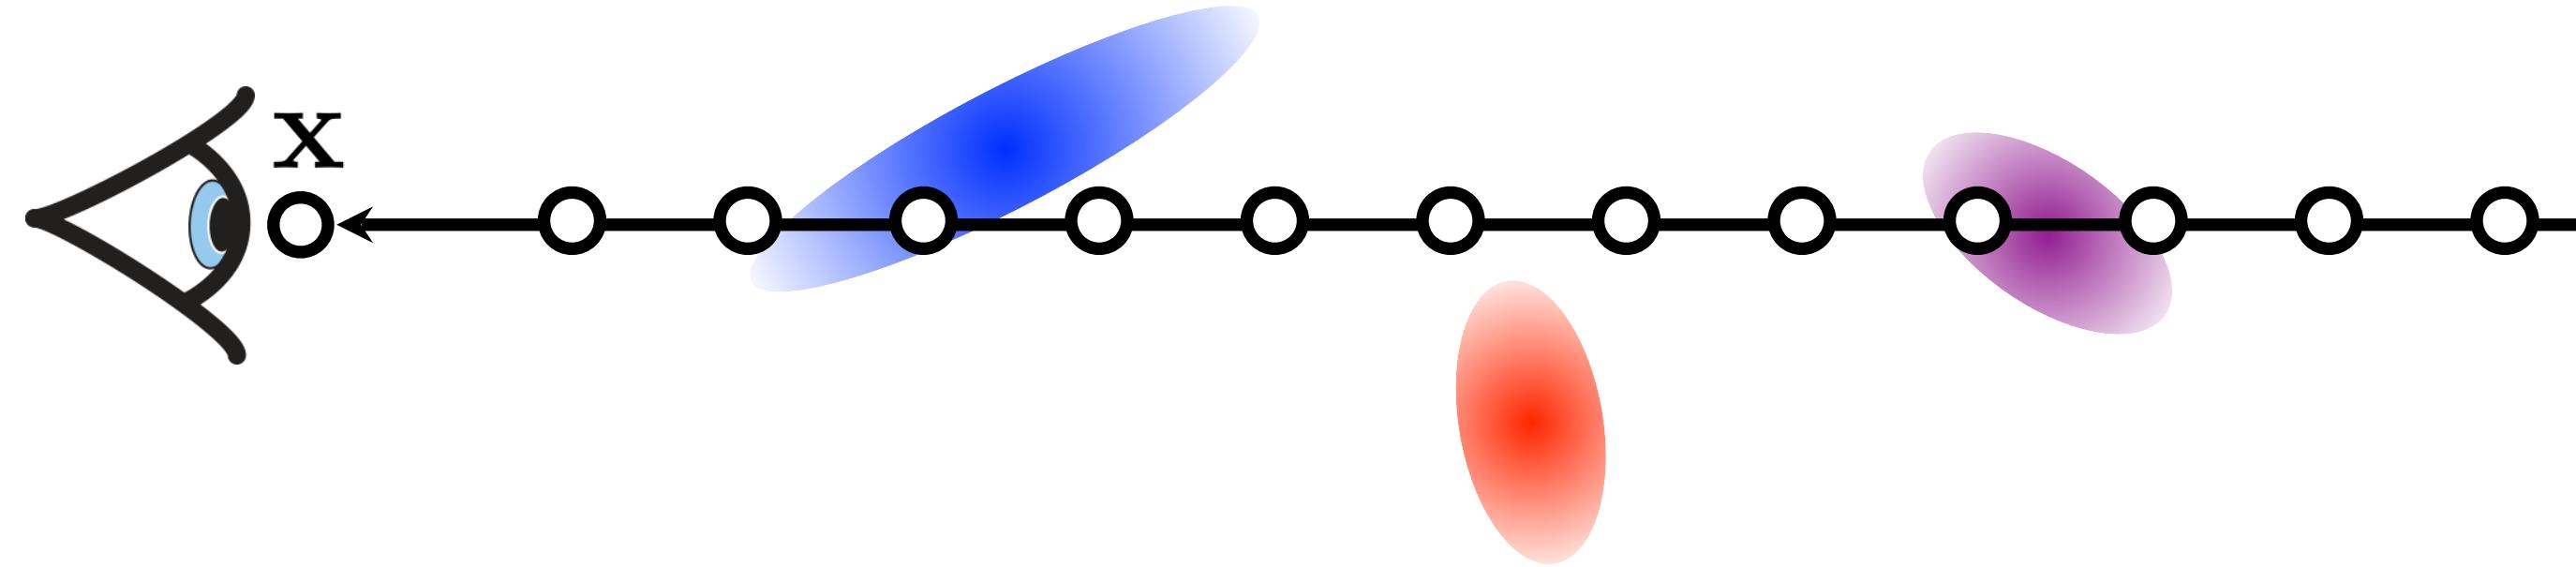


~~1. Draw samples along the ray~~

(wasteful – we know where the primitives are!)

2. Aggregate their contributions to render

Rendering Primitives (e.g. Gaussians)



1. Find primitives that affect the ray/pixel
(can lead to significant speedups in rendering)
2. Aggregate their contributions to render

Differentiable Primitive Rendering

Rasterization

{

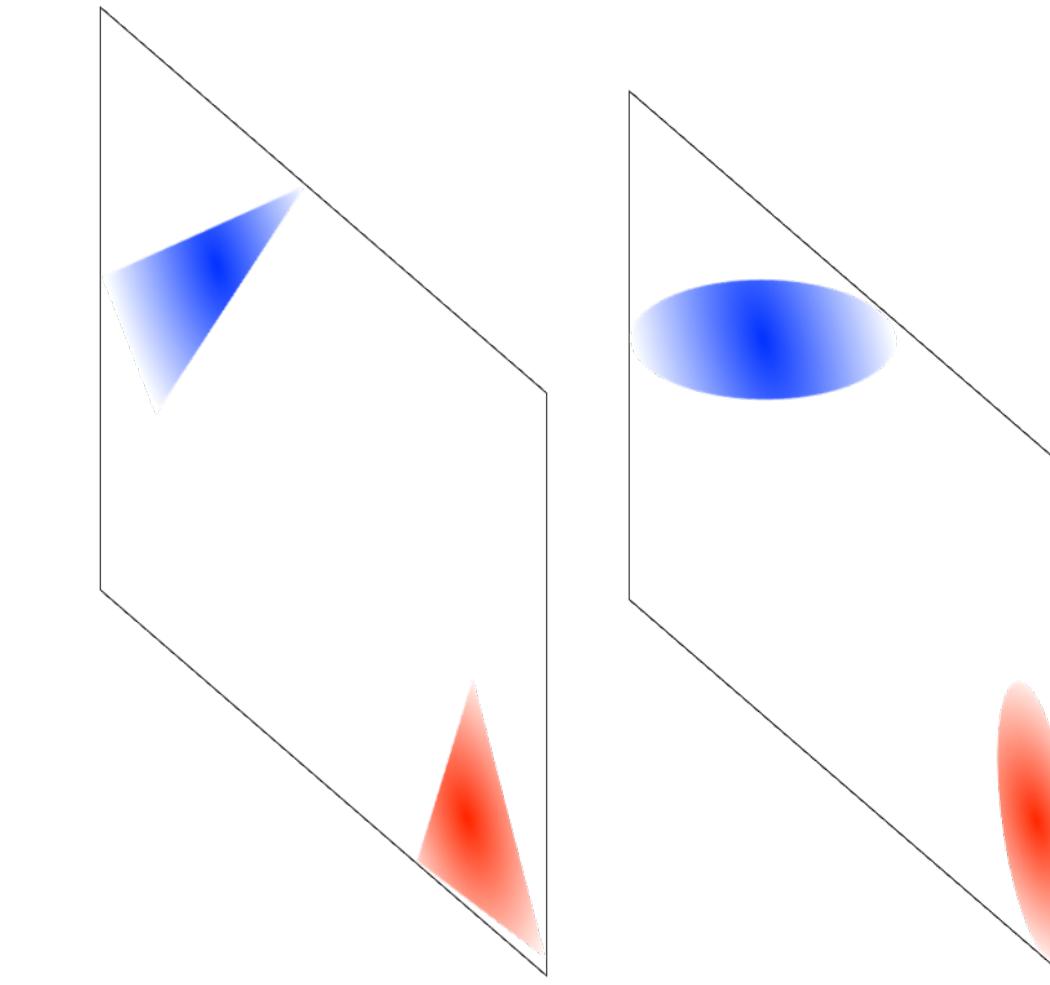
```
def render(primitives, camera):
    """
    # Initialize a rasterization data structure
    # (records influencing primitives for each pixel)
    """

    for primitive in primitives:
        prim2d = project(primitive, camera)
        """
        # Update rasterization data structure
        """

    for pixel in camera.grid:
        """
        # Aggregate appearance from influencing primitives
        """
```

Blending

{



Differentiable Gaussian Rendering

**What is the representation
of a 3D gaussian?**

**How to project to 2D
and rasterize?**

**How to model/aggregate
appearance?**

Rasterization

{

```
def render(gaussians, camera):
    """
    # Initialize a rasterization data structure
    # (records influencing primitives for each pixel)
    """

    for gaussian in gaussians:
        gauss2d = project(gaussian, camera)
        """
        # Update rasterization data structure
        """
```

Blending

{

```
for pixel in camera.grid:
    """
    # Aggregate appearance from influencing gaussians
    """
```

Differentiable Gaussian Rendering

**What is the representation
of a 3D gaussian?**



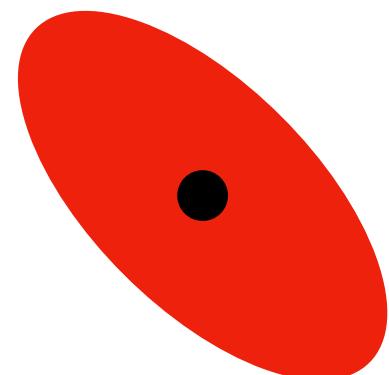
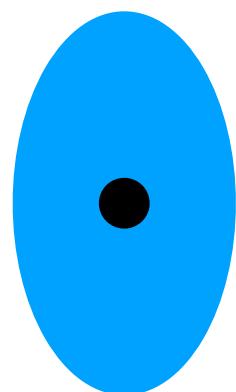
**How to project to 2D
and rasterize?**



**How to model/aggregate
appearance?**

Differentiable Gaussian Rendering

**What is the representation
of a 3D gaussian?**



**How to project to 2D
and rasterize?**

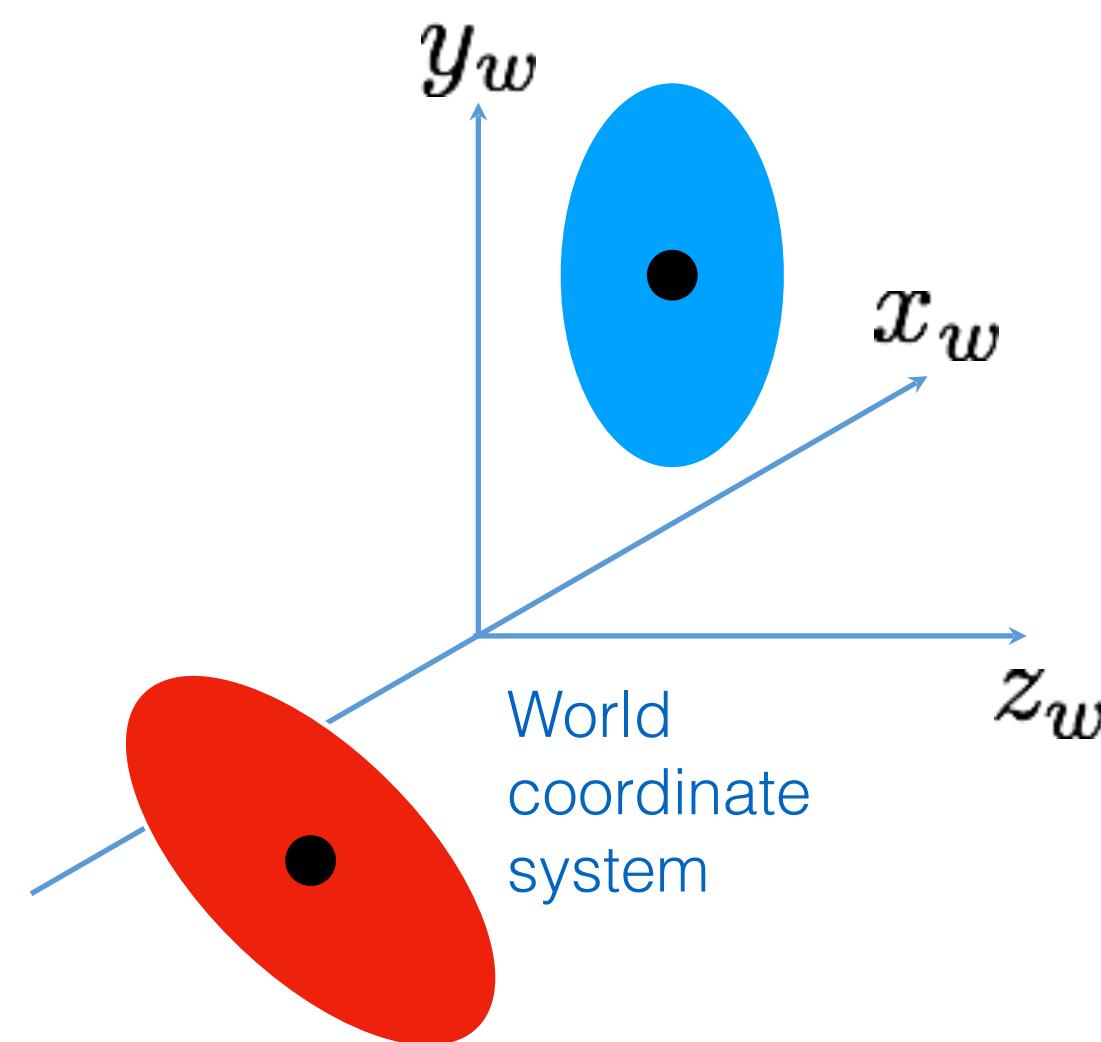
$$\mathcal{G}_{\mathbf{V}}(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|\mathbf{V}|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T \mathbf{V}^{-1} (\mathbf{x}-\mathbf{p})}$$

Factorize as scale and rotation: $\mathbf{V} = RSS^T R^T$

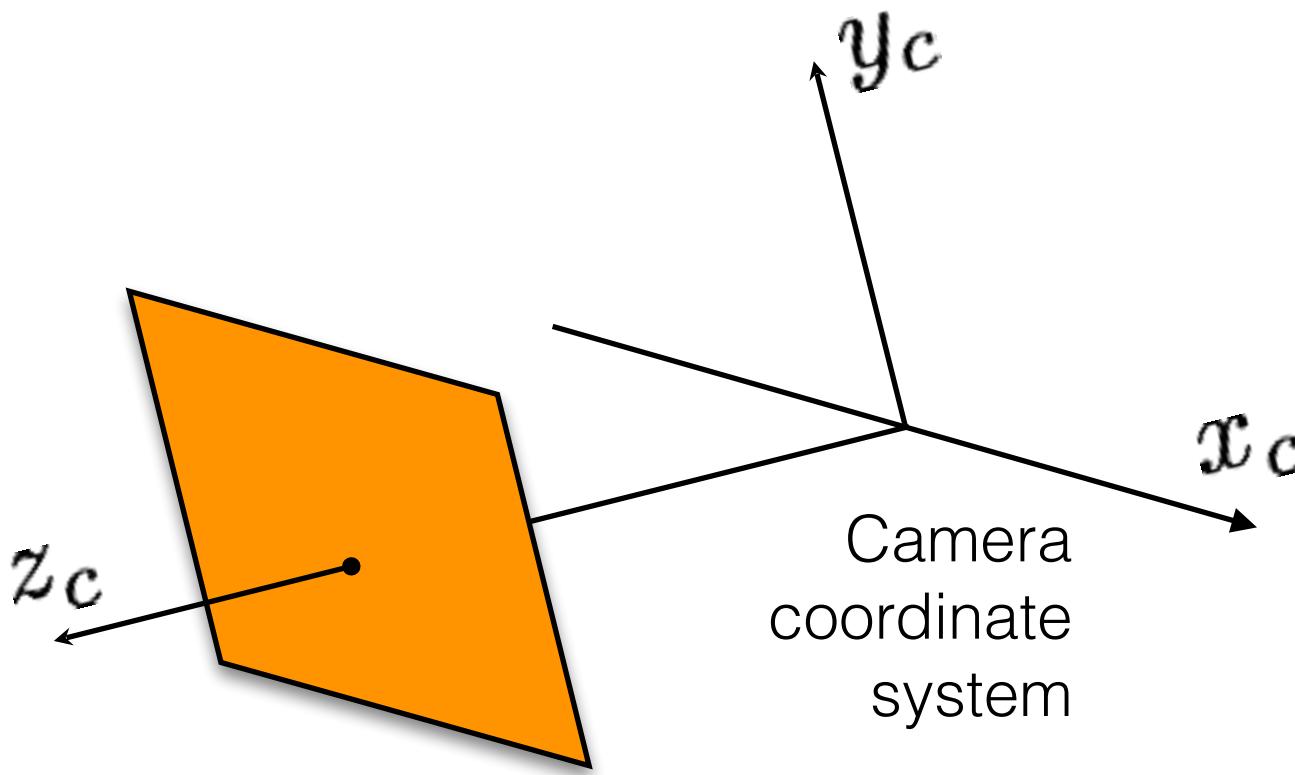
Each gaussian also has an opacity and view-dependent color
(via SH coefficients): α, \mathbf{c}

Differentiable Gaussian Rendering

What is the representation
of a 3D gaussian?



$$\mathbf{p}, \mathbf{R}, \mathbf{S}$$



$$\mathbf{p}', \mathbf{R}', \mathbf{S}'$$

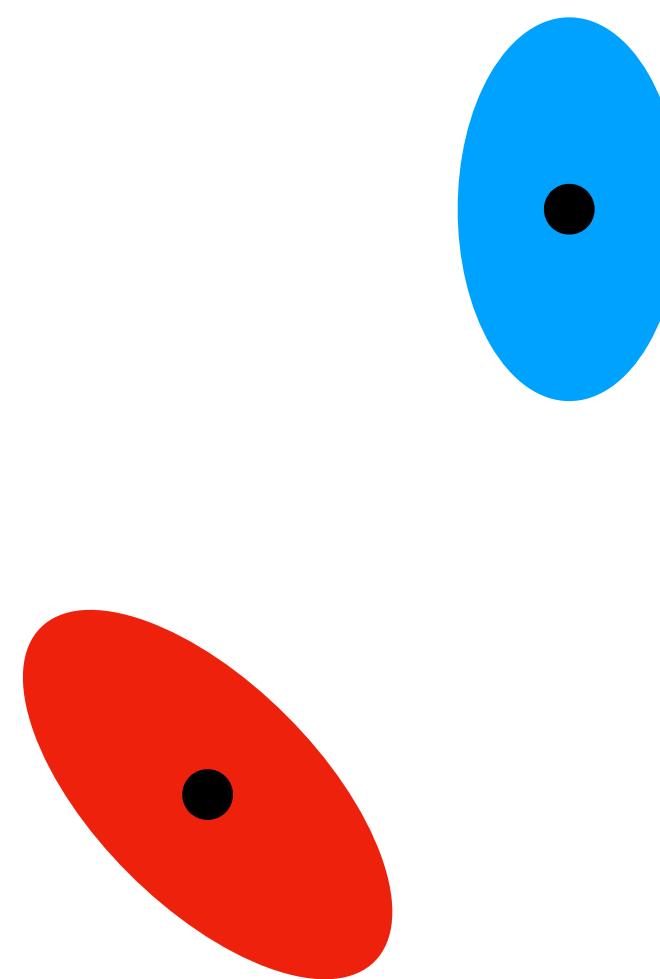
How to project to 2D
and rasterize?

How to model/aggregate
appearance?

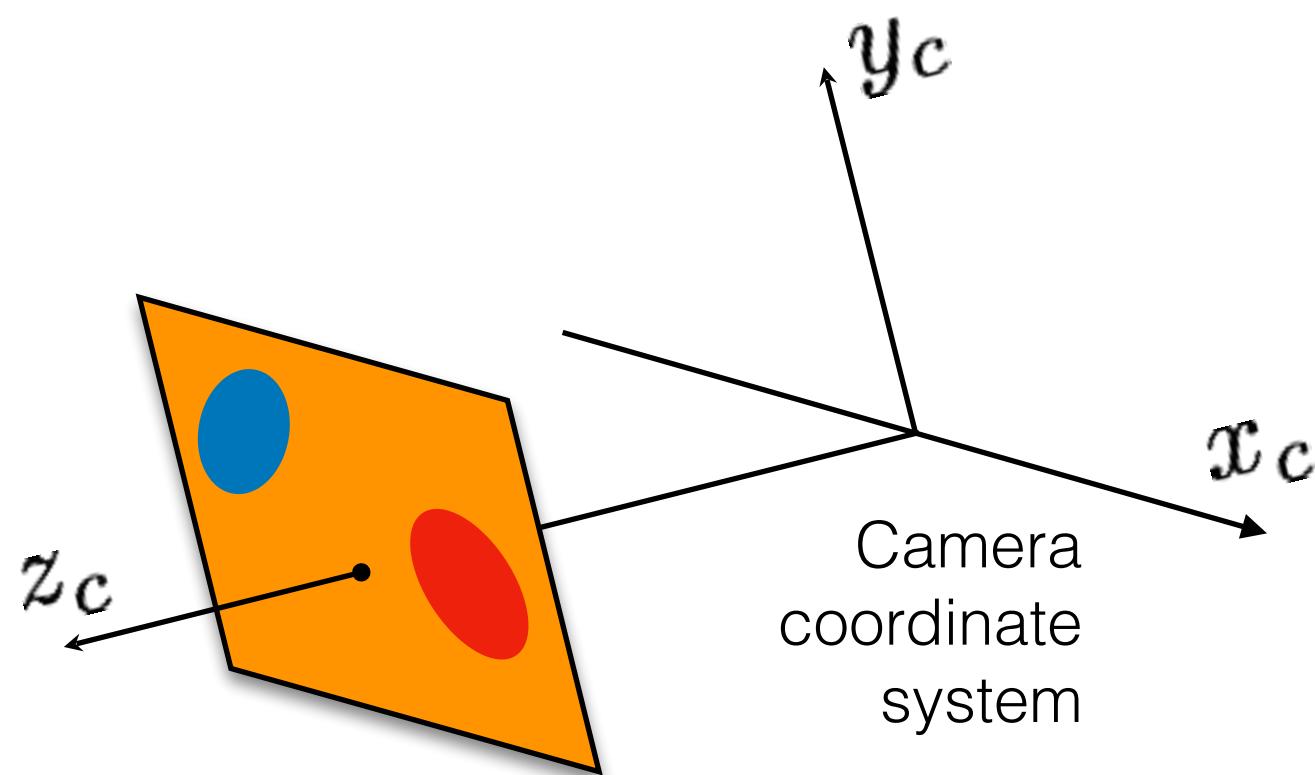
We can use the camera
extrinsics to transform each 3D
gaussian to the camera frame

Differentiable Gaussian Rendering

What is the representation
of a 3D gaussian?



$$\mathbf{p}', \mathbf{R}', S$$



Q: What is the image-space
projection of a 3D gaussian?

A: Can approximate as a 2D gaussian!
(EWA Volume Splatting. Zwicker et. al., 2001)

How to project to 2D
and rasterize?

How to model/aggregate
appearance?

$$\pi(\mathbf{x}) = \mathbf{u}$$

π : Projection function for
mapping 3D points to pixels

$$z \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

2D mean: $\mu_{2D} = \pi(\mu_{3D})$

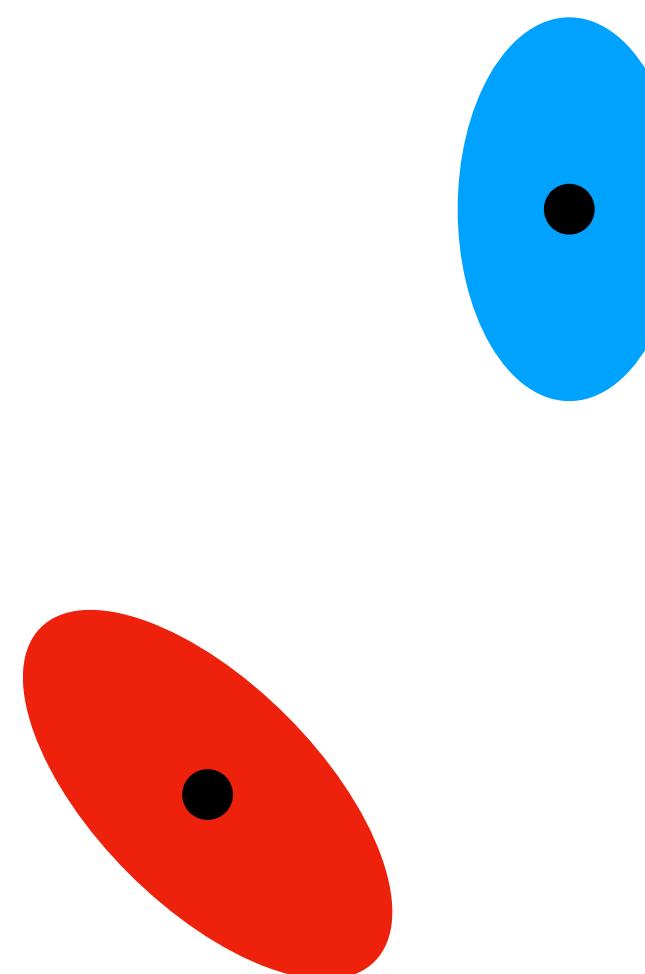
2D covariance:

$$J = \frac{\partial \pi}{\partial \mathbf{x}}(\mu_{3D})$$

$$\Sigma_{2D} = J \Sigma_{3D} J^T$$

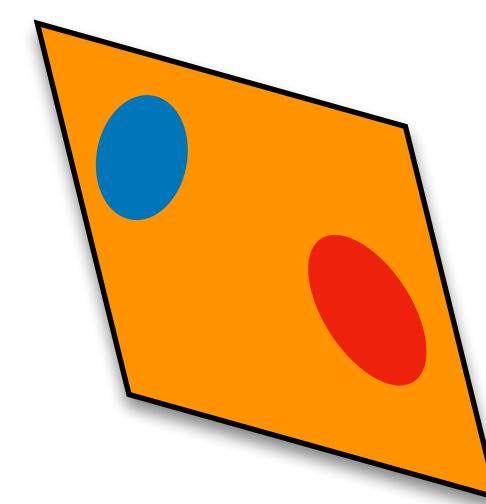
Differentiable Gaussian Rendering

What is the representation
of a 3D gaussian?



$$\begin{aligned}\mu_{2D} &= \pi(\mu_{3D}) \\ \Sigma_{2D} &= J \Sigma_{3D} J^T\end{aligned}$$

How to project to 2D
and rasterize?



How to model/aggregate
appearance?

1. Sort gaussians from closest to furthest
from the camera

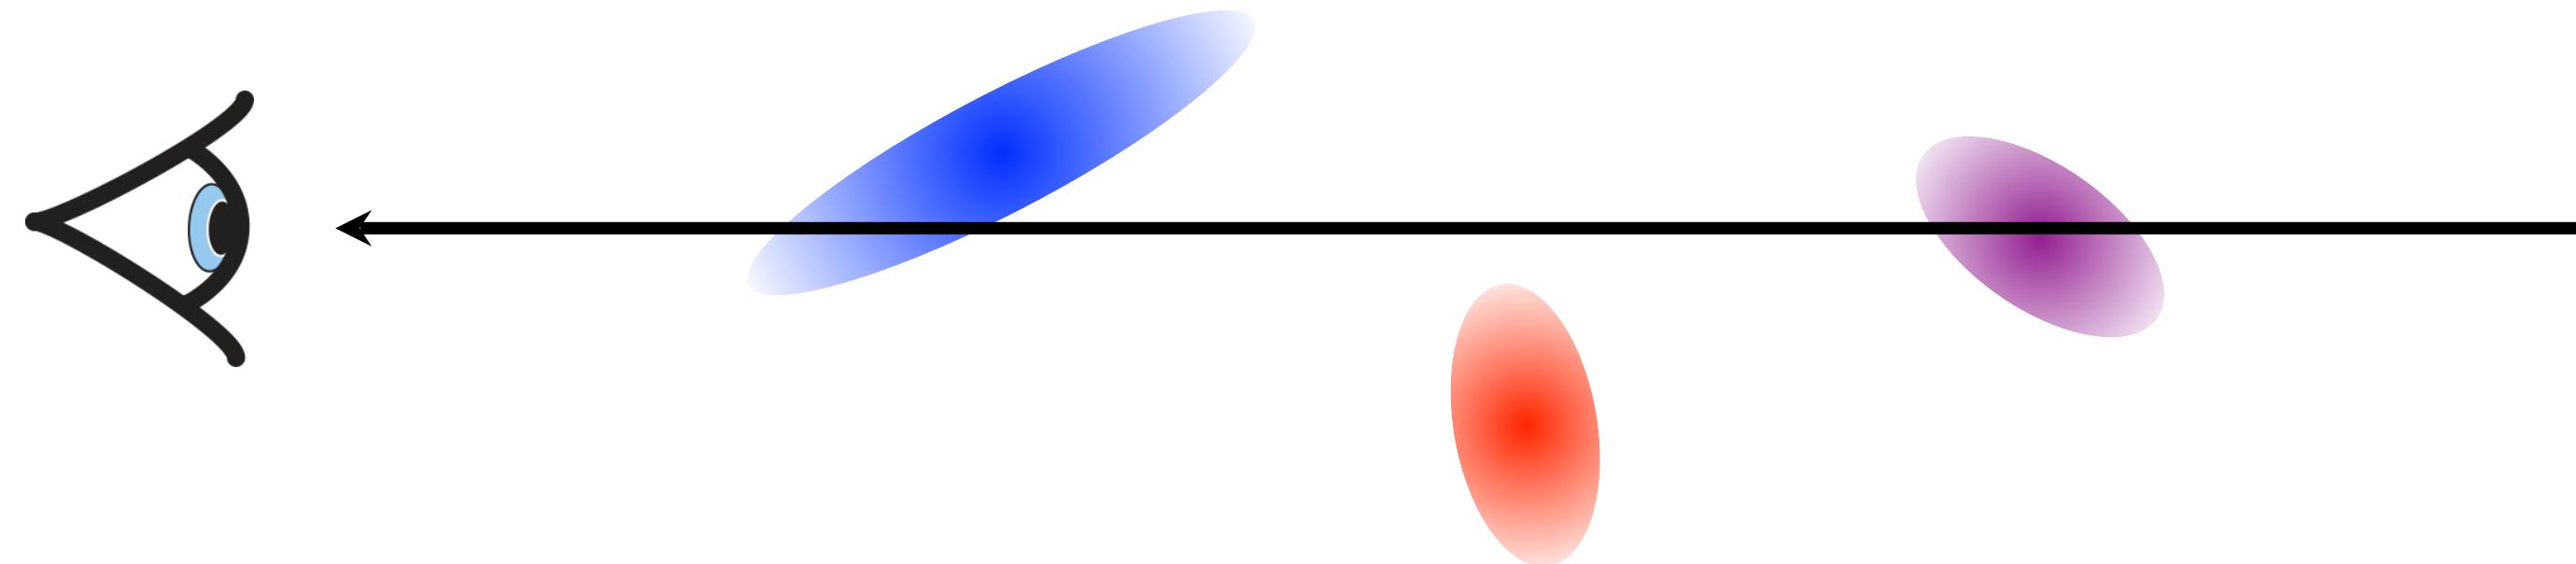
2. For each pixel \mathbf{u} , compute opacity for
each gaussian \mathcal{G}_k :

$$\bar{\alpha}_k = \alpha_k \frac{e^{-(\mathbf{u}-\mu_{2D}^k)^T (\Sigma_{2D}^k)^{-1} (\mathbf{u}-\mu_{2D}^k)}}{2\pi |\Sigma_{2D}^k|^{0.5}}$$

(In practice, can rasterize ‘blocks’ instead of entire image as not all
gaussians influence all blocks)

Differentiable Gaussian Rendering

What is the representation
of a 3D gaussian?



How to project to 2D
and rasterize?

How to model/aggregate
appearance?

Compute per-gaussian weights based on opacities of current and previous gaussians:

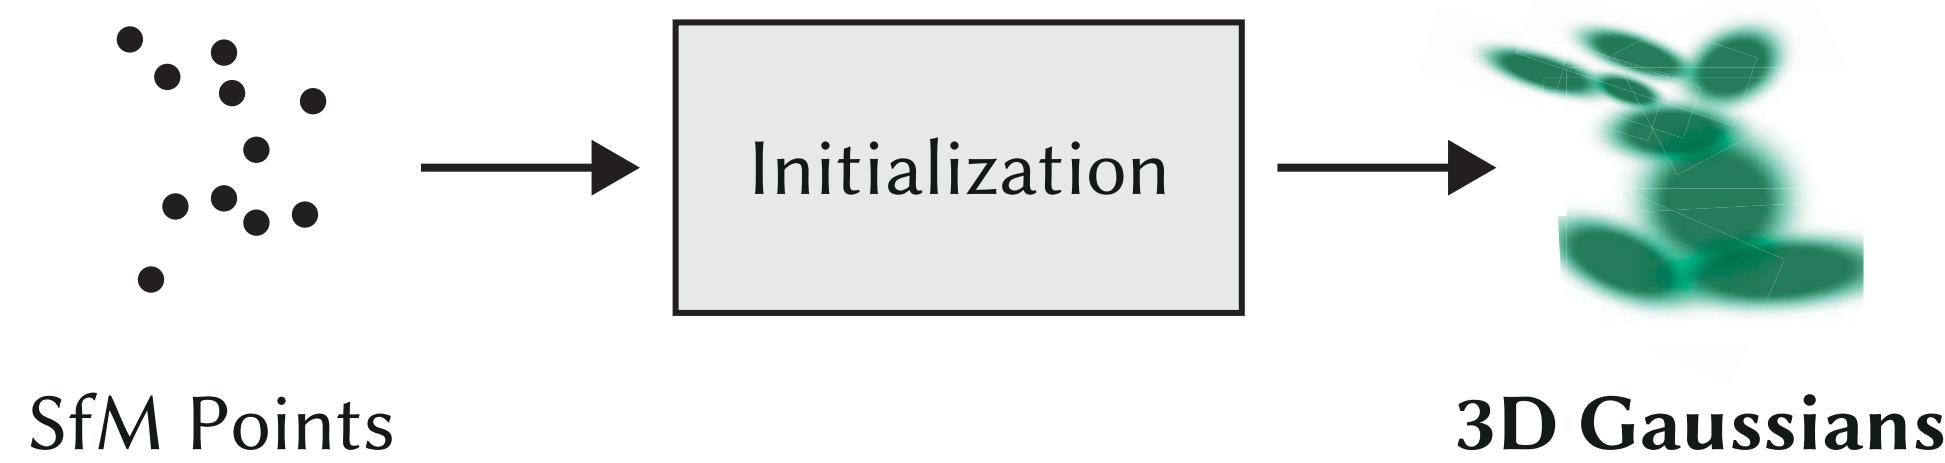
$$w_k = \bar{\alpha}_k \prod_{j=1}^{k-1} (1 - \bar{\alpha}_j)$$

Use per-gaussian SH coefficients and ray direction to get view-dependent color \mathbf{c}_k

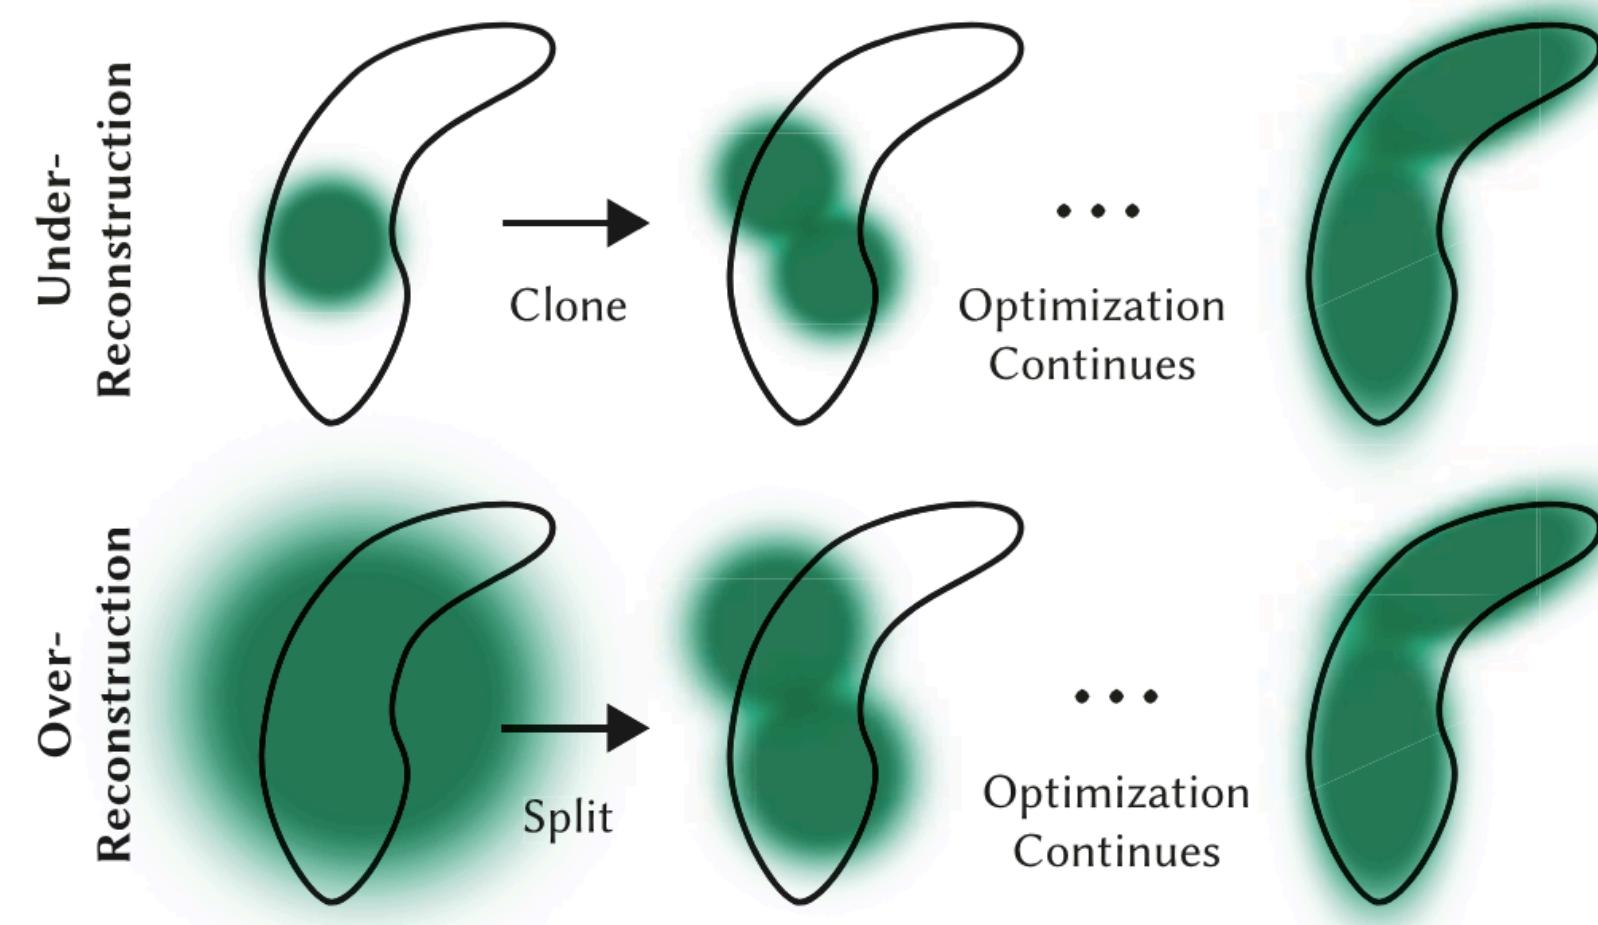
Aggregate to obtain pixel color:

$$\mathbf{c} = \sum_k w_k \mathbf{c}_k$$

Gaussian Splatting: Bells and Whistles



Initialize with sparse point cloud from SfM



Split/clone gaussians based on heuristics

Menu Views Capture

► 3D Gaussians

► Camera Point view



▼ Metrics

57.56 (17.37 ms)

VSync On





▼ Metrics
89.37 (11.19 ms)





This video contains a voice-over

3D Gaussian Splatting for Real-Time Radiance Field Rendering

SIGGRAPH 2023
(ACM Transactions on Graphics)

Bernhard Kerbl*



Georgios Kopanas*



Thomas Leimkühler



max planck institut
informatik

George Drettakis



* Denotes equal contribution

Bonus: NeRF Song



Questions?